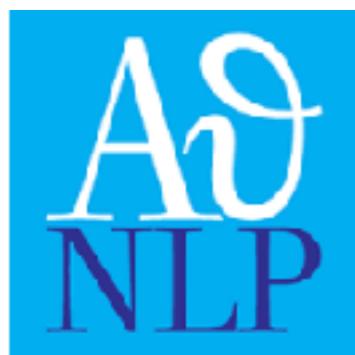
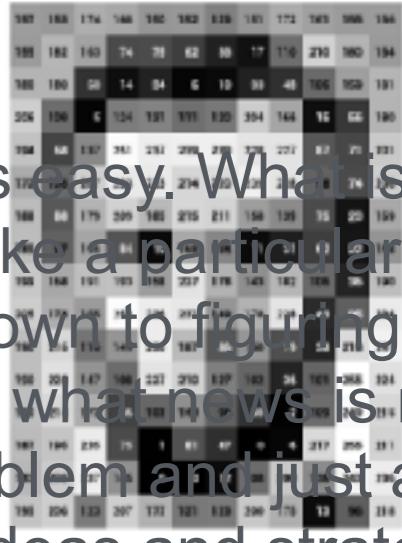


# Encoder-Decoder Models

Barbara Plank  
ITU, Copenhagen, Denmark



September 21, 2019,  
#AthNLP2019 Athens, Greece



## Motivation

chicken flock's market is down for participants say the unflocks are fine. In fact, the actual price of

eggs in the country's market has fallen for the only other Republiagan government to have

competitive reelection races taking place according to Cook Re

is Pat McCrory 8 who hasn't made a p

statistics. However, there are 34 Republican

members of the House who face competitive races, according to Cook The

one to have endorsed Trump is Tom

the incumbent from New York's Congressional District, a Republi

Read more: Stocks Basics: What Causes Stock Prices To Change? | Investopedia http://

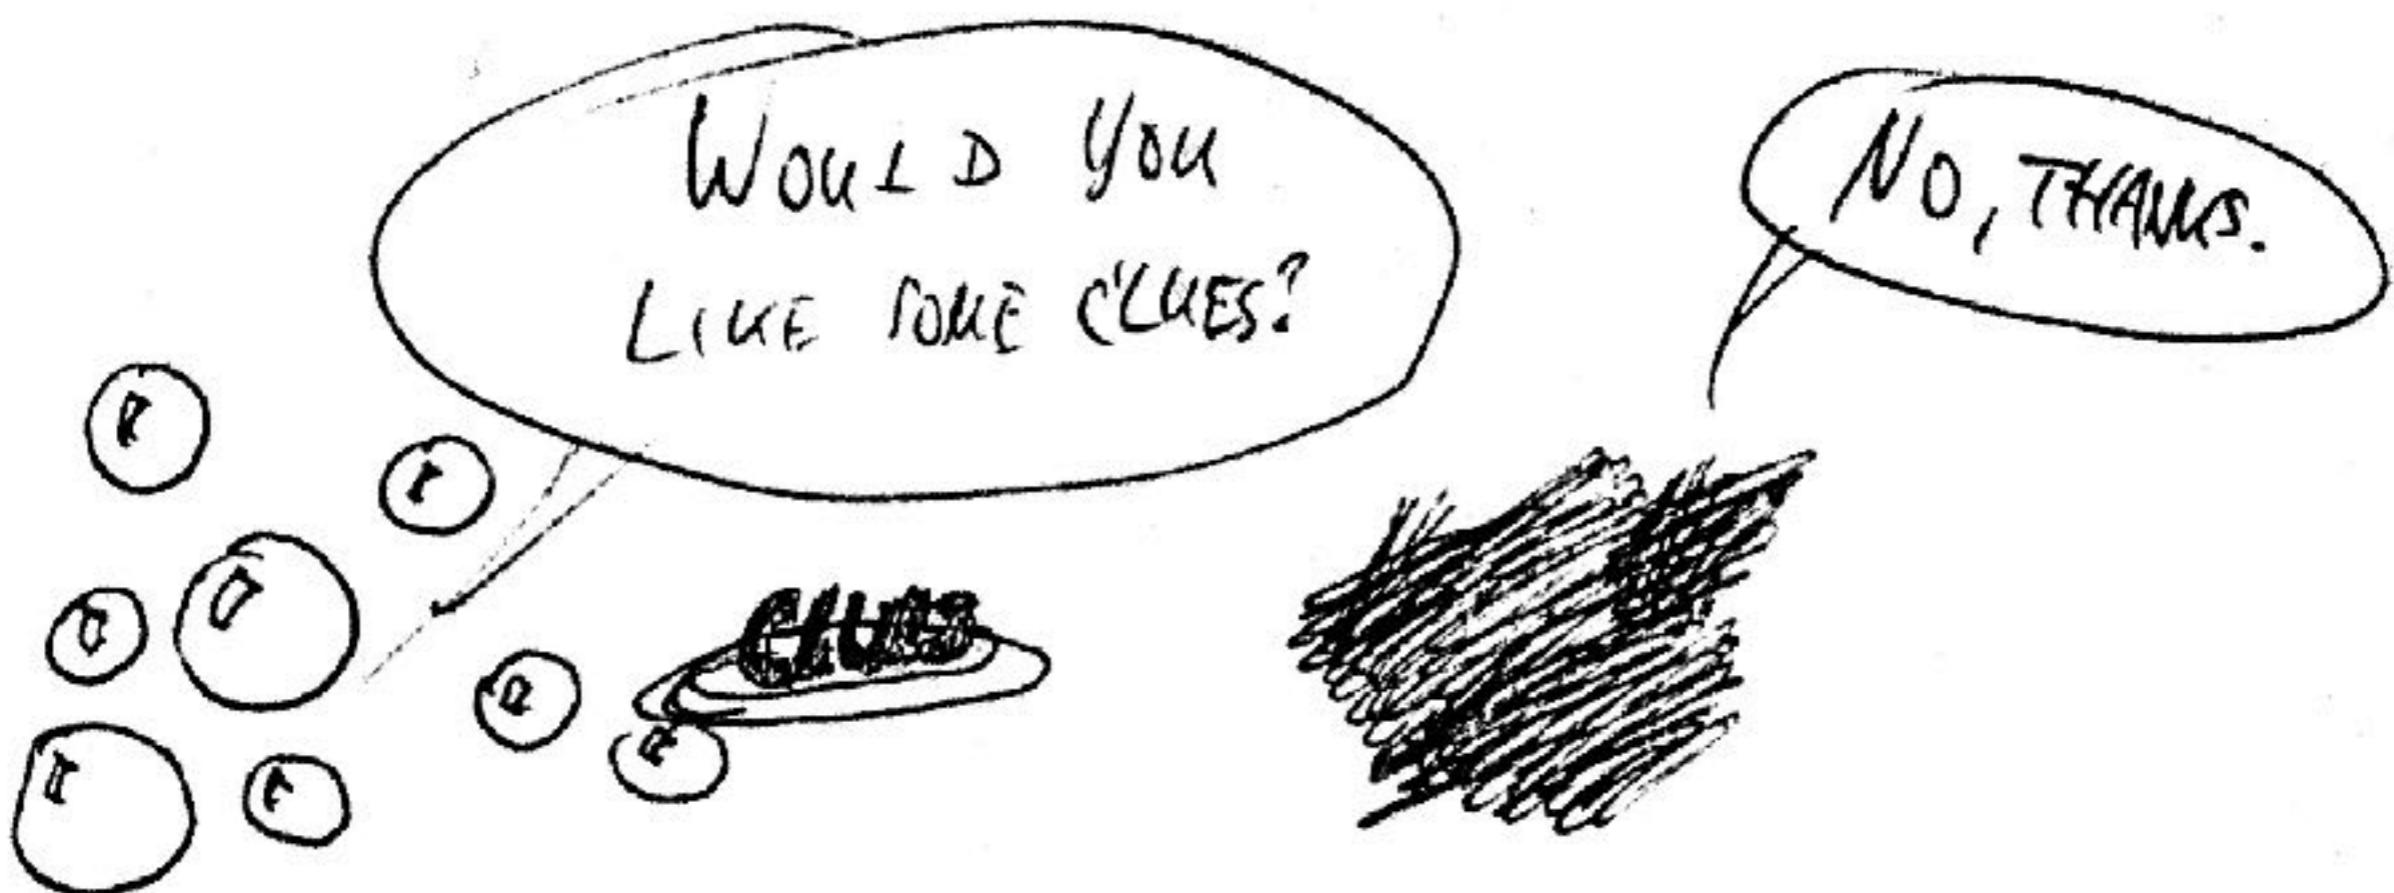
Inspired by slides from S.Gouws & D. Hovy.

2

# Challenge: Language is ambiguous

Galactic bubbles offer clues to **dark matter**

Space Daily - 22 hours ago



# Challenge: Language is productive

How to sunny-day Saturday in Seattle:

- ✓ pop out of bed and fling open the drapes
- ✓ brew coffee ☕ and grab your **go-cup**
- ✓ get outside asap
- ✓ dog walk, hike, run, bike, kayak, sail
- ✓ 🍺 soak up the sun in your favorite beer garden
- ✓ 😊 👍 ✓



2



22



# Sequence Modeling Tasks

- In some applications, we want to **condition** on sequential data to make a prediction

*AUTHOR ATTRIBUTES*

Female

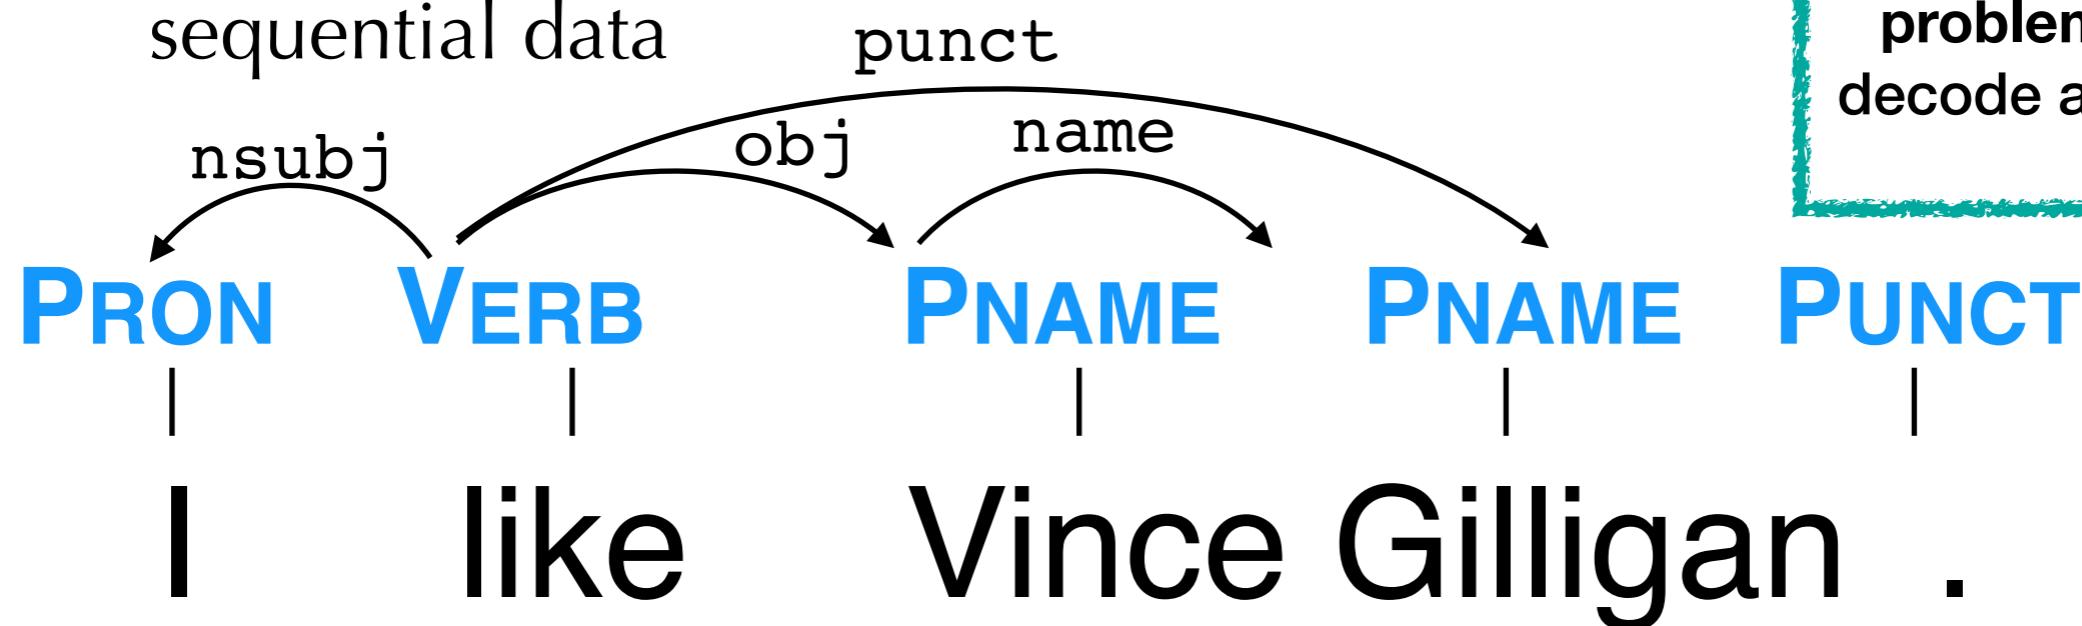
30+

**condition problems:**  
how to represent a  
sequence?

*SENTIMENT ANALYSIS*

positive

- In other applications, we want to **generate** sequential data



**generation**  
problems: how to  
decode a sequence?

**A step back...  
How did the field evolve?**

NLP



# Machine Learning

- ▶ Early approaches in NLP: **symbolic & rule-based**
- ▶ In the late 80s: development of annotated corpora (especially the well-known Penn Treebank Wall Street Journal)

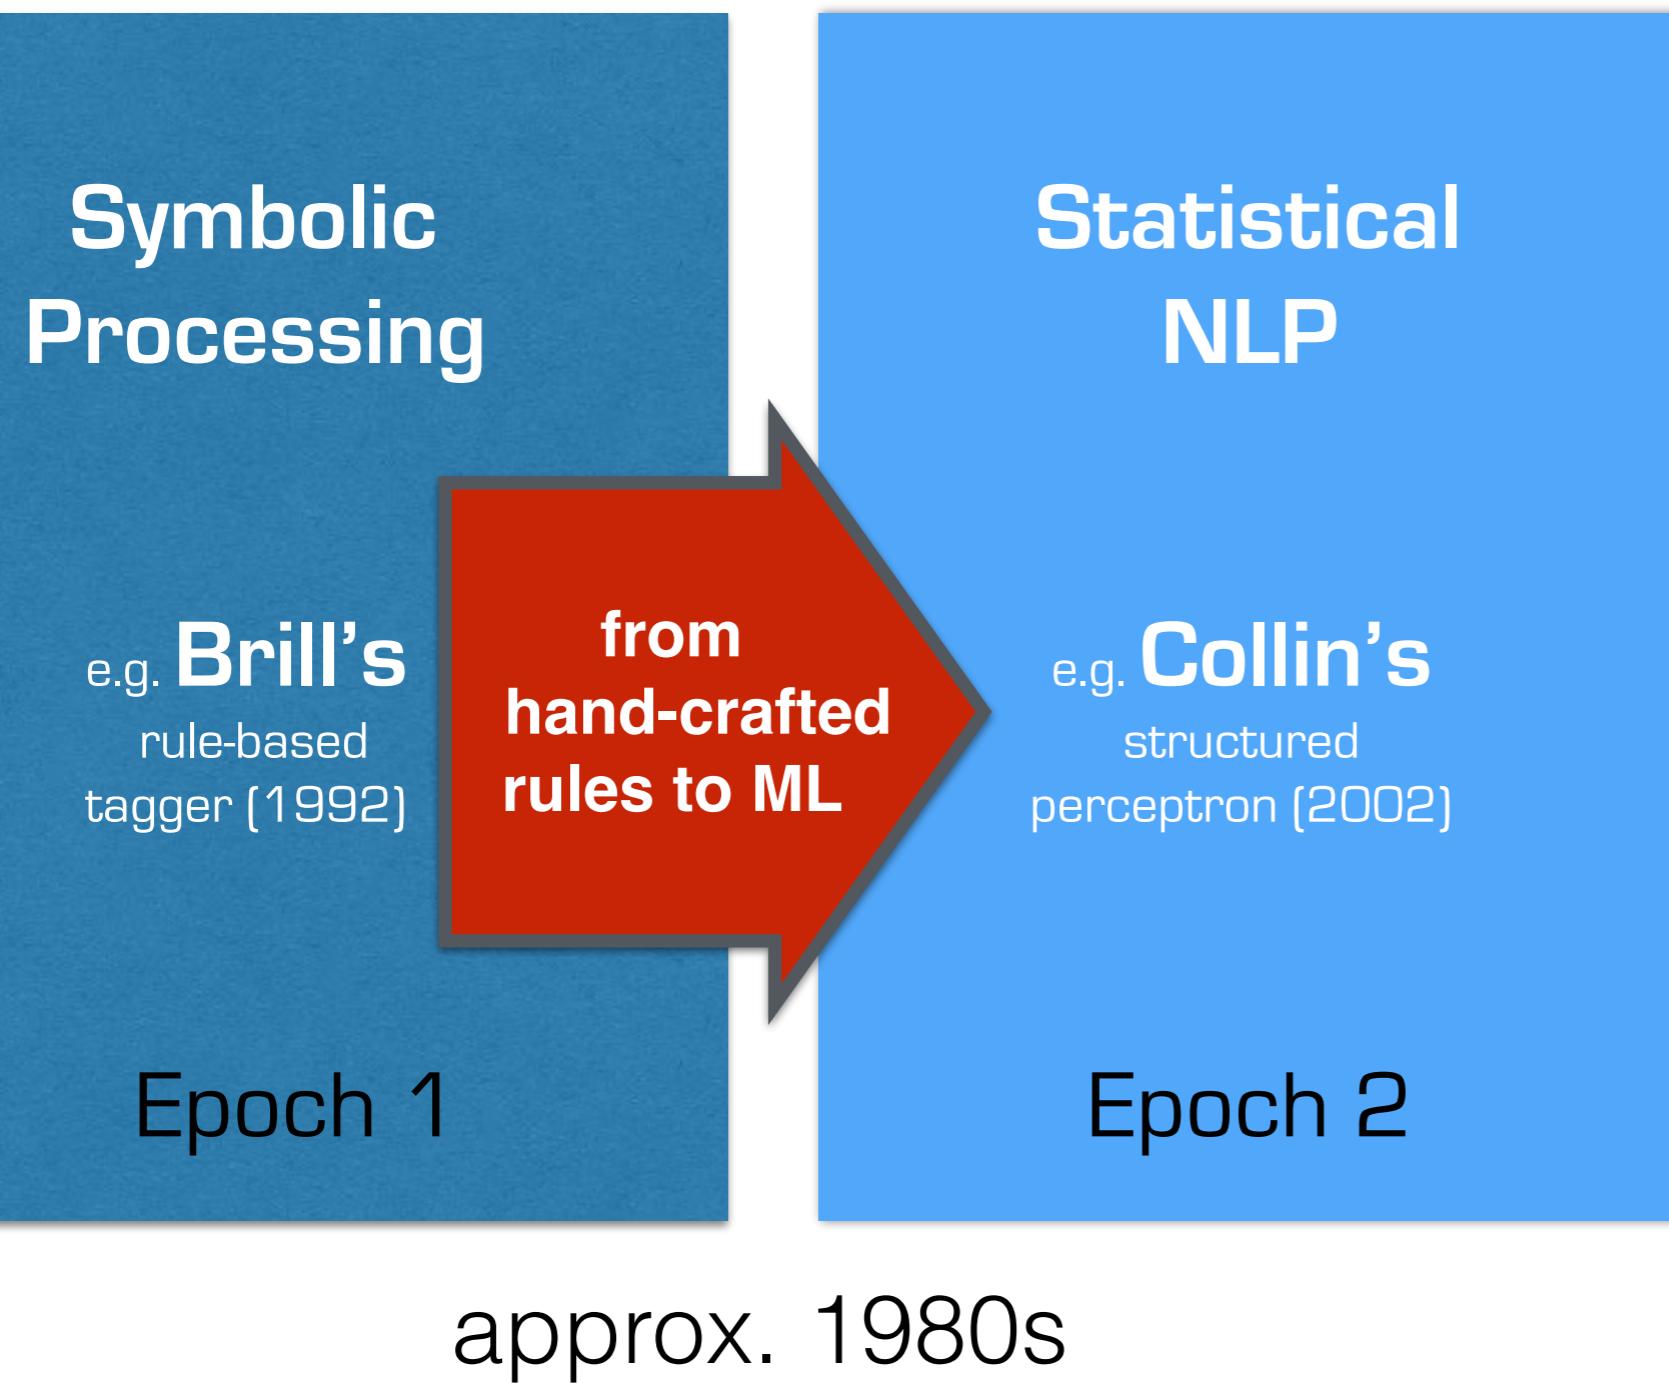


- ▶ And corresponding emergence of **statistical approaches**
- ▶ Common evaluation corpora and measures pushed the field

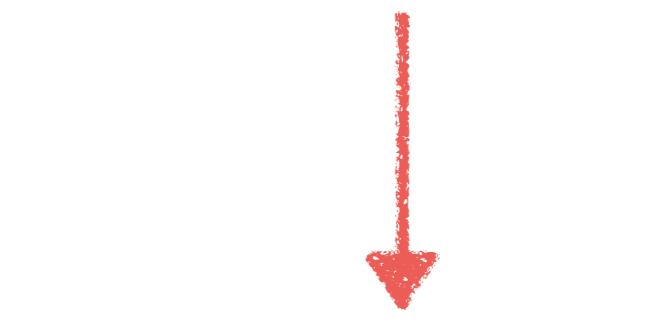
80s



# First big jump: statistical learning



**x:** the dog barks  
if prev\_w = DET and ...  
tag=NOUN



**x:**

1	0	..	0	1	0
---	---	----	---	---	---

$w_i=\text{dog}$        $w_{i-1}=\text{the}$

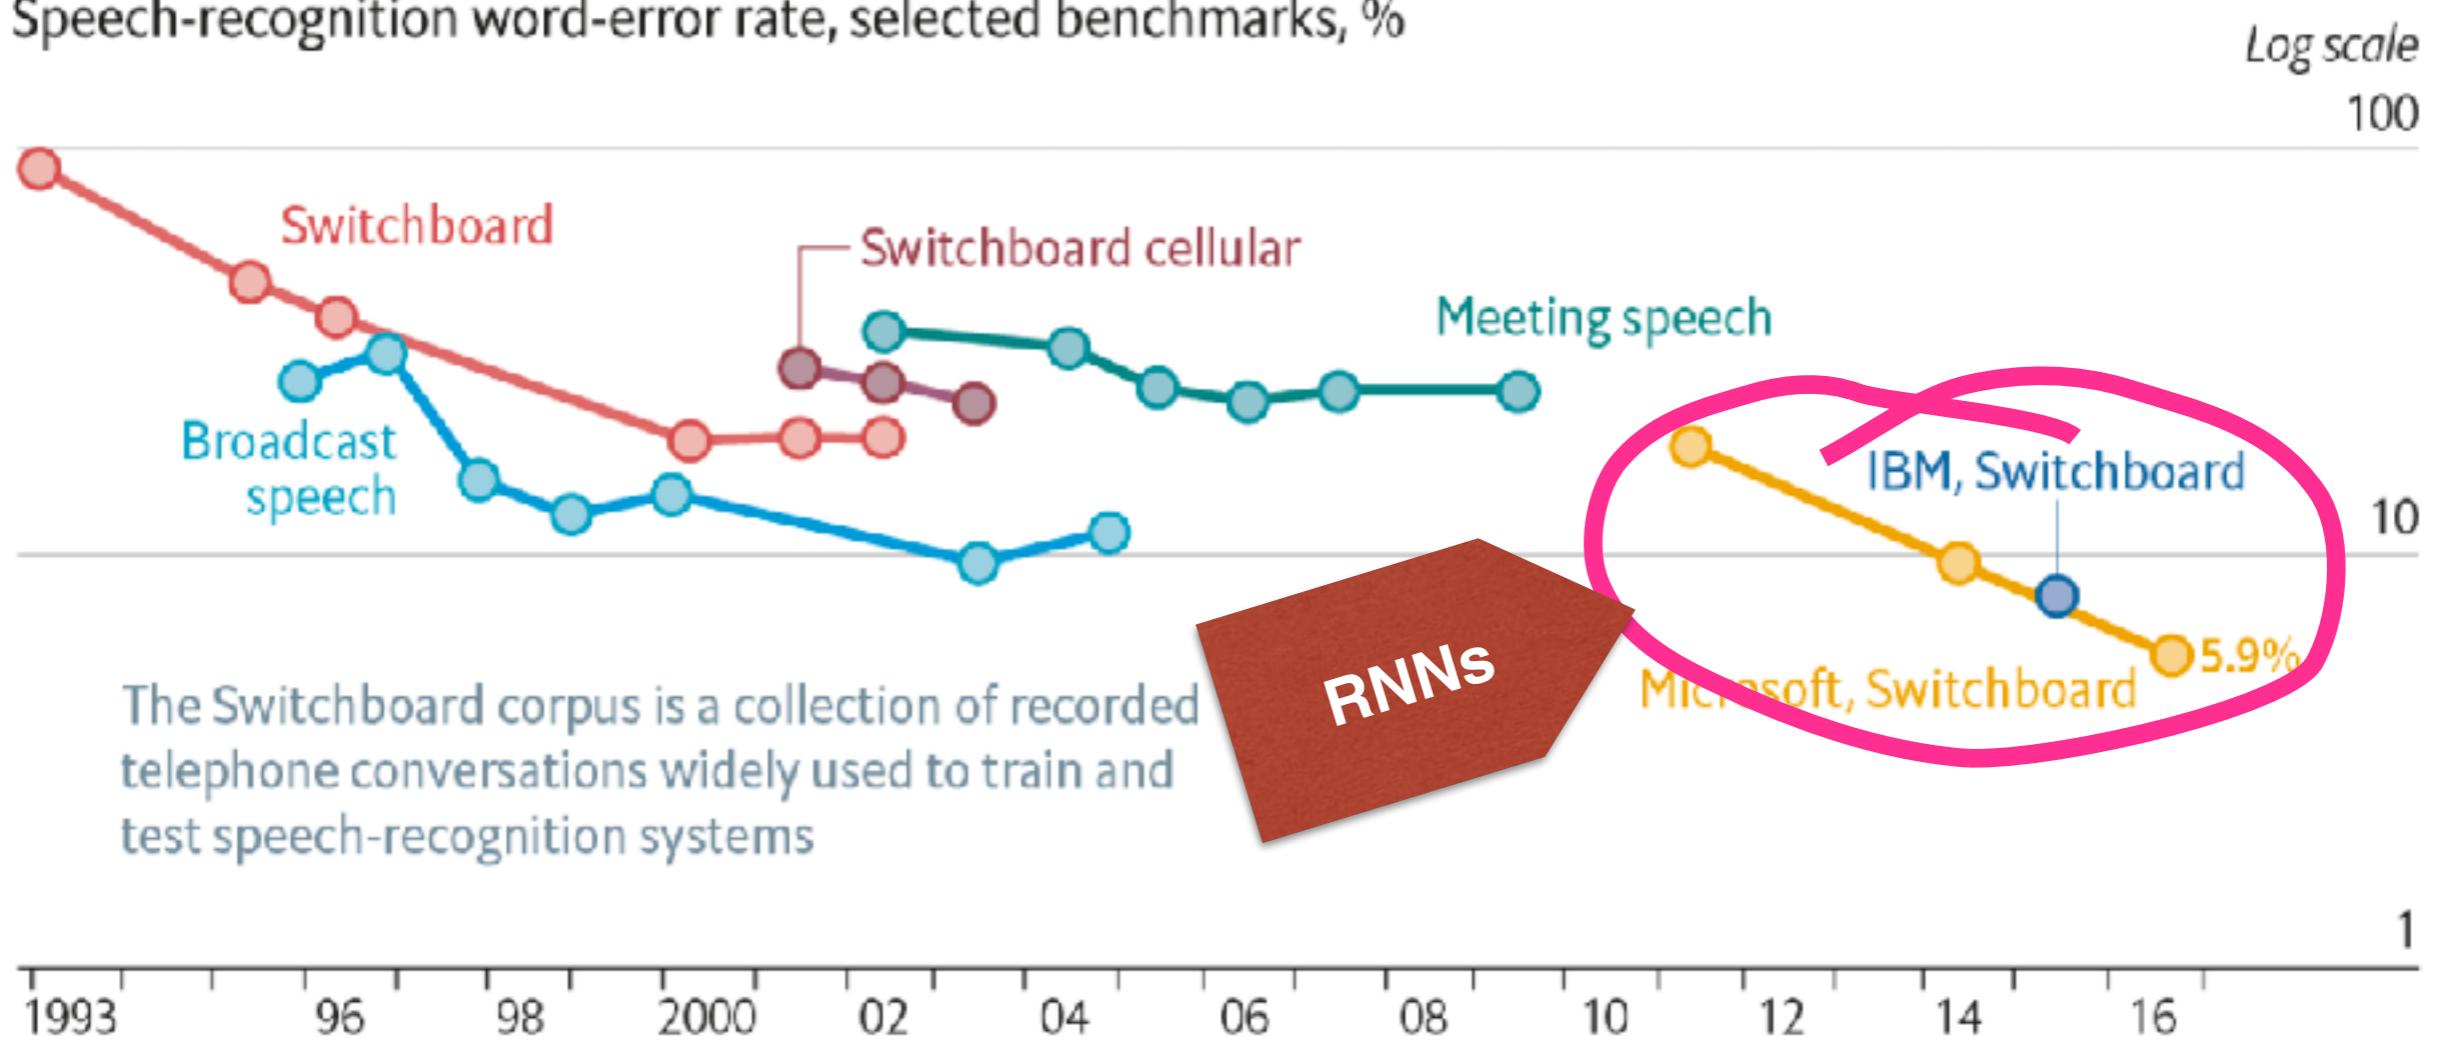
classic **sparse** 'n-hot' encoding

# **The emergence of deep learning (in NLP)**

# In Speech Recognition

## Loud and clear

Speech-recognition word-error rate, selected benchmarks, %



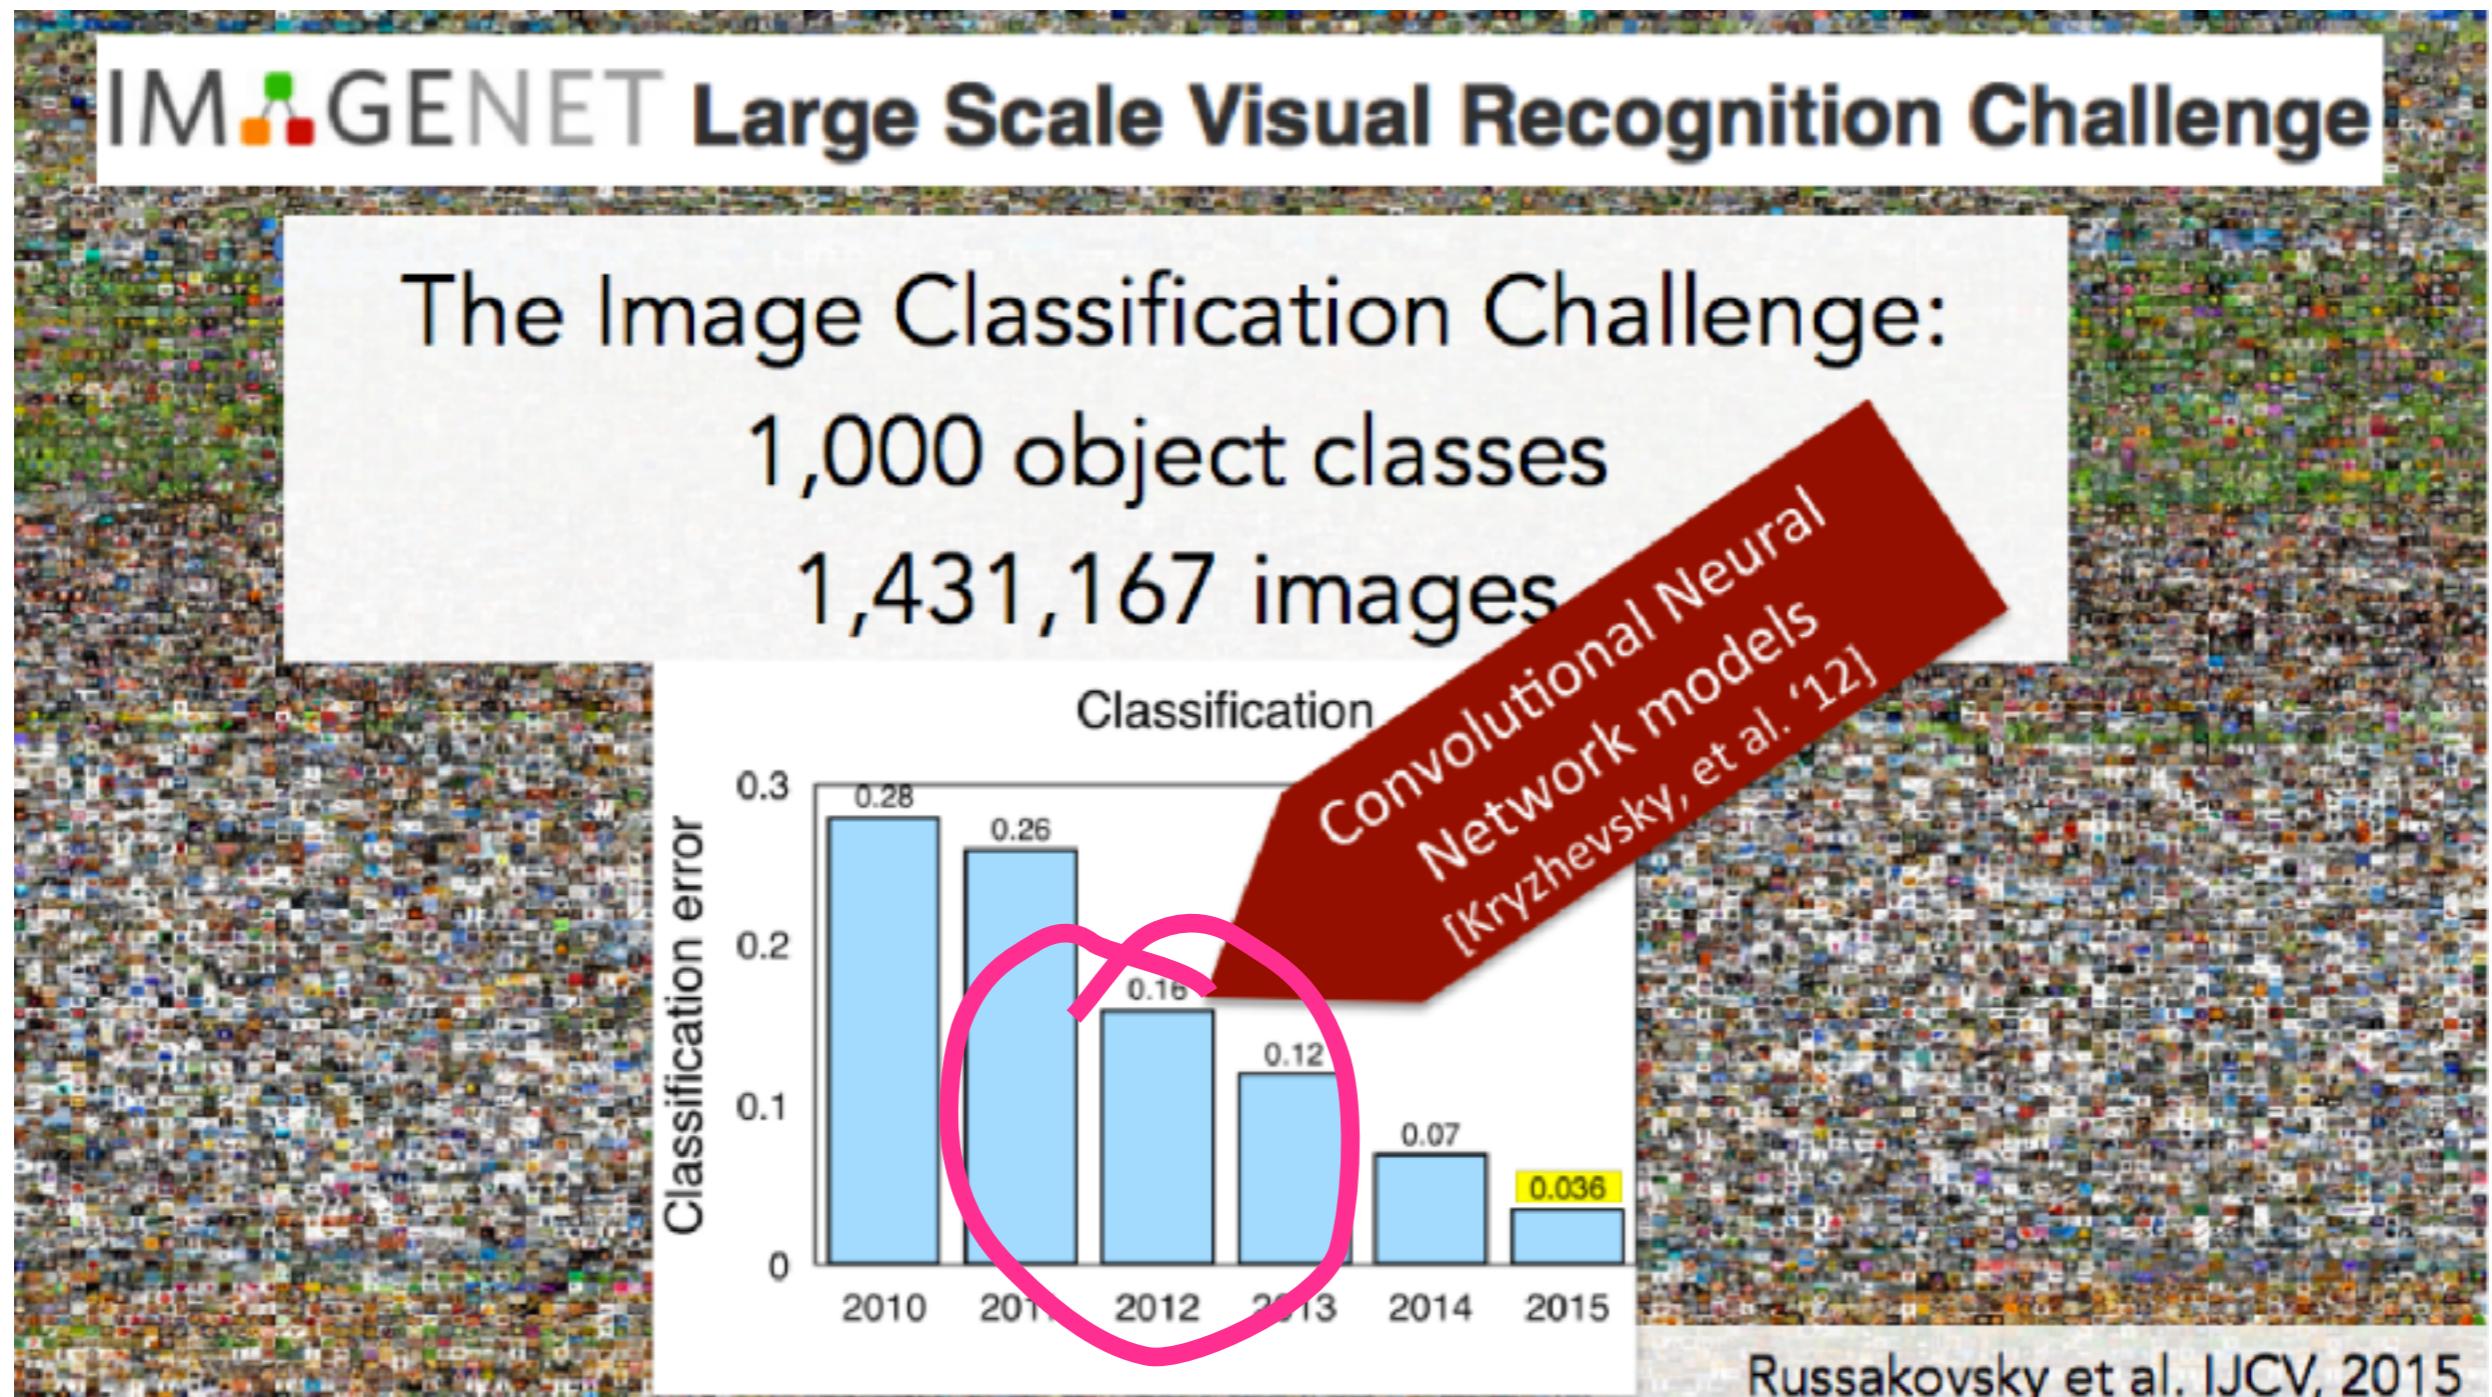
Sources: Microsoft; research papers

(Source: The Economist)

2010



# In Computer Vision



(src: slide by Fei-Fei Li)

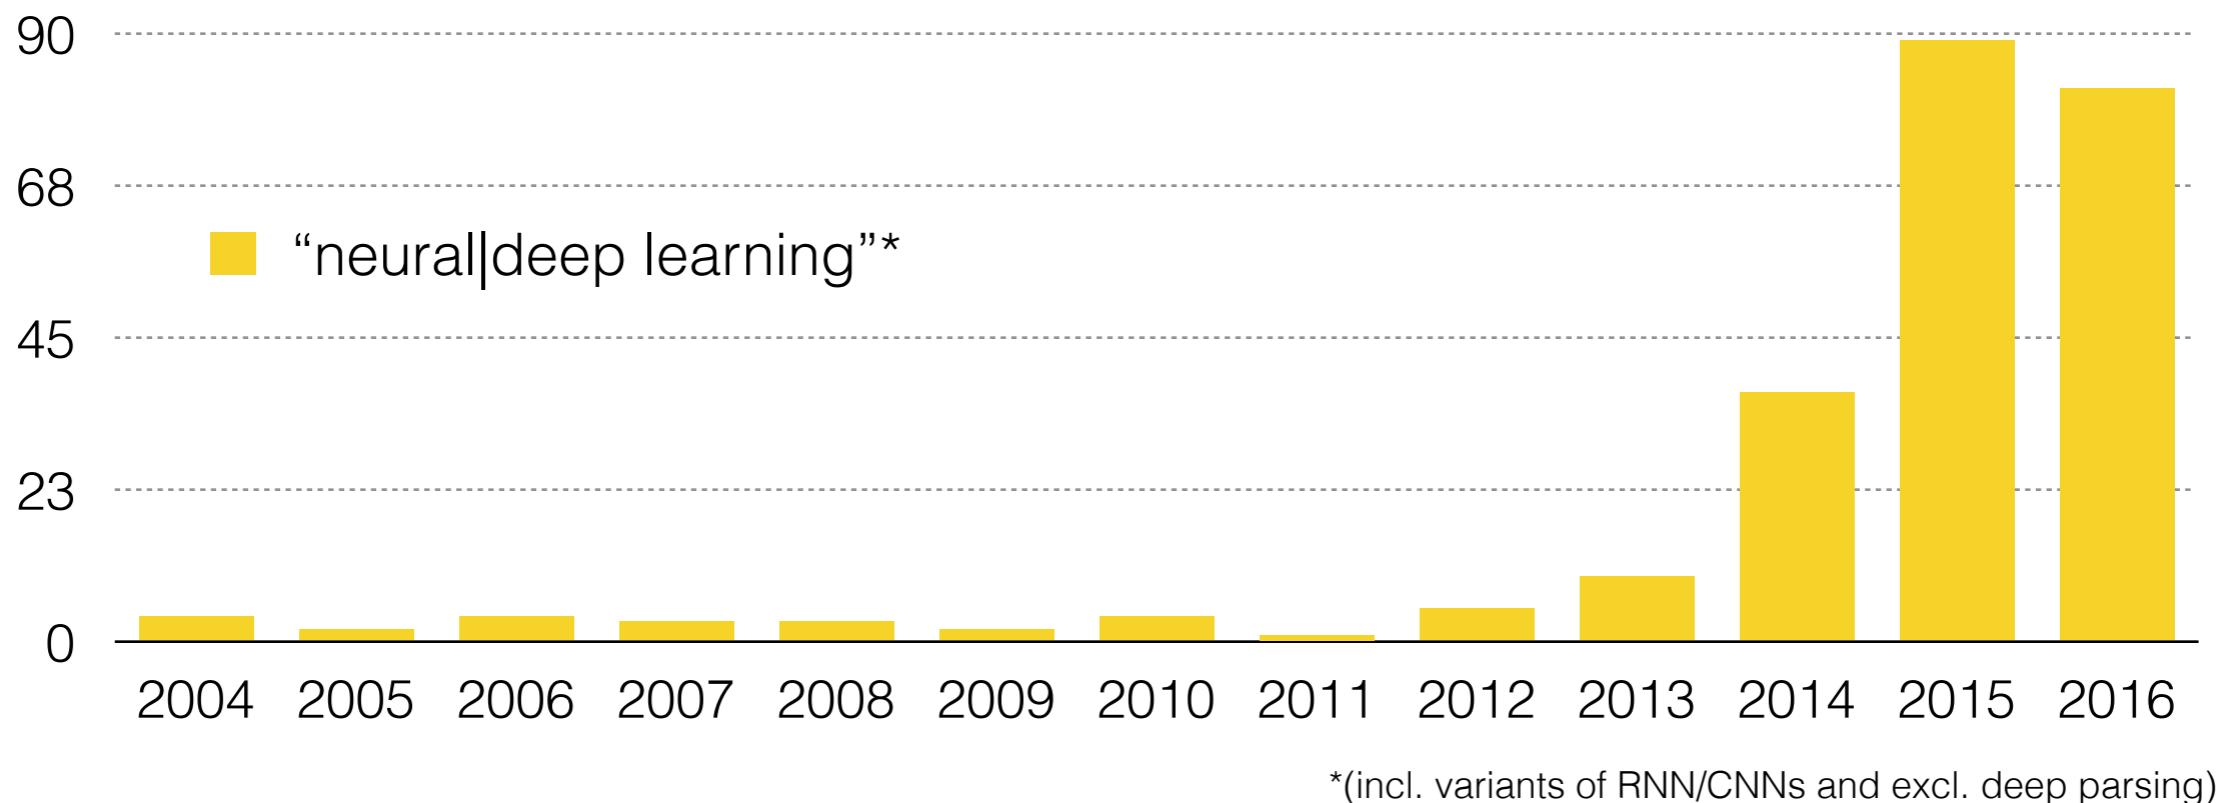
2012



# Papers: Deep learning in NLP

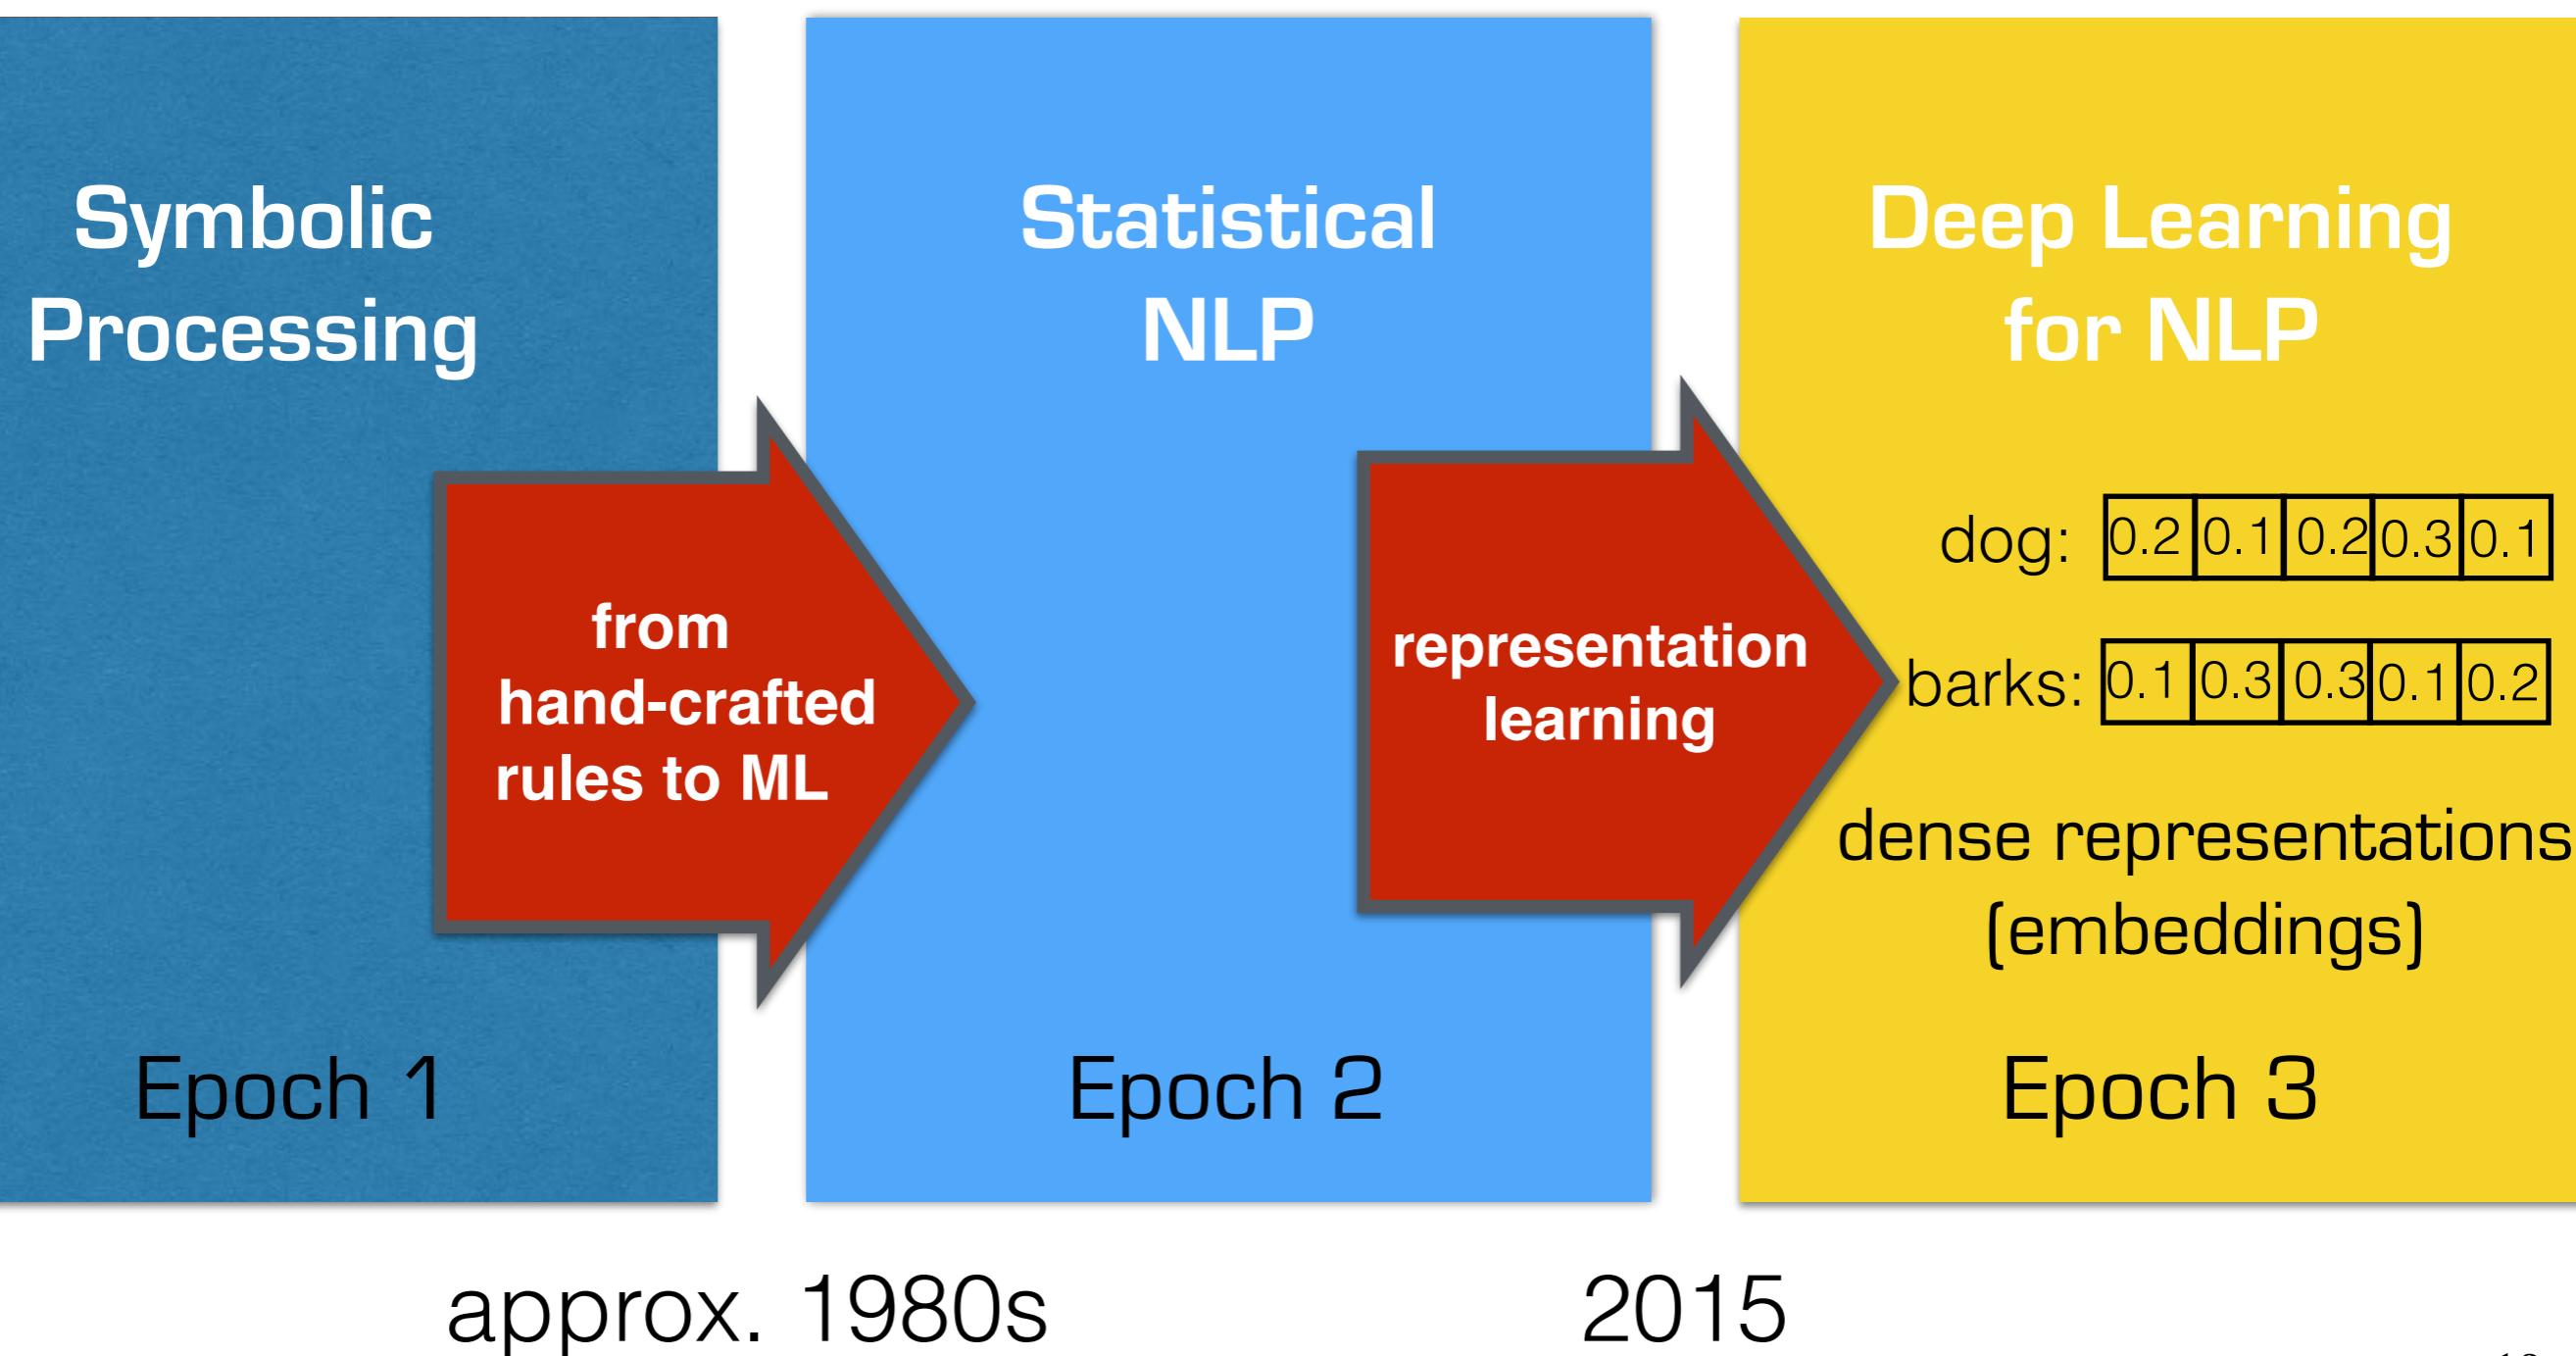


**“2015 seems like the year when the full force of the tsunami hit the major NLP conferences”**  
—Chris Manning (2015)



**Titles of papers in ACL anthology (from 2004)**

# NLP ❤ Deep Learning



# Overview

*foundations*



Motivation,  
Brief History, Overview



Back to the roots:  
Language Models



Feedforward NNs  
(FFNNs)

*representations*



What's the input?  
Representations



Contextualized  
Embeddings, Fine-tuning

*beyond FFNNs*

Convolutional Neural  
Networks (CNNs)



Recurrent Neural  
Network (RNNs)

Advanced RNNs,  
Decoders

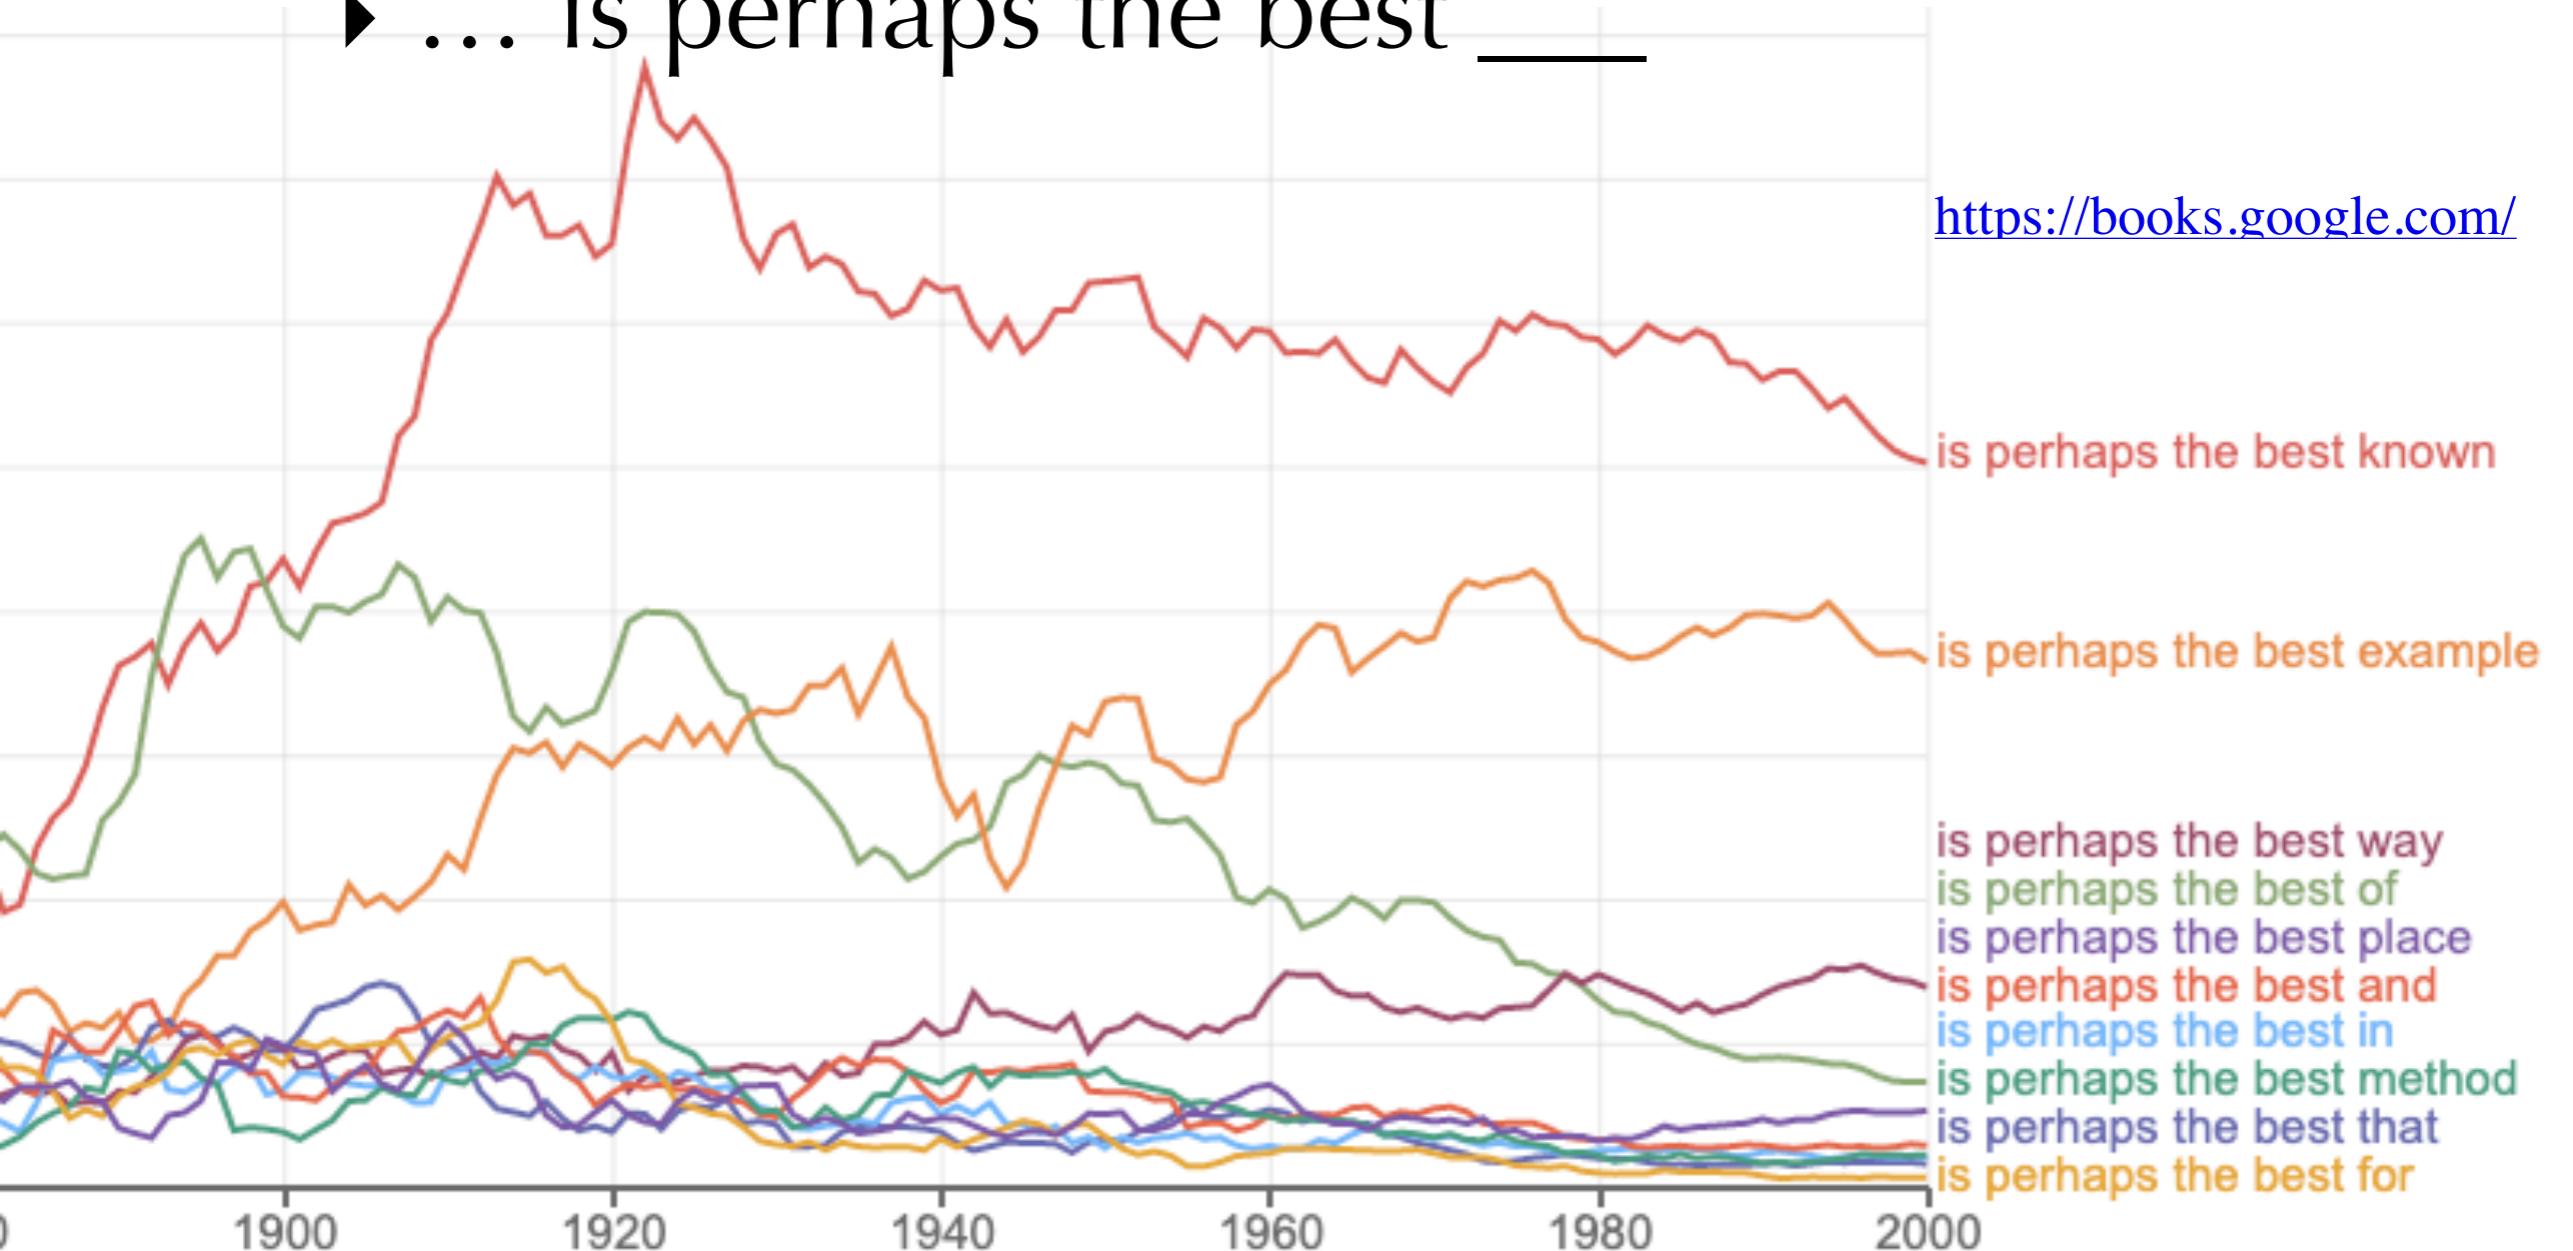
Attention? Attention!

# Predicting the next word: A Simple (?) Exercise

- ▶ [www.mentimeter.com](https://www.mentimeter.com)  
Room: (see code)

# More examples

- ▶ Recurrent Neural \_\_
- ▶ ... is perhaps the best \_\_

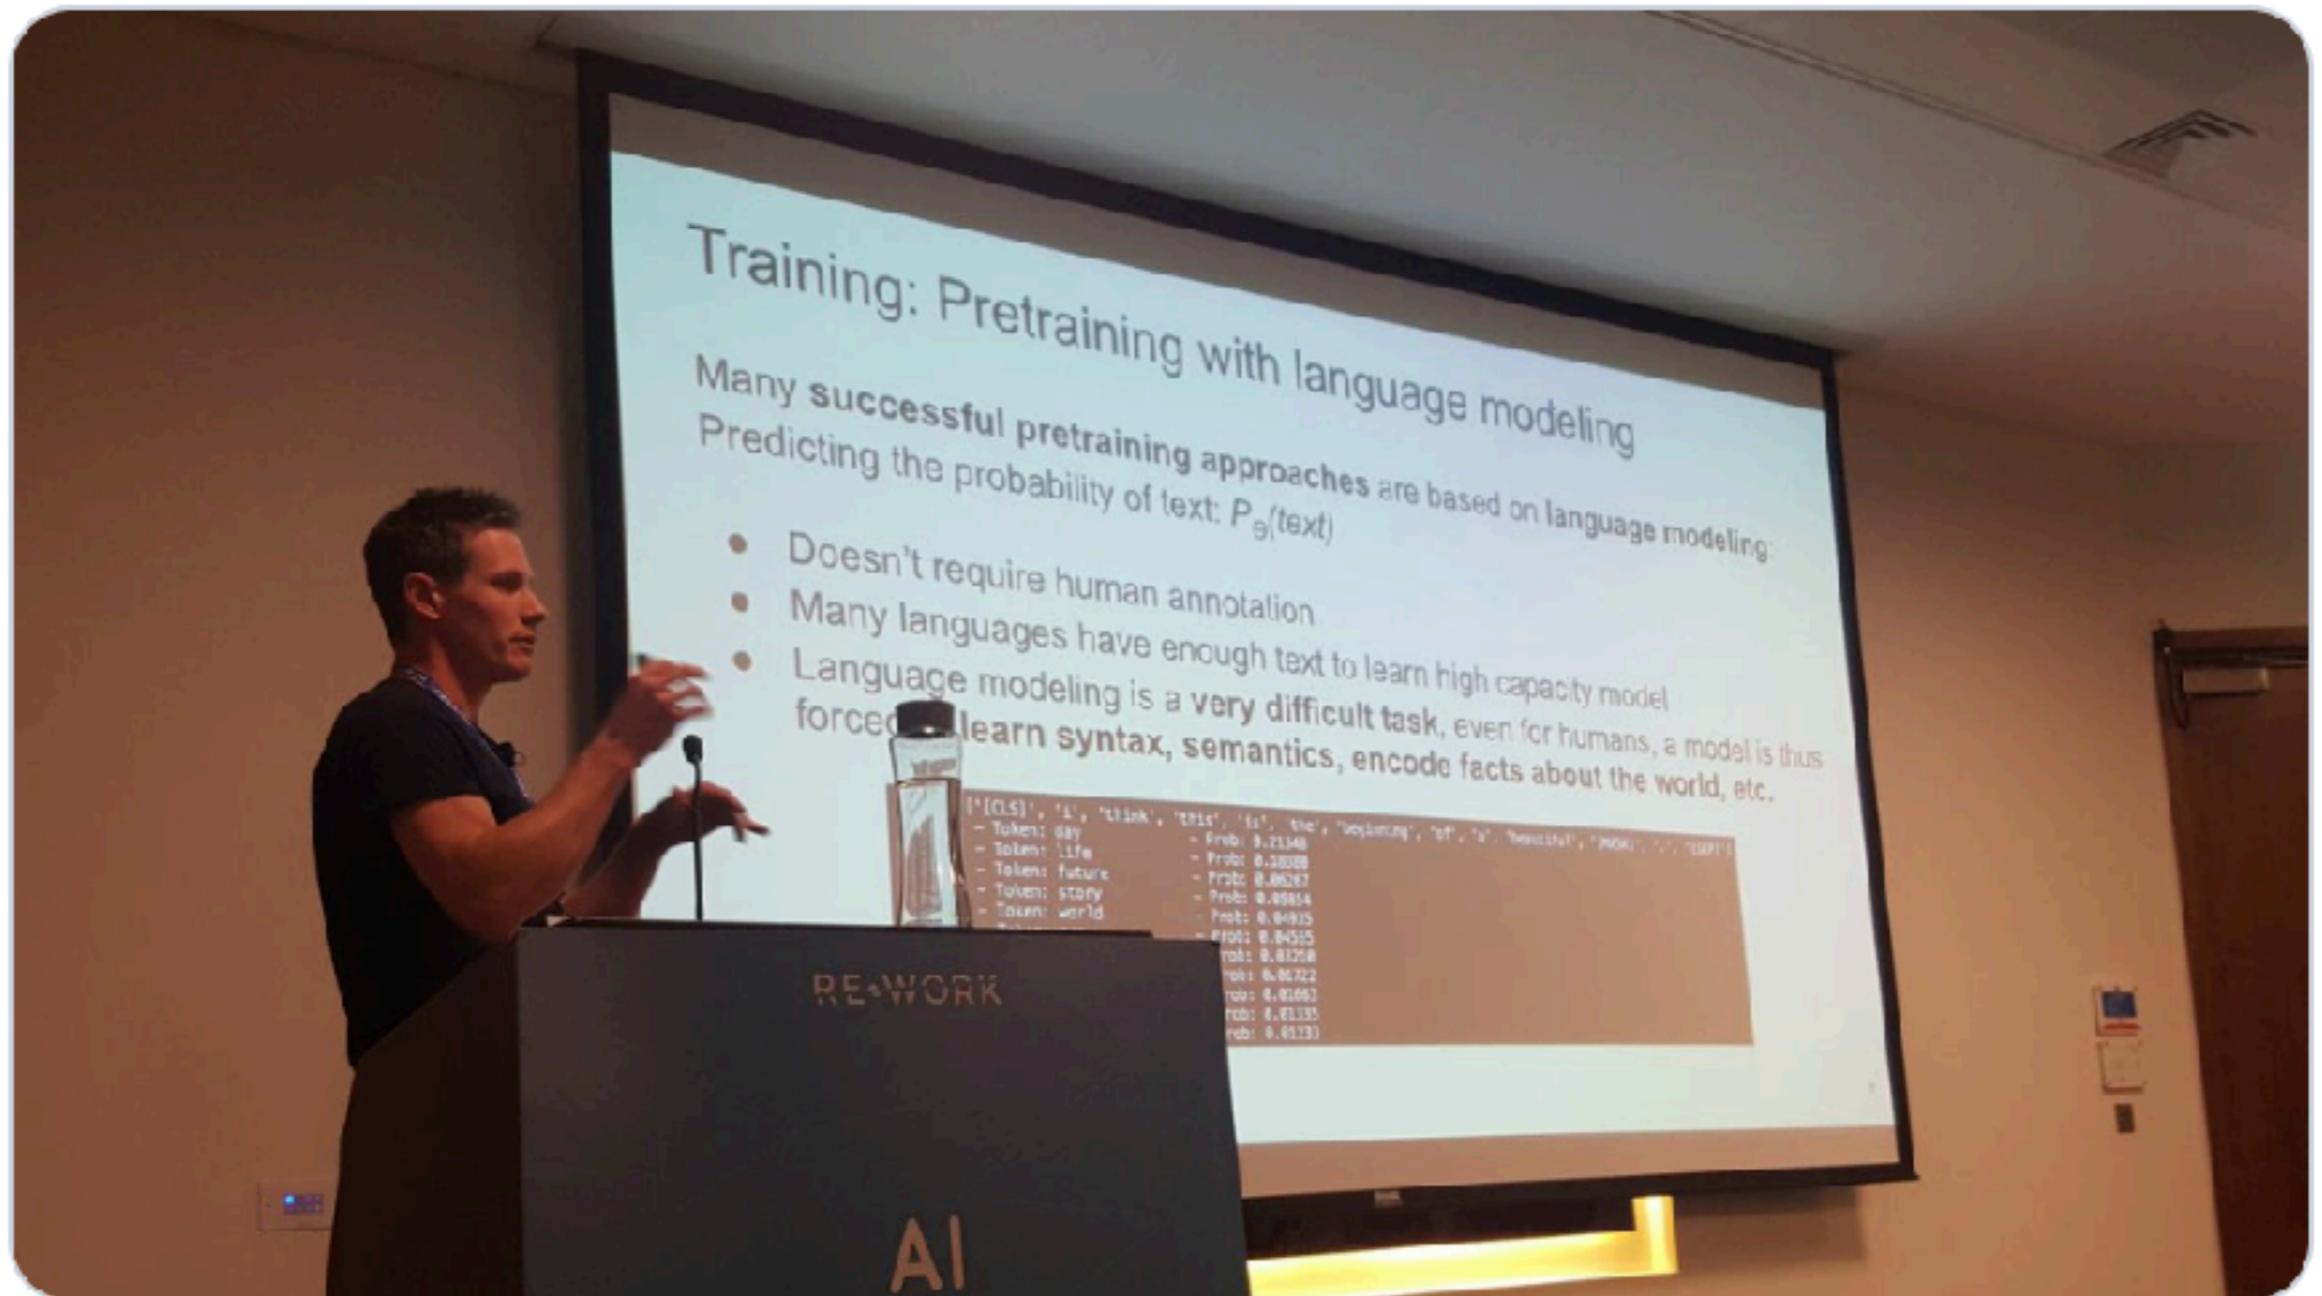


# Why care about LMs?



**Verena Rieser** @verena\_rieser · 3m

Great talk by **@Thom\_Wolf** on transfer learning **@reworkAI** !



[https://twitter.com/verena\\_rieser/status/1174694748310953984](https://twitter.com/verena_rieser/status/1174694748310953984)



**So let's look deeper at LMs:  
from traditional LMs to  
contextualized embeddings**

# What is a Language Model (LM)?

- ▶ A computational model that can be used to either of the following two tasks is called a Language Model (LM):
  - ▶ to compute the probability of a text\*

$$P(\text{today is a great day}) = ??$$

- ▶ to compute the probability of the next word

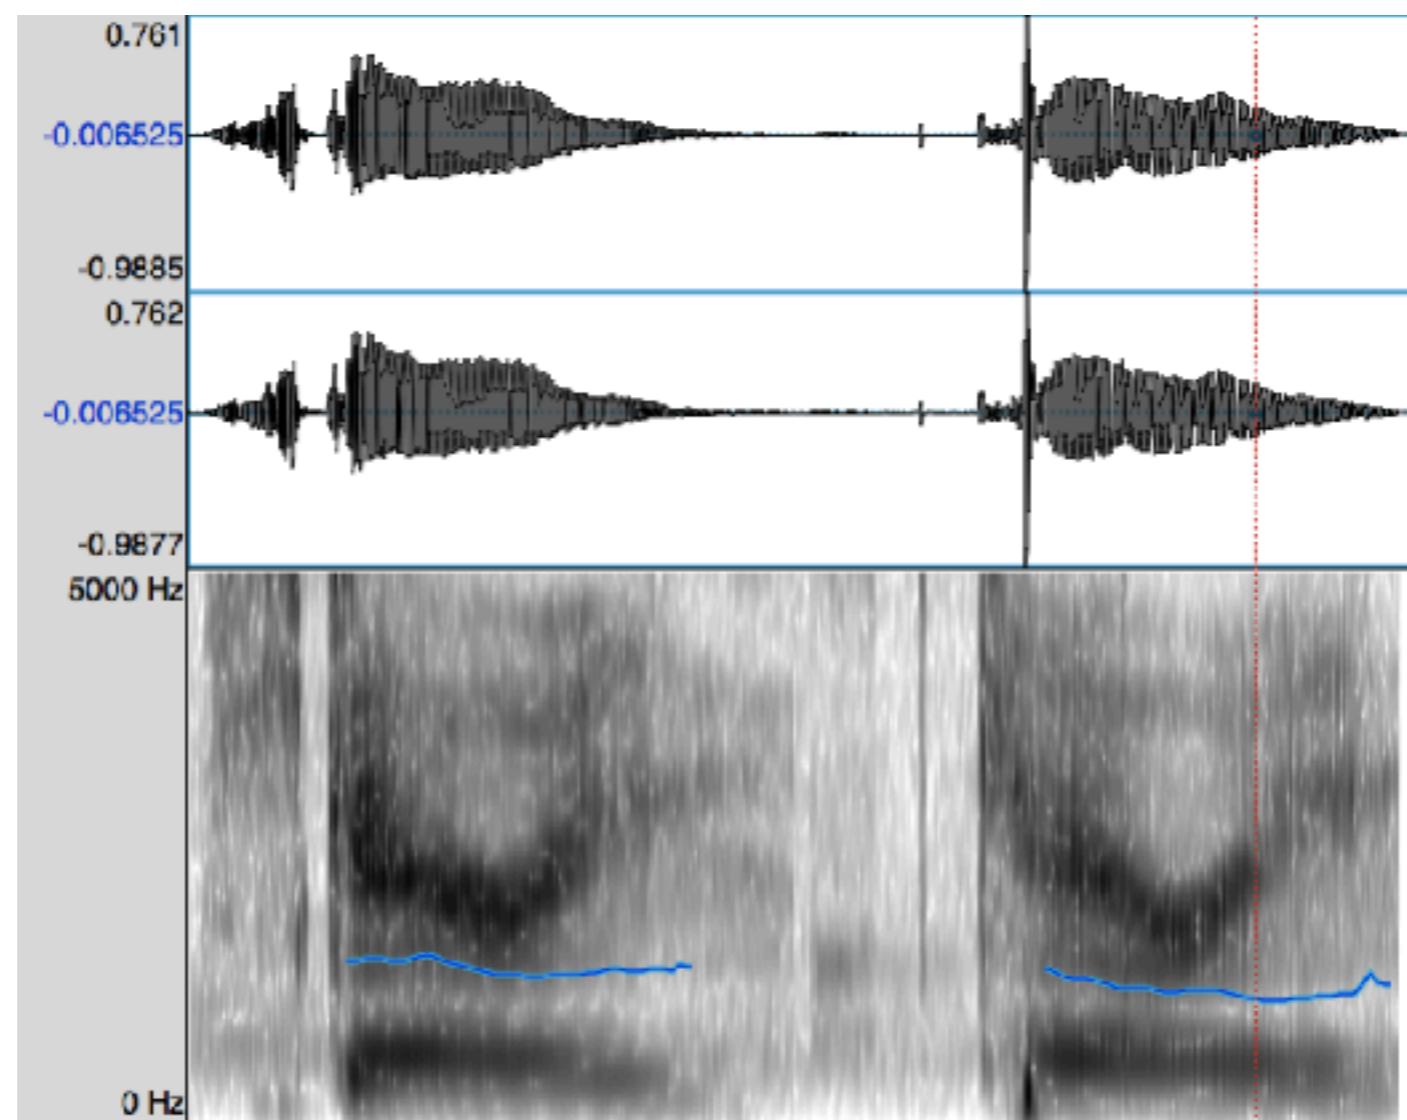
$$P(\text{day} \mid \text{today is a great}) = ??$$

\* (can be a text, sentence, phrase,...)

# **Why? Example Use Cases**

# Speech Recognition

$P(\text{where is the nearest beach}) > P(\text{where is the nearest breach})$



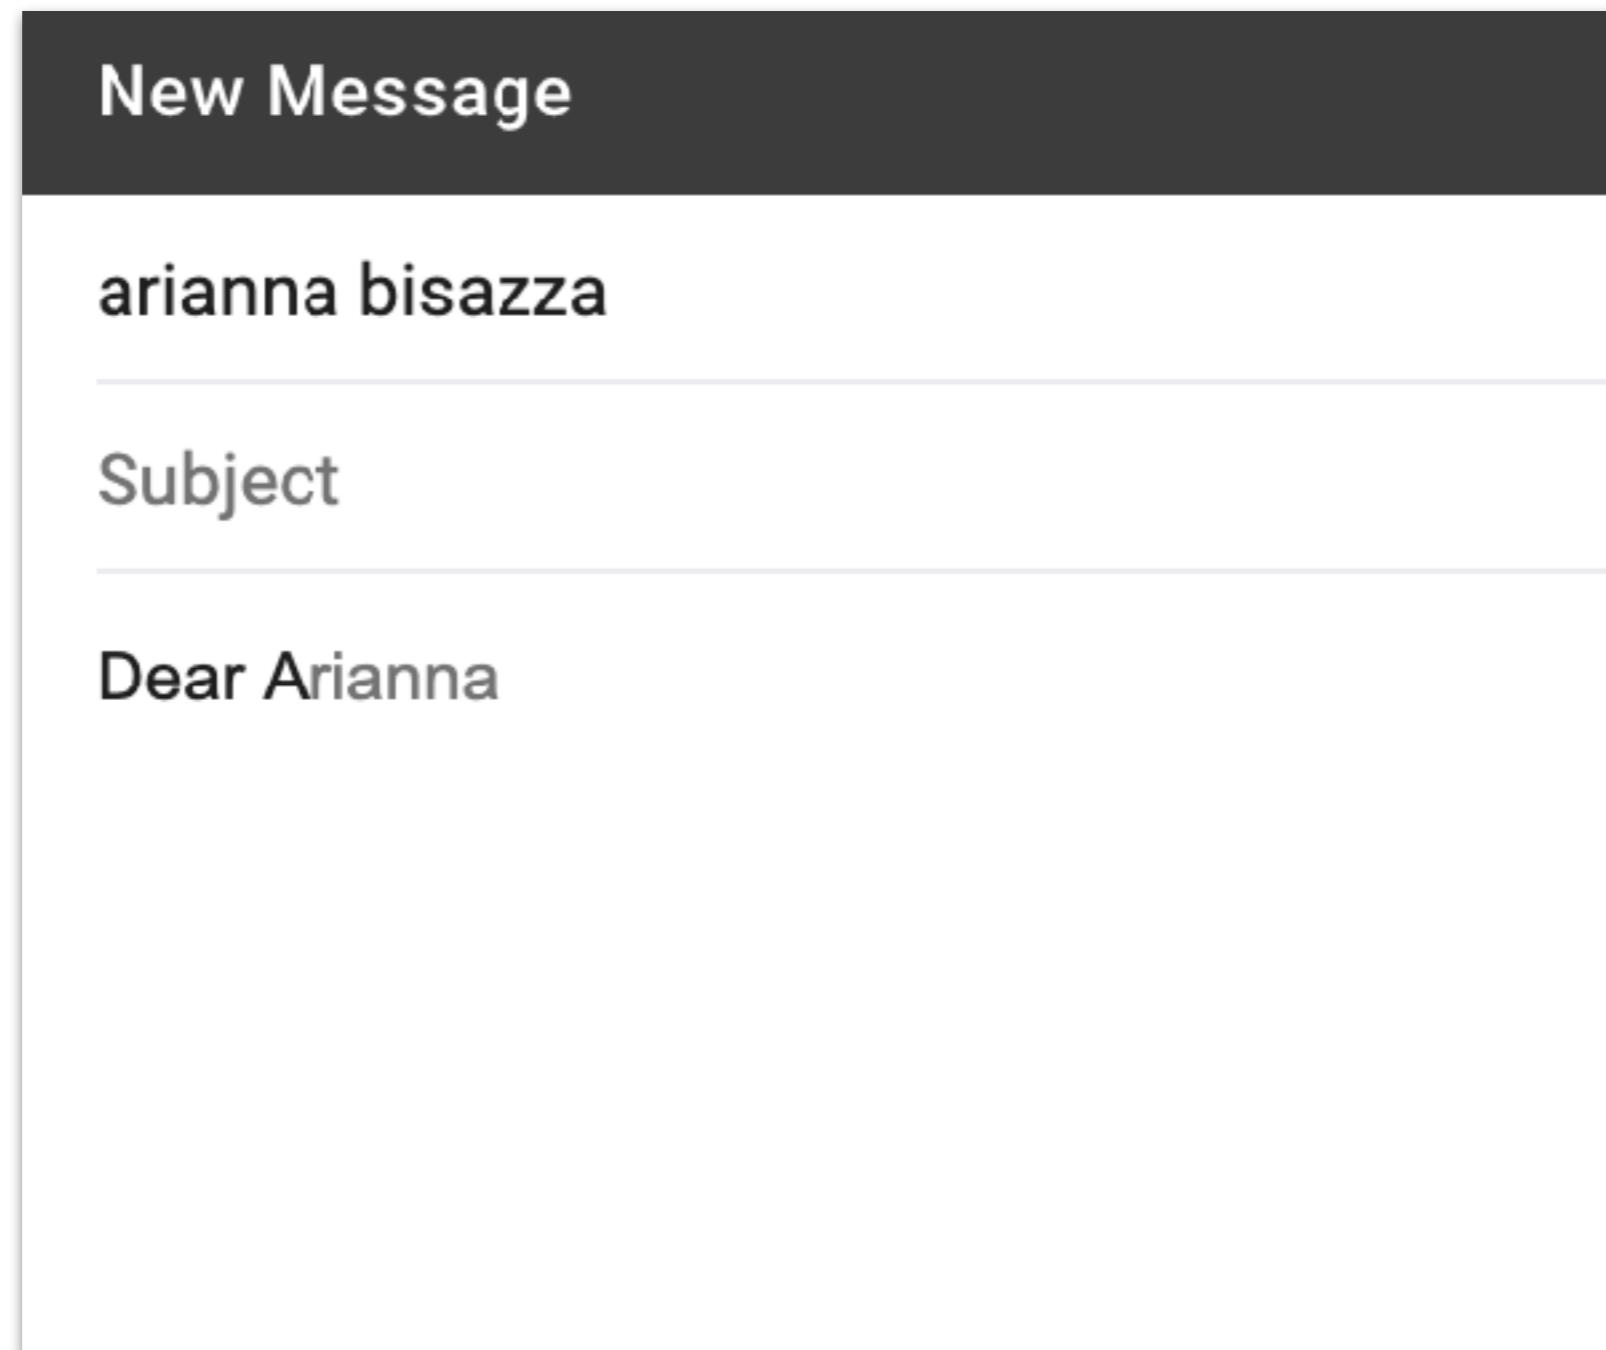
# Spelling Correction



# You probably use a LM every day...

-  what is the mo
-  what is the mo - Google Search
-  what is the most **spoken language in the world**
-  what is the most **played game in the world**
-  what is the most **dangerous animal in the world**
-  what is the most **expensive car in the world**
-  what is the moon made of

# You probably use a LM every day...

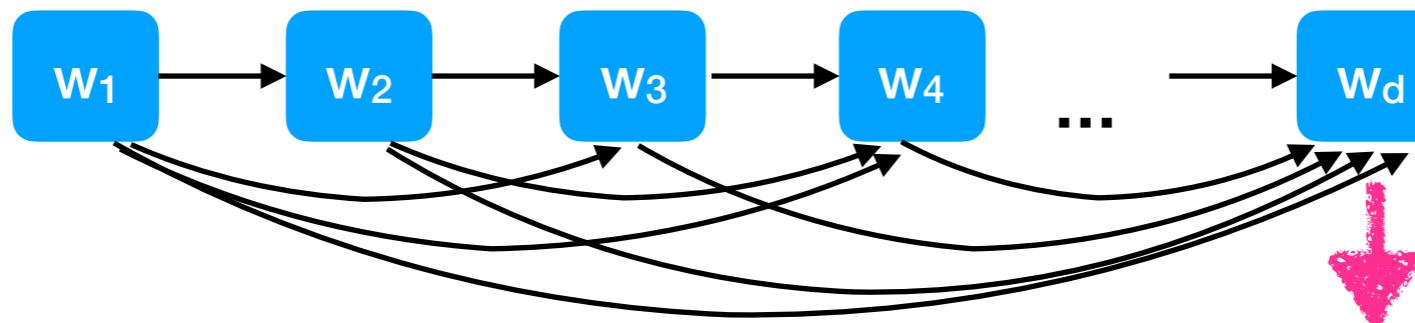


# A Language Model - Formally

- ▶ Given a sequence of words:  $(w_1, \dots, w_d)$
- ▶ LM models the probability:  $P(w_1, \dots, w_d)$
- ▶ Without loss of generality (**Chain Rule**):

$$\begin{aligned} P(w_1, \dots, w_d) &= P(w_1)P(w_2 | w_1)P(w_3 | w_1, w_2)\dots \\ &= P(w_1) \prod_{i=2}^d P(w_i | \boxed{w_1, \dots, w_{i-1}}) \end{aligned}$$

*history*



$P(\text{Athens} | \text{an awesome summer school this year is in})$

# Markov assumption

- ▶ A common assumption in sequence modeling is to make the **Markov assumption**:

$$P(x_1, \dots, x_d) = \prod_{i=1}^d P(x_i | x_1, \dots, x_{i-1}) \approx \prod_{i=1}^d P(x_i | x_{i-(n-1)}, \dots, x_{i-1})$$

$$\begin{aligned} p(\mathbf{w}) = & p(w_1) \times \\ & p(w_2 | w_1) \times \\ & p(w_3 | \cancel{w_1}, w_2) \times \\ & p(w_4 | \cancel{w_1}, \cancel{w_2}, w_3) \times \\ & \dots \end{aligned}$$

*Markov: forget "distant" past  
Valid for language? No...*

**Is it practical?** *Often!*

**n-th order Markov assumption:**  
**history of n-1 words**



Adapted from Chris Dyer

26

# How to learn a LM?

- ▶ **(Pre-deep learning) era:** Learn an **n-gram** Language Model
- ▶ **n-gram:** a chunk of consecutive words
  - ▶ n=2 (bigram): “to buy”, “buy a”, “a house”...
  - ▶ n=3 (trigram): “to buy a”, “buy a house”,...
- ▶ Key method: **collect statistics** of n-grams from a corpus to estimate the parameters of the model (maximum likelihood)

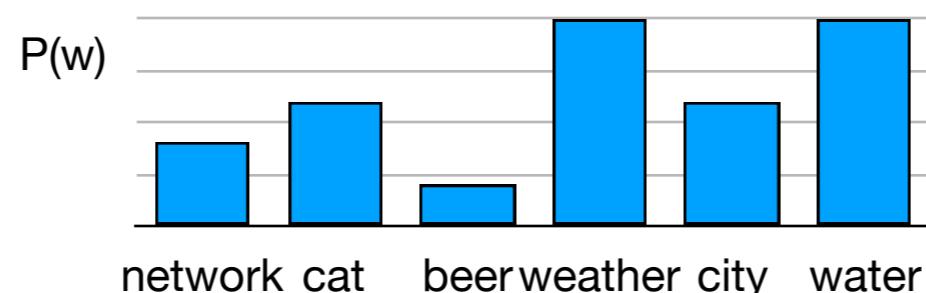
# Unigram LM (1<sup>st</sup> order)

- n=1, 1<sup>st</sup> order Markov assumption, history (n-1): 0

$$P(w_1, \dots, w_d) = P(w_1)P(w_2)P(w_3)\dots$$

$$= \prod_{i=1}^d P(w_i)$$

$P(\text{started})$       *unigram LM*



# Bigram Language Model

- $n=2$ , 2<sup>nd</sup> order Markov assumption, history  $(n-1): 1$

$$P(w_1, \dots, w_d) = P(w_1)P(w_2 | w_1)P(w_3 | w_2)\dots$$

$$= P(w_1) \prod_{i=2}^d P(w_i | w_{i-1})$$

$P(\text{started} | \text{has})$  *bigram LM*



# Higher-order LMs

- ▶ A **bigram** model conditions on the previous word ( $n=2$ ; or: a window of 2 words)

$$P(x_i | x_{i-1})$$

- ▶ A **trigram** model uses a history of 2 words ( $n=3$ , history 2)

$$P(x_i | x_{i-2}, x_{i-1})$$

- ▶ E.g. a 5-gram LM

# Sparsity problems with n-gram LMs

C=count()

**Sparsity problem 1:** If “its water is so w” never occurred in the corpus: prob for w is 0!

What can we do? **Smoothing** (add small count to every w)

$$P(w \mid \text{its water is so}) = \frac{C(\text{its water is so } w)}{C(\text{its water is so})}$$

**Sparsity problem 2:** If “its water is so” never occurred: prob for any w is 0.

What can we do? **Backoff** (condition on lower-level n-grams)

**In general:** Increasing n-gram size makes sparsity problem worse.

# Further issues with n-gram LMs

- ▶ What about similar words?
  - ▶ *she bought a bicycle*
  - ▶ *she purchased a bicycle*



cannot share strength  
among similar words

- ▶ Long-distance dependencies?
  - ▶ *for programming she yesterday purchased her own brand new laptop*
  - ▶ *for running she yesterday purchased her brand new sportswatch*

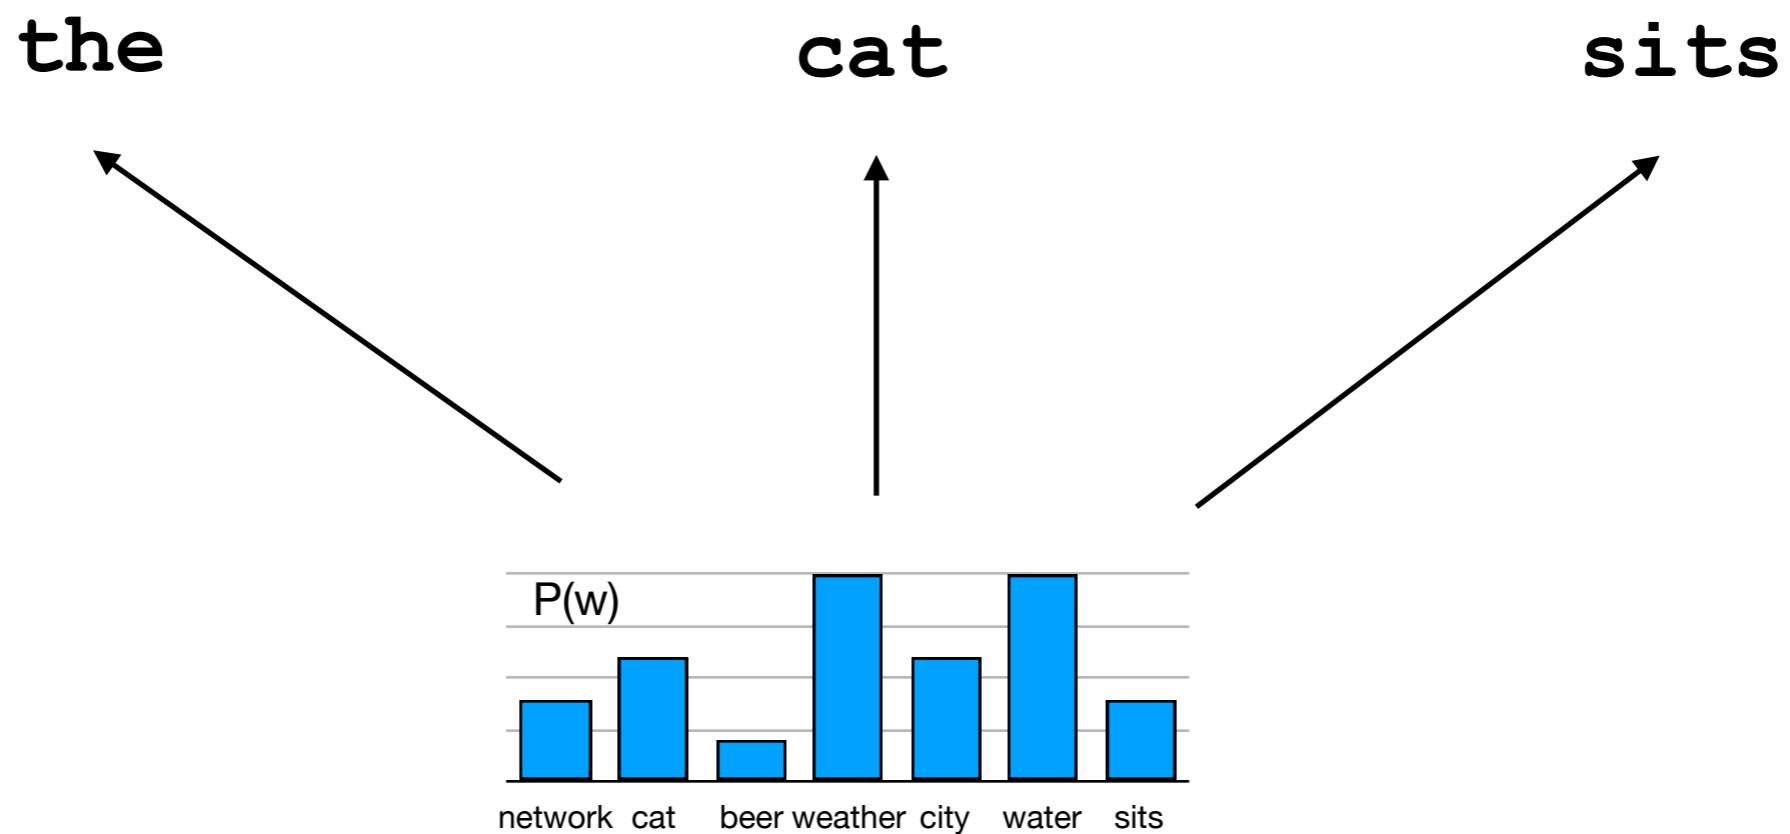


cannot handle long-  
distance dependencies

# **Generating text from a LM**

# Sample from a unigram LM

- ▶ We can sample **incrementally** from a Language Model, one word at a time



# Sample from a unigram LM

- We can sample **incrementally** from a Language Model, one word at a time

*word salad?*

*sequences / sequential data!*

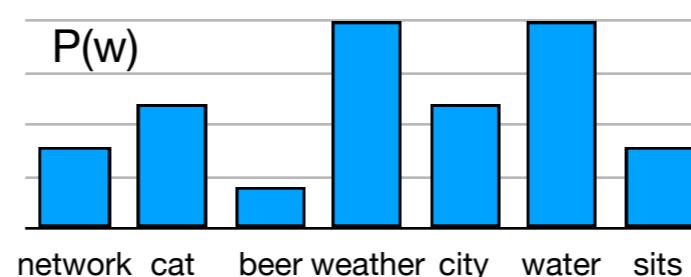
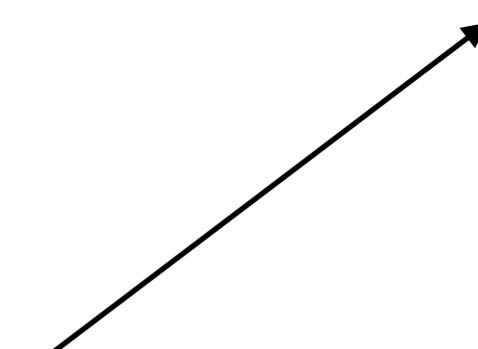
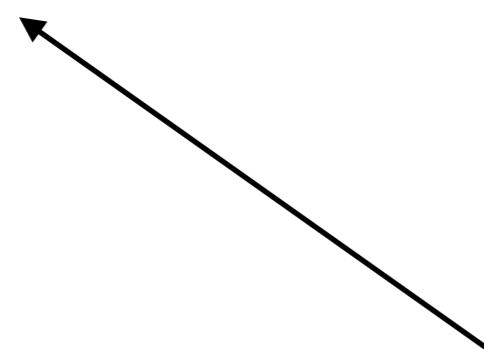
**cat**



**sits**



**the**



# Outlook: Why RNNs are so great for Language

- ▶ No more Markov assumptions
- ▶ Great fit for *sequences / sequential data*
- ▶ Arbitrary length input

the → cat → sits

the cat sits there

the sleepy cat sits there

the sleepy cat which chased the dog sits

# How to learn a *neural* LM?

- ▶ Language model task:

- ▶ input: sequence of words:  $(w_1, \dots, w_d)$
- ▶ output: probability of next word  $P(w_{t+1} | w_t, \dots, w_2, w_1)$

- ▶ An Early (deep learning era) solution: a window-based n-gram neural language model (Bengio et al., 2003)

[A neural probabilistic language model](#)

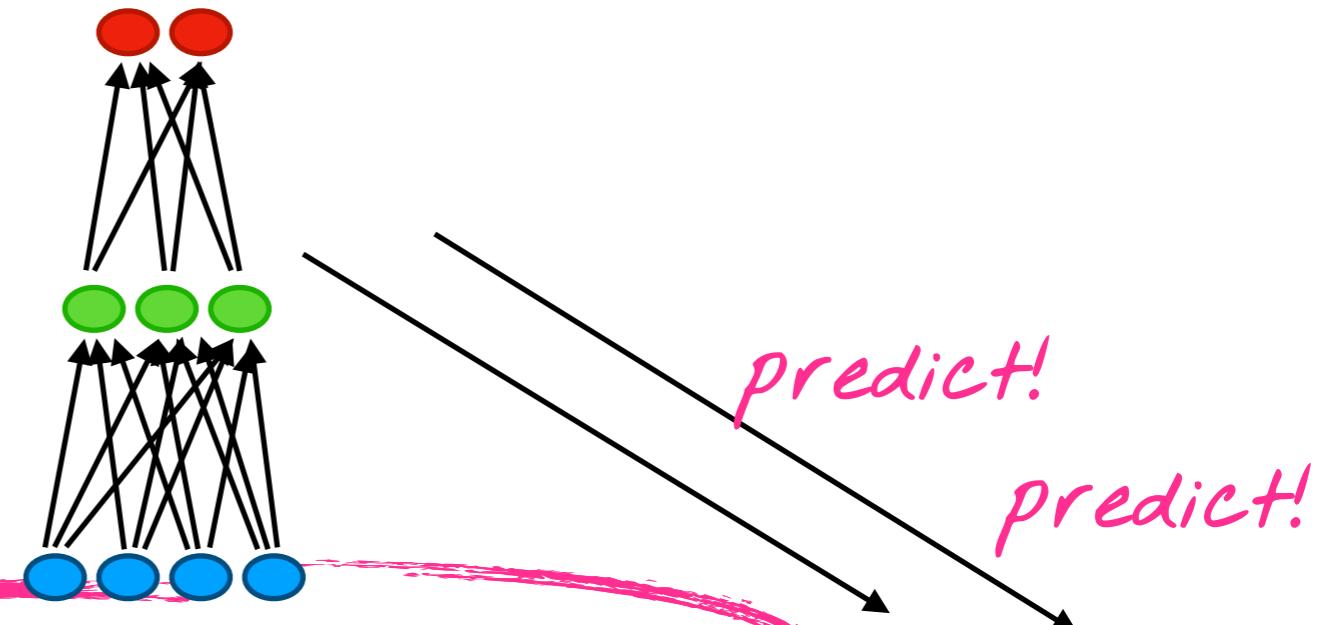
[Y Bengio, R Ducharme, P Vincent, C Jauvin - Journal of machine learning ..., 2003 - jmlr.org](#)

A goal of statistical language modeling is to learn the joint probability function of sequences of words in a language. This is intrinsically difficult because of the curse of dimensionality: a word sequence on which the model will be tested is likely to be different from all the word sequences seen during training. Traditional but very successful approaches based on n-grams obtain generalization by concatenating very short overlapping sequences seen in the training set. We propose to fight the curse of dimensionality by learning a distributed ...

☆ 99 Cited by 5046 Related articles All 57 versions ▾

feed-forward neural network

# Window-based neural LM via FFNN



~~As the clock rang the students opened the~~

*discard*

*fixed window of n words*

# Overview

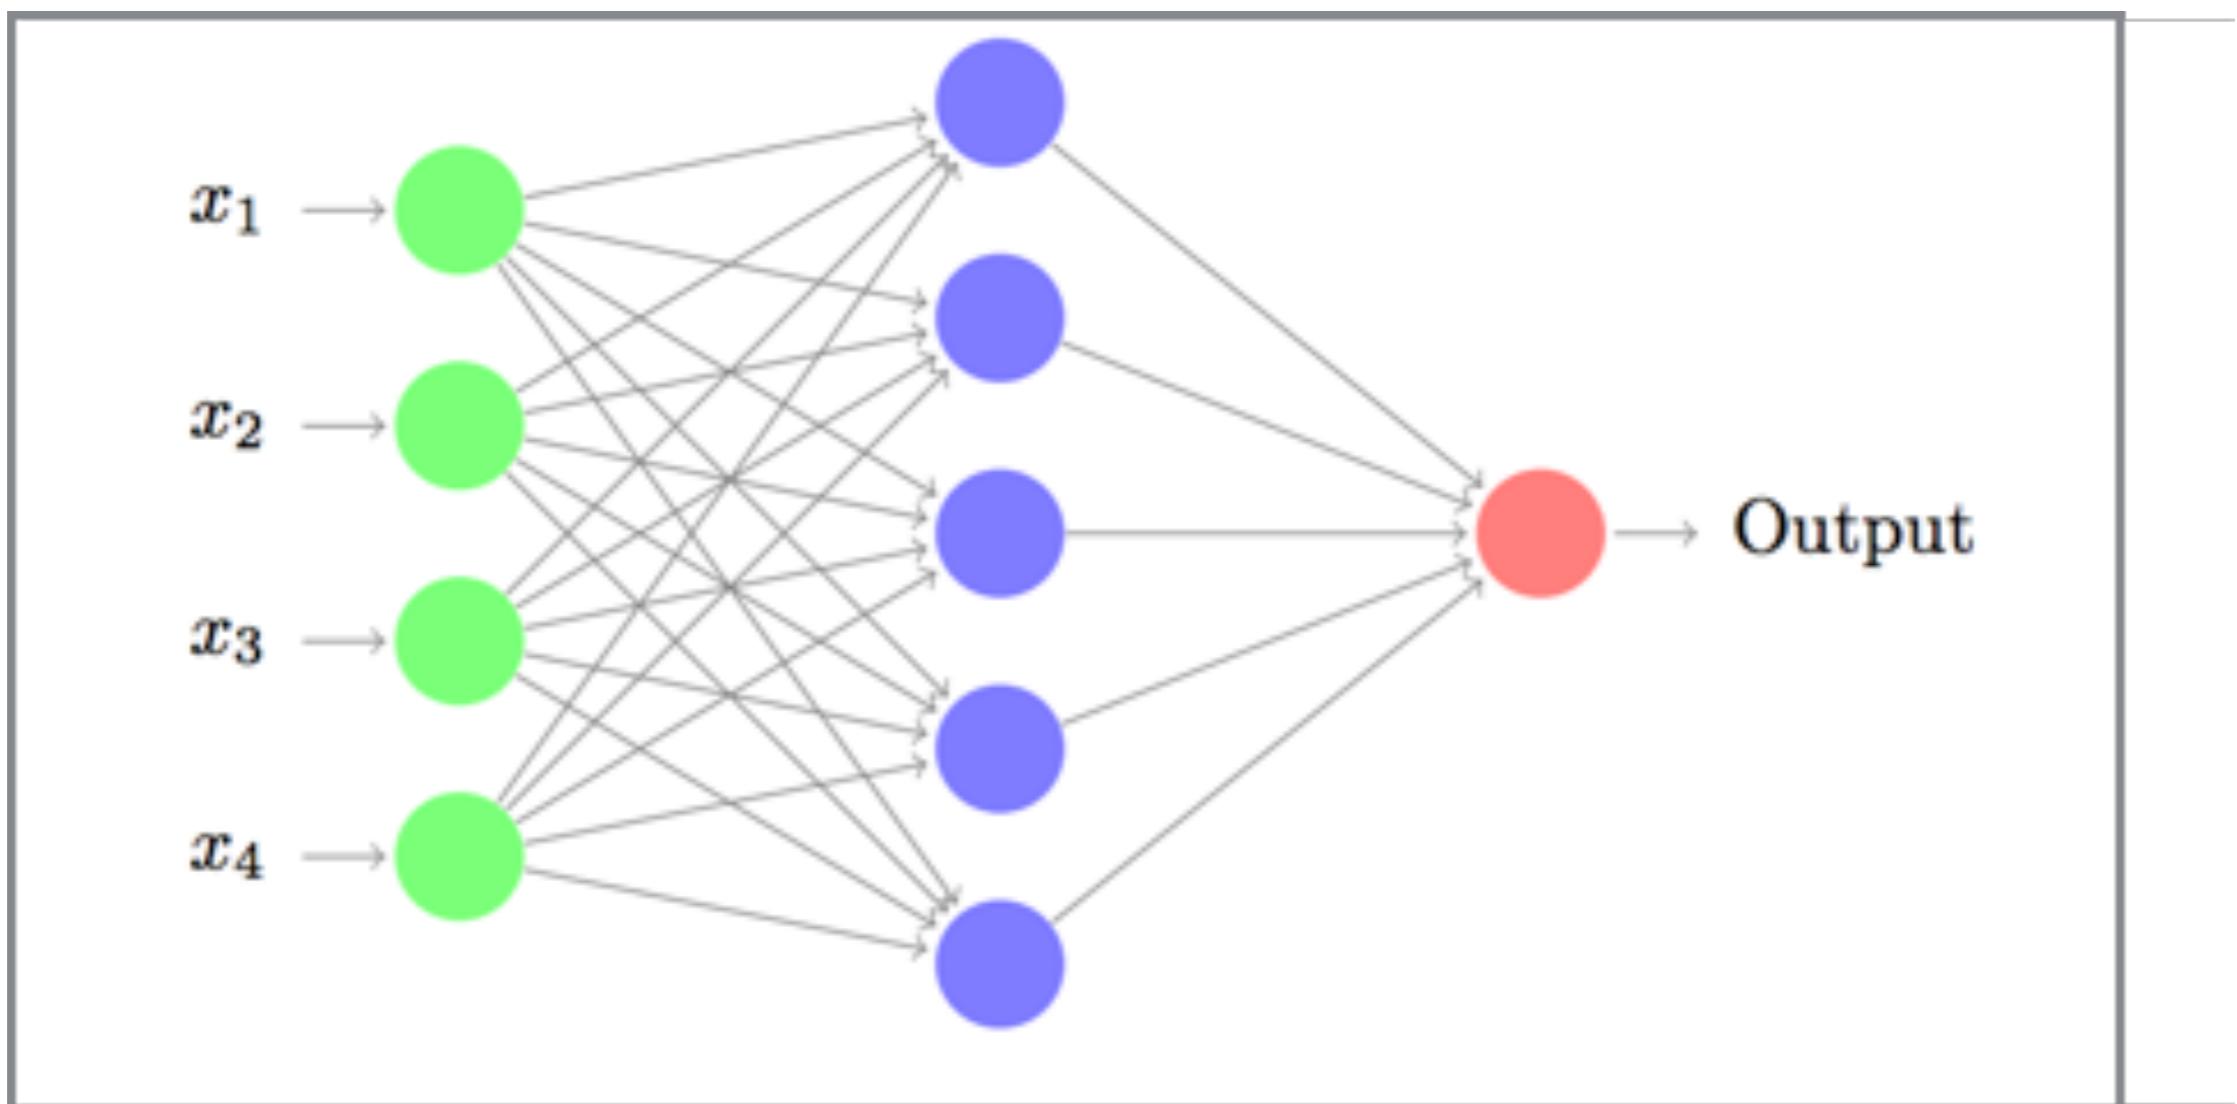
*foundations*



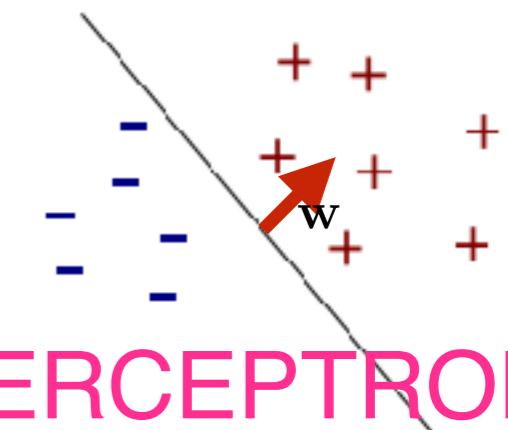
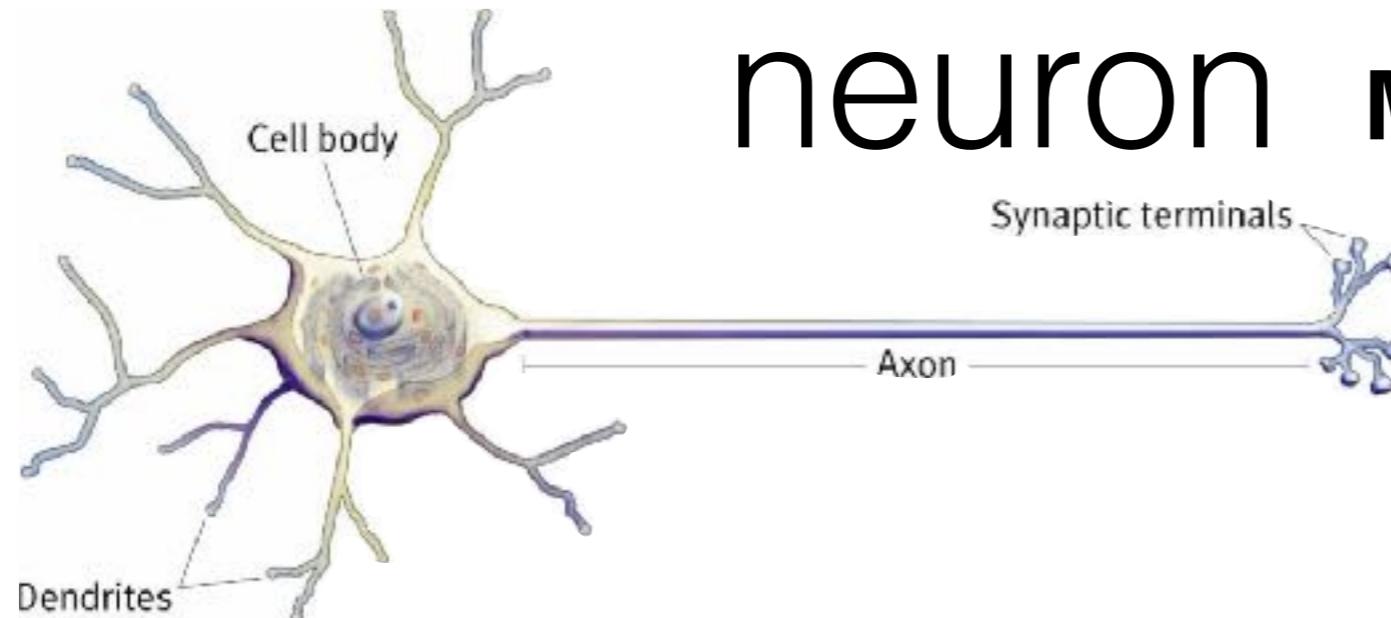
# Feedforward Neural Network (FFNN)

- ▶ After this **recap** you should:
  - ▶ connect the different views on FFNNs
  - ▶ refresh ourselves on how to represent input in NLP

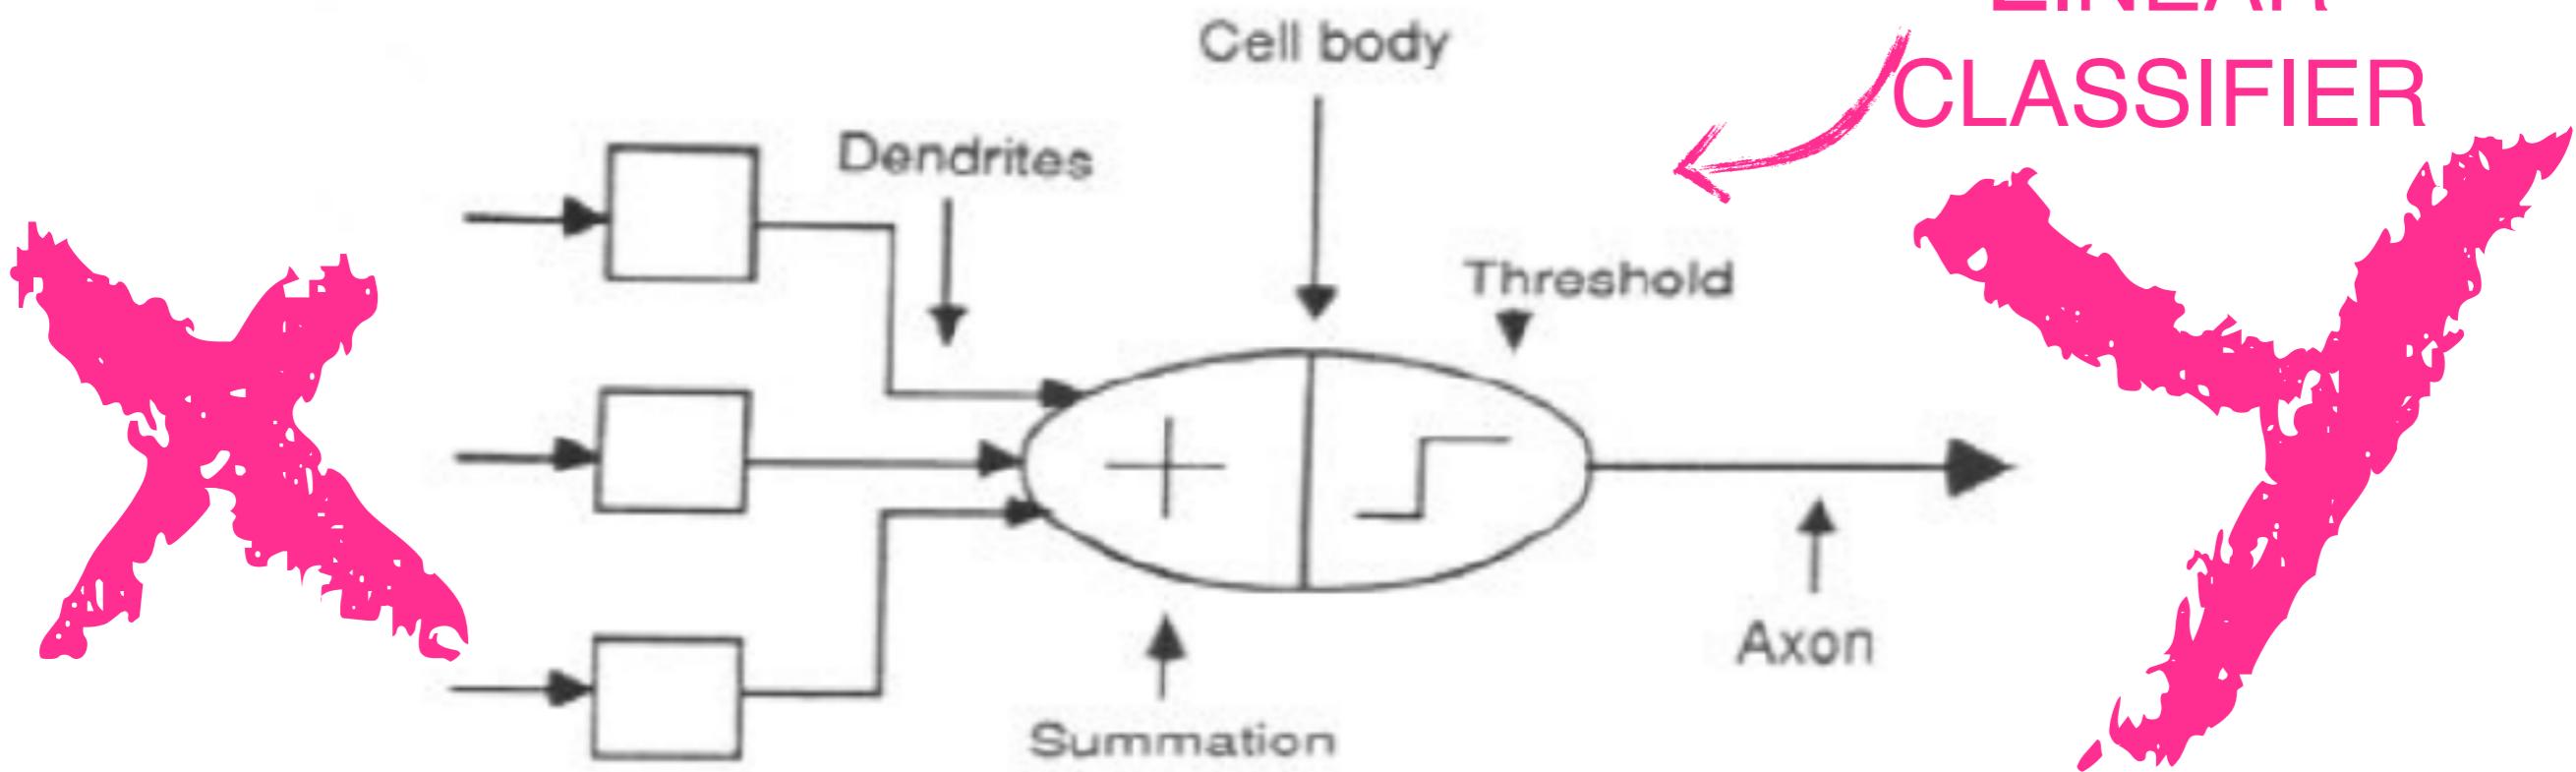
# Neural Network



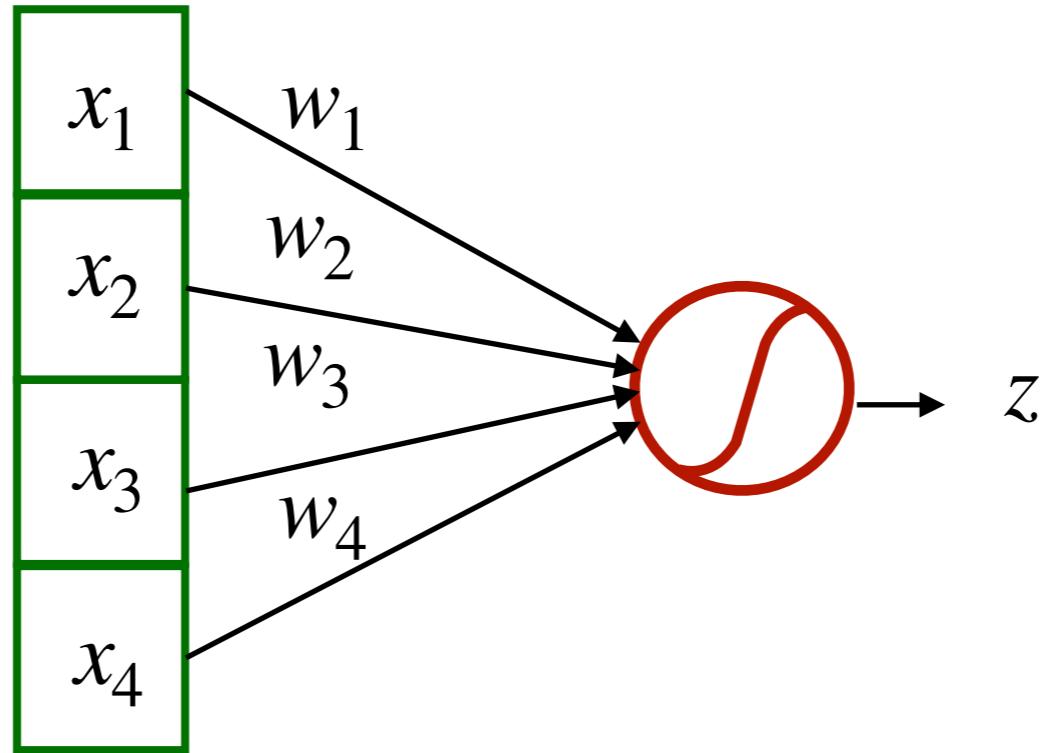
# From biological to artificial neuron McCulloch & Pitt (1943)



PERCEPTRON:  
LINEAR  
CLASSIFIER



# A single neuron

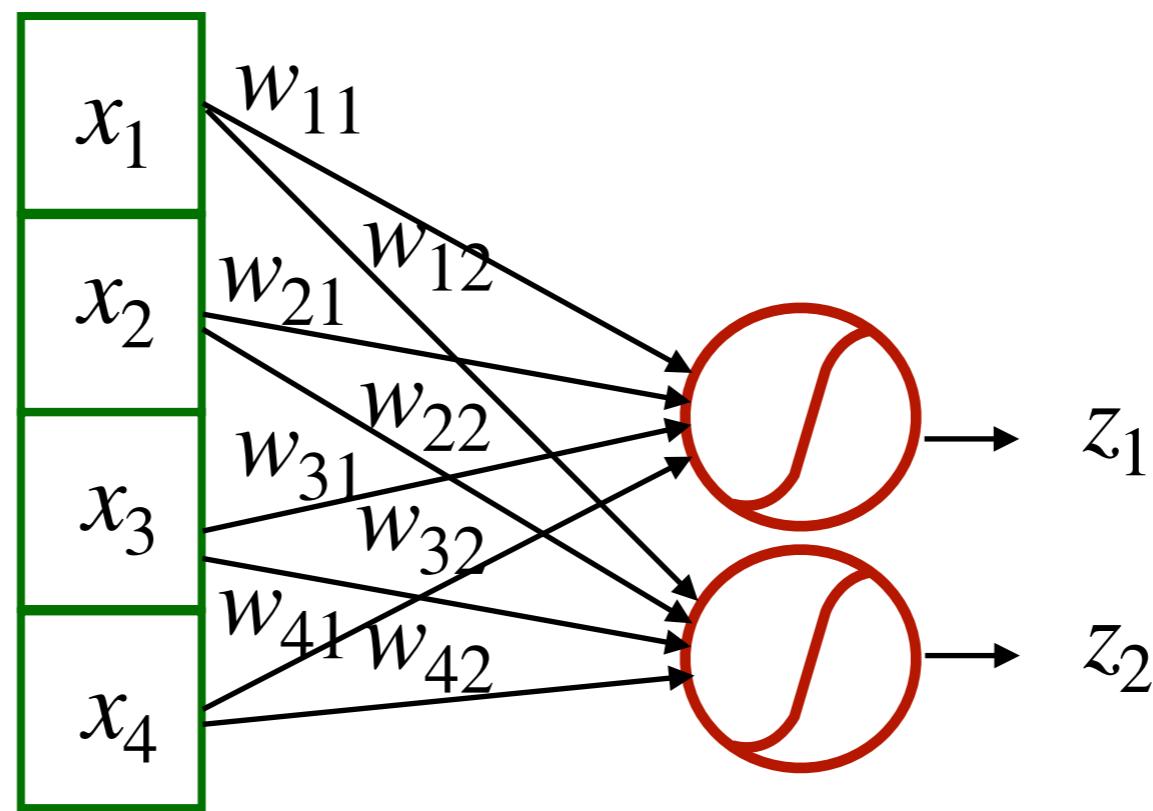


$$\begin{aligned} z &= \sigma(x_1w_1 + x_2w_2 + x_3w_3 + x_4w_4 + b) \\ &= \sigma(\mathbf{x} \cdot \mathbf{w} + b) \quad \text{with } \mathbf{x}, \mathbf{w} \in \mathbb{R}^4 \end{aligned}$$

(bias node omitted in visualization)

# Multiple neurons

$$f_{\theta} : \mathbb{R}^4 \rightarrow \mathbb{R}^2$$



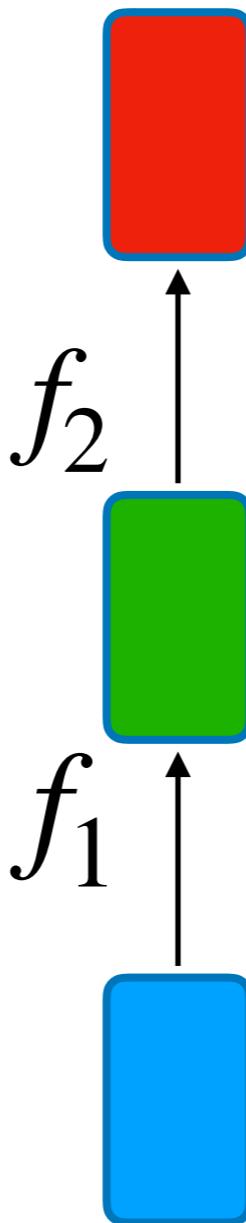
$$z_1 = \sigma(\mathbf{x} \cdot \mathbf{w}_1 + b_1)$$

$$z_2 = \sigma(\mathbf{x} \cdot \mathbf{w}_2 + b_2)$$

# FFNN: abstract & functional view

$$\hat{y} = \text{softmax}(f_2(f_1(x)))$$

- ▶ Each layer is a **function**, acts on the **output** of the layer below (input)
- ▶ Final output: **cascade** of functions
- ▶ Given the true output  $y$ , we compute the **error**  $E_r$



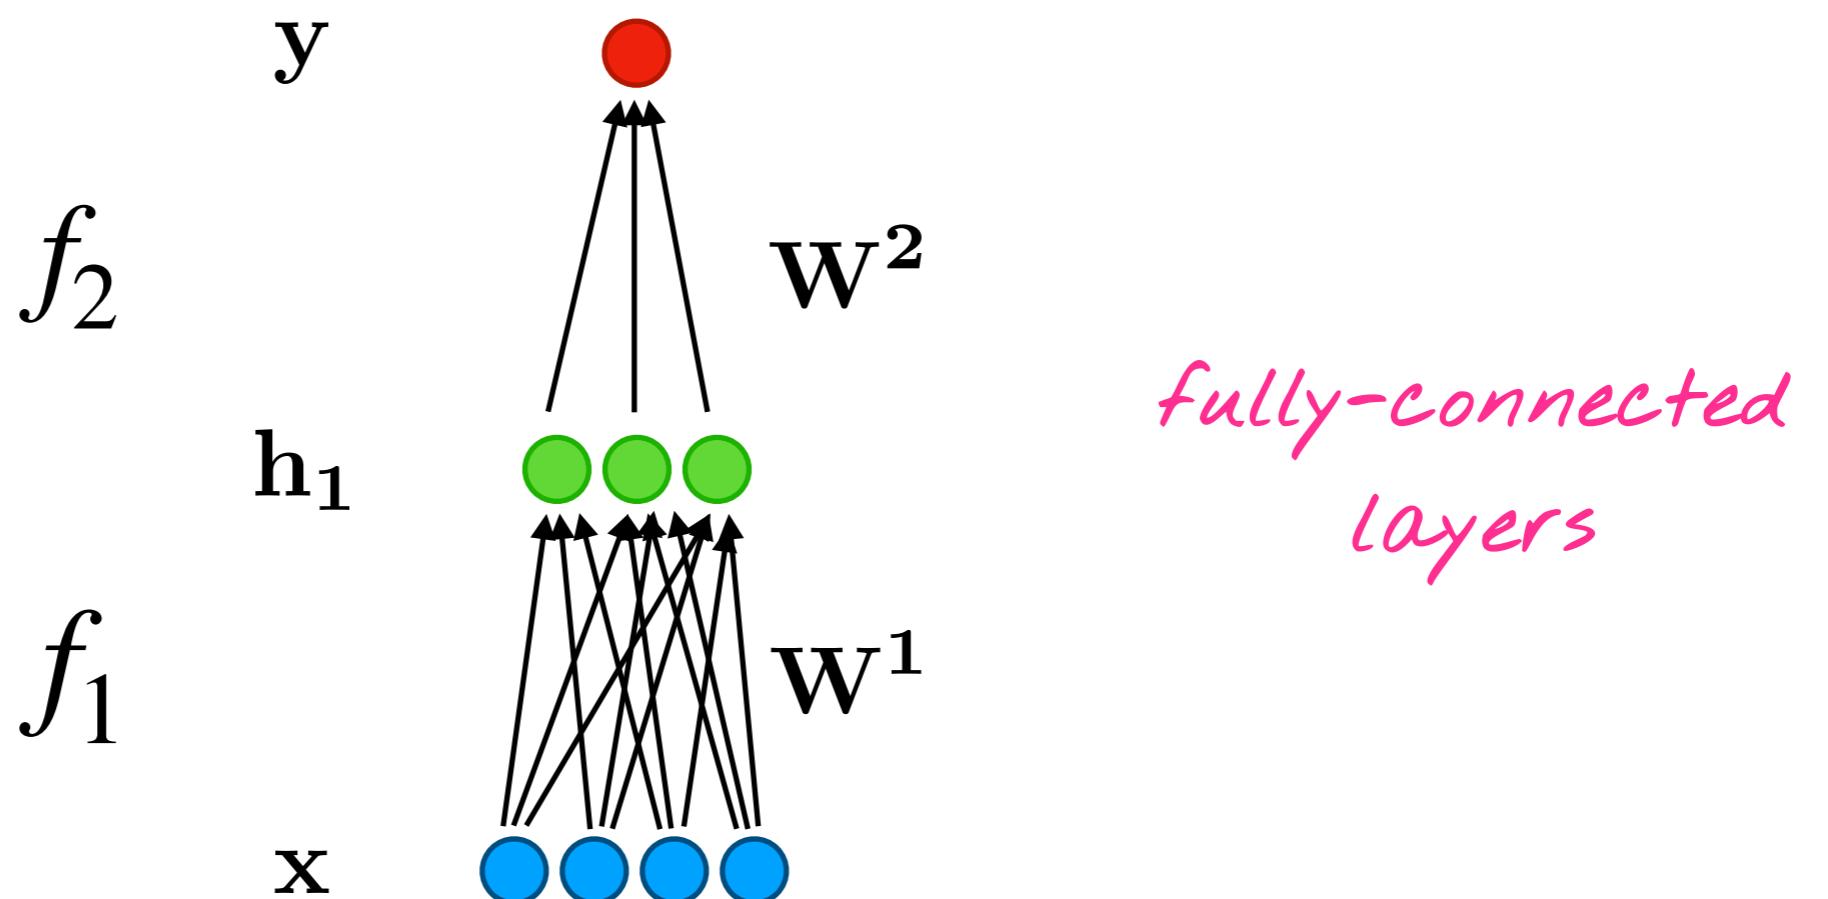
$$error = \text{loss}(y, \hat{y})$$

- ▶ Using the chain rule, we can compute the derivative (**gradient**) of the Error wrt any of the intermediate layer weights (Lecture 1)

“vanilla” Neural Network

# FFNN: Graphical view

$$\hat{y} = \text{softmax}(f_2(f_1(x)))$$



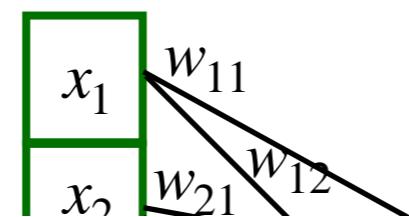
# Multiple neurons: Vectorization

$$z_1 = \sigma(\mathbf{x} \cdot \mathbf{w}_1 + b_1)$$

$$z_2 = \sigma(\mathbf{x} \cdot \mathbf{w}_2 + b_2)$$



$$\mathbf{z} = \sigma(\mathbf{x} \cdot \mathbf{W} + \mathbf{b})$$

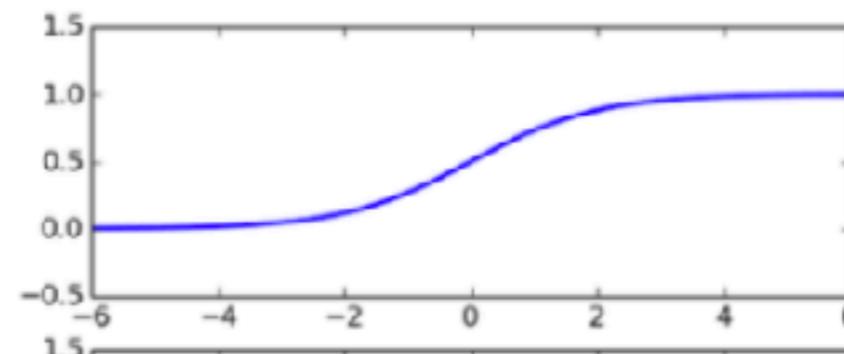


linear projection followed  
by non-linearity

$$f_{\theta} : \mathbb{R}^4 \rightarrow \mathbb{R}^2$$

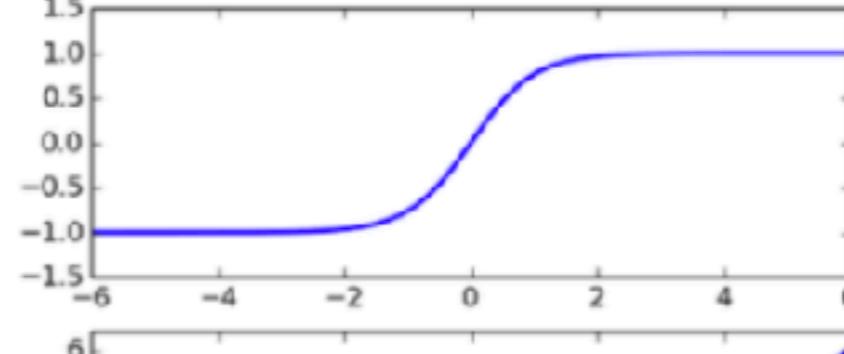
$\sigma$  is the layer's  
**(non-linear)**  
activation  
function

$$\mathbf{x} = \begin{matrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{matrix}$$



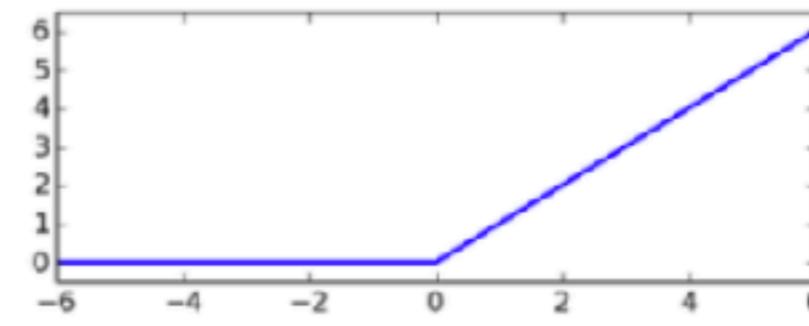
Sigmoid

$$\phi(z) = \frac{1}{1 + e^{-z}}$$



Hyperbolic Tangent

$$\phi(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

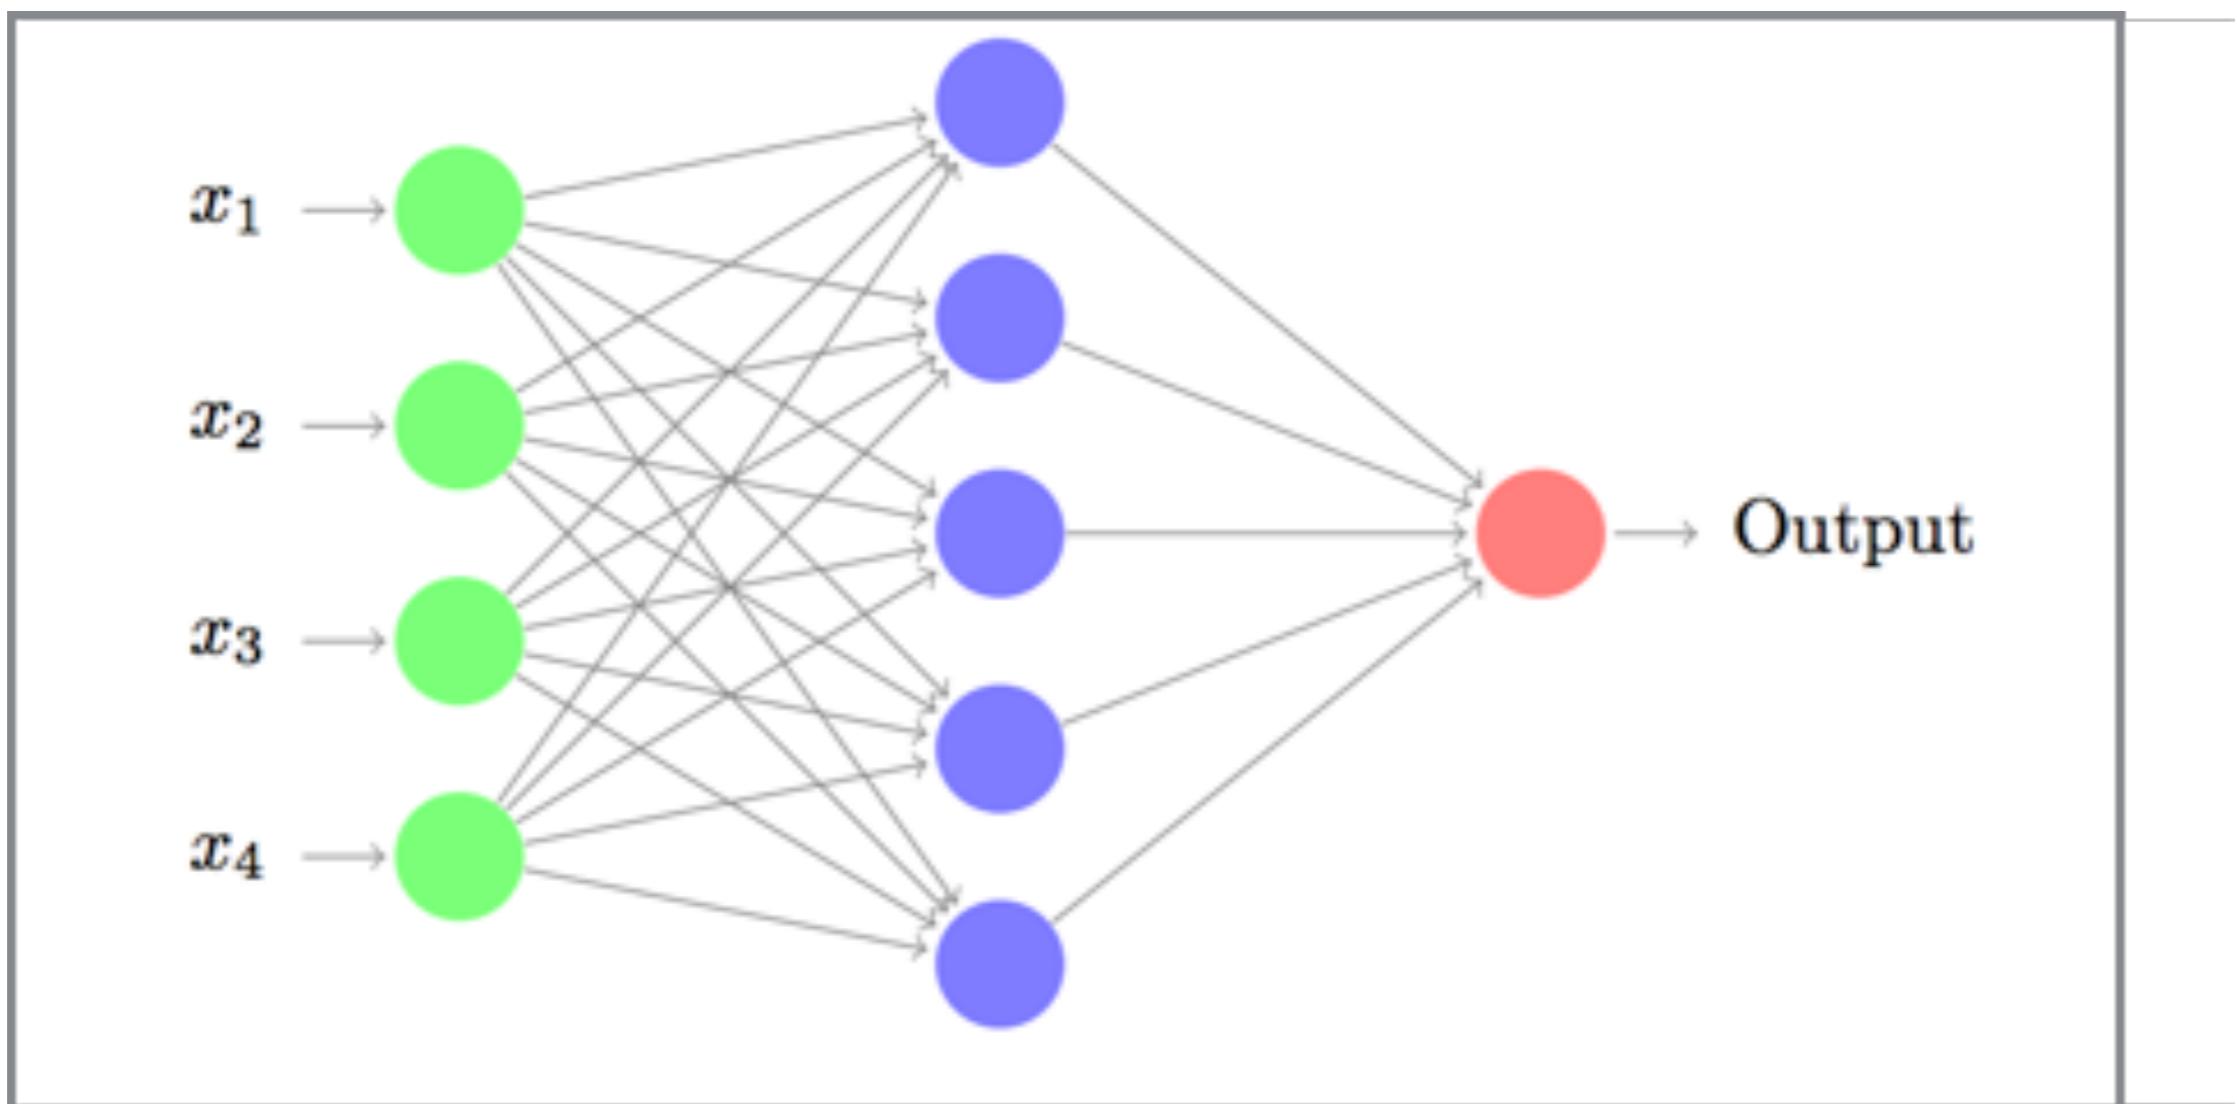


Rectified Linear

$$\phi(z) = \begin{cases} 0 & \text{if } z < 0 \\ z & \text{if } z \geq 0 \end{cases}$$

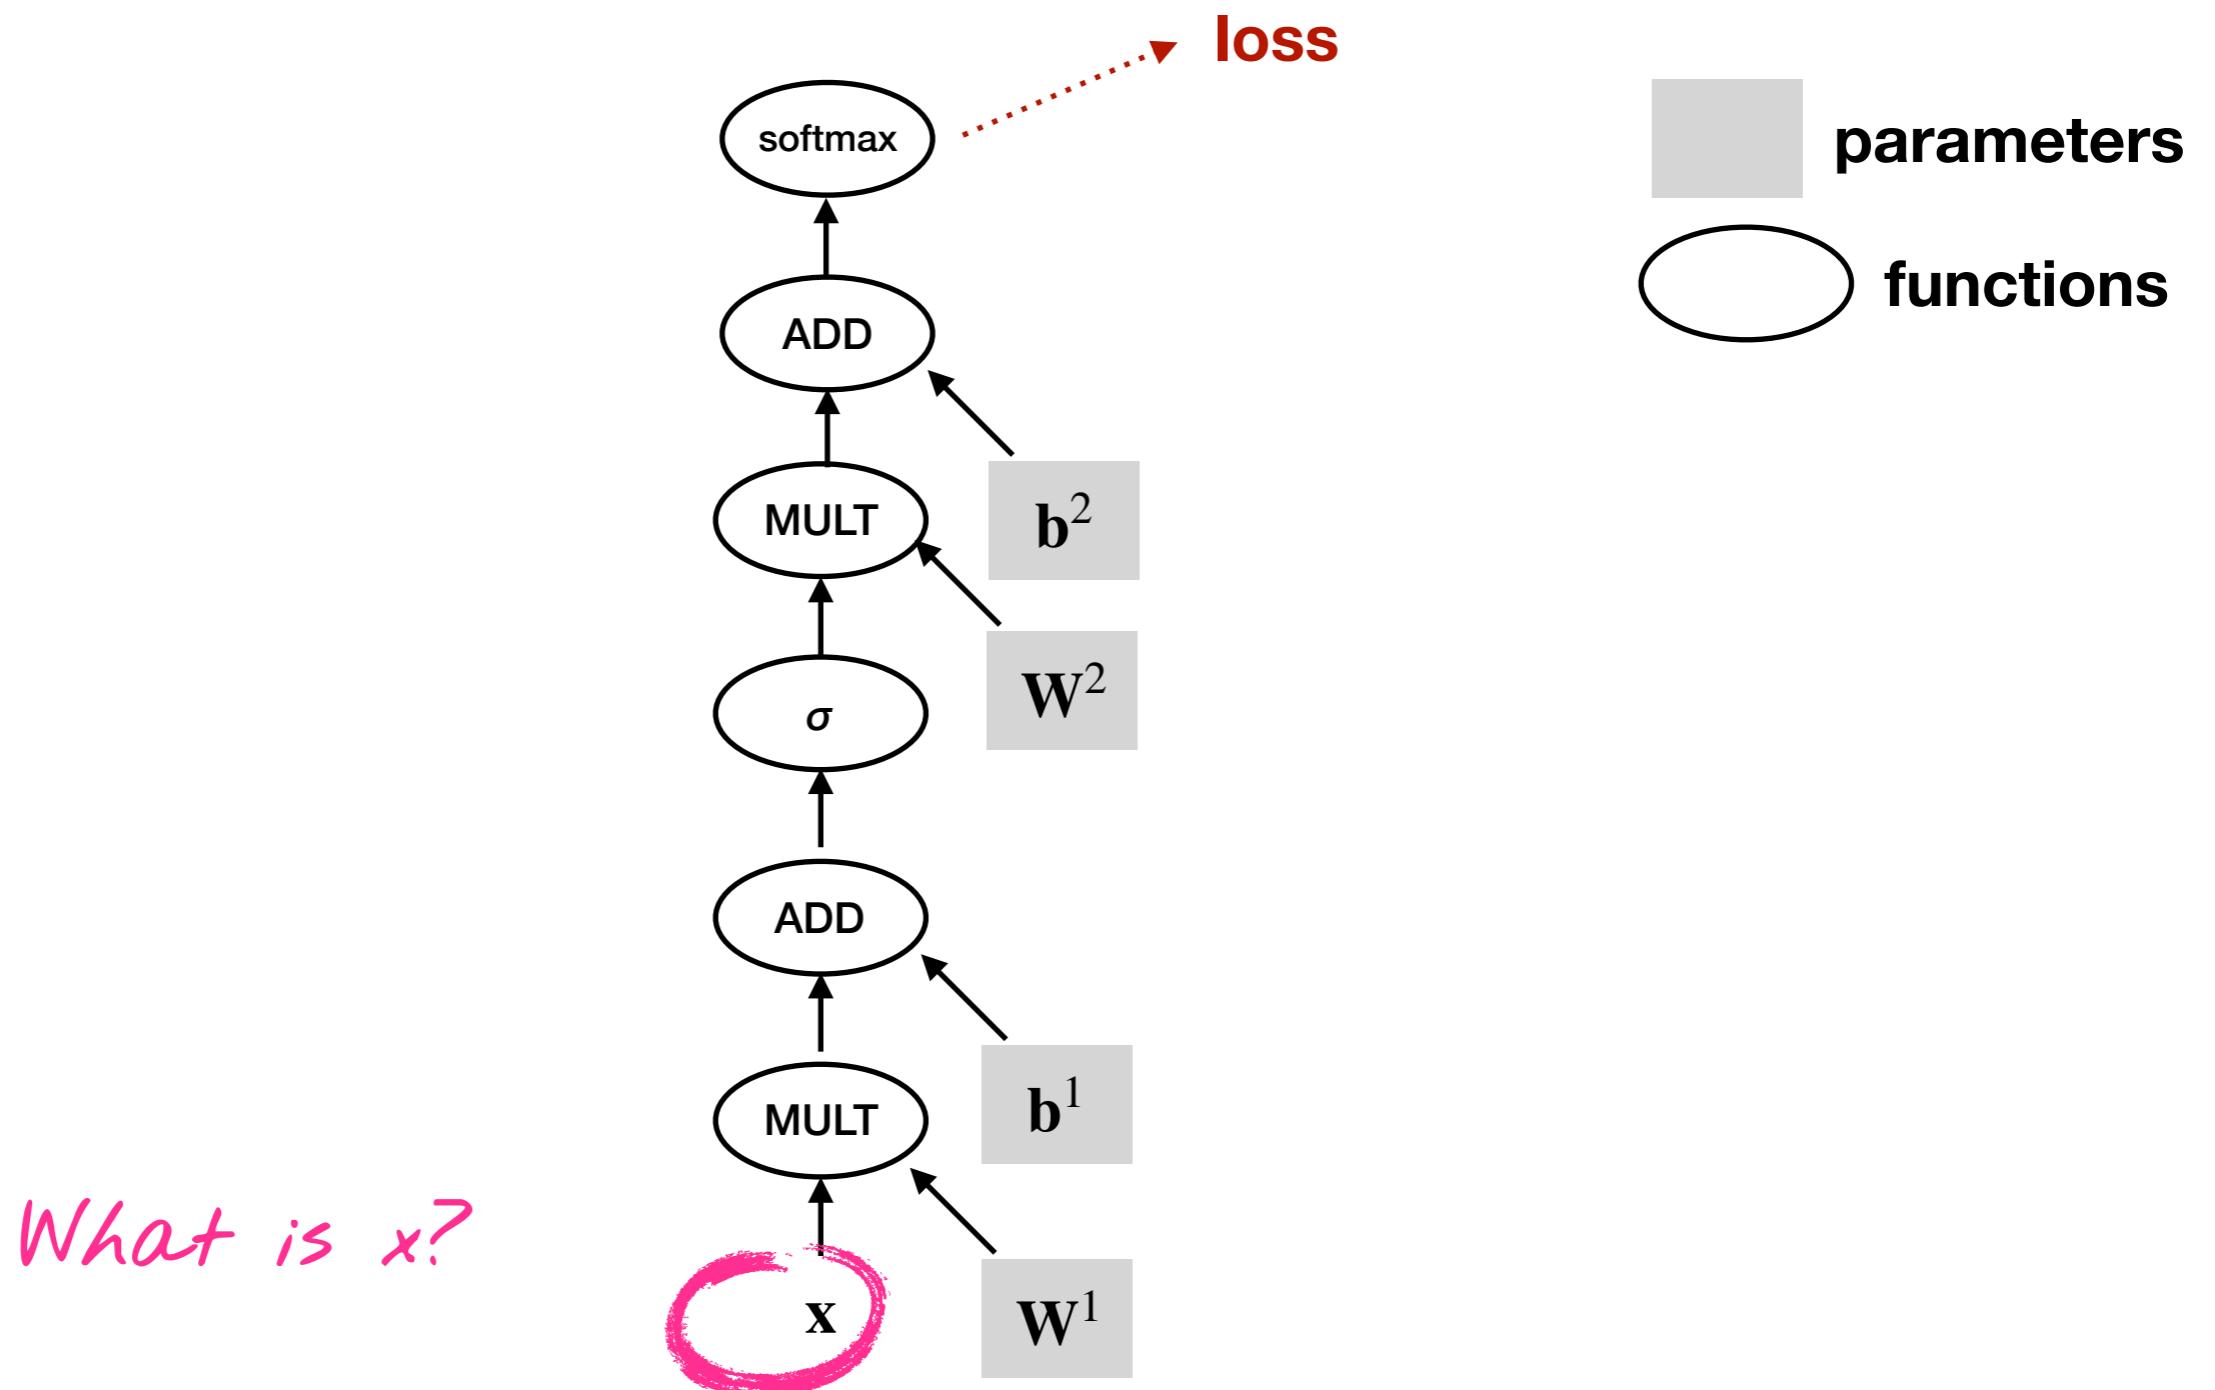
From Hughes and Correll 2016

# Connecting the views: FFNNs (MLPs)



$$NN_{MLP1}(\mathbf{x}) = g(\mathbf{x}\mathbf{W}^1 + \mathbf{b}^1)\mathbf{W}^2 + \mathbf{b}^2$$

# Computation Graph View



$$NN_{MLP1}(\mathbf{x}) = g(\mathbf{x}\mathbf{W}^1 + \mathbf{b}^1)\mathbf{W}^2 + \mathbf{b}^2$$

# Overview

*foundations*



*representations*



# What does 'eienskappe' mean?

# Keyword in Context (KWIC)

# Distributional Hypothesis

"You shall know a word by the  
company it keeps"

(Firth, J. R. 1957:11)

# The company it keeps

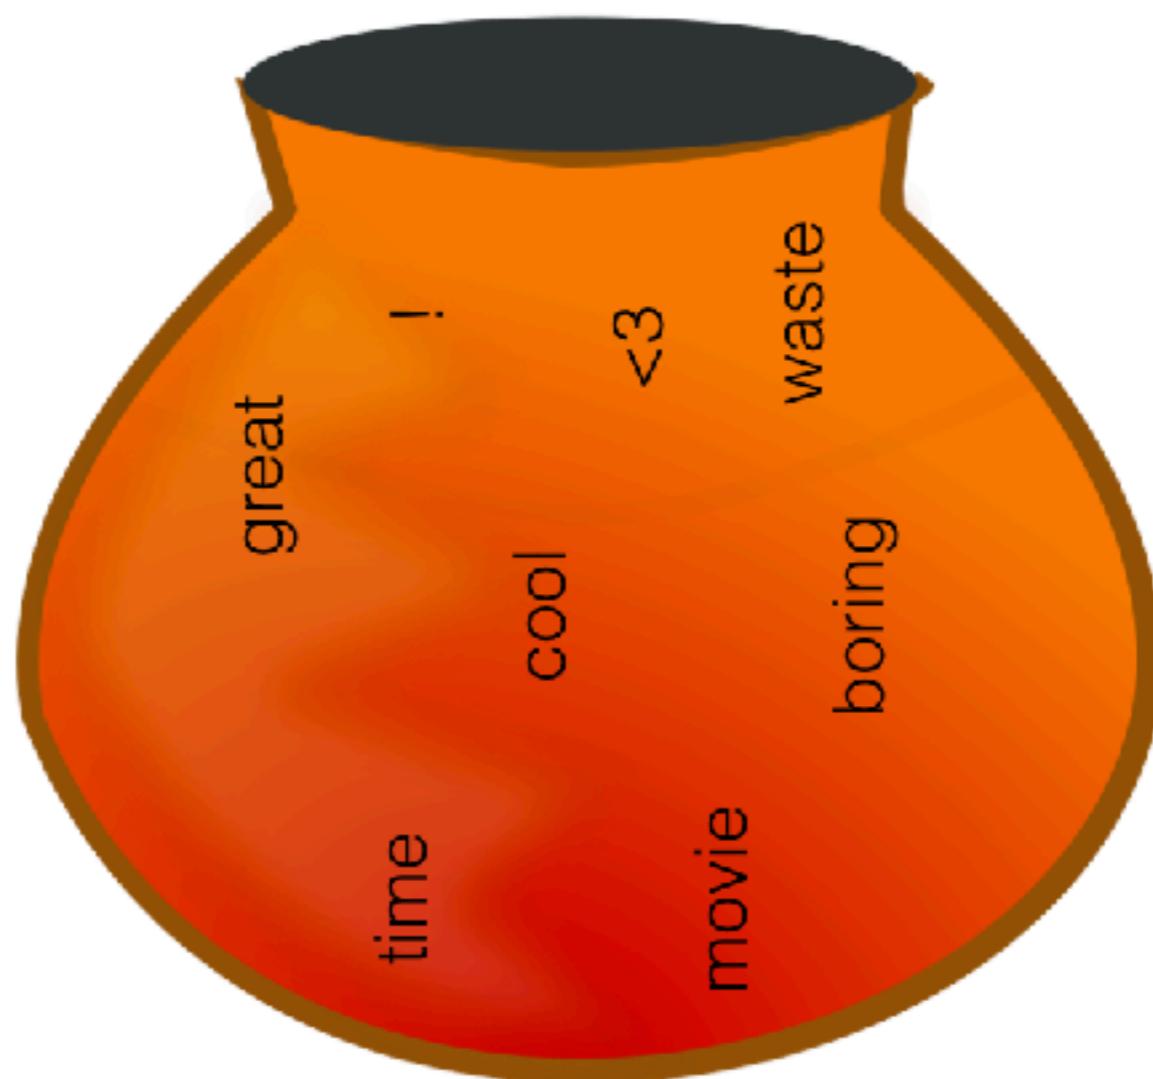
- ▶ Key idea in NLP: the meaning of a word is represented by the words which occur frequently close to it
- ▶ One of the most successful ideas in NLP
- ▶ Nowadays, we talk about **representations**

# What are good representations?

- ▶ Representations are **distinct**
- ▶ Similar words (or units) have **similar** representations

# Traditional sparse text encoding: BOW

bag-of-words



# Sparse binary text encoding: BOW

this is great !

cool <3

a waste of time

great movie

boring just boring

this    is    great    -    cool    <3    boring    a    waste    of    time    movie    just

1	1	1	1																
---	---	---	---	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--

					1	1													
--	--	--	--	--	---	---	--	--	--	--	--	--	--	--	--	--	--	--	--

										1	1	1	1						
--	--	--	--	--	--	--	--	--	--	---	---	---	---	--	--	--	--	--	--

					1													1	
--	--	--	--	--	---	--	--	--	--	--	--	--	--	--	--	--	--	---	--

										1									1
--	--	--	--	--	--	--	--	--	--	---	--	--	--	--	--	--	--	--	---

**n-hot encoding**

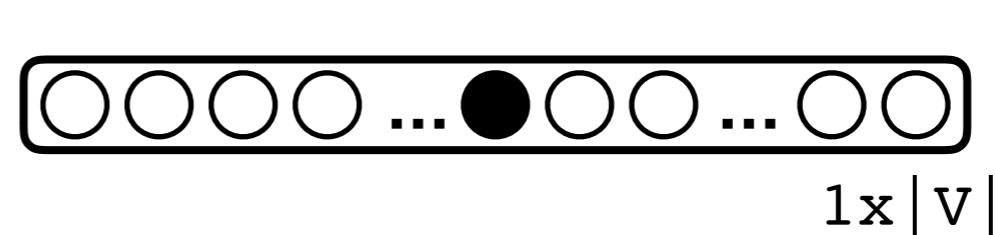
# One-hot encoding

- ▶ Sparse high-dimensional vector of dimension  $|V|$  (=size of vocabulary)

**Symbol** (word, char,...)

**yellow**

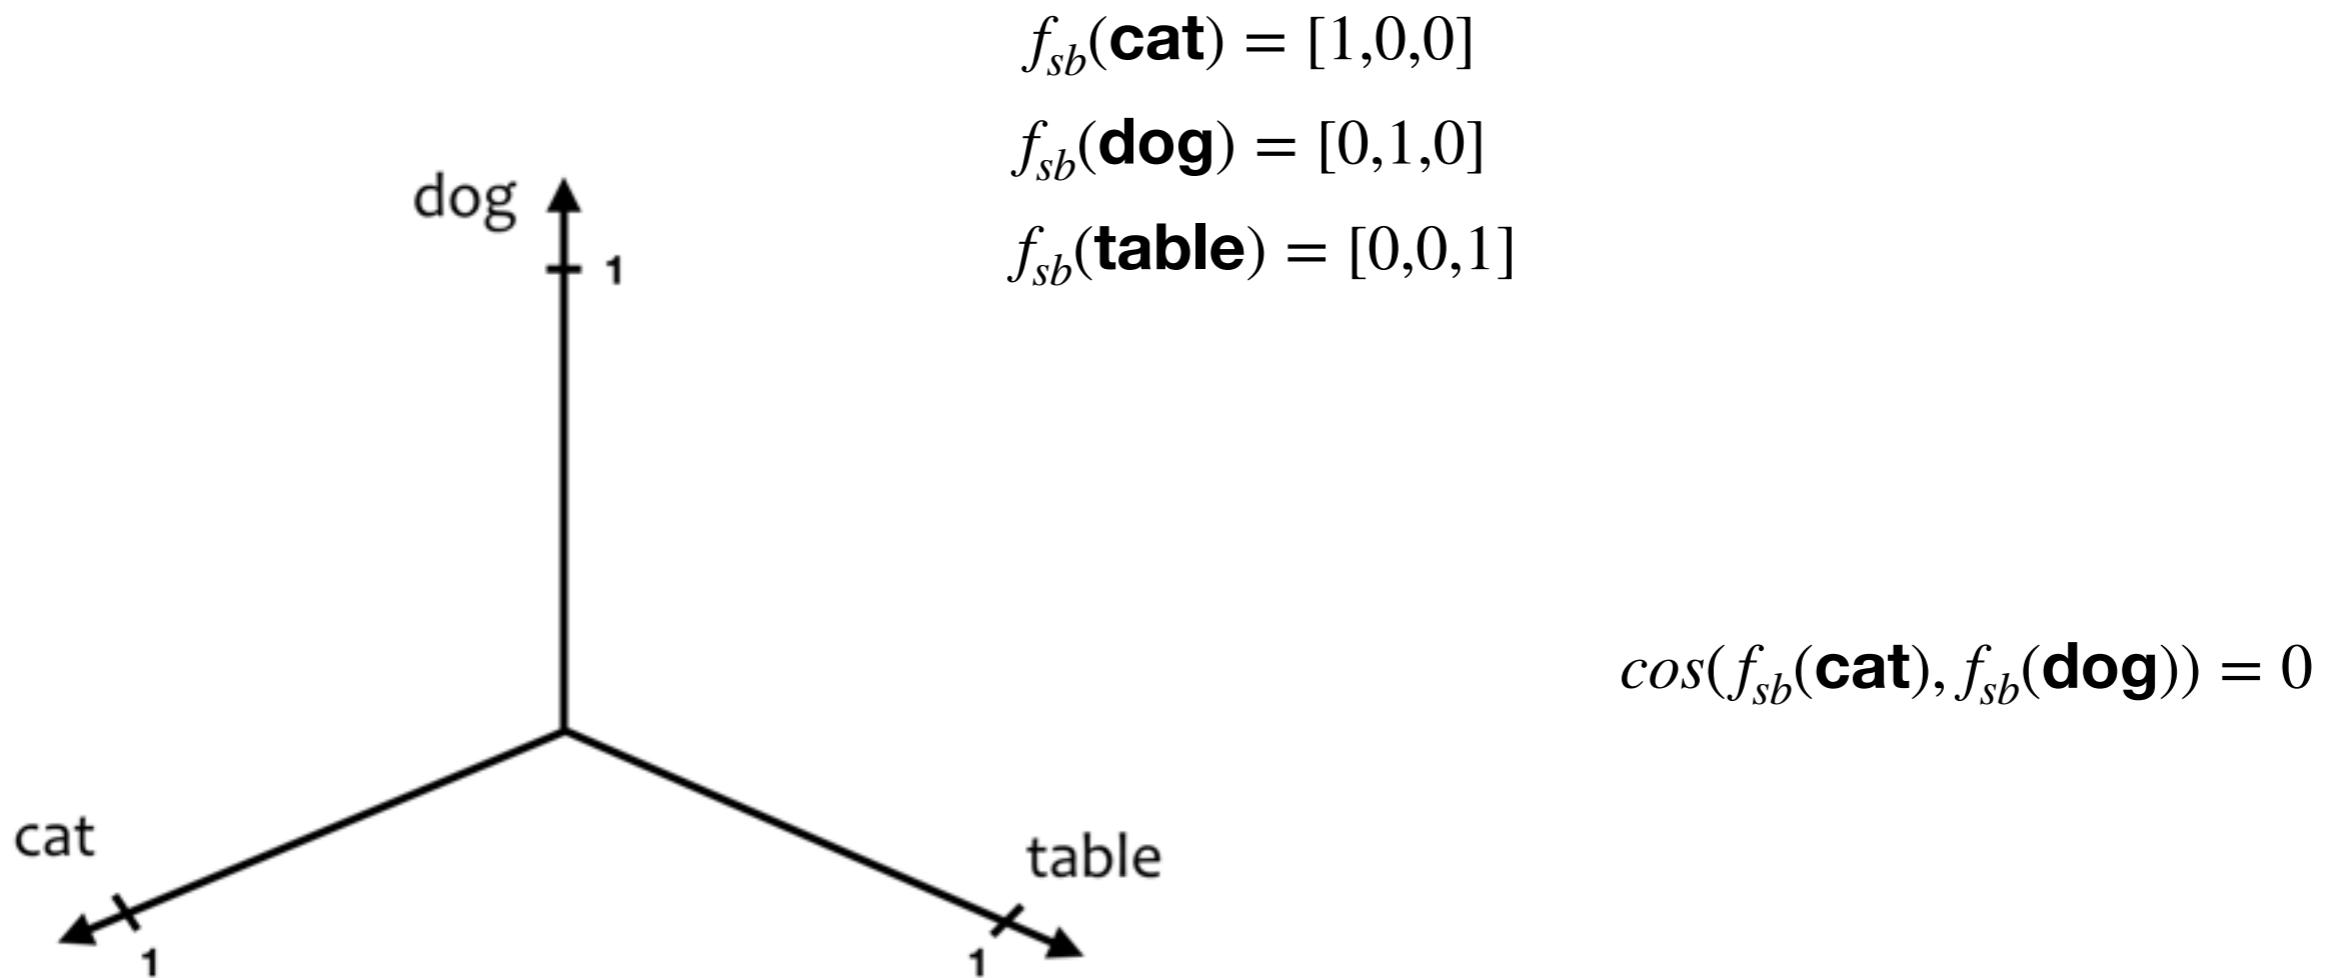
**one-hot vector**  
(length  $V$ , one entry is 1)



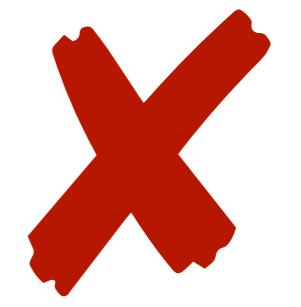
# One-hot encoding: Sparse binary repr.

- sb: sparse binary representation

$$\mathbb{V} = \{\mathbf{cat}, \mathbf{dog}, \mathbf{table}\}$$



# Sparse binary representations

- ▶ Representations are distinct 
- ▶ Similar symbols have similar representations 
- ▶ Despite of this, n-hot representations are often very **powerful** for text classification.

**From sparse high-dim  
to continuous low-dim**

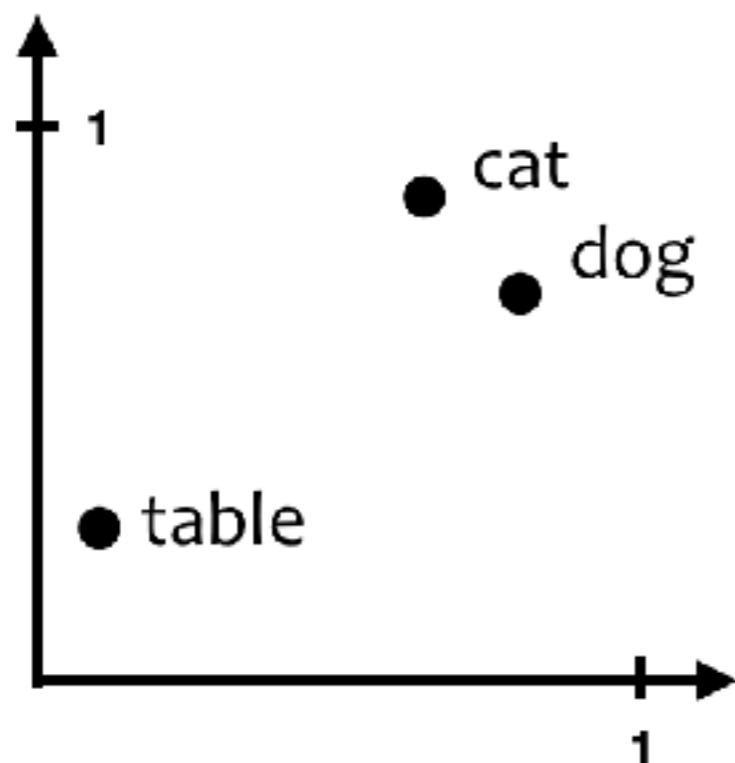
# Dense continuous: Embeddings

- ▶ “Embed” symbol  $f_{dc}(w) \mapsto \mathbb{R}^d$  in dense low-dimensional space ( $d \ll |V|$ )
- ▶ Dimensionality  $d$  (hyperparameter)

$$V = \{\text{cat}, \text{dog}, \text{table}\}$$

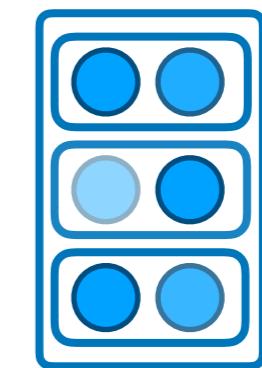
$$d = 2$$

$$E \in \mathbb{R}^{3 \times 2}$$



$$\begin{aligned}f_{sb}(\text{cat}) &= [0.7, 0.8] \\f_{sb}(\text{dog}) &= [0.75, 0.6] \\f_{sb}(\text{table}) &= [0.1, 0.15]\end{aligned}$$

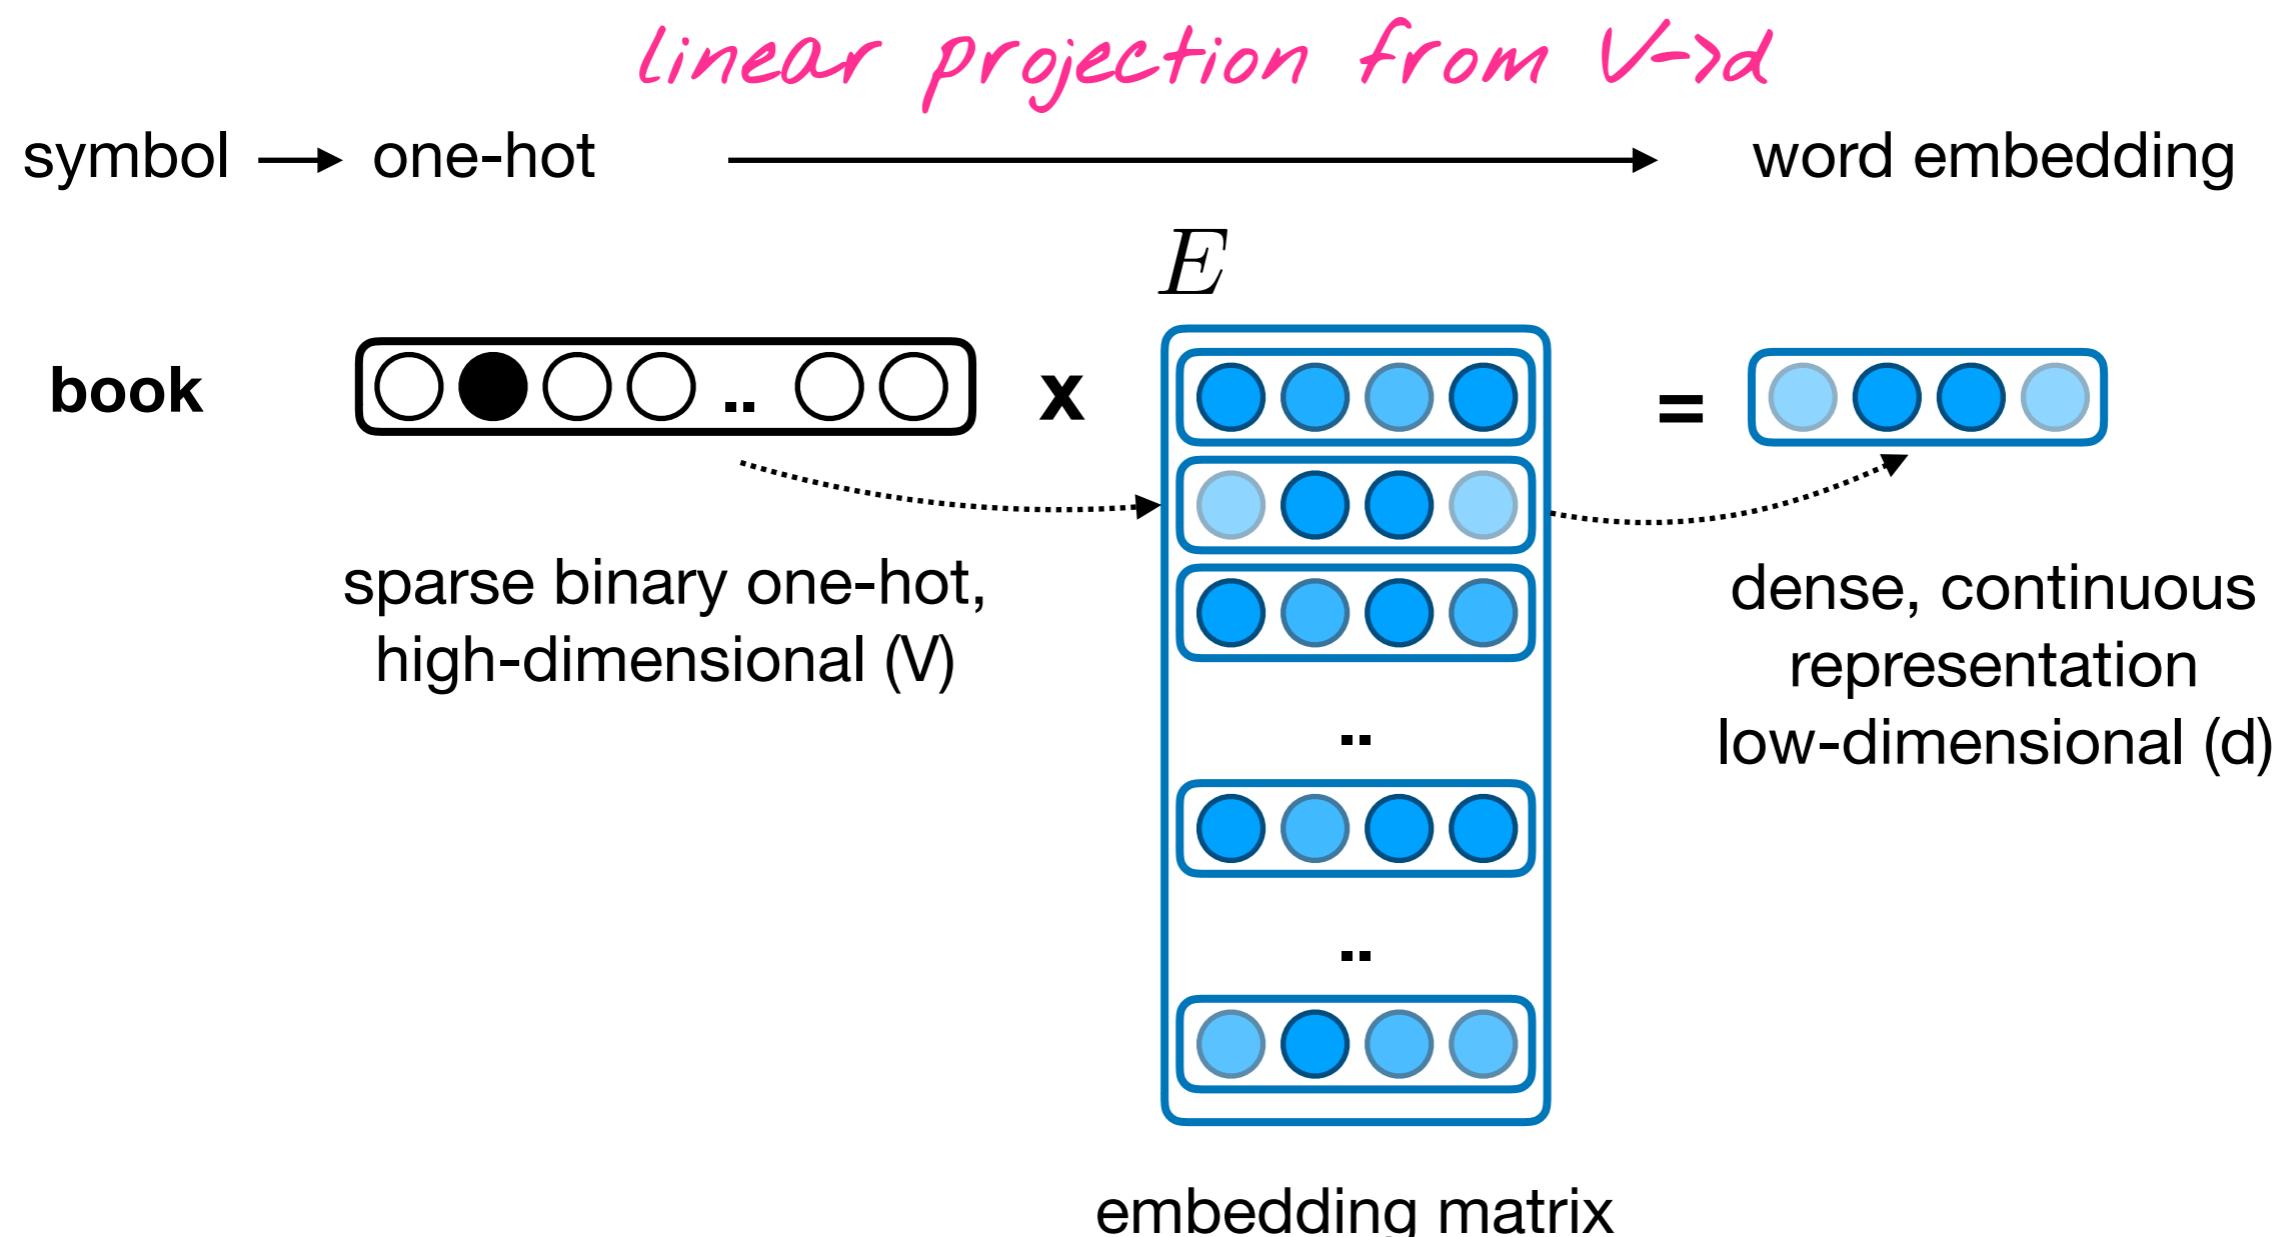
$$E$$



$|V| \times d$   
embedding matrix

Note:  $d < |V|$

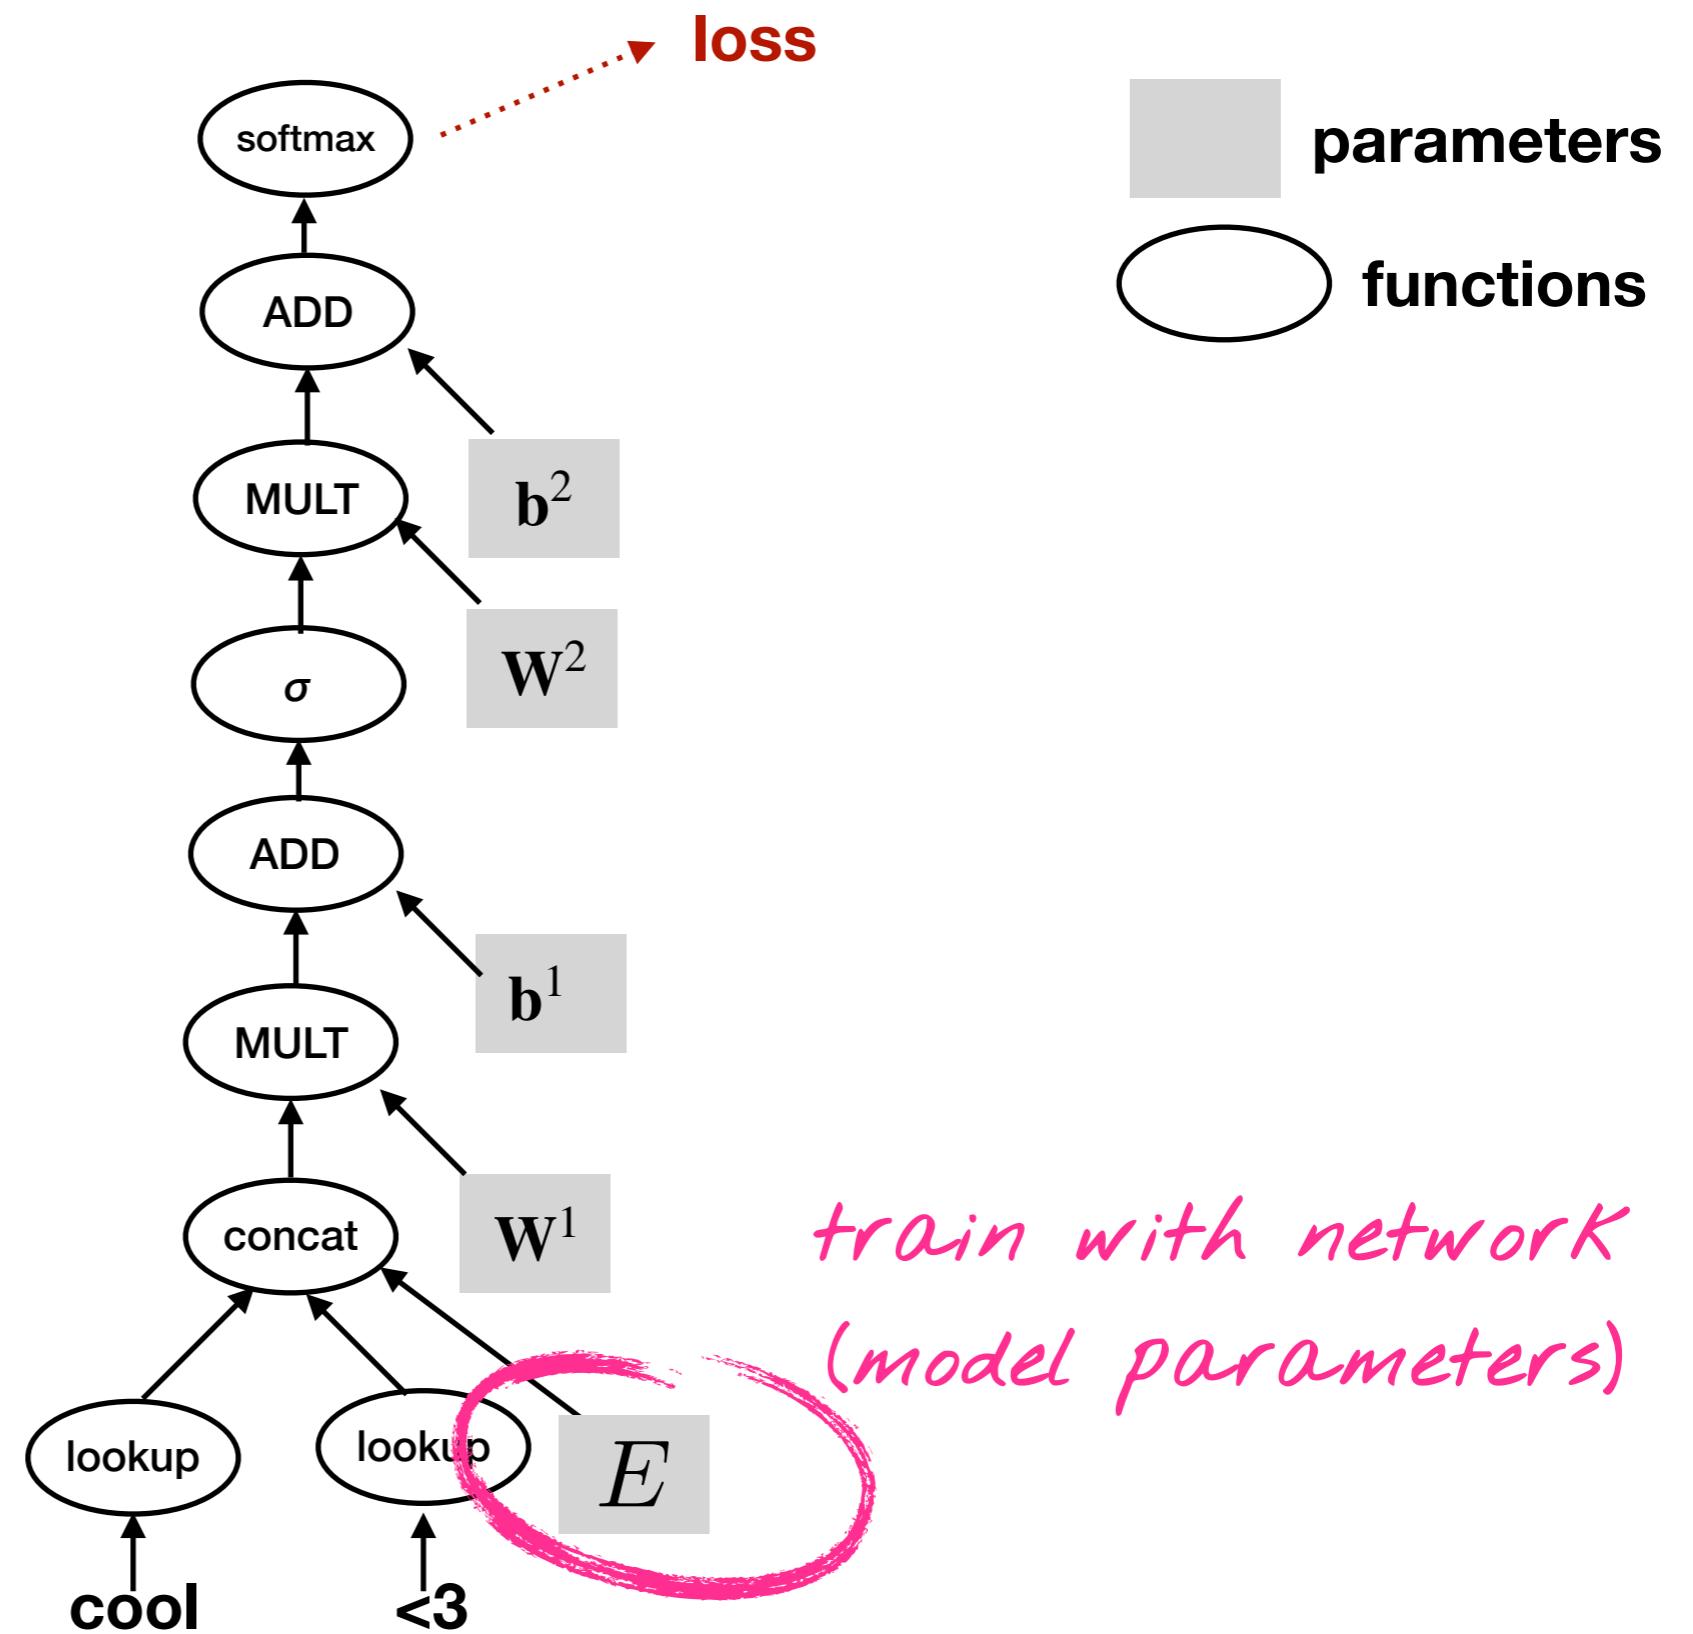
# Lookup: Representing a symbol



# In general, the **neural** way for extracting features:

- ▶ Extract **core linguistic features**  $f_1, \dots f_n$
- ▶ **Define** a **vector** for each feature (lookup Embedding table)
  - ▶ Can train representation E together with the network

# Computational Graph View



# Dense continuous text encodings

this is great !

cool <3

this

is

great

!



cool

<3



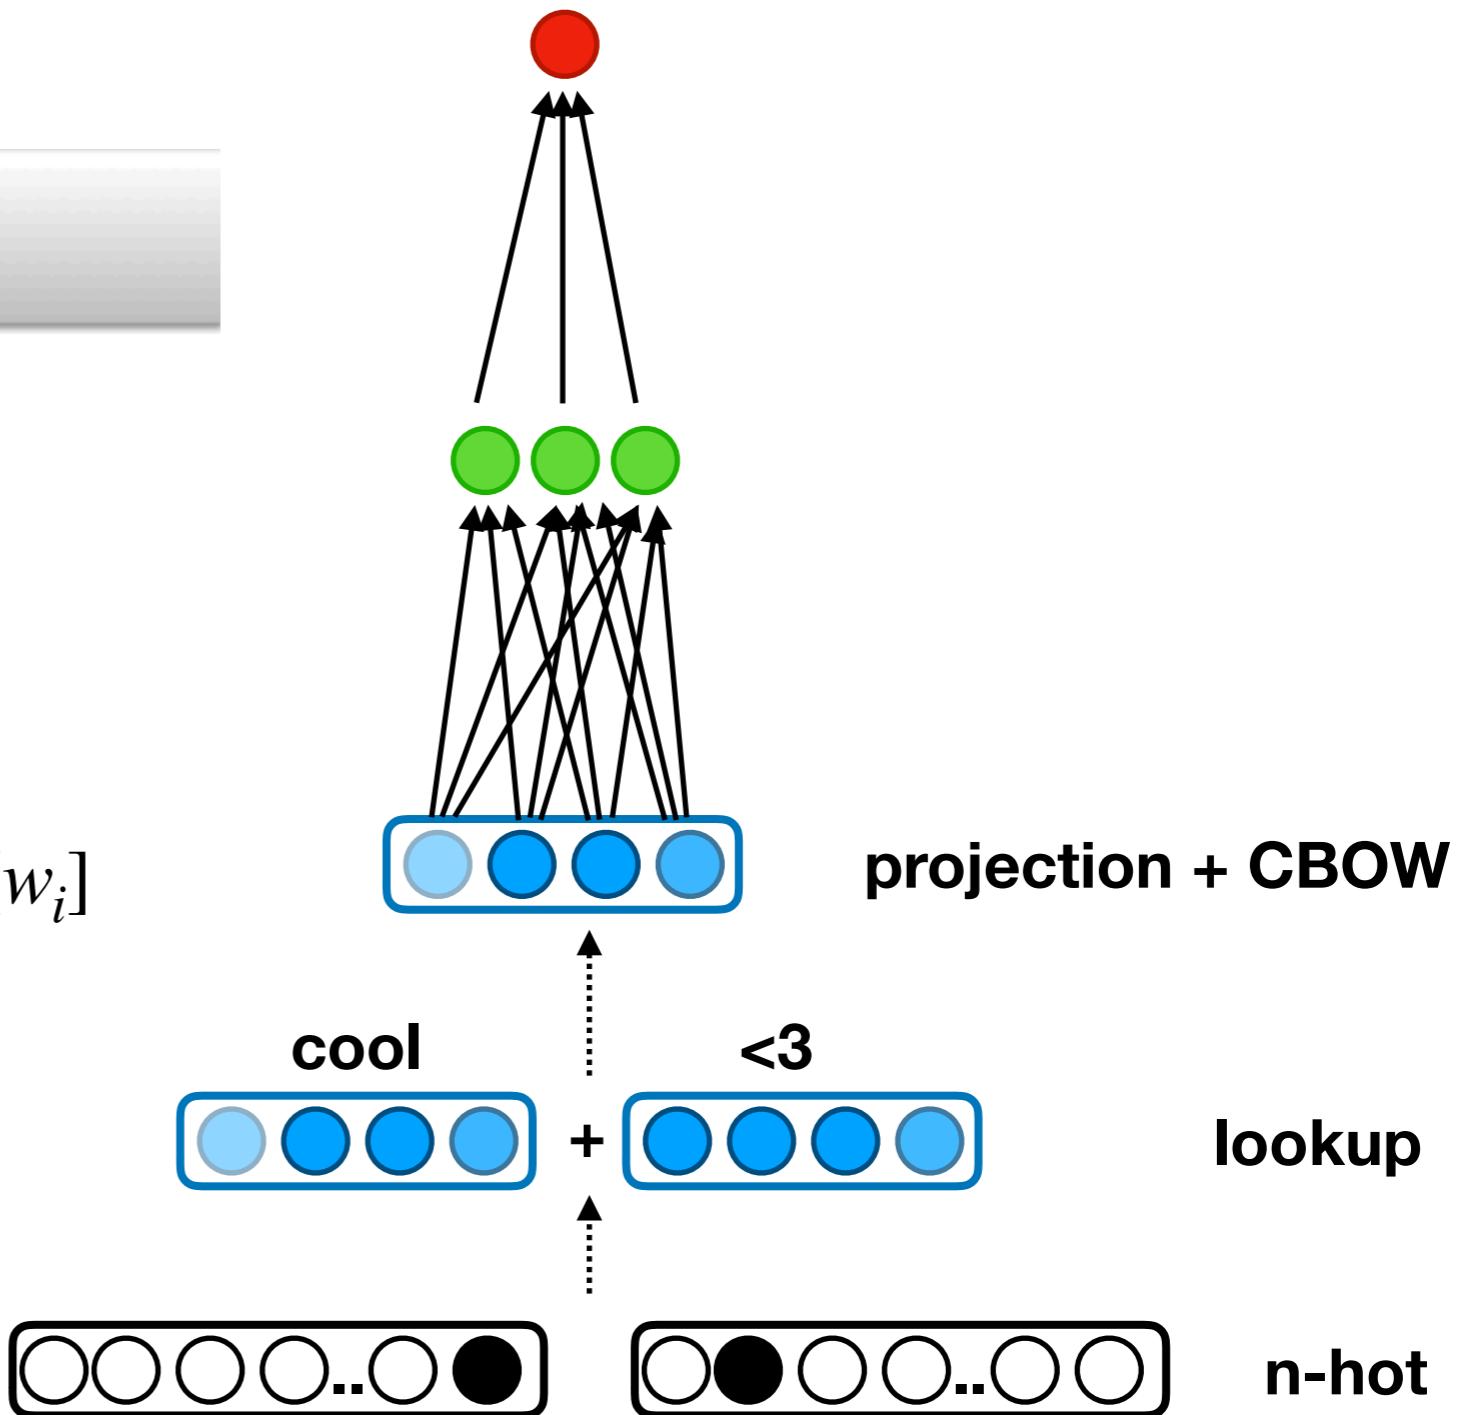
How to combine word embeddings?

# Dense continuous text encoding: e.g. continuous BOW (CBOW)

Example input document 1:

cool <3

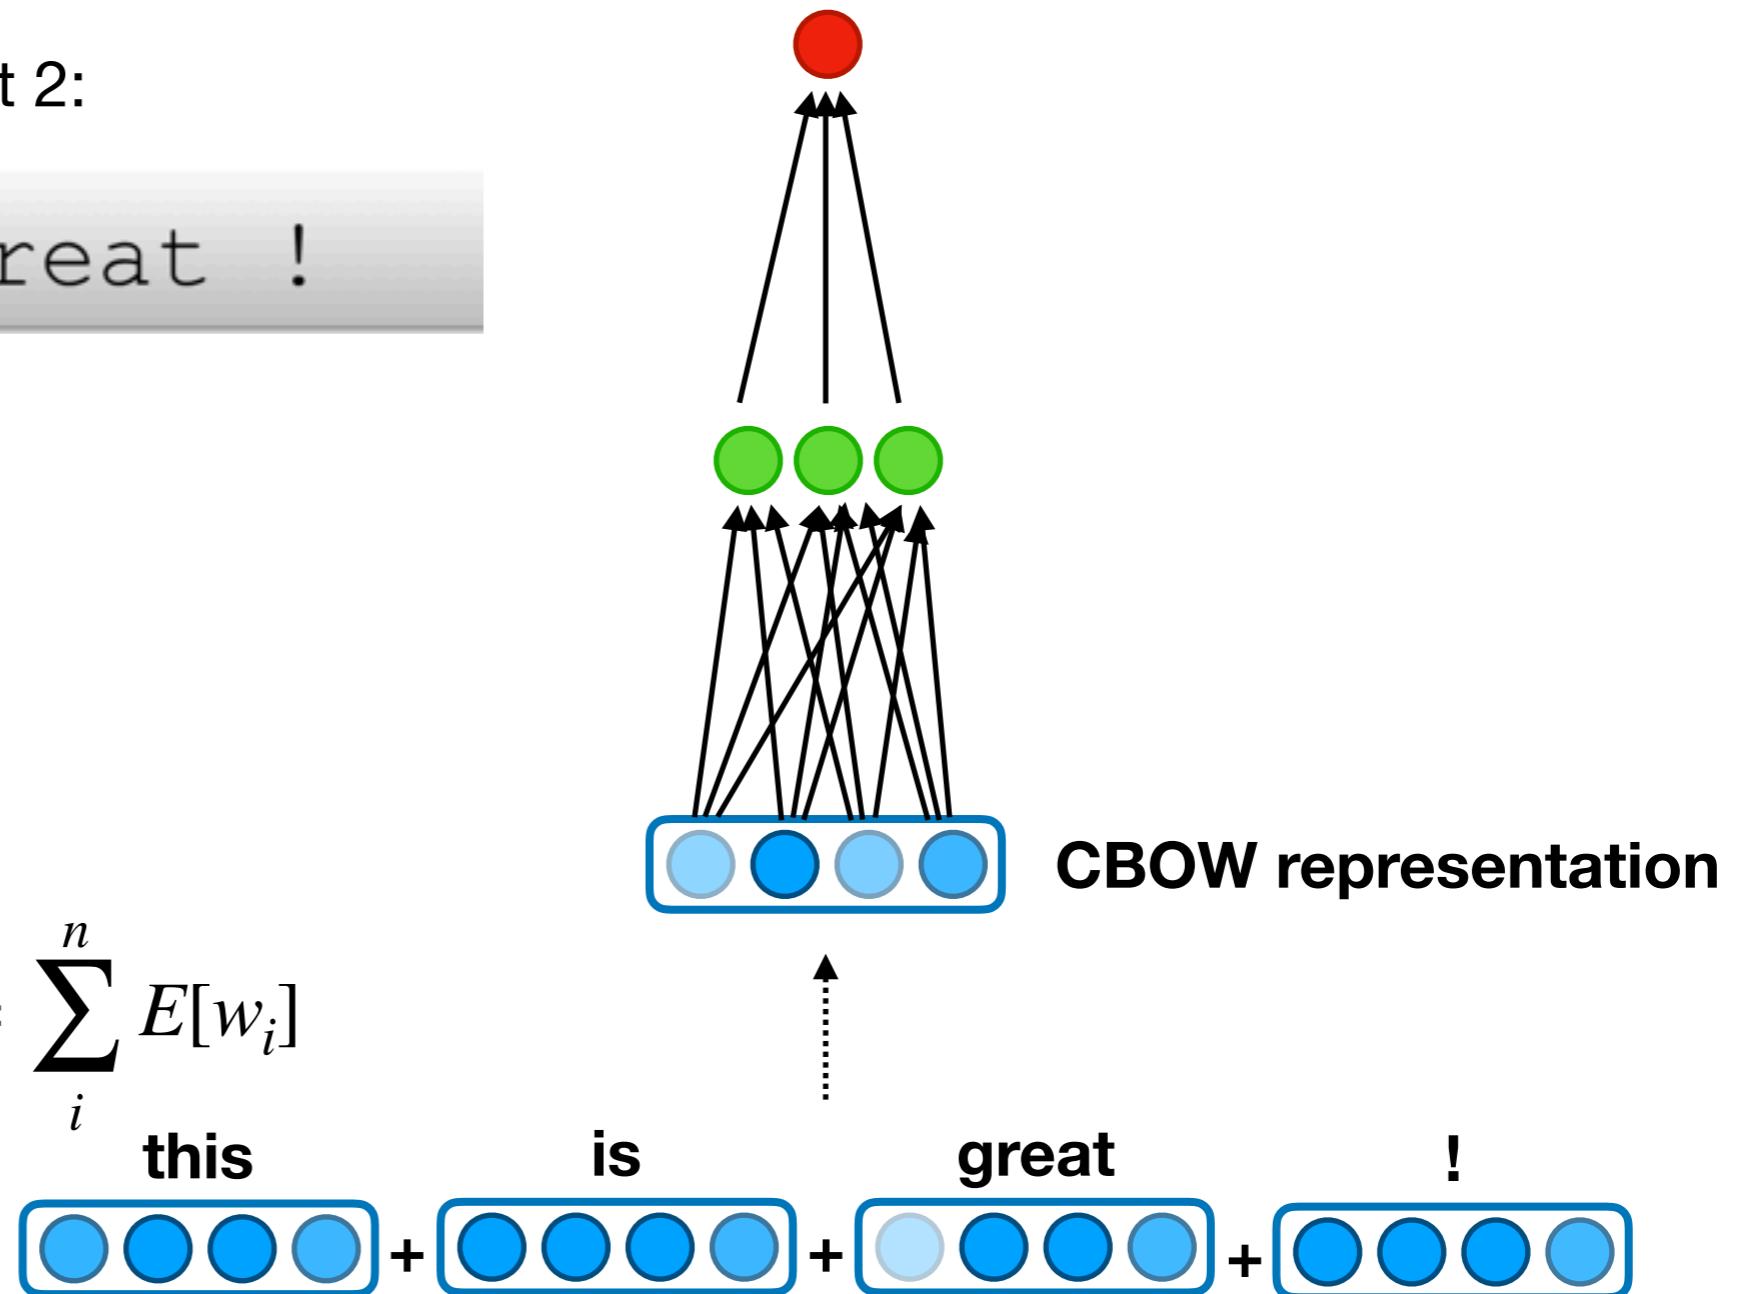
$$\text{CBOW}(w_i, \dots, w_n) = \sum_i^n E[w_i]$$



# Dense continuous text encoding: e.g. continuous BOW (CBOW)

Example input document 2:

this is great !



# Limitation of BOW

- ▶ What's the biggest limitation of BOW/CBOW?
  - ▶ Similar to unigram model: It **disregards the order of items** (e.g. words in a sequence)
  - ▶ Example:
    - ▶ “it was not good, it was actually terrible”
    - ▶ “it was not terrible, it was actually good”
- ▶ A simple solution?

# Possible Improvement

- ▶ Bag of n-grams
  - ▶ “not good”, ...
- ▶ Problems:
  - ▶ Parameter explosion (BOW/n-hot) or even more averaging (CBOW)
  - ▶ No sharing between similar words & n-grams

# Where to get task-specific $E$ from? From Scratch vs Pre-trained

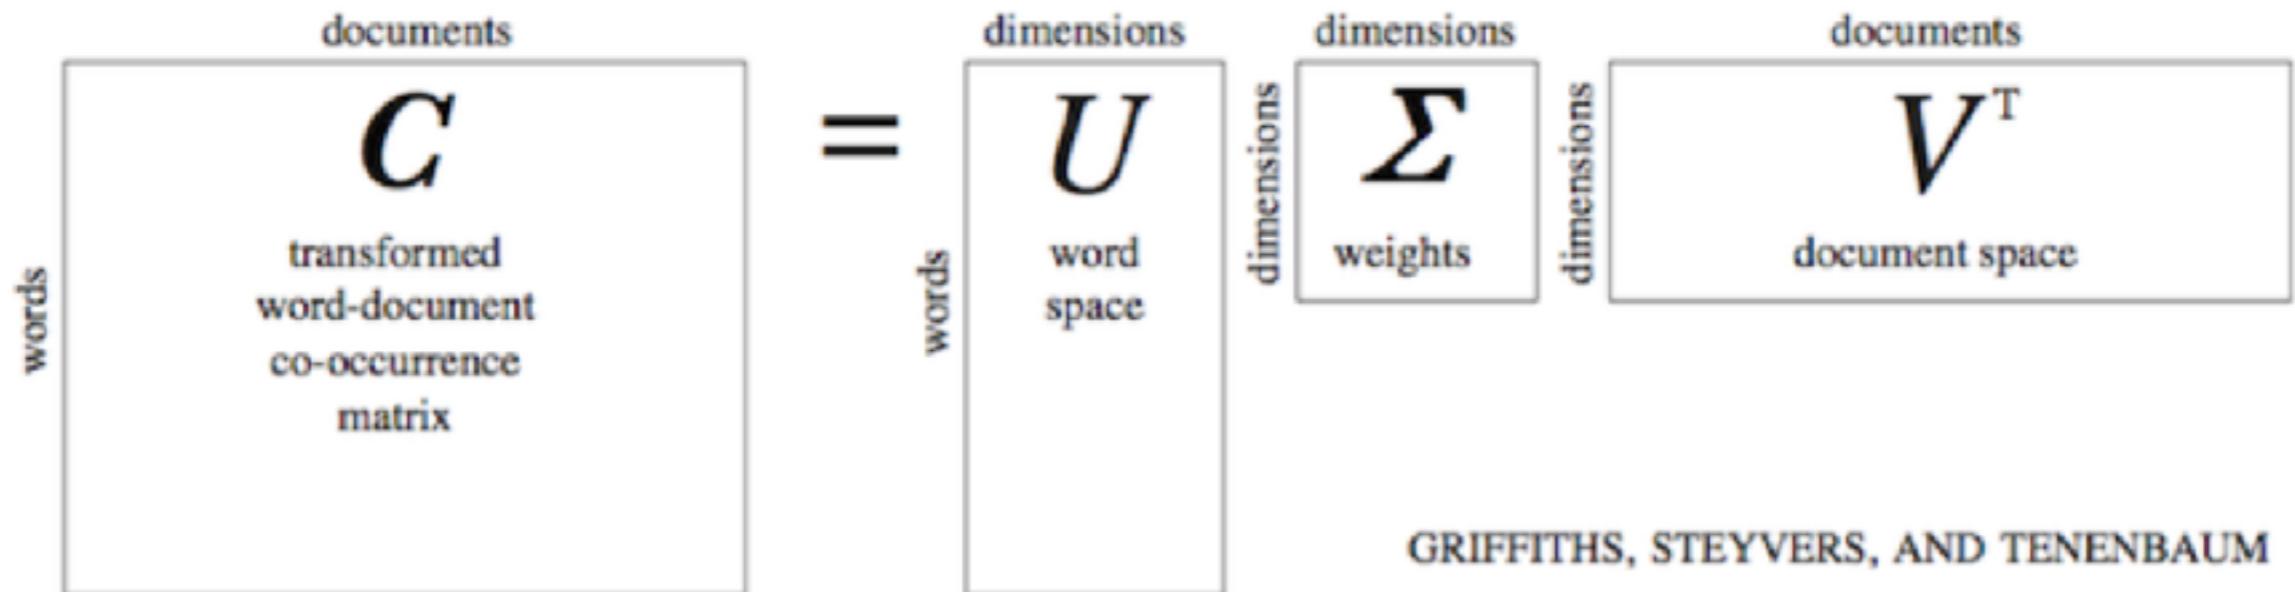
- ▶ Embedding layer  $E$ :
  - ▶ trained with network from **scratch** - task-specific
  - ▶ **initialized** with off-the-shelf pre-trained word embeddings (e.g., Glove, Polyglot, fastText)
- ▶ Pre-trained embedding **initialization** typically leads to performance gains. Why?
  - ▶ train on **more words**
  - ▶ implicitly **more data**
- ▶ Ways to obtain off-the-shelf embeddings? (word vector space representation?)

# Embeddings: New? No!

- ▶ **Two major methods:**
  - ▶ **Count!** (pre-deep learning method, aka “word vector space models”)
  - ▶ **Predict!** (core idea underlying word2vec - lecture 1)

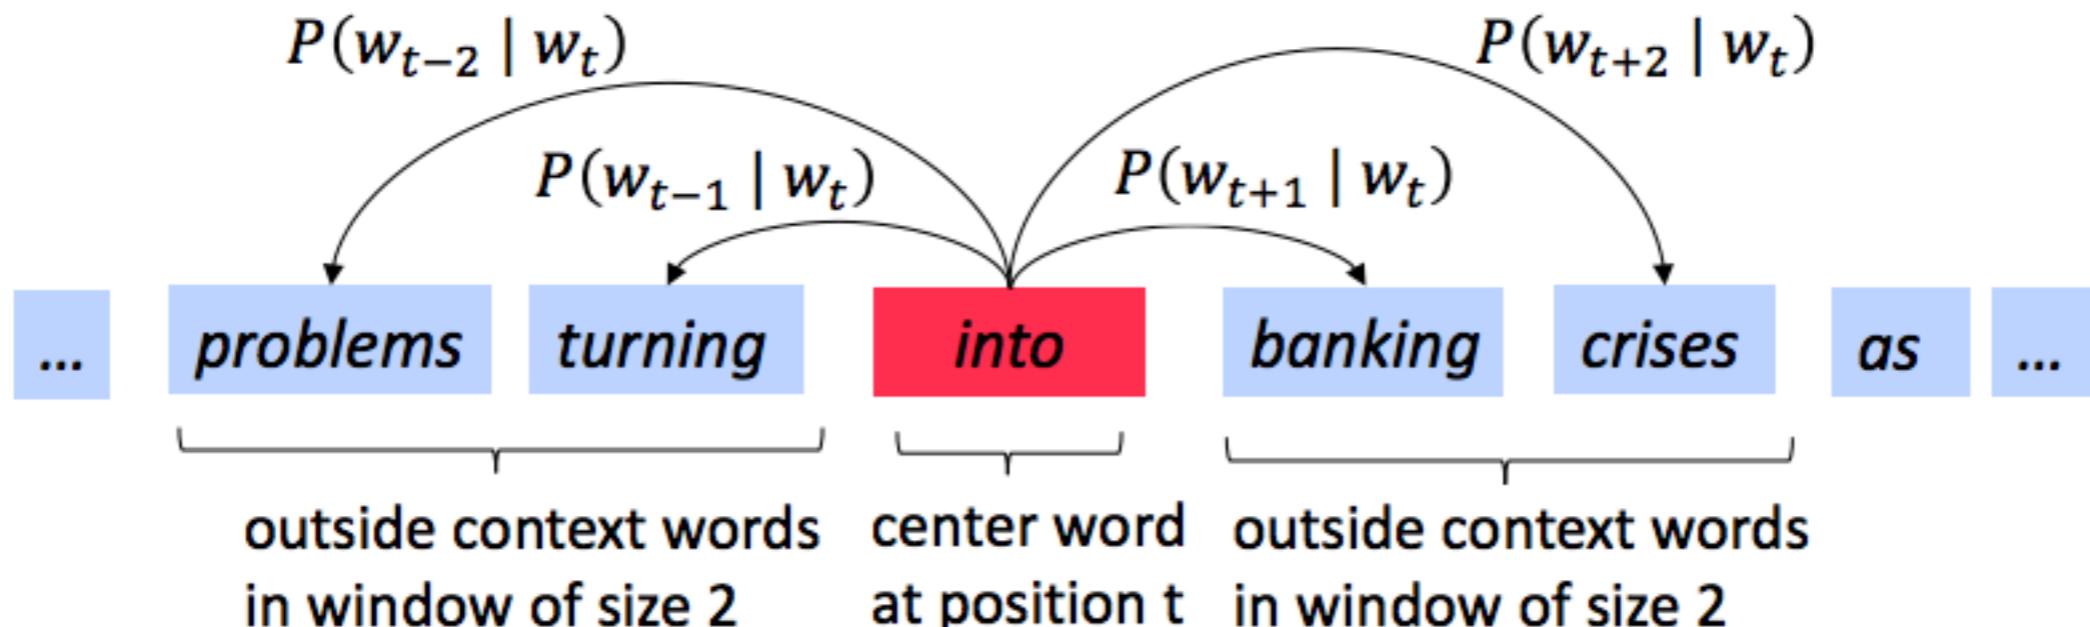
# Count-based methods

- ▶ Represent the “company” of a word in terms of a word **co-occurrence-matrix**, get the statistics (counts)
- ▶ E.g. **Latent Semantic Analysis** (LSA) (Deerwester et al., 1990)  
SVD decomposition over co-occurrence matrix to reduce to lower-dimensional space (matrix  $U$  where  $\text{dim} < |\text{docs}|$ )



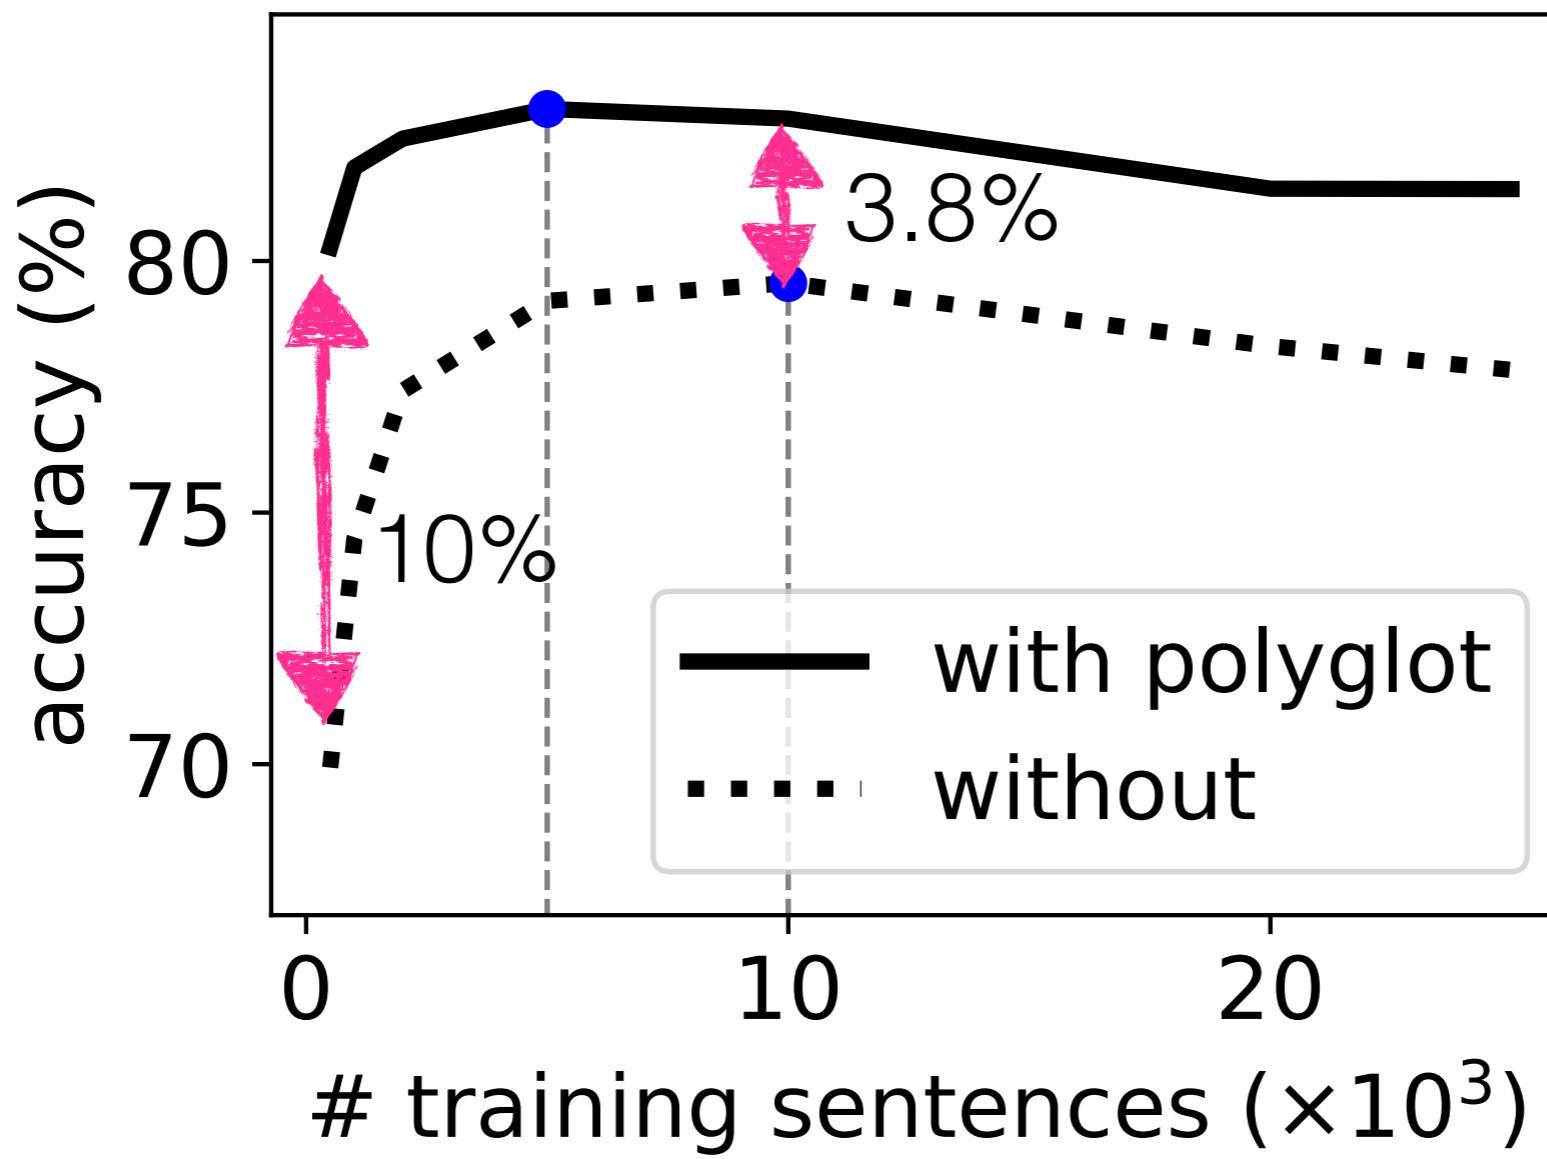
# Prediction-based methods

- ▶ **Key idea:** predict the context of a word (instead of capturing co-occurrence statistics in matrix C) to directly learn the low-dimensional word vector representation
- ▶ **Word2vec** (family of methods) Mikolov et al. (2013)  
[Lecture 1 by Ryan] - scales well to large data



# Example: Cross-lingual POS tagging- Word embedding initialization

(Plank & Agic, 2018)



Mean over 21 languages

# Overview

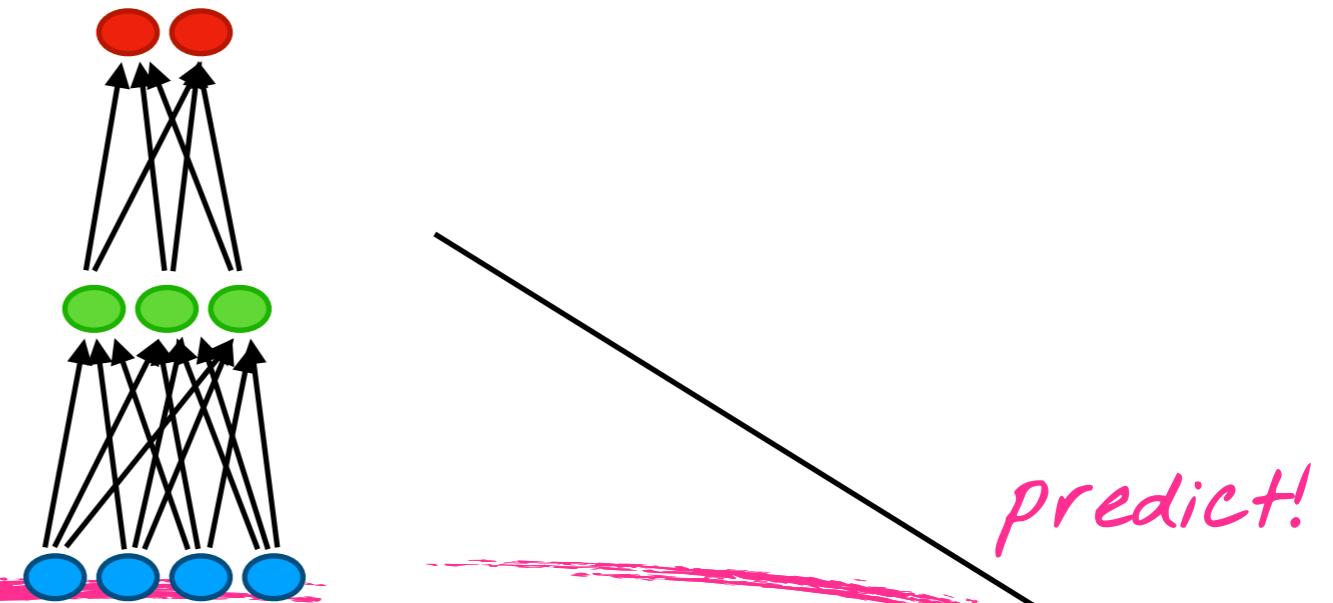
*foundations*



*representations*



# Window-based neural LM via FFNN



~~As the clock rang the students opened the~~

*fixed window of n words*

# A fixed-window based neural LM

output distribution

$$\hat{y} = \text{softmax}(Uh + b_2) \in \mathbb{R}^{|V|}$$

hidden layer

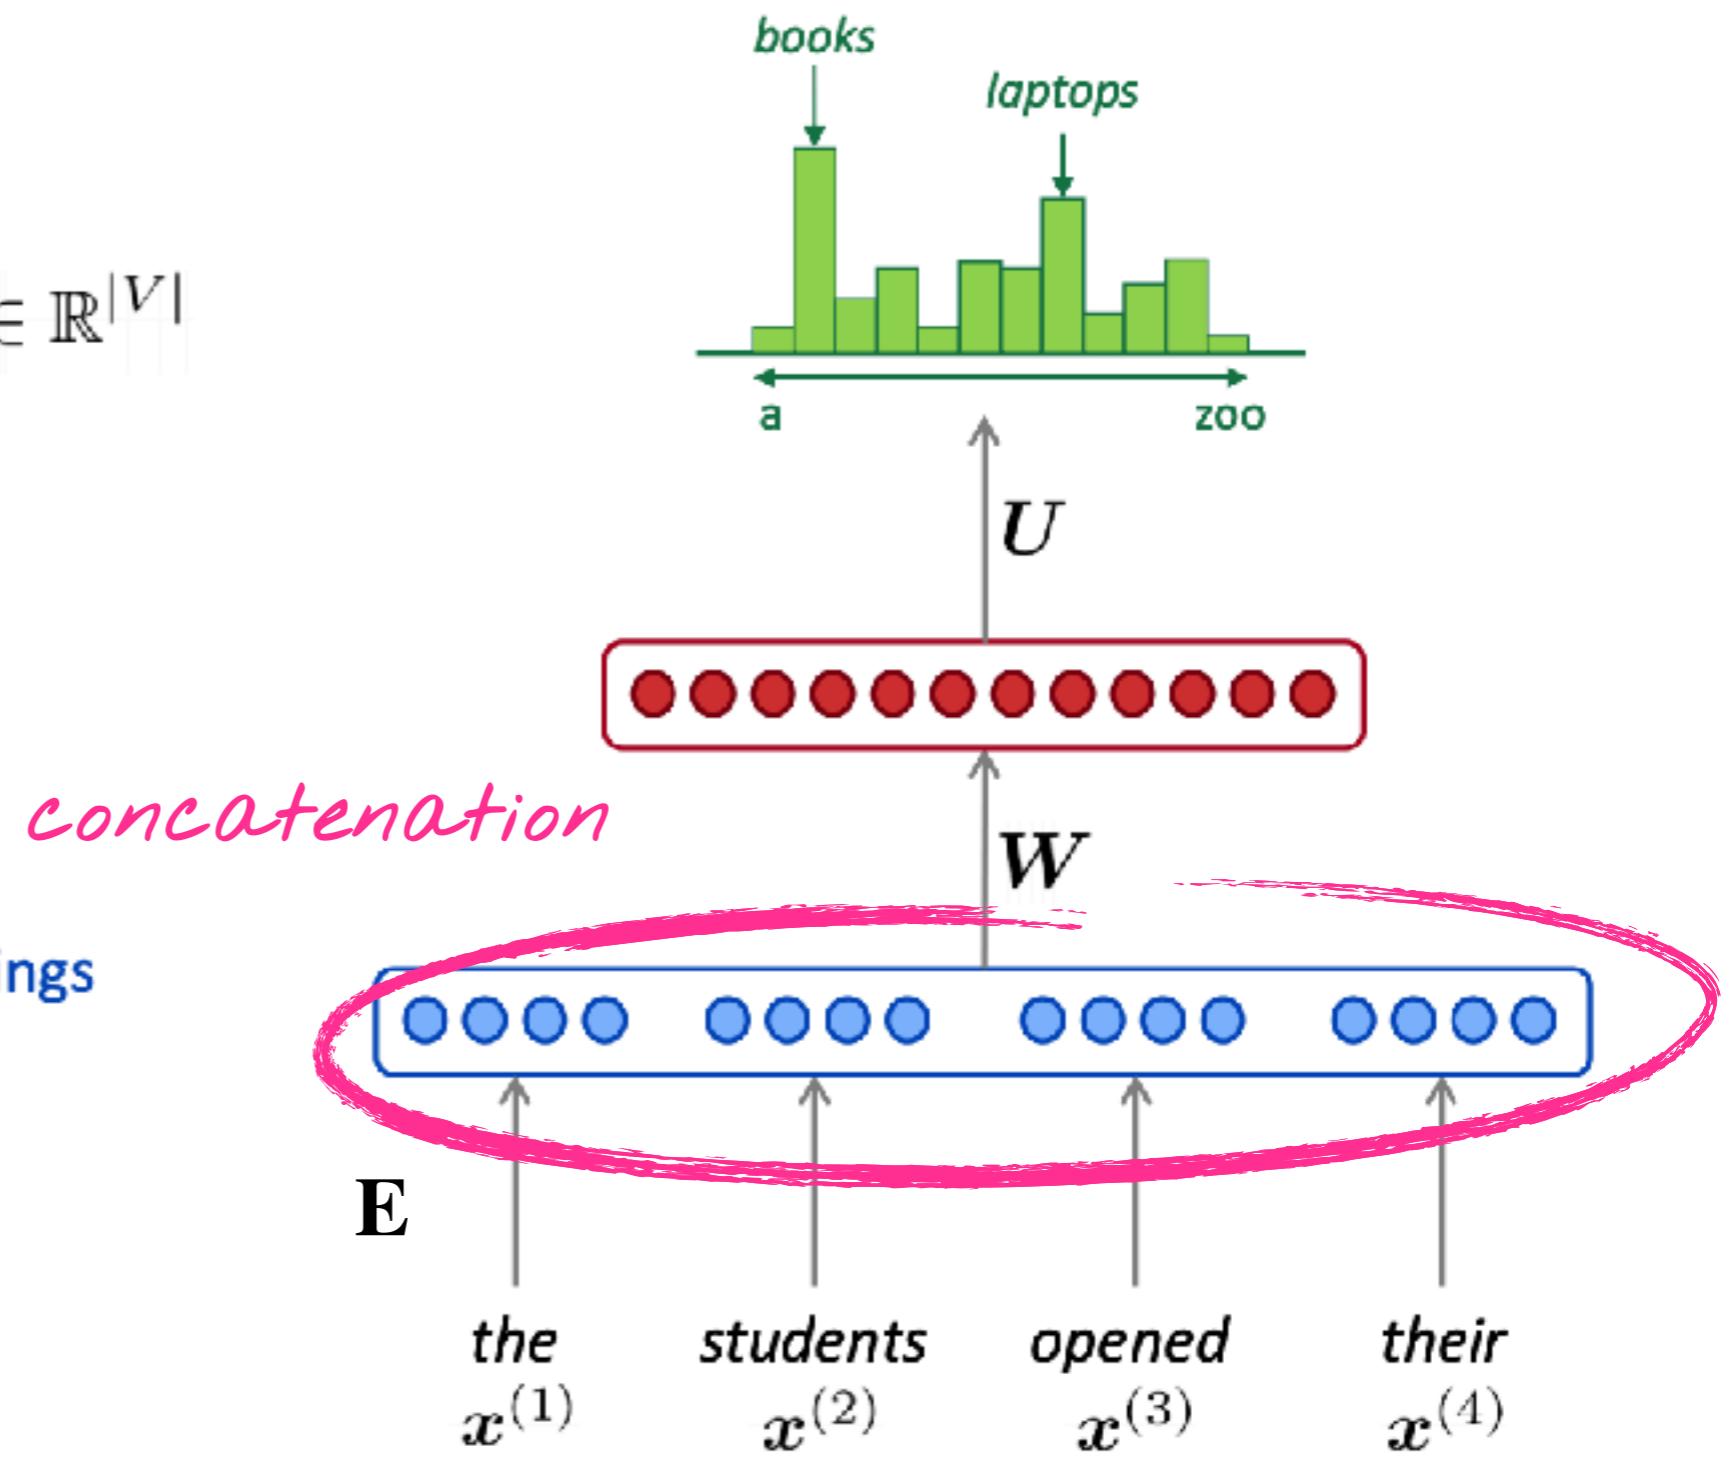
$$h = f(We + b_1)$$

concatenated word embeddings

$$e = [e^{(1)}; e^{(2)}; e^{(3)}; e^{(4)}]$$

words / one-hot vectors

$$x^{(1)}, x^{(2)}, x^{(3)}, x^{(4)}$$



# Training the Neural n-gram LM

- ▶ Iteratively move the n-gram window through a very large corpus to predict the next word at each time step
- ▶ Cross-entropy loss (negative log-likelihood):

$$L = -\log p(w_t | w_{t-1} \dots w_{t-n+1})$$

- ▶ Note: typically very large vocabulary (softmax)
  - ▶ Workaround: negative sampling (lecture 1)

# What about these issues?

- ▶ Can it handle similar words?
  - ▶ she *bought* a bicycle
  - ▶ she *purchased* a bicycle
- ▶ Long-distance dependencies?
  - ▶ for *programming* she yesterday purchased her own brand new *laptop*
  - ▶ for *running* she yesterday purchased her brand new *sportswatch*



# Tips for unknown words

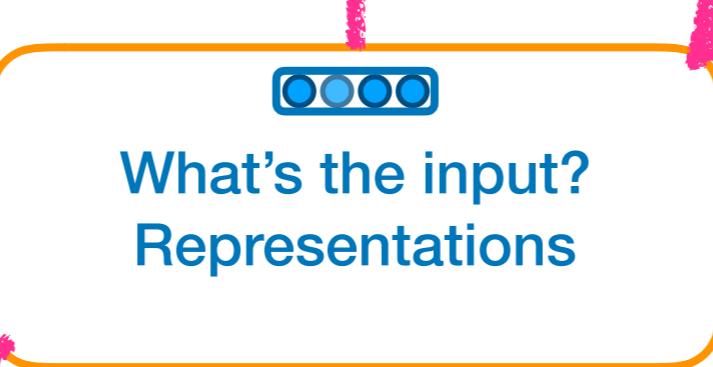
- ▶ Simplest solution:
  - ▶ Train an <UNK> word vector, e.g., map rare words to UNK (count < threshold)
- ▶ Problem:
  - ▶ Conflates a long tail into the same vector representation
  - ▶ Subword representations (character-level models) to the rescue!
  - ▶ More on these later (after we have seen CNNs)

# Overview

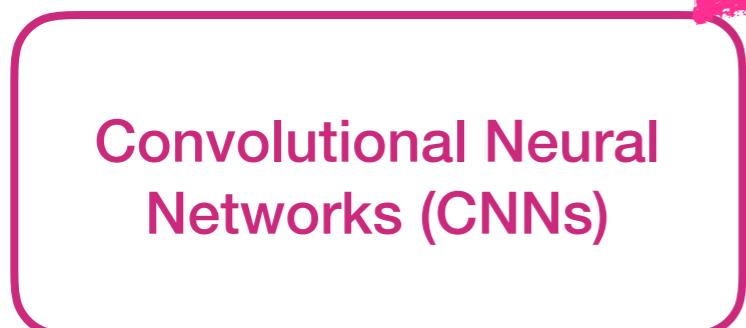
*foundations*



*representations*



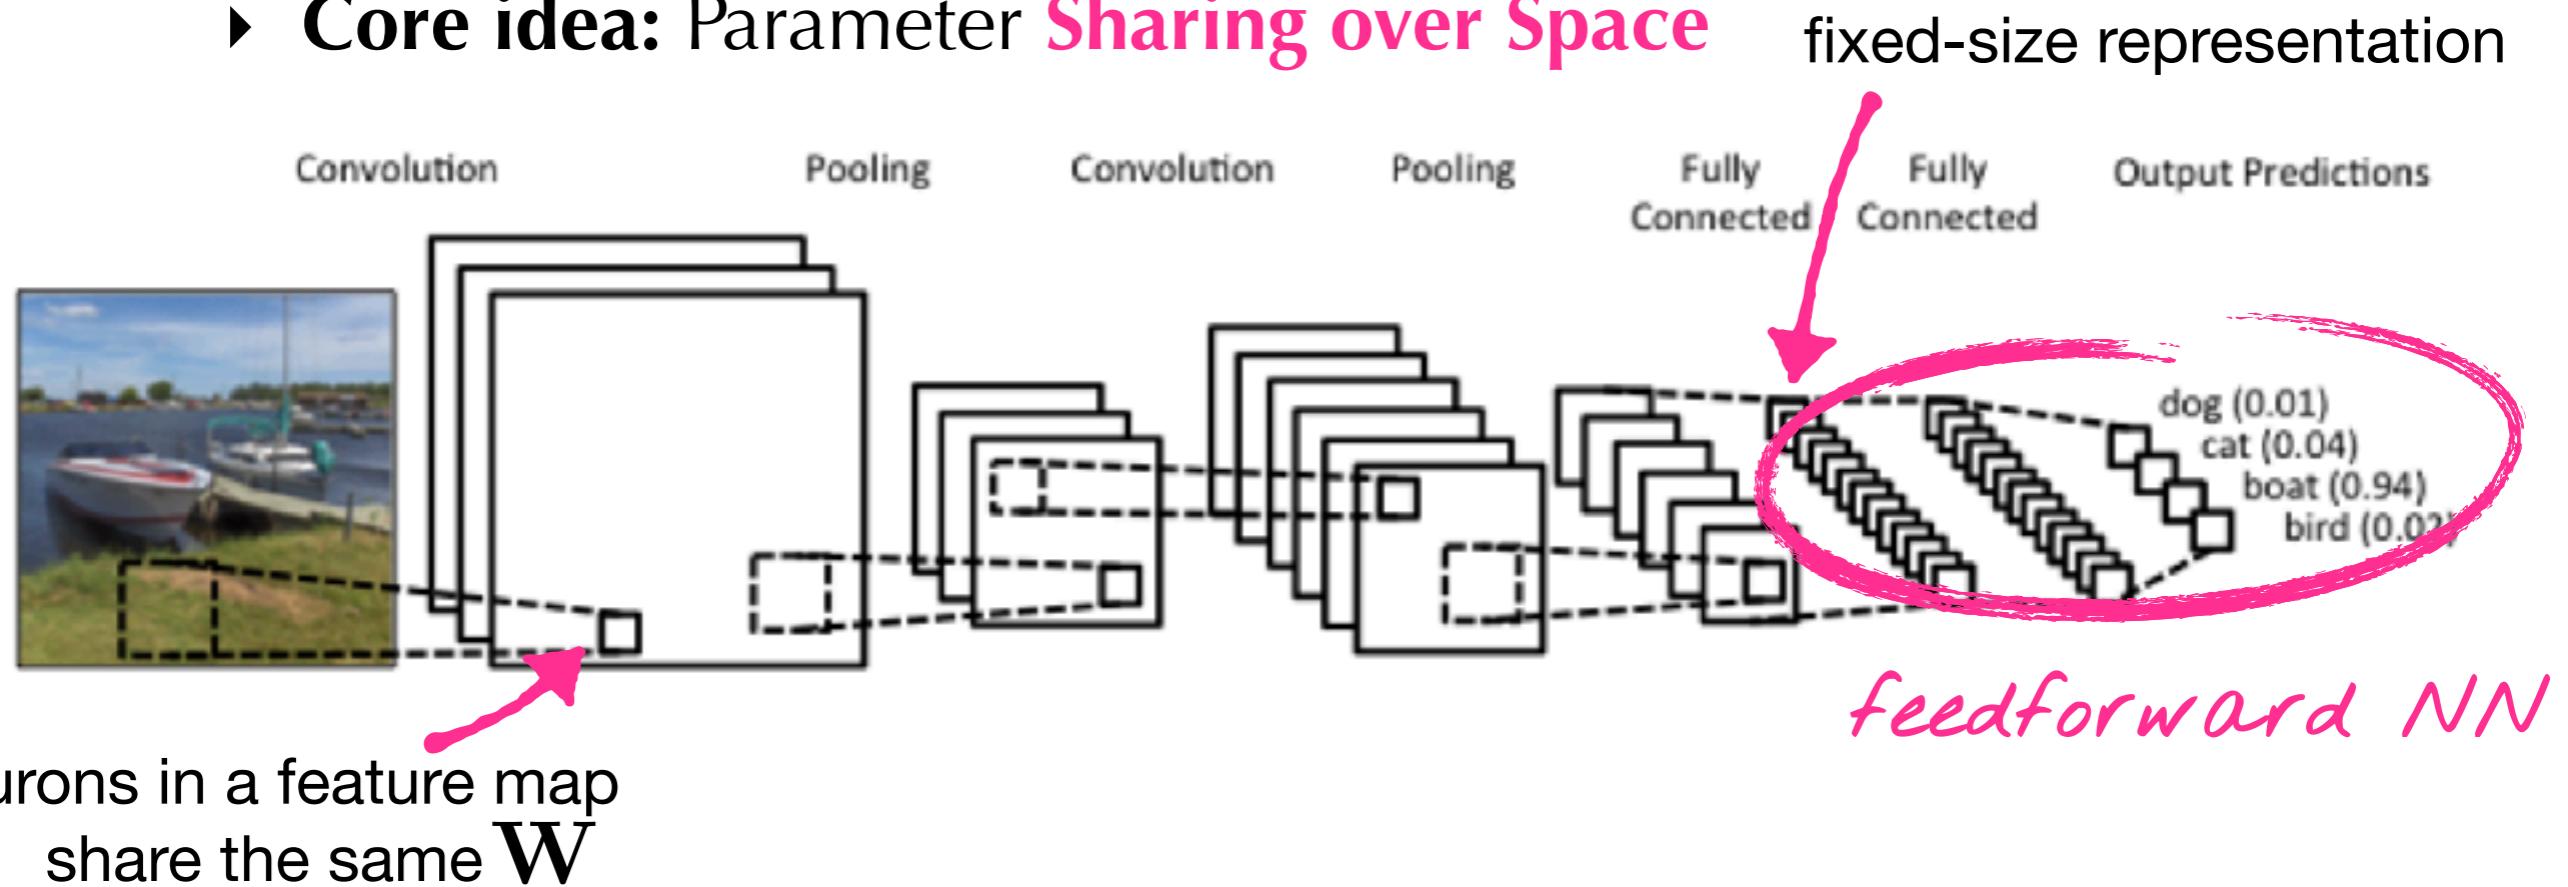
*beyond FFNNs*



# **CNNs / Convnets**

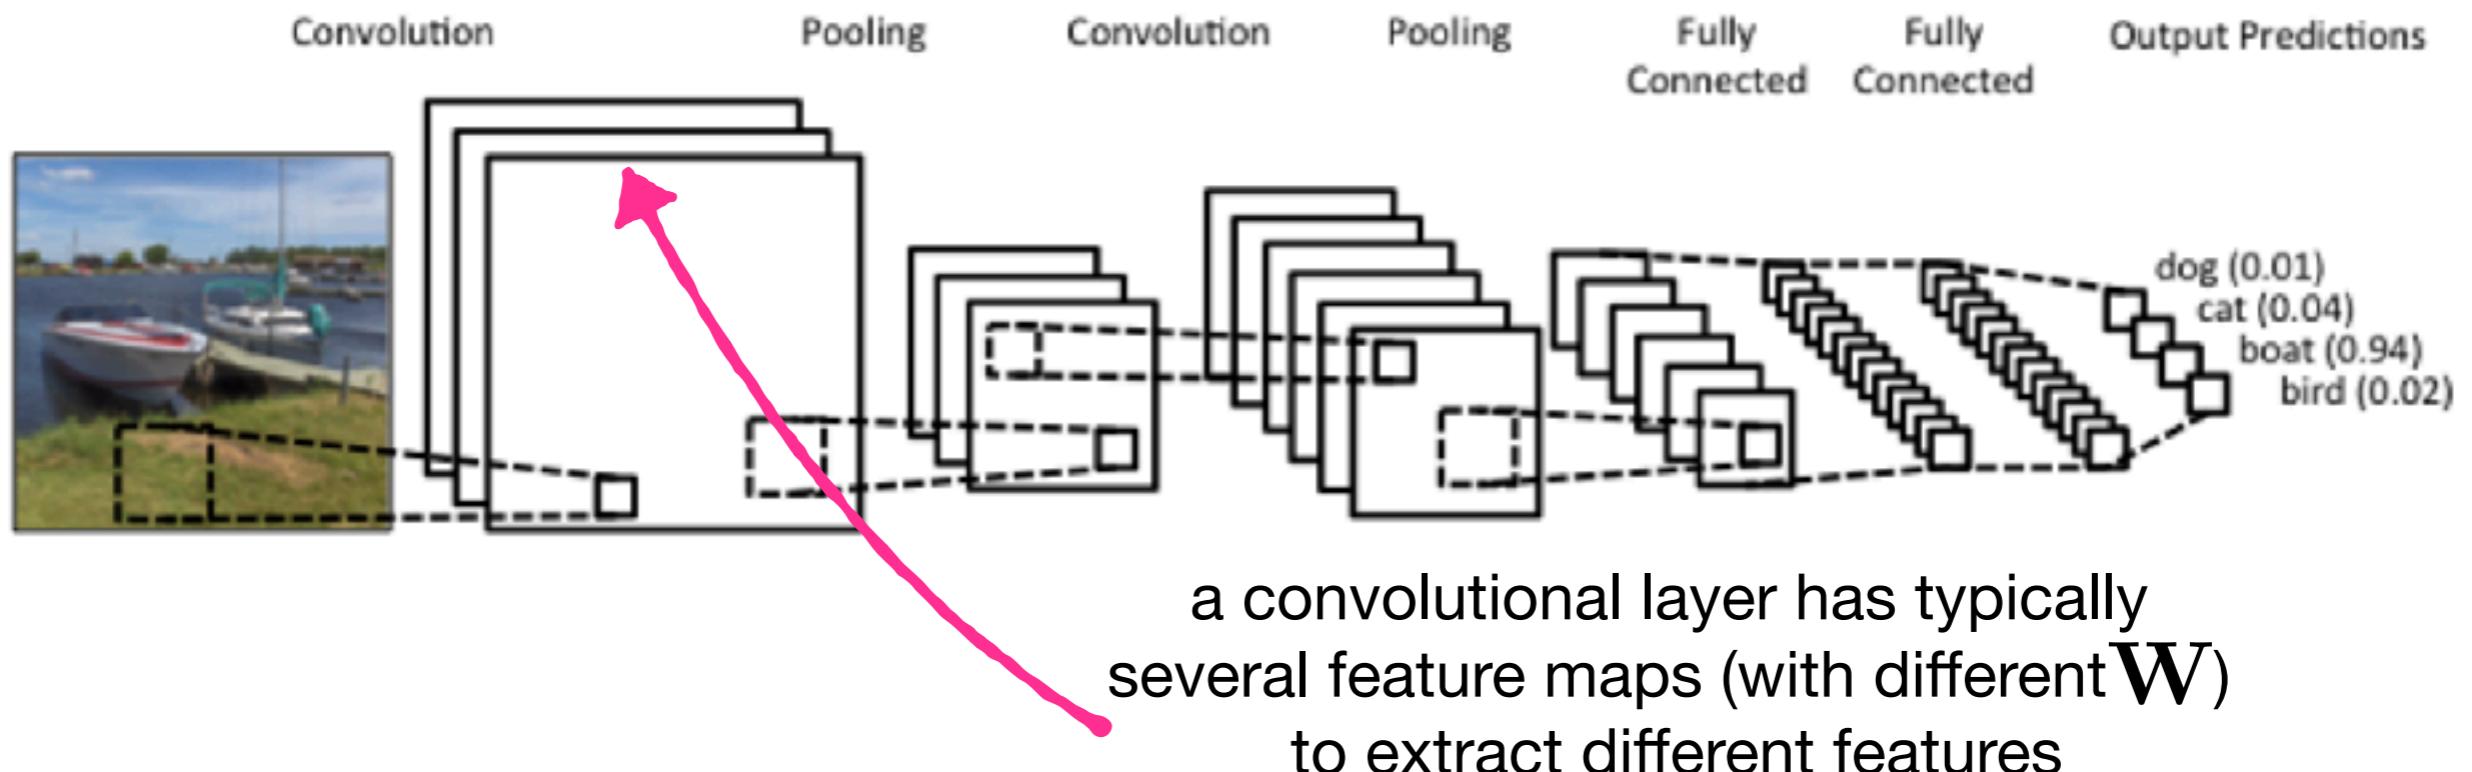
# Convolutional Neural Network (CNN)

- ▶ Neural networks for processing data with a **grid-like typology** (LeCun & Bengio, 1995)
- ▶ **Can handle arbitrary-length inputs** and reduce them down to a **fixed size vector** representation
- ▶ **Core idea:** Parameter **Sharing over Space**

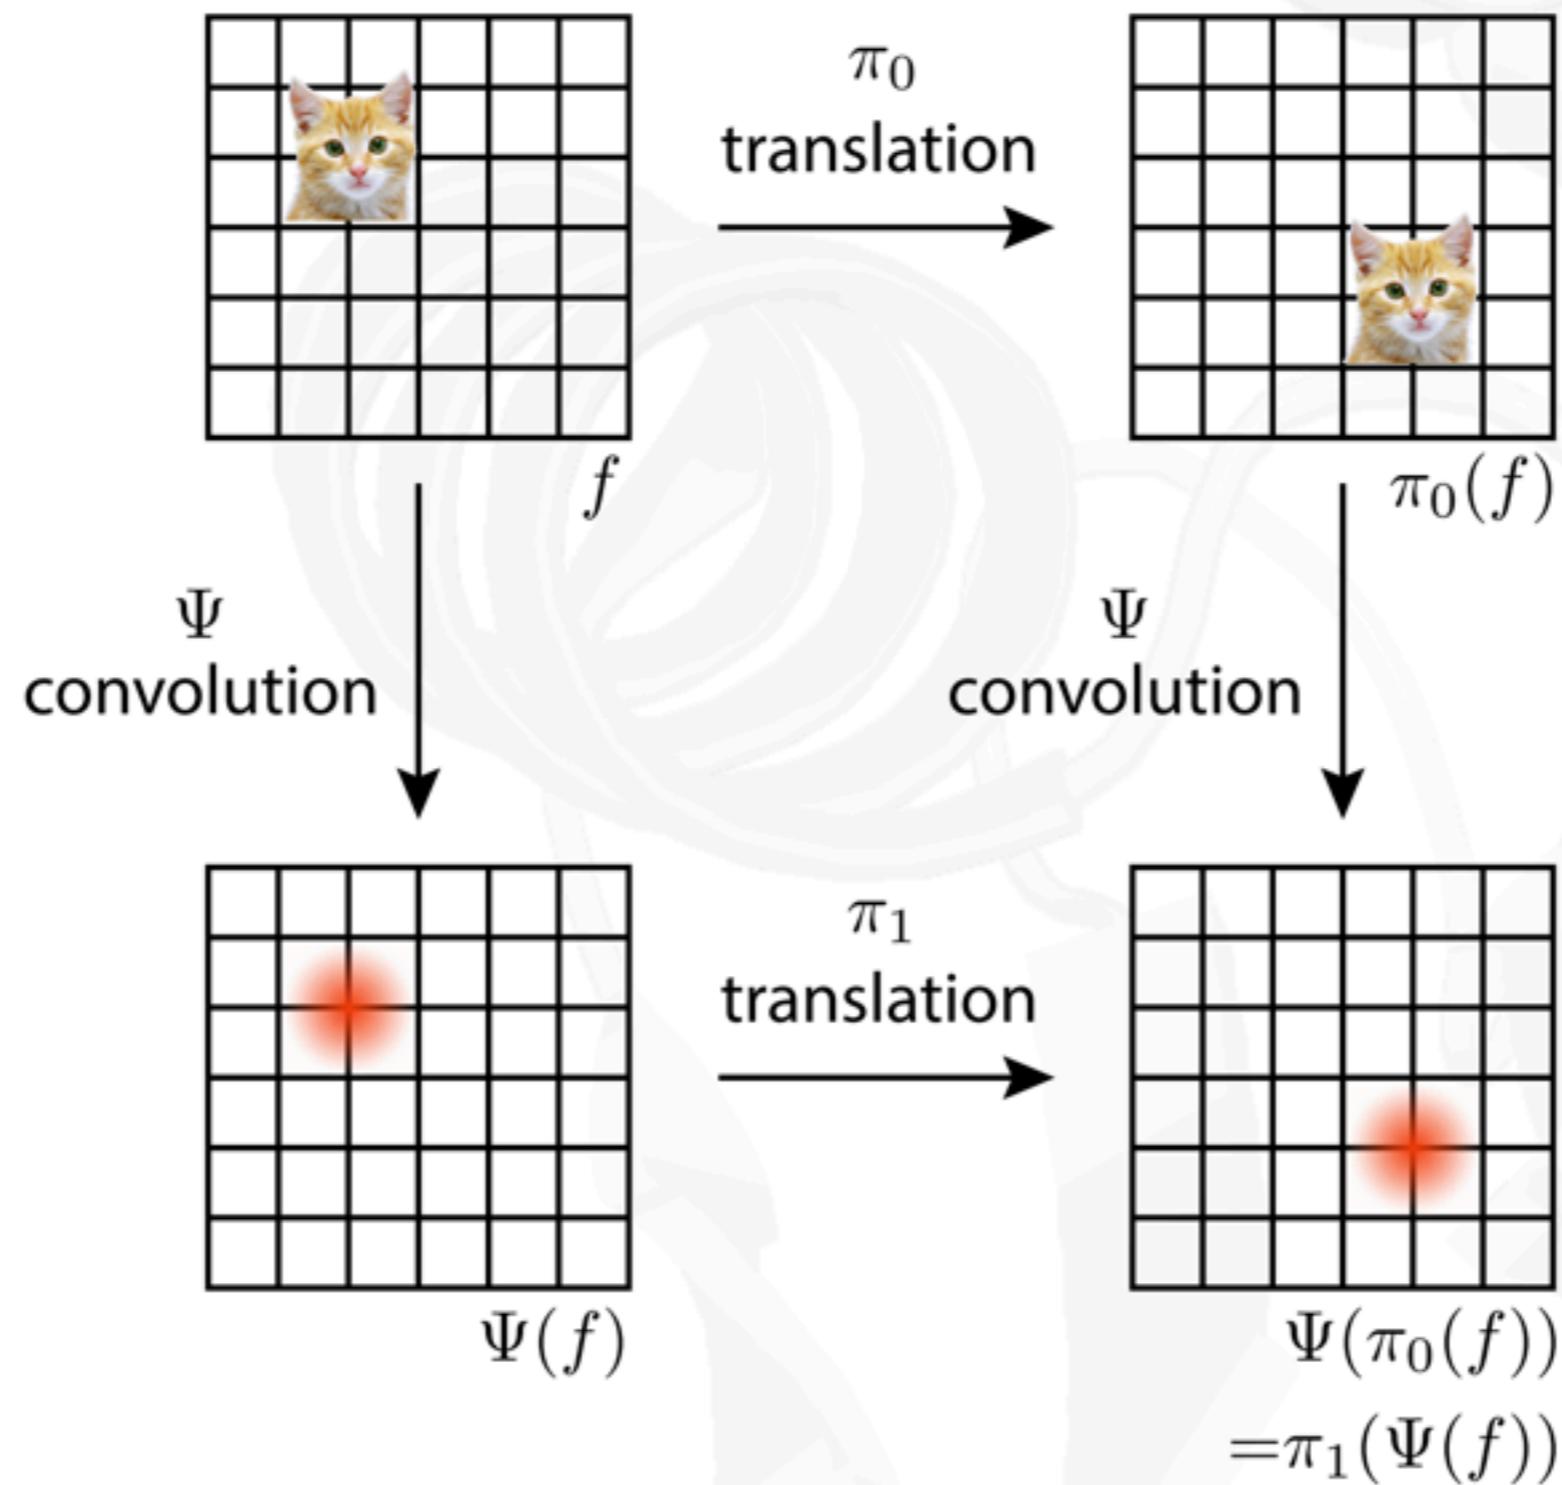


# What are CNNs - Terminology

- ▶ CNNs use **convolutions** over the input (convolution + pooling)
  - ▶ Each convolution applies filters (or kernels; often several hundreds of them) and combines their results via **pooling** (to reduce the resolution of the feature map and the sensitivity of the output to shifts and distortion)



# Translational equivalence



# Example of a 2D convolution

- ▶ Filter (kernel) of size 3x3
  - ▶ “to identify indicative local predictors” (Goldberg, 2015)

1 <small>x1</small>	1 <small>x0</small>	1 <small>x1</small>	0	0
0 <small>x0</small>	1 <small>x1</small>	1 <small>x0</small>	1	0
0 <small>x1</small>	0 <small>x0</small>	1 <small>x1</small>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved  
Feature

Source: [http://deeplearning.stanford.edu/wiki/index.php/Feature\\_extraction\\_using\\_convolution](http://deeplearning.stanford.edu/wiki/index.php/Feature_extraction_using_convolution)

# Convolution - Filter (Kernel) example

- ▶ Imagine a 1d input vector

*What is this Kernel doing?*

- ▶  $f: [10, 50, 60, 10, 20, 40, 30]$

**computing a moving average**

- ▶  $g: [1/3, 1/3, 1/3]$

$$(f * g)(i) = \sum_{j=1}^m g(j) \cdot f(i - j + m/2)$$

- ▶ Let's compute the value at position  $h(3)$

$$[10, \boxed{50, 60, 10}, 20, 40, 30]$$

$$[ 0, \boxed{0/3, 1/3, 1/3}, 10, 30, 0 ]$$

$$50 * \frac{1}{3} + 60 * \frac{1}{3} + 10 * \frac{1}{3} = 40$$

$$h(3) = 40$$

$$h(4) = 30$$

# Convolutions for Text

Collobert et al. (2011); Kim (2014)

# Types of convolution

She likes strong coffee

**narrow/“valid” convolution**

She likes  
likes strong  
strong coffee

< /s > She likes strong coffee < /s >

**wide/“same” convolution**  
(padded)

# CNN on Text

$k=2$  window length

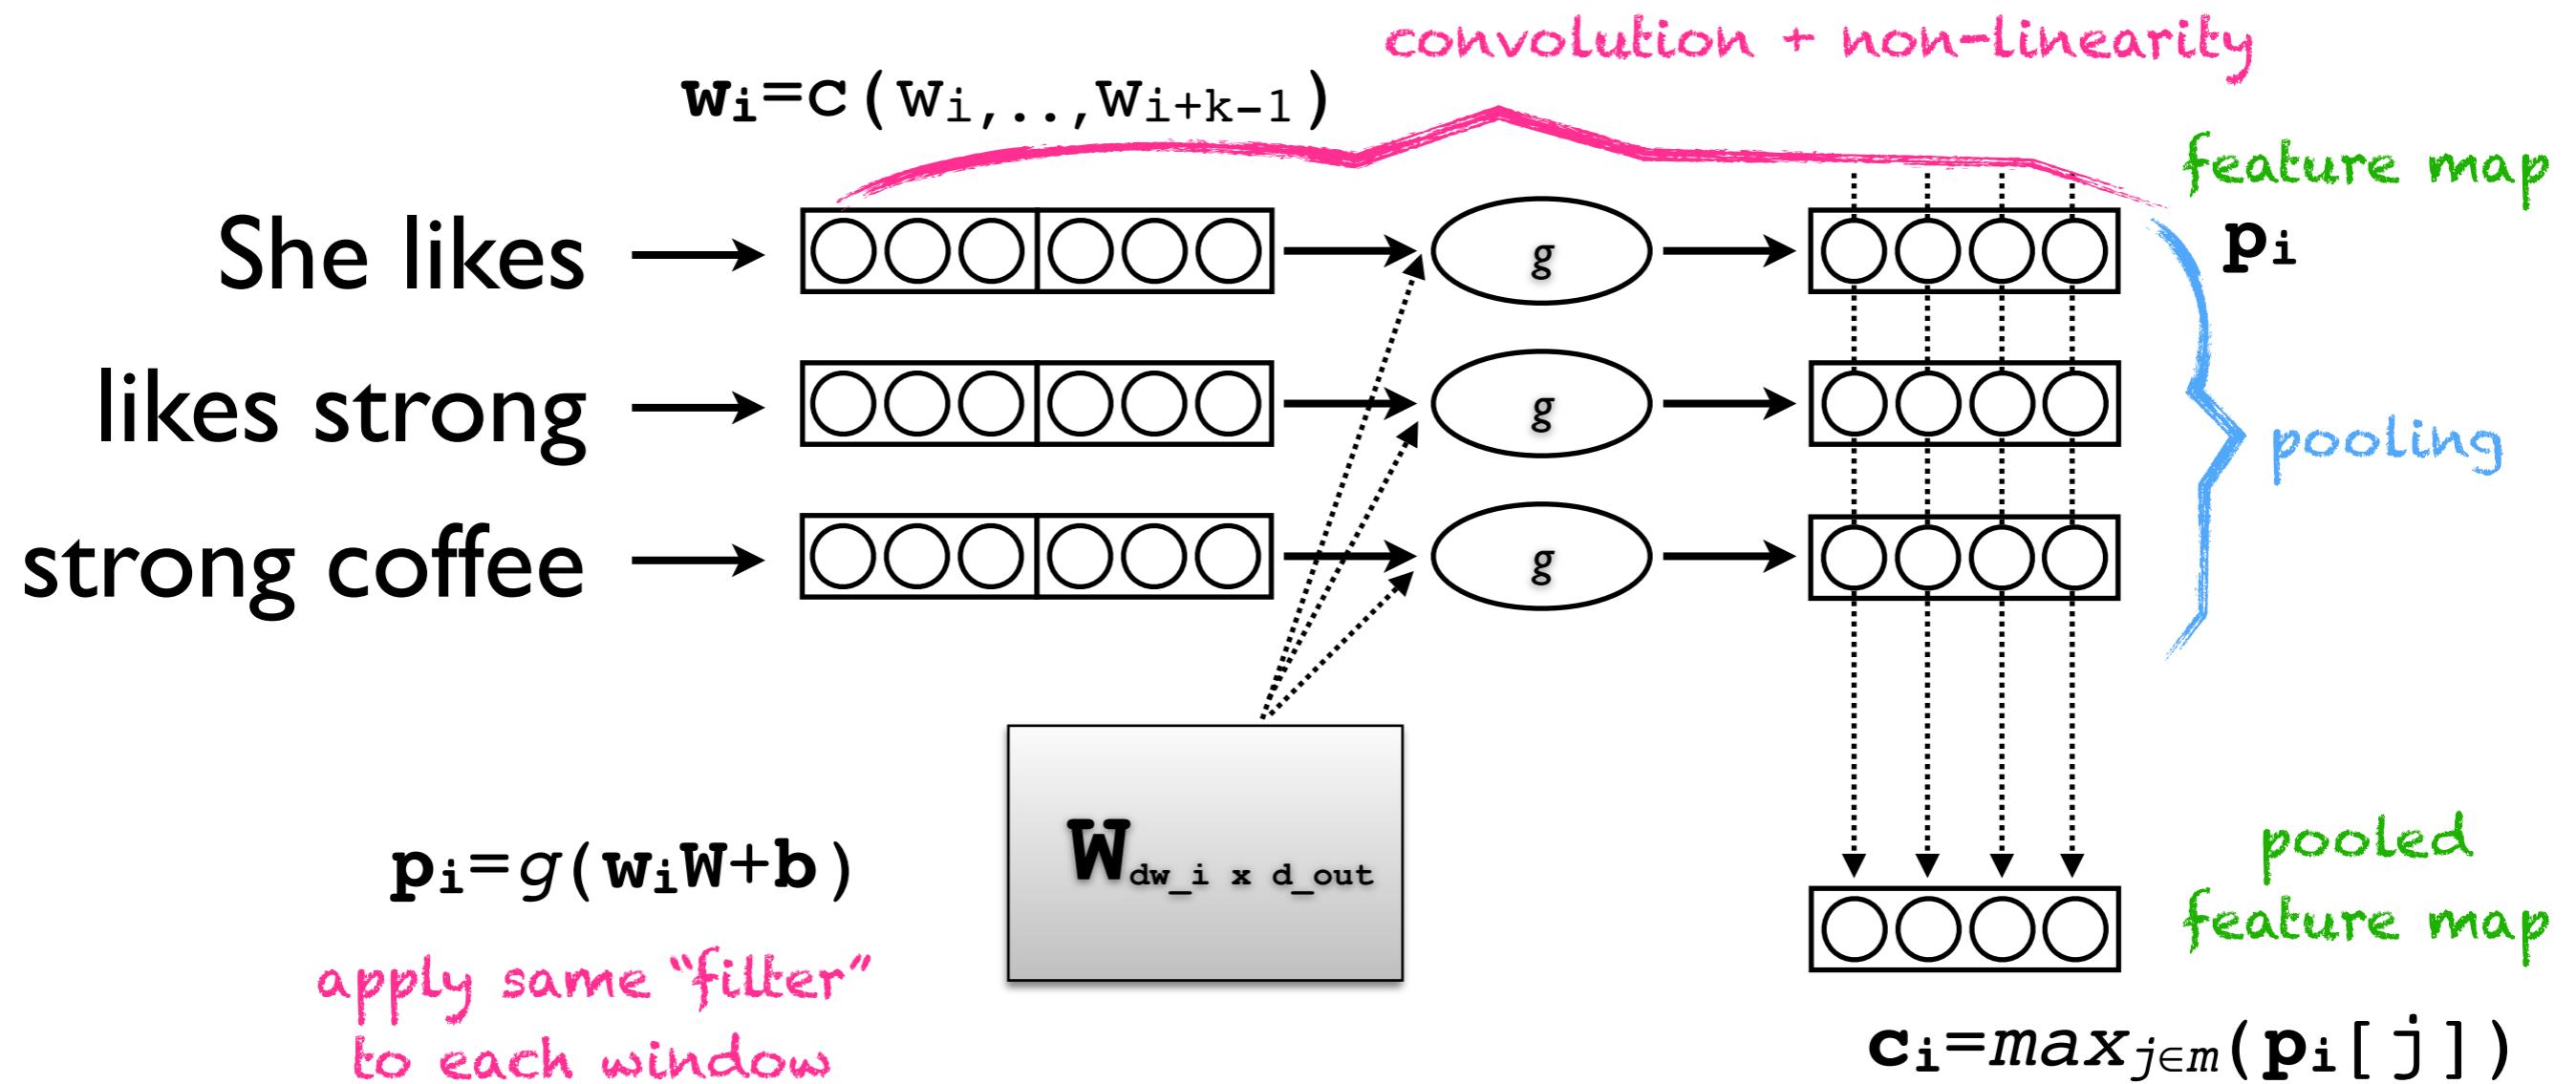
$d_{in}=3$  embedding dim

$d_{w_i}=kd_{in}=6$  filter width

$d_{out}=4$  conv output dim

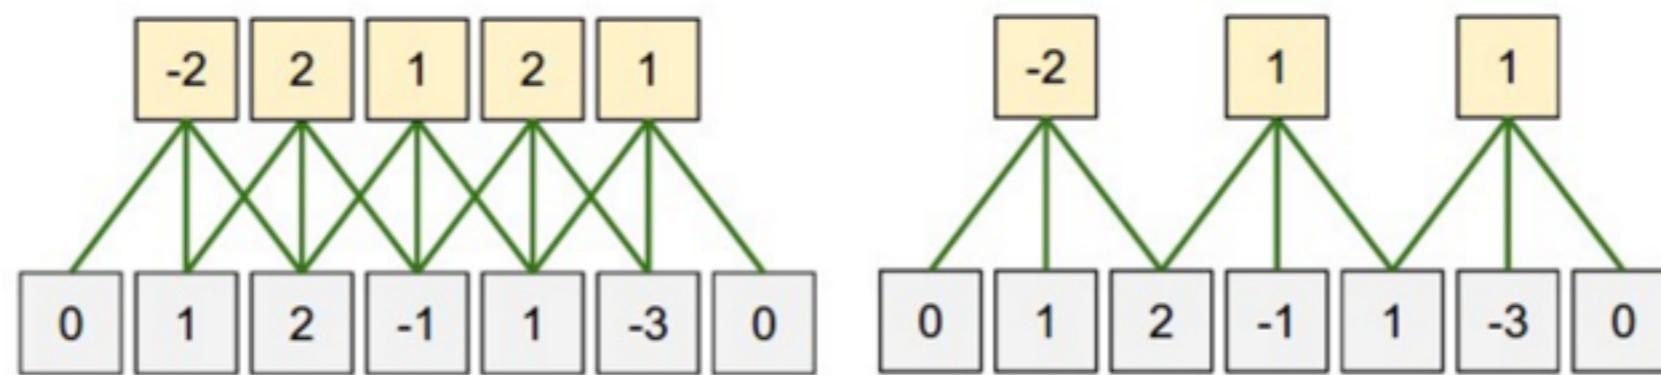
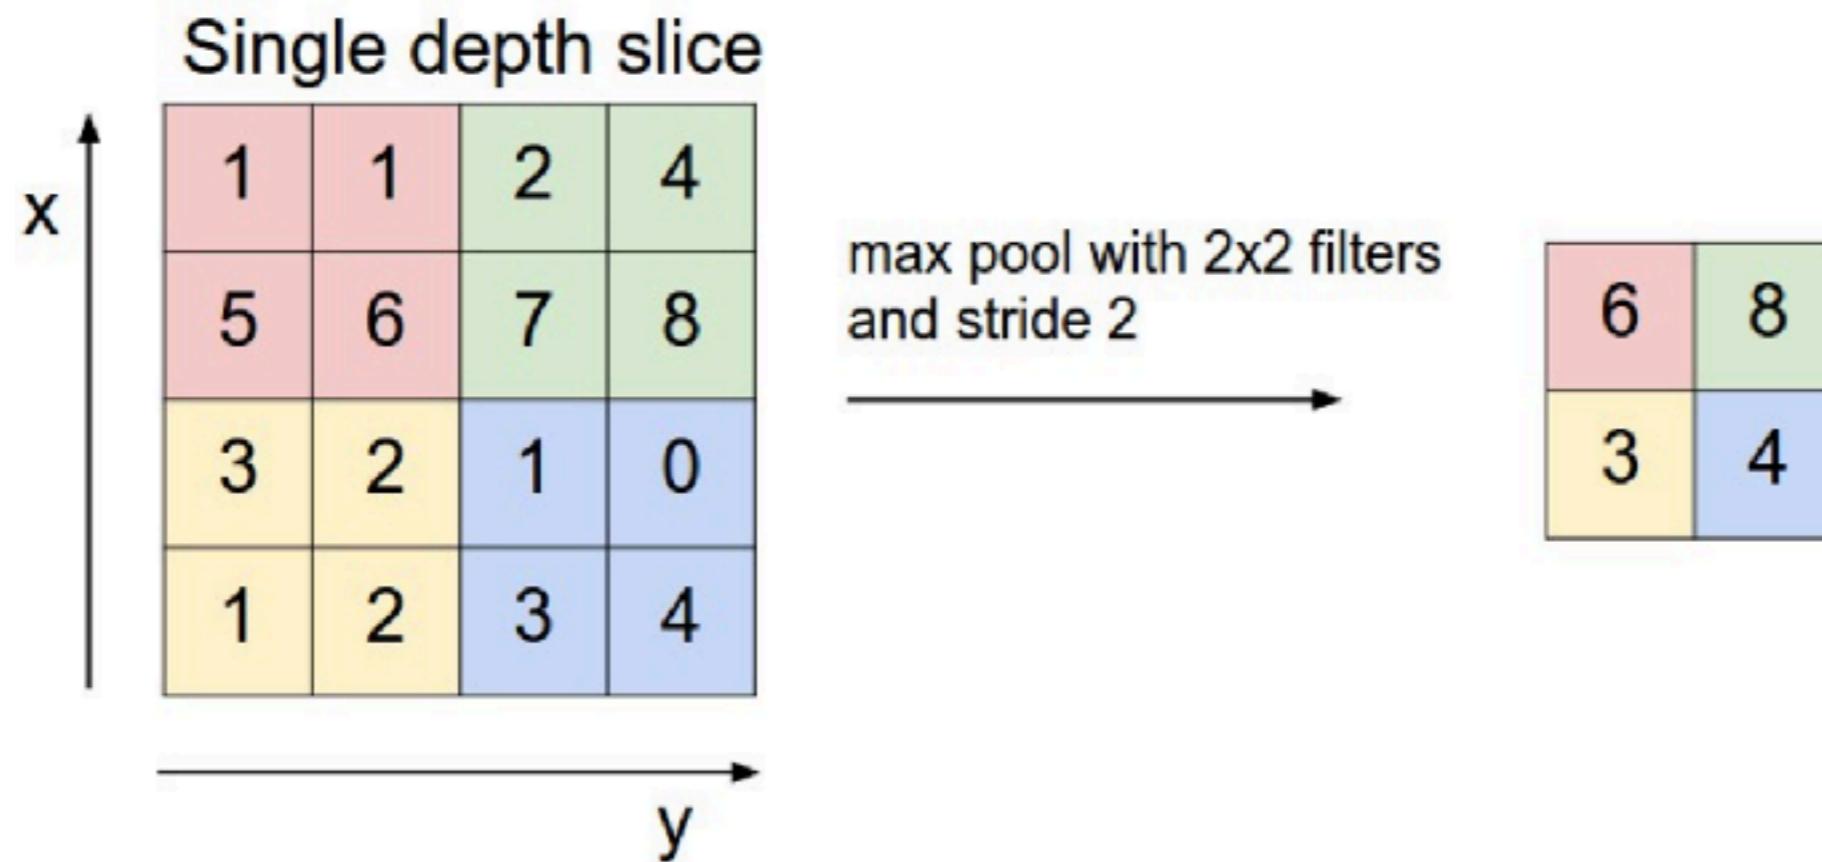
$n=4$  input length

She likes strong coffee



**“soft” n-grams**

# Stride

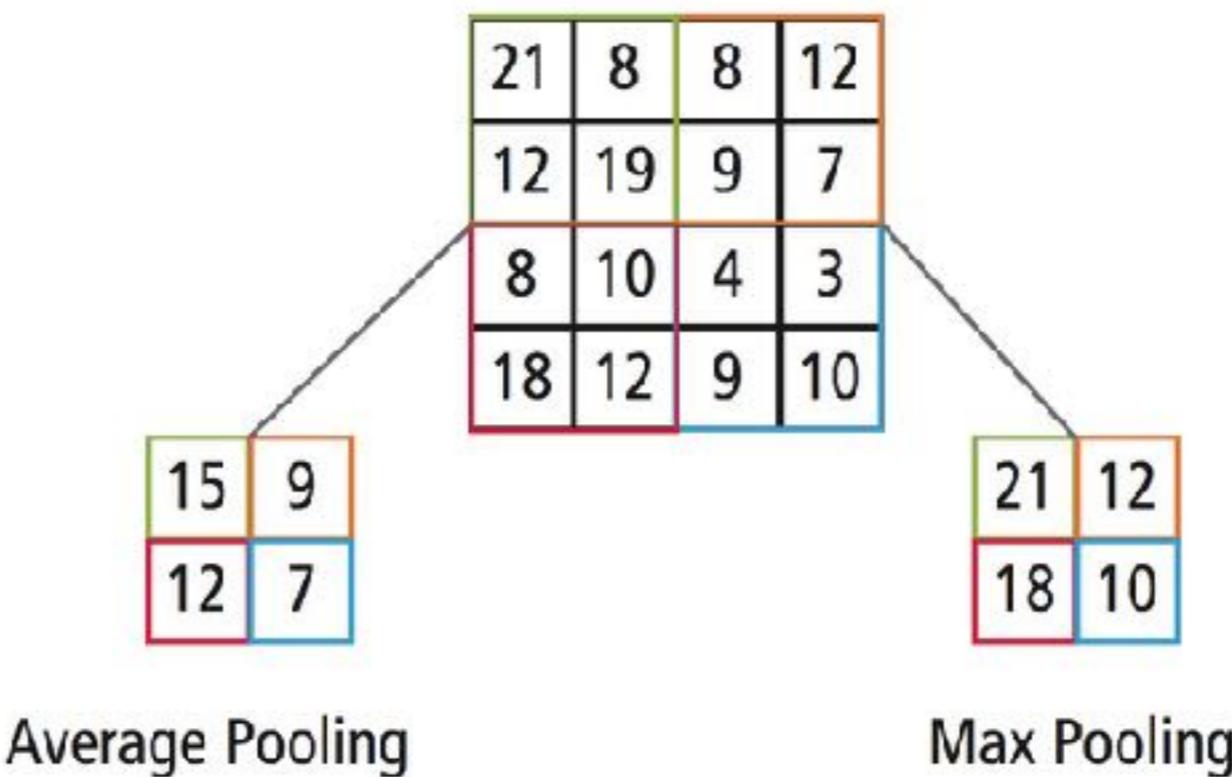


<http://d3kbpzbmcynnmx.cloudfront.net/wp-content/uploads/2015/11/Screen-Shot-2015-11-05-at-2.18.38-PM.png>

<http://d3kbpzbmcynnmx.cloudfront.net/wp-content/uploads/2015/11/Screen-Shot-2015-11-05-at-10.18.08-AM.png>

# Types of pooling (1/3)

- ▶ **Max pooling:** “Did you see this feature anywhere in the range?” (most common)  $\mathbf{c}_i = \max_{j \in m} (\mathbf{p}_i[j])$
- ▶ **Average pooling:** “How prevalent is this feature over the entire range”  $\mathbf{c}_i = \frac{1}{m} \sum_m \mathbf{p}_i$

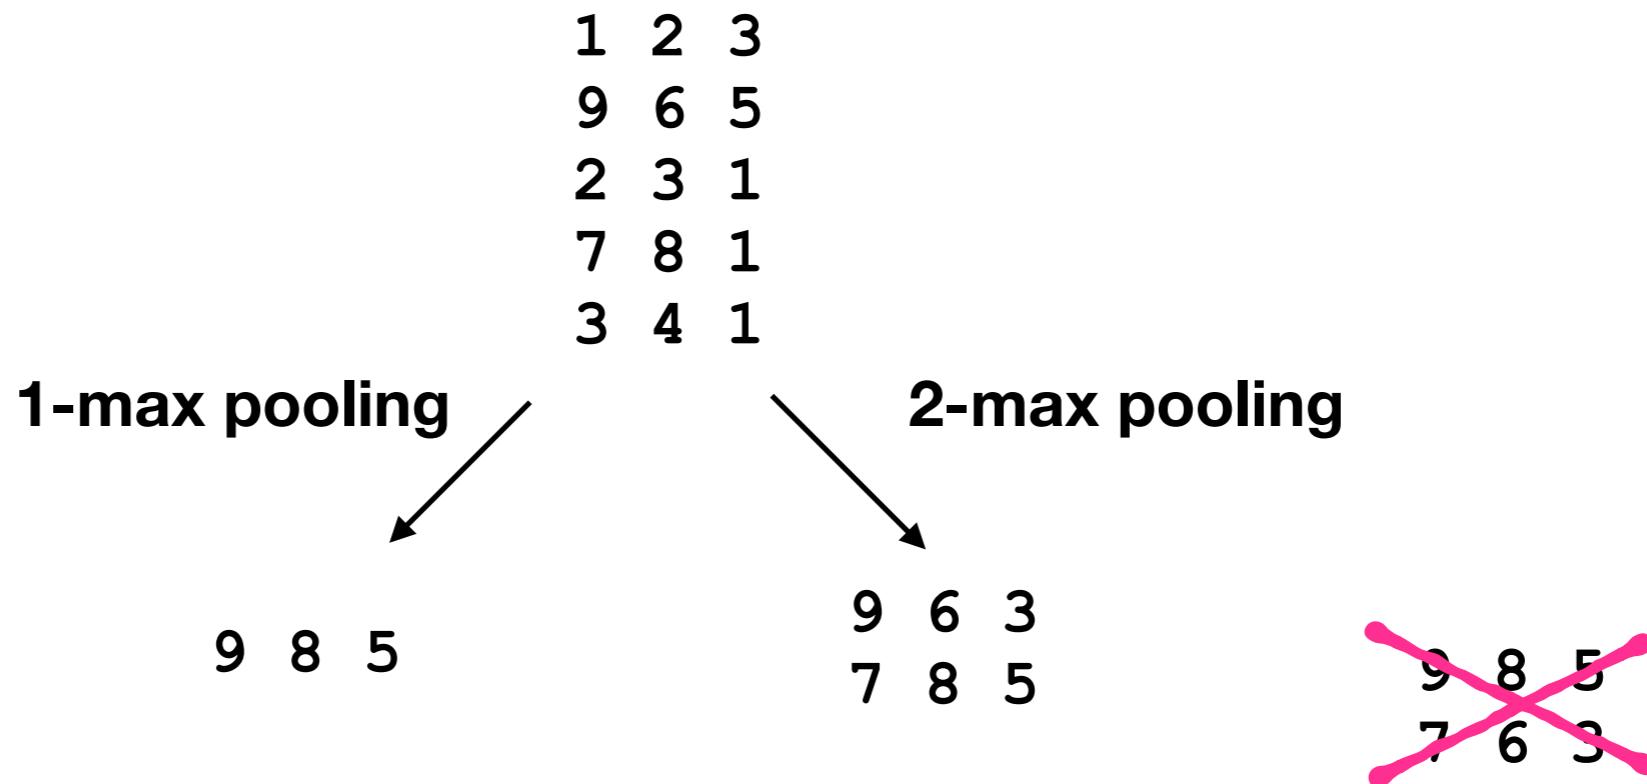


Average Pooling

Max Pooling

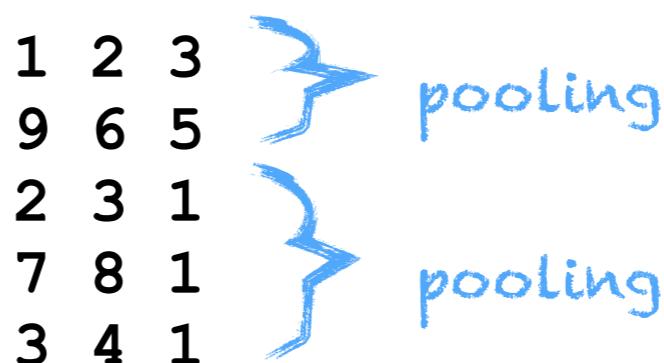
# Types of pooling (2/3)

- ▶ **k-Max pooling**: “Did you see this feature up to k times?” (Kalchbrenner et al., 2014)
  - ▶ retain top k values in each dimension instead of only the best one, while preserving the order in which they appear



# Types of pooling (3/3)

- ▶ **Dynamic pooling:** “Are some parts more informative?” (Johnson & Zhang, 2015)
  - ▶ split  $\mathbf{p}_i$ 's into separate groups based on domain knowledge and apply max-pooling to each region/group
  - ▶ e.g. initial sentences more predictive for news topic classification (Johnson & Zhang, 2015)



# CNNs for Text Classification (Kim, 2014)

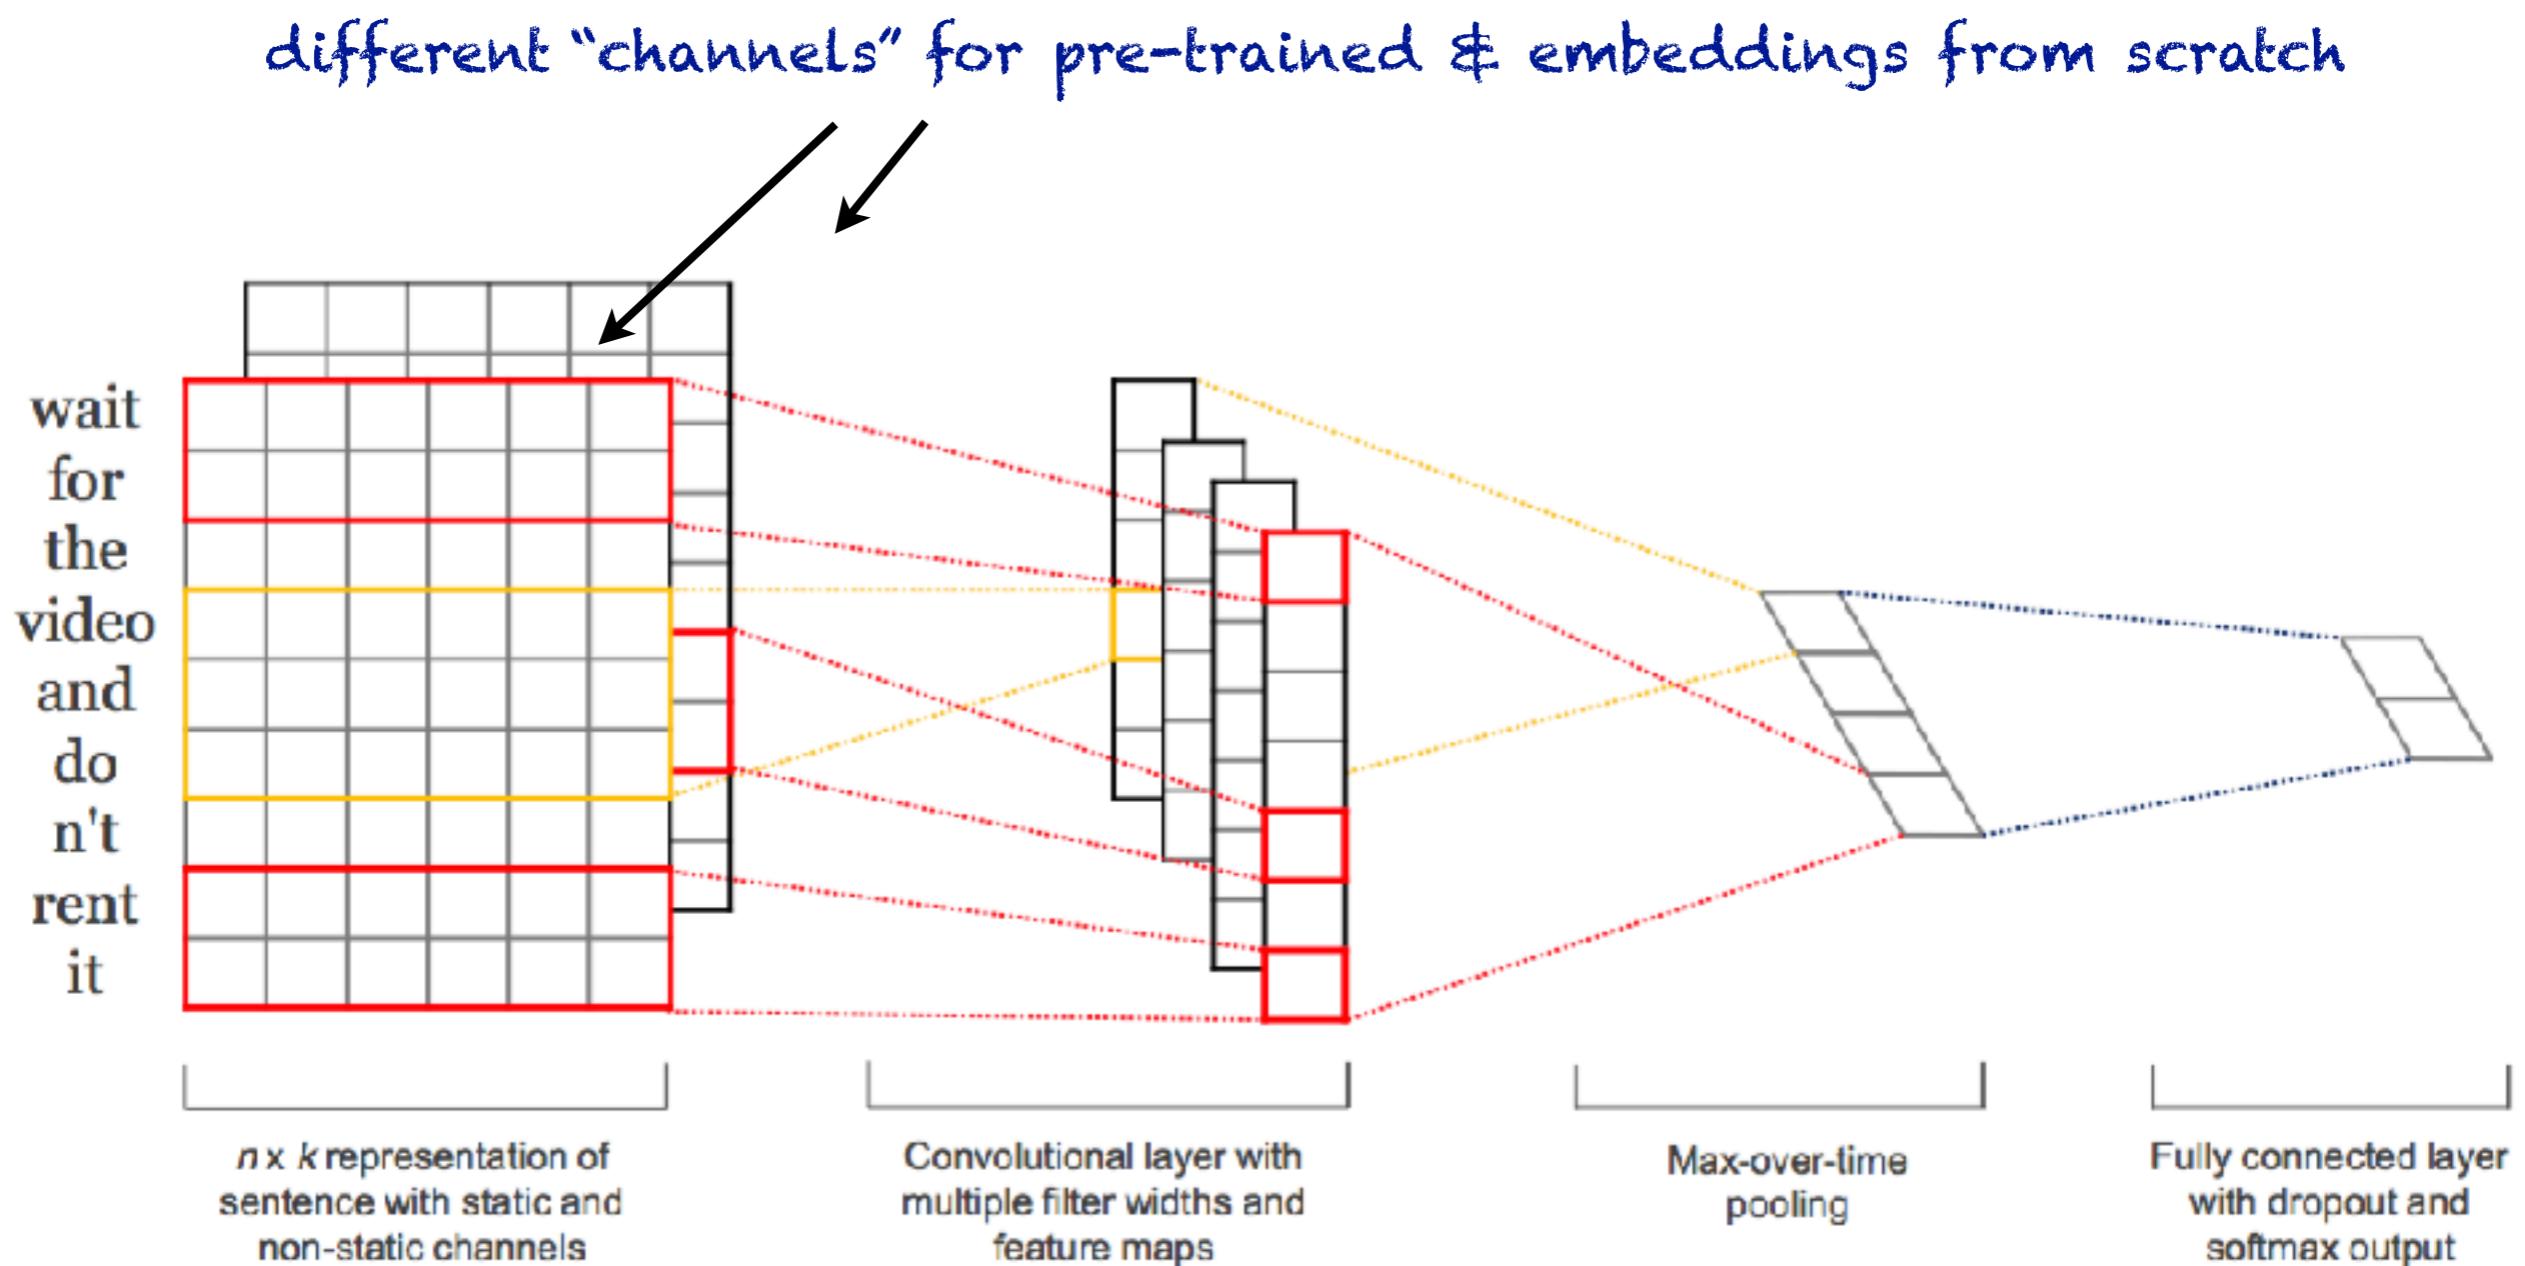


Figure 1: Model architecture with two channels for an example sentence.

# CNNs - Interim summary

- ▶ **Main idea:** apply the same parametrized function over all n-grams in the sequence.
- ▶ This creates a series of m vectors, each representing a particular n-gram in the sequence
- ▶ The representation is sensitive to the identity and order of the words in the n-gram, but the same representation will be extracted for a n-gram regardless of its position in the sequence

# **Two advances in CNNs**

# Stacked convolutions

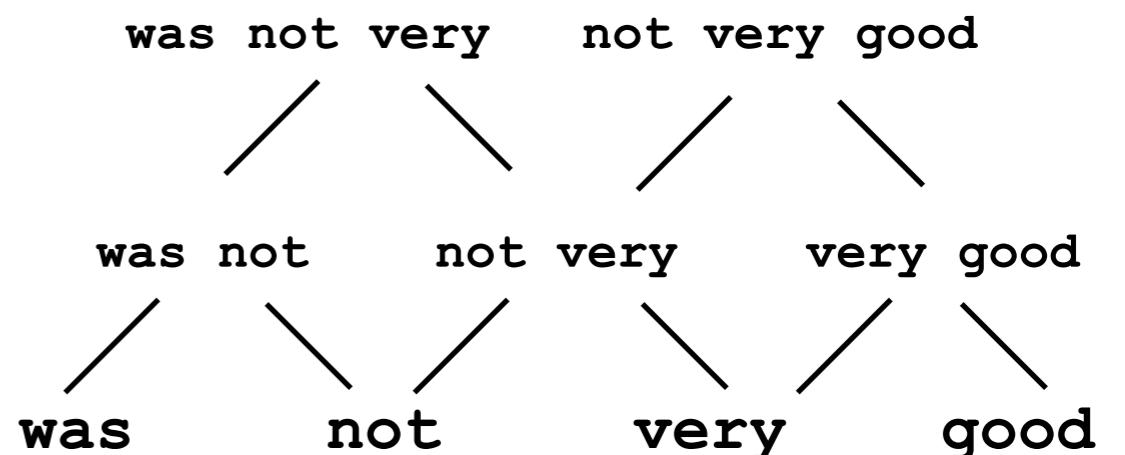
- ▶ **Hierarchical convolutions:** apply a sequence of  $r$  convolutions that feed into each other
- ▶ Resulting vectors capture increasingly larger windows (“receptive fields”)

$$\mathbf{p}_{1:m_1}^1 = \text{CONV}^k 1_{U^1, b^1}(\mathbf{w}_1 : \mathbf{n})$$

$$\mathbf{p}_{1:m_2}^2 = \text{CONV}^k 2_{U^2, b^2}(\mathbf{p}_1 : \mathbf{m}_1)$$

...

$$\mathbf{p}_{1:m_r}^r = \text{CONV}^k r_{U^r, b^r}(\mathbf{p}_{r-1} : \mathbf{m}_{r-1})$$



# Dilated convolutions

(Strubell et al., 2017; Kalchbrenner et al., 2016; Yu and Koltun, 2016)

- ▶ Each layer in the hierarchy has a stride size of  $k-1$ 
  - ▶ speed gains over RNNs

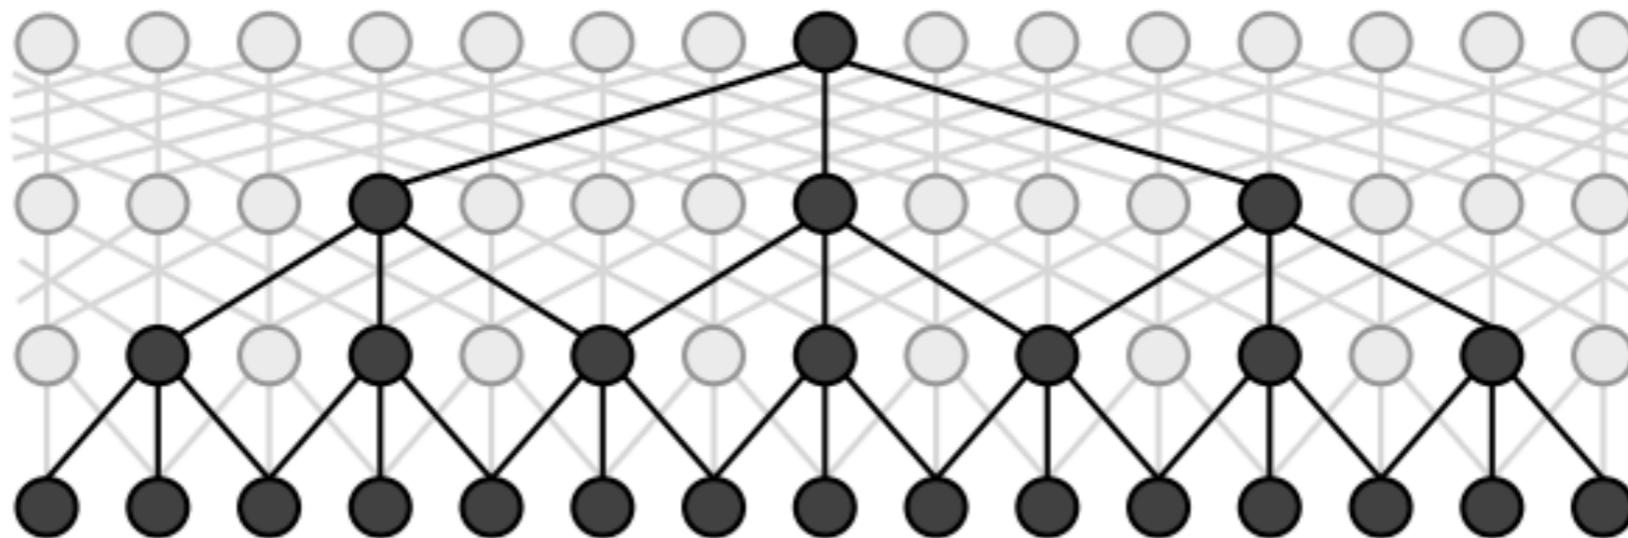


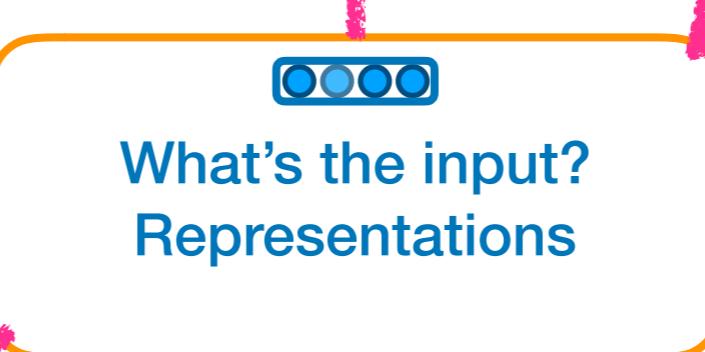
Figure 1: A dilated CNN block with maximum dilation width 4 and filter width 3. Neurons contributing to a single highlighted neuron in the last layer are also highlighted.

# Overview

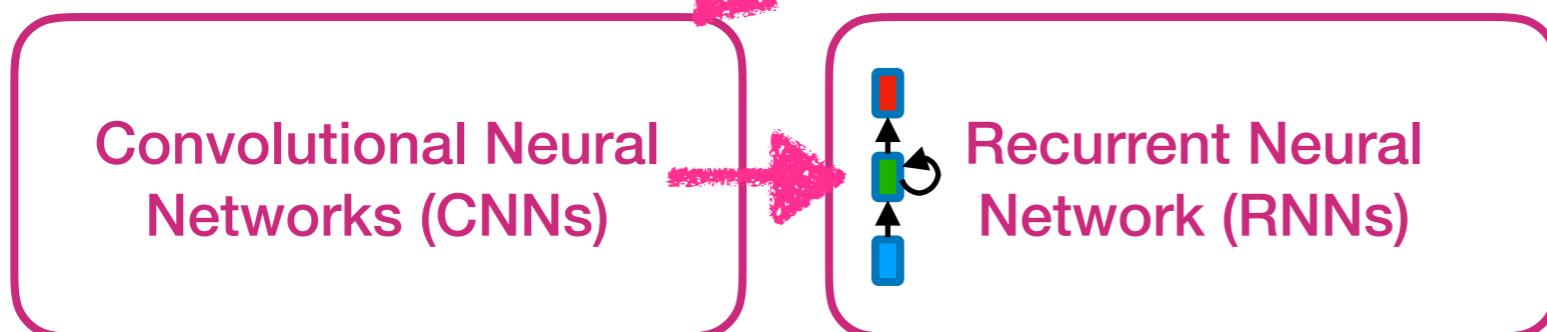
*foundations*



*representations*



*beyond FFNNs*



# BREAK

# RNNs

# Recurrent Neural Networks (RNNs)

Elman, 1990

- ▶ RNNs (and their variants) are **one of the most powerful and widespread architectures to date**
- ▶ From J.Schmidhuber's homepage:



speech recognition. They learn through gradient descent and / or evolution or both. Compare the RNN Book Preface. LSTM is getting popular: Google, Apple, Microsoft, Facebook, IBM, Baidu, and many other companies use LSTM RNNs to improve large vocabulary speech recognition, machine translation, language identification / time series prediction / text-to-speech synthesis, etc.

# Recurrent Neural Networks (RNNs)

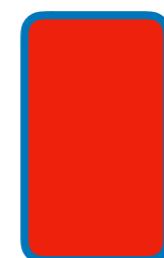
- ▶ Can handle **arbitrary length inputs** (just like CNNs or a FFNN with a CBOW input representation)
  - ▶ Unlike CBOW, they **model the order** in the sequence
  - ▶ Unlike vanilla CNNs, they **can deal with long-distance dependencies** (especially the gated RNN variants)
- ▶ Do **not need to make the Markov assumption**
  - ▶ Opens up for a family of models:  
***Conditioned generation models***

*Key Idea:*

# Recurrent Neural Networks (RNNs)

“vanilla” Neural Network

$$\hat{y} = Wh + b$$



$$h = g(Vx + b)$$

RNN

RNNs have an internal  
“memory” (state)  
which is updated as the  
sequence is read

$$\hat{y}_t = Wh_t + b$$



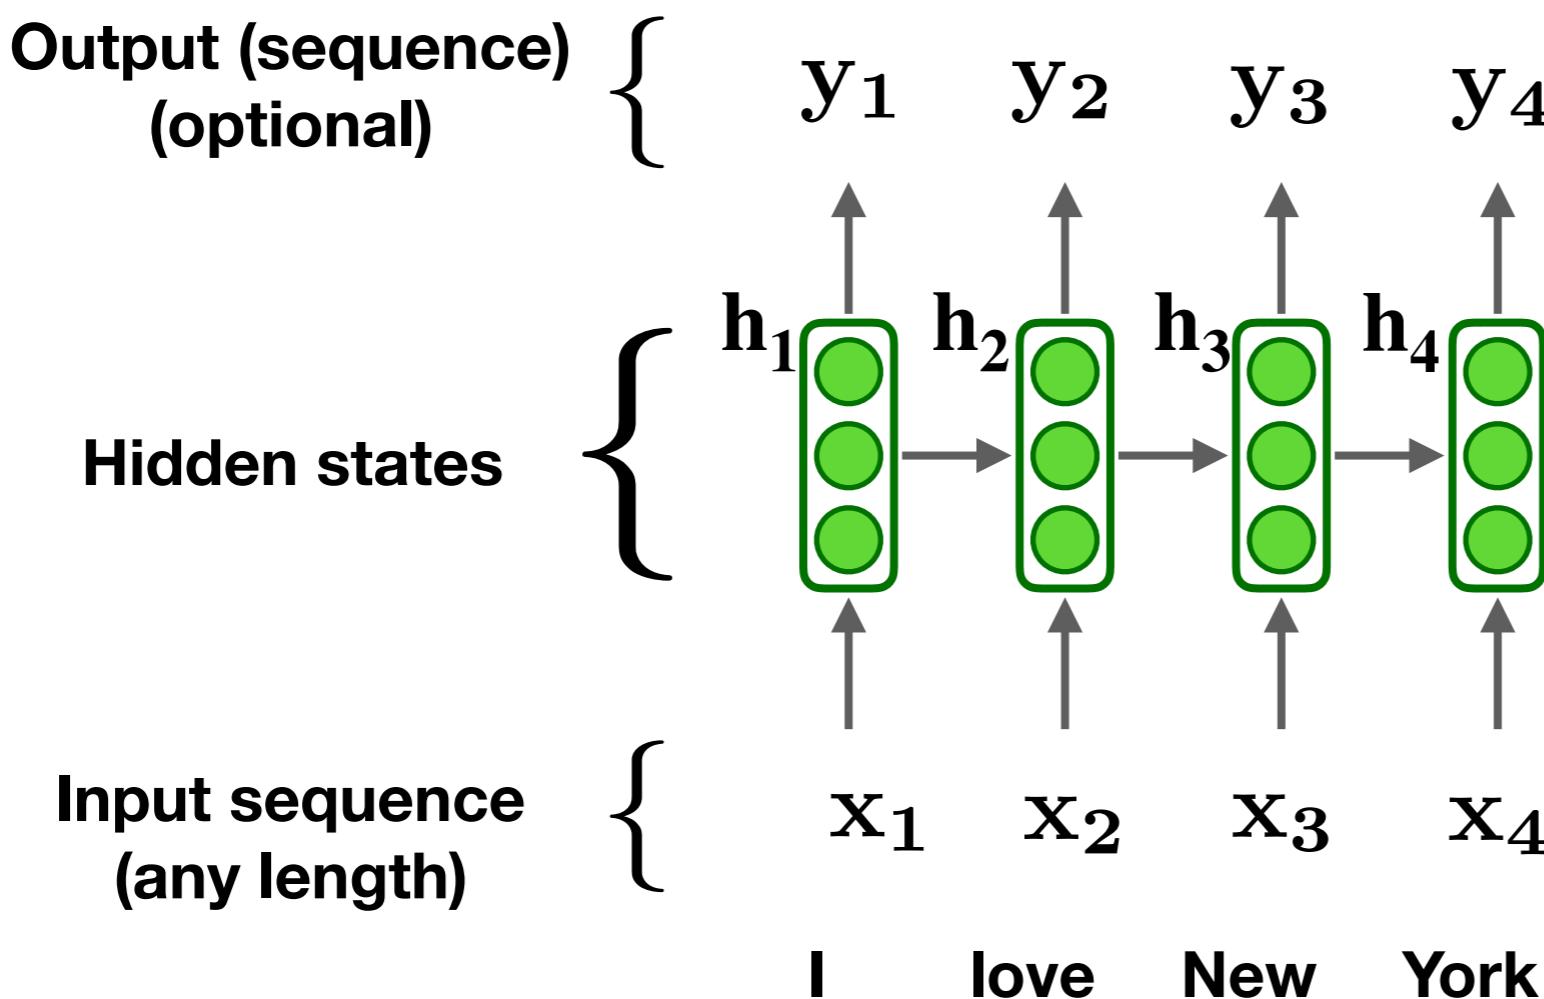
$$h = g(Vx_t + Uh_{t-1} + c)$$

*recurrence /  
memory / state*

# Recurrent Neural Networks

A family of recurrent NN architectures

$$\mathbf{h} = g(\mathbf{V}\mathbf{x}_t + \mathbf{U}\mathbf{h}_{t-1} + \mathbf{c})$$



Core idea:  
Parameter  
Sharing over  
Time  
(=apply  $\mathbf{U}$   
repeatedly)

**Before we dig into details**

- The RNN abstraction**

# Count the number of 1s



Example from Cho (2015)

# Count the number of 1s

```
def add1(el,s):  
    if el==1: return s+1  
    else: return s
```

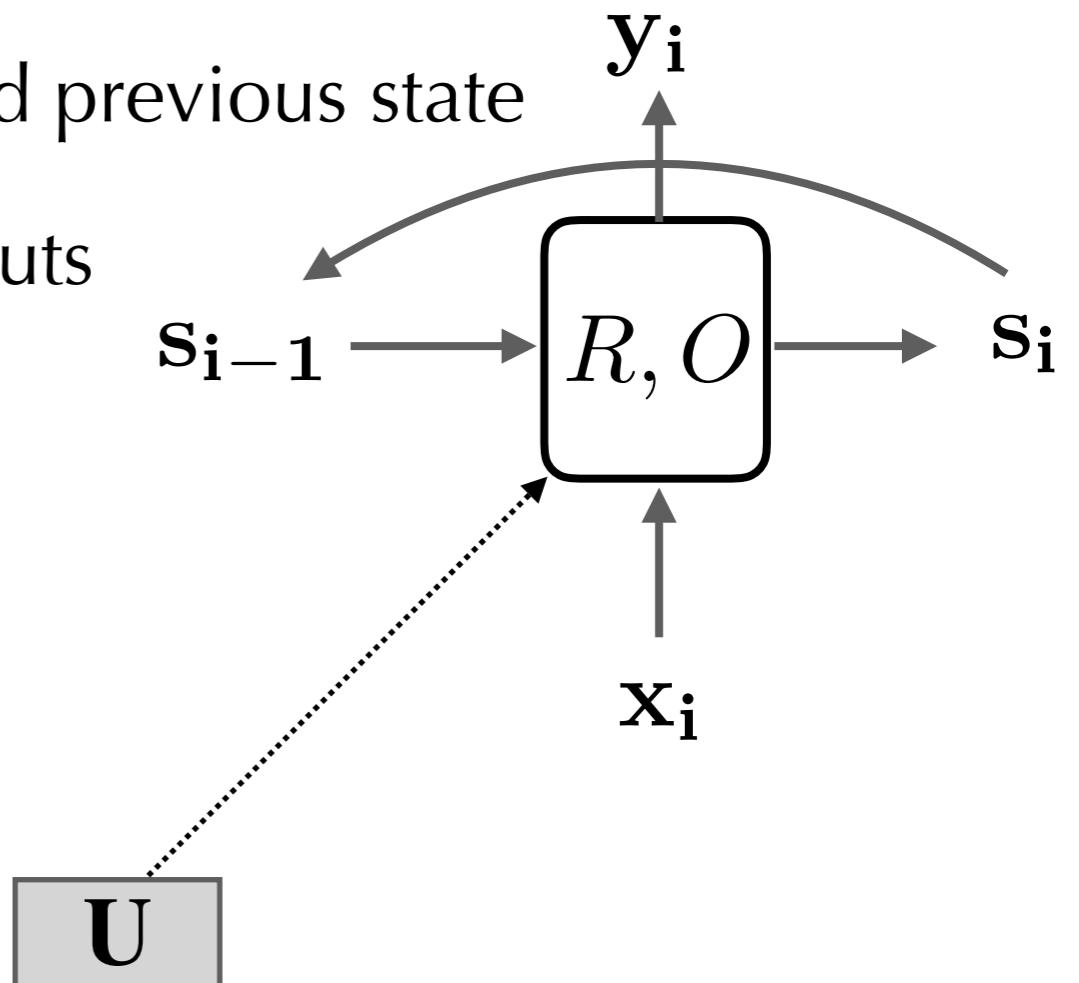
**Two important components:**

- memory  $s$
- function  $add1$  is applied to each symbol in the input *one at a time* to update the memory

```
v=[ 0 , 1 , 0 , 0 , 1 , 1 ]  
s=0  
for el in v:  
    s=add1(el,s)  
print("count(1):", s)
```

# The RNN abstraction

- ▶ Input sequence of vectors:  $\mathbf{x}_{1:n}$
- ▶ Start state:  $s_0$
- ▶  $RNN(s_0, \mathbf{x}_{1:n})$  consists of two functions:
  - ▶ function  $R$  consumes input and previous state
  - ▶ function  $O$  maps states to outputs



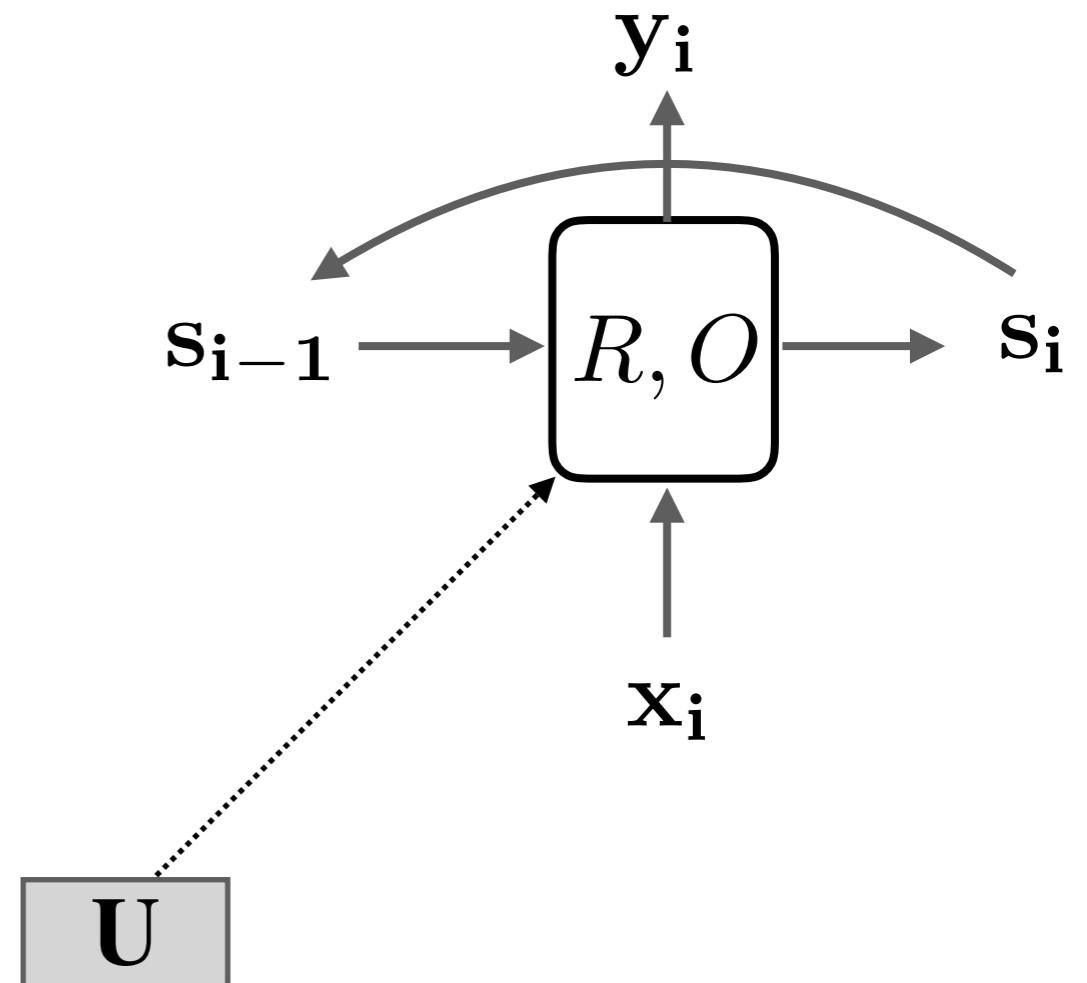
(Graphical illustration - recursive - from Yoav Goldberg's primer, 2015)

# The RNN abstraction - More formally

$$RNN(\mathbf{s}_0, \mathbf{x}_{1:n}) = \mathbf{s}_{1:n}, \mathbf{y}_{1:n}$$

$$\mathbf{s}_i = R(\mathbf{s}_{i-1}, \mathbf{x}_i)$$

$$\mathbf{y}_i = O(\mathbf{s}_i)$$



(Graphical illustration - recursive - from Yoav Goldberg's primer, 2015)

# RNN: Unrolled over time

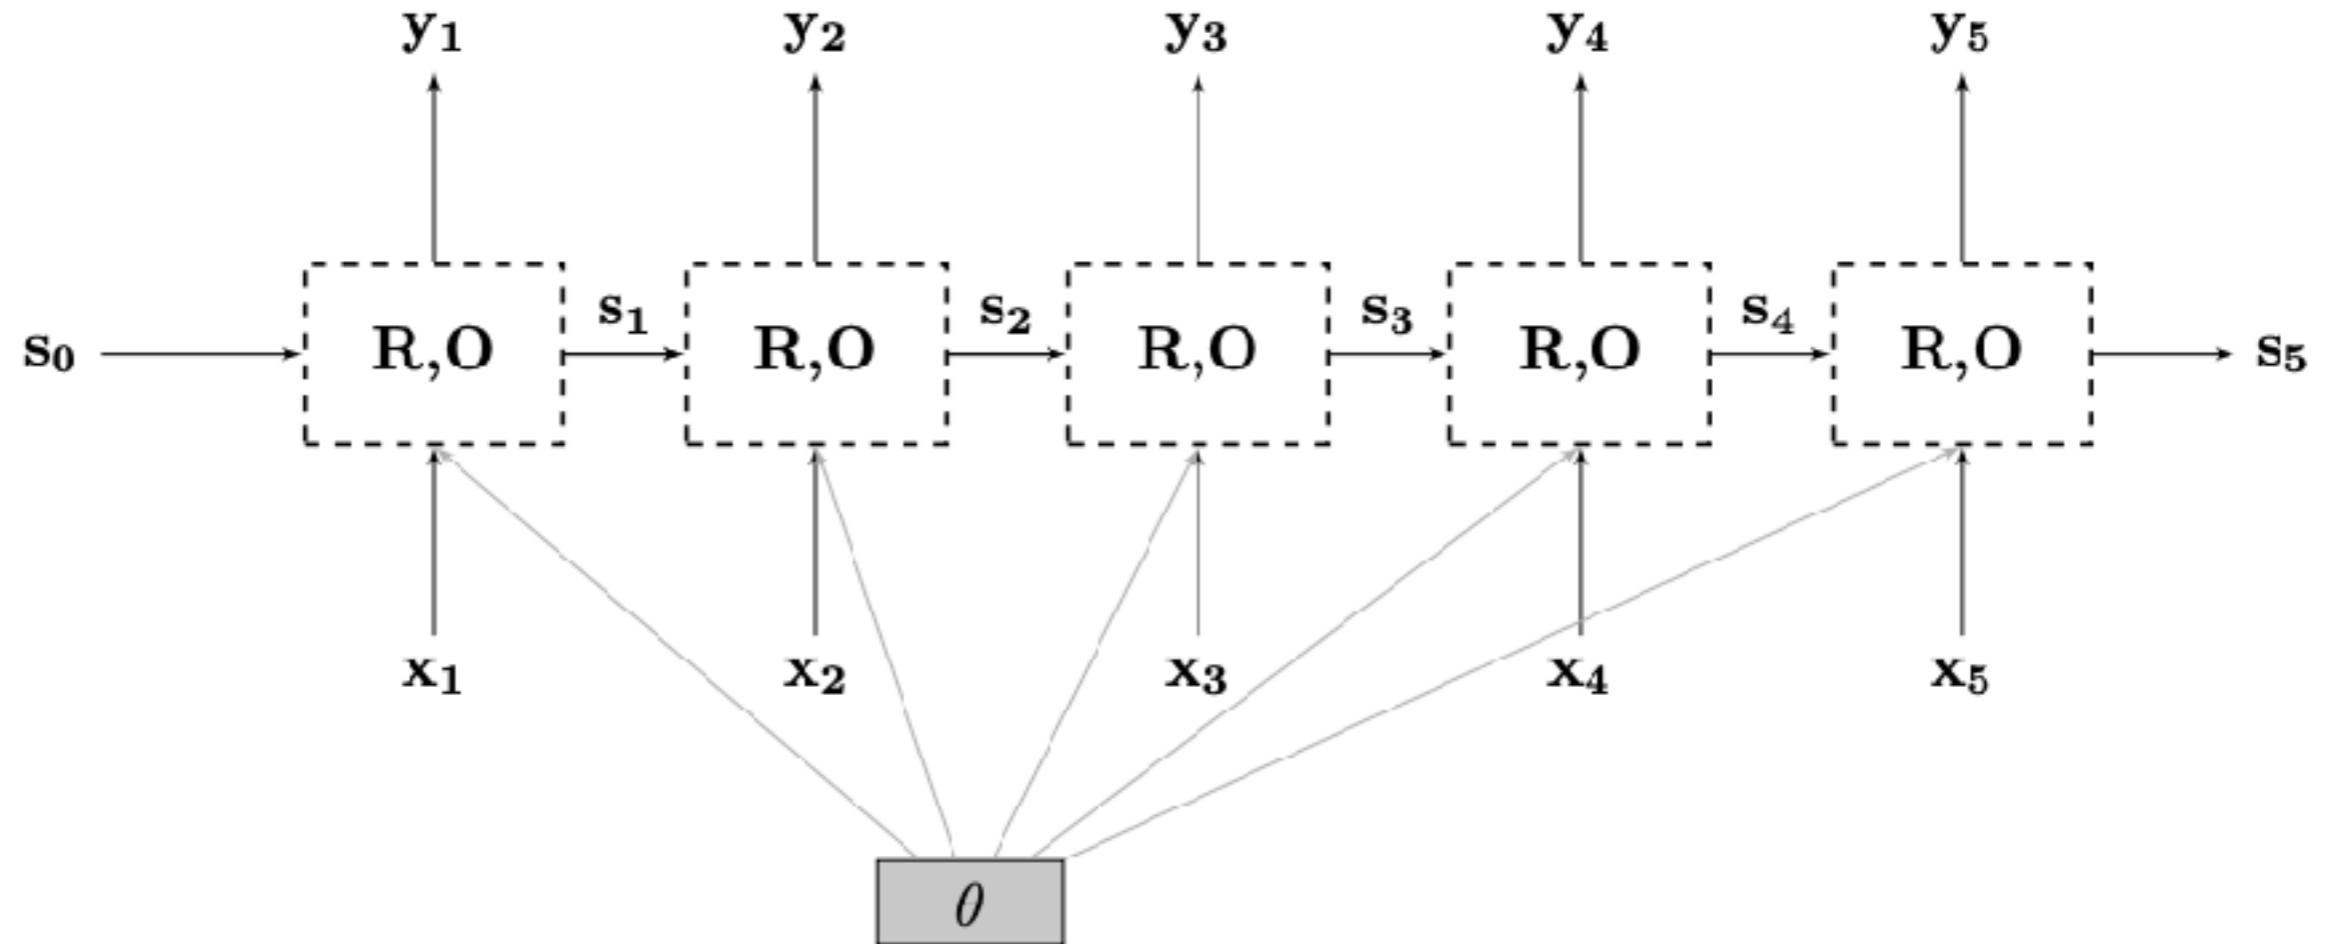


Figure 6: Graphical representation of an RNN (unrolled).

(Graphical illustration - recursive - from Yoav Goldberg's primer, 2015)

# Expansion at time step 4

$$\mathbf{s}_4 = R(\mathbf{s}_3, \mathbf{x}_4)$$

$$= R(\overbrace{R(\mathbf{s}_2, \mathbf{x}_3)}^{\mathbf{s}_3}, \mathbf{x}_4)$$

$$= R(R(\overbrace{R(\mathbf{s}_1, \mathbf{x}_2)}^{\mathbf{s}_2}, \mathbf{x}_3), \mathbf{x}_4)$$

$$= R(R(R(\overbrace{R(\mathbf{s}_0, \mathbf{x}_1)}^{\mathbf{s}_1}, \mathbf{x}_2), \mathbf{x}_3), \mathbf{x}_4)$$

# Training a RNN, parameter tying

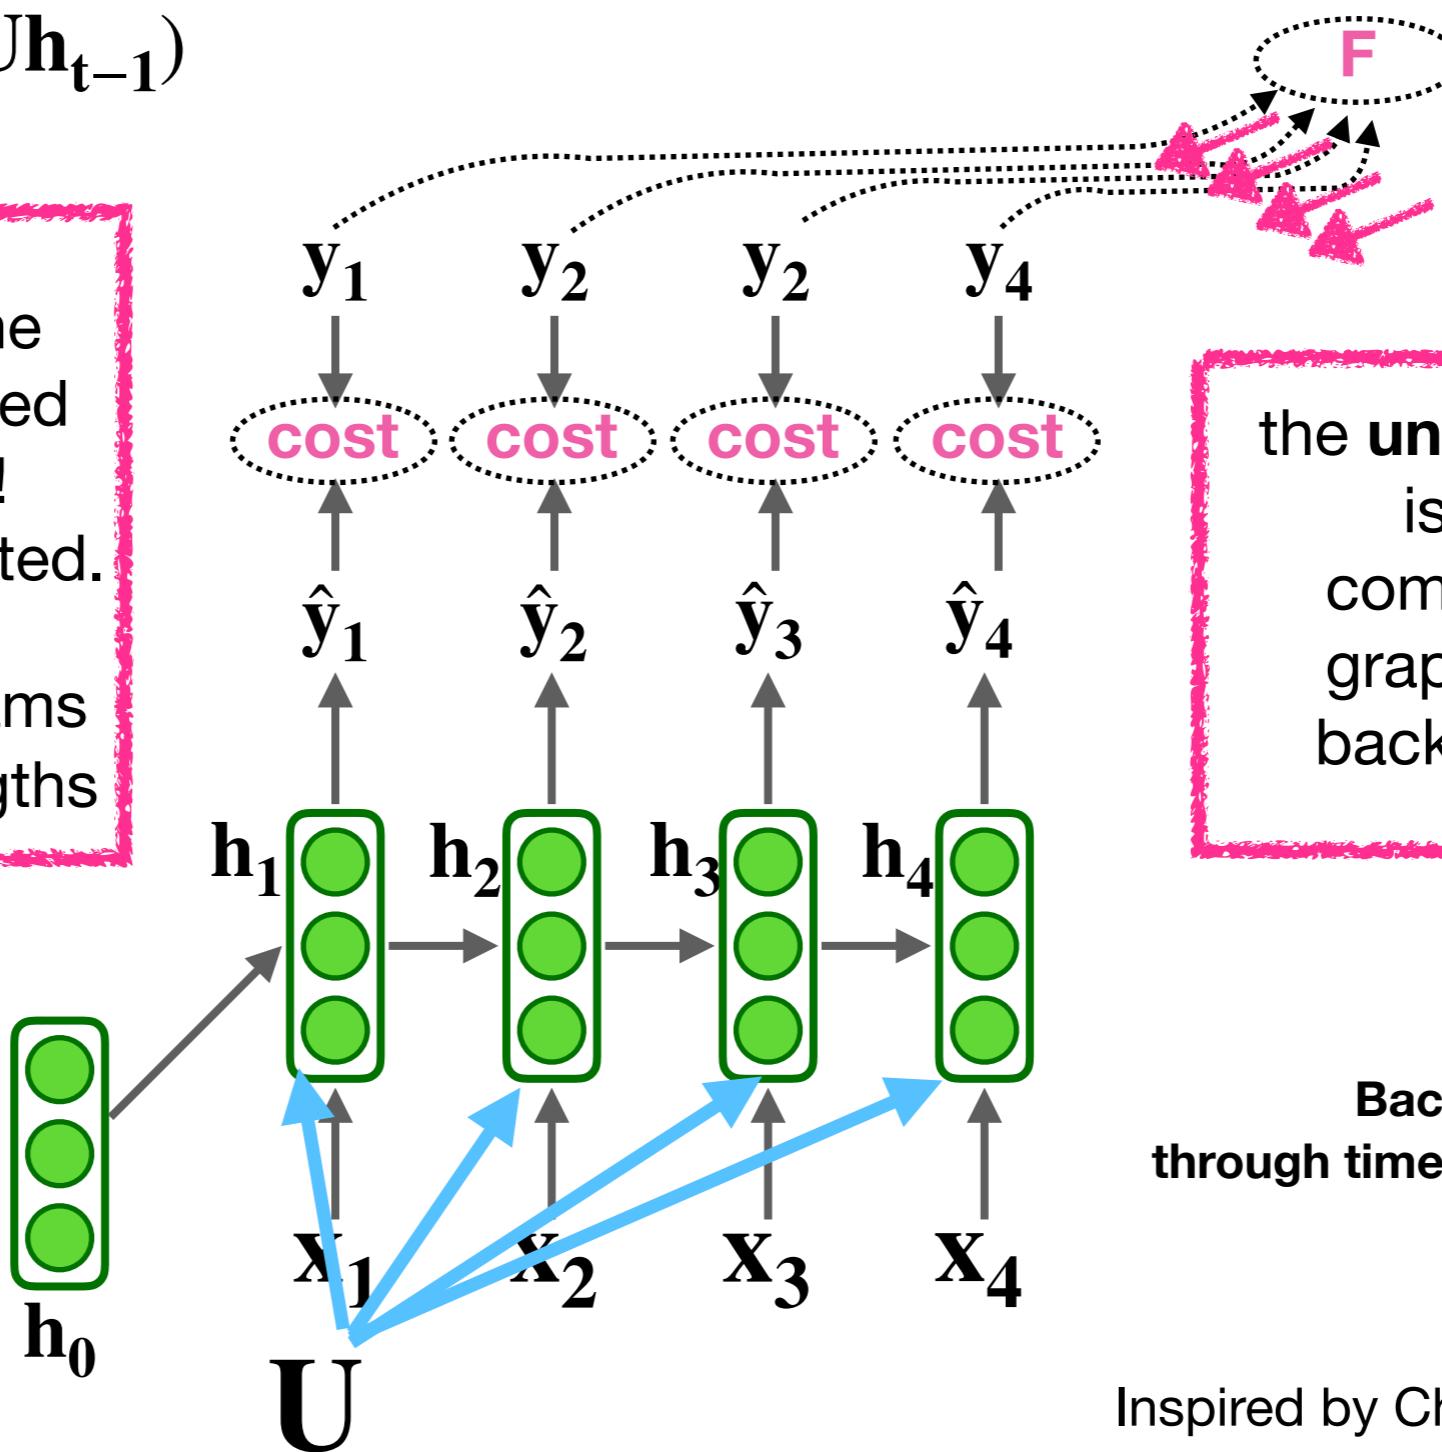
$$\hat{y}_t = Wh_t + b$$

$$h = g(Vx_t + Uh_{t-1})$$

Parameter **tying**: the parameters are shared across time steps!

Derivatives accumulated.

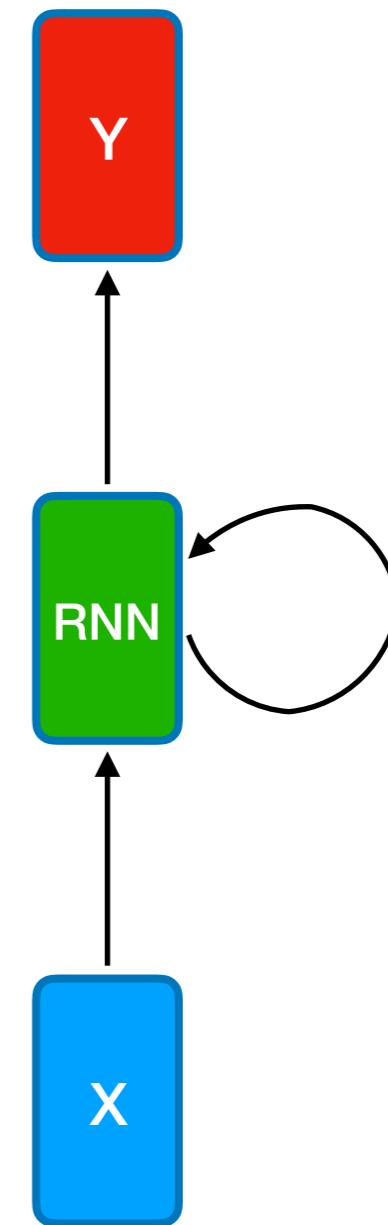
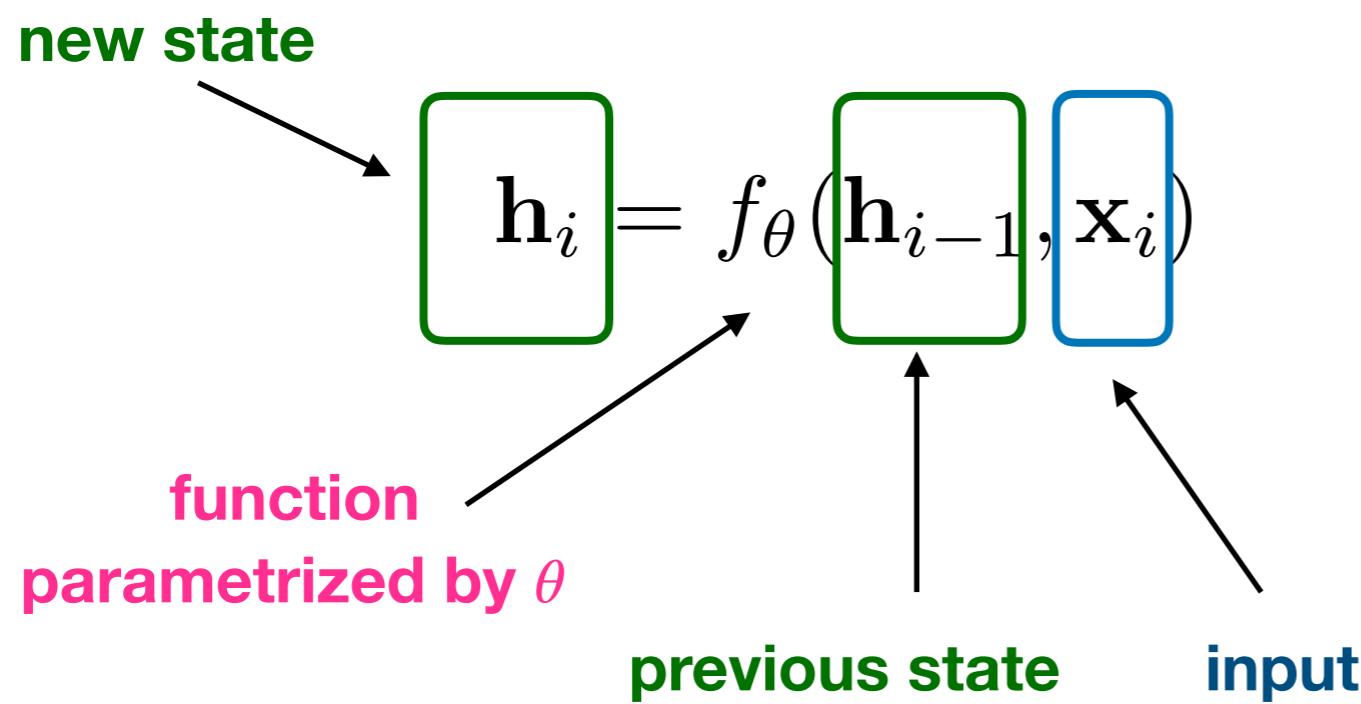
**Pros:** - reduce #params  
- model arbitrary lengths



Backpropagation through time (BPTT, Werbos, 1990).

# A closer look: inside an RNN

- We process a sequence  $\mathbf{x}$  by applying a recurrence formula at every time step  $i$ :



# Vanilla RNN

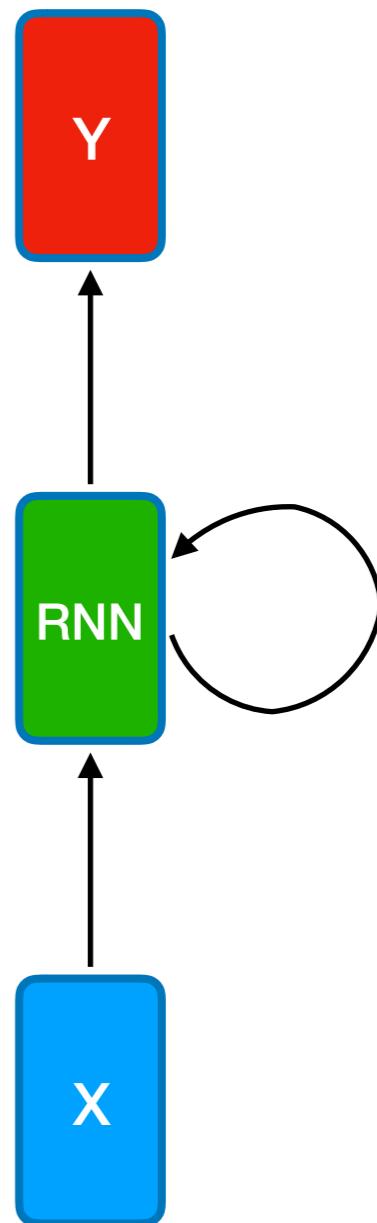
- ▶ Simple vanilla RNN (Elman, 1990)

RNN: example instantiation  
of function  
parametrized by  $\theta$

$$\hat{y}_t = Wh_t + b$$

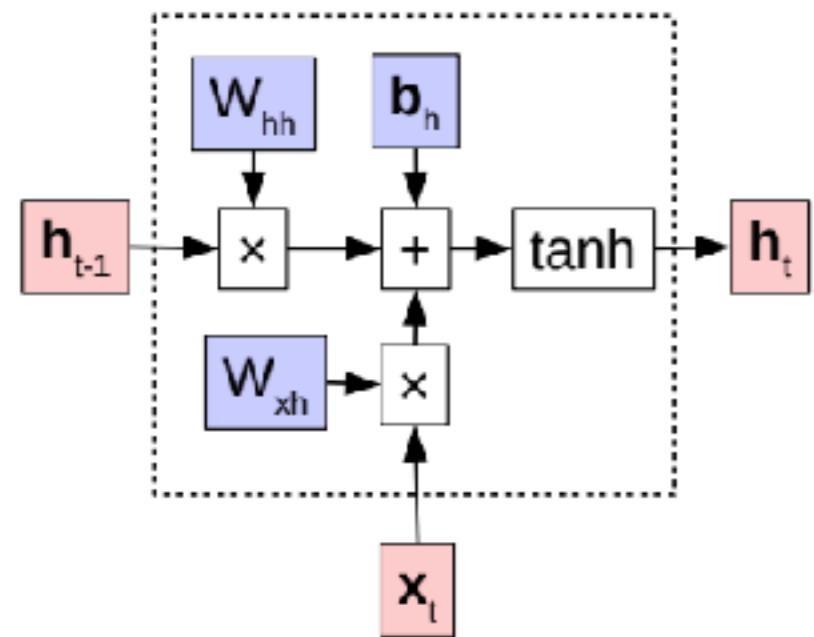
$$h = g(Vx_t + Uh_{t-1})$$

$$h_i = f_\theta(h_{i-1}, x_i)$$

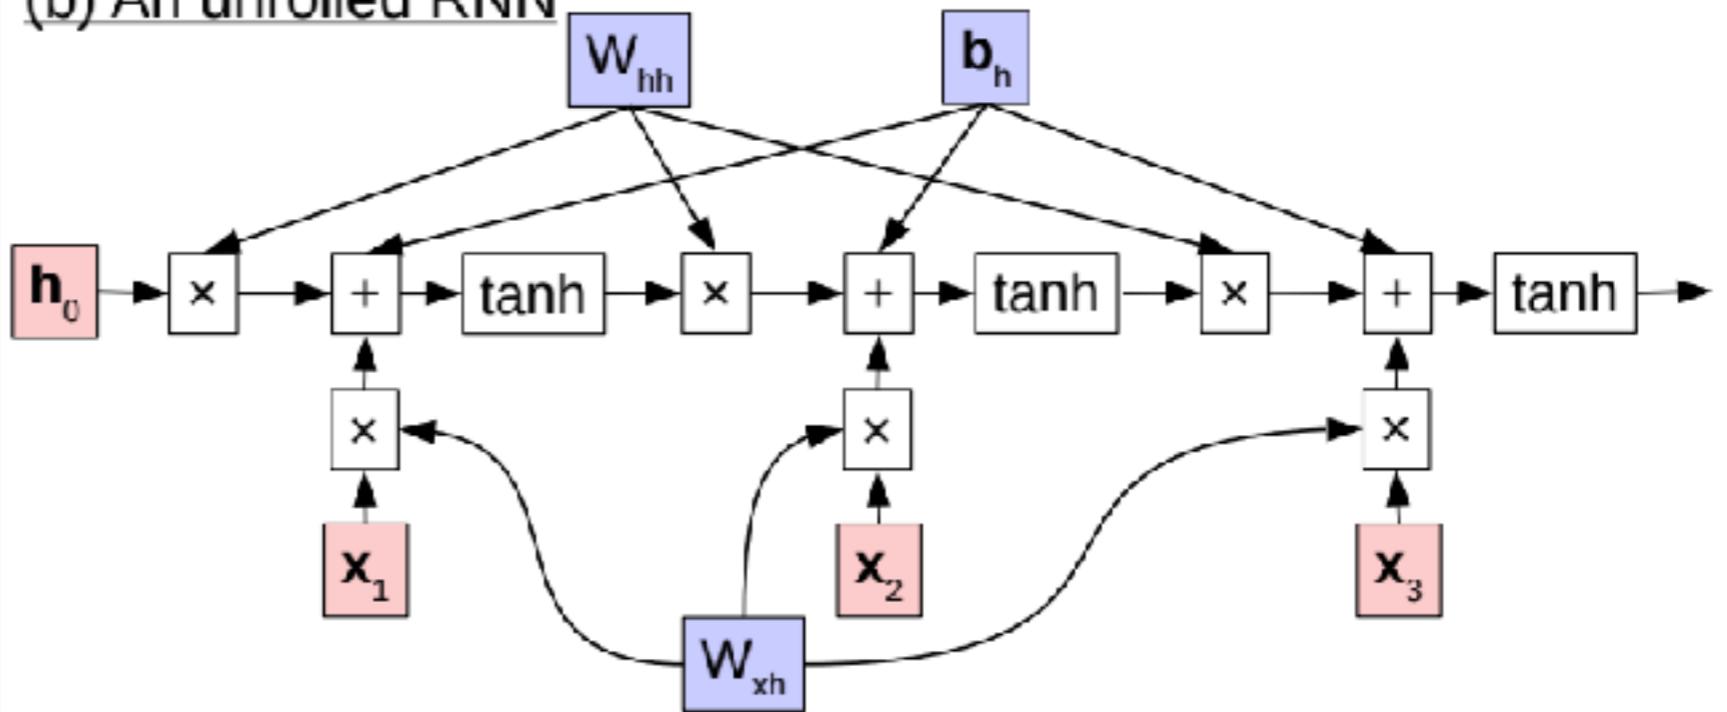


# Summary of Views:

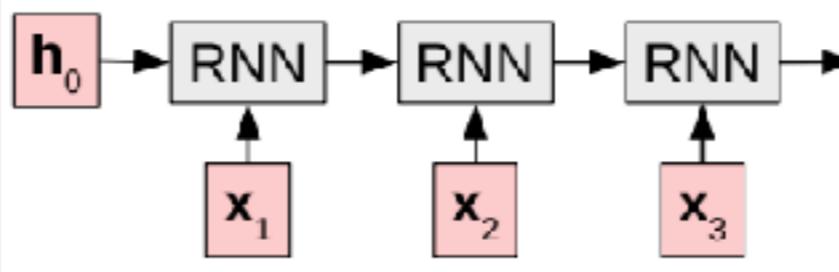
(a) A single RNN time step



(b) An unrolled RNN



(c) A simplified view

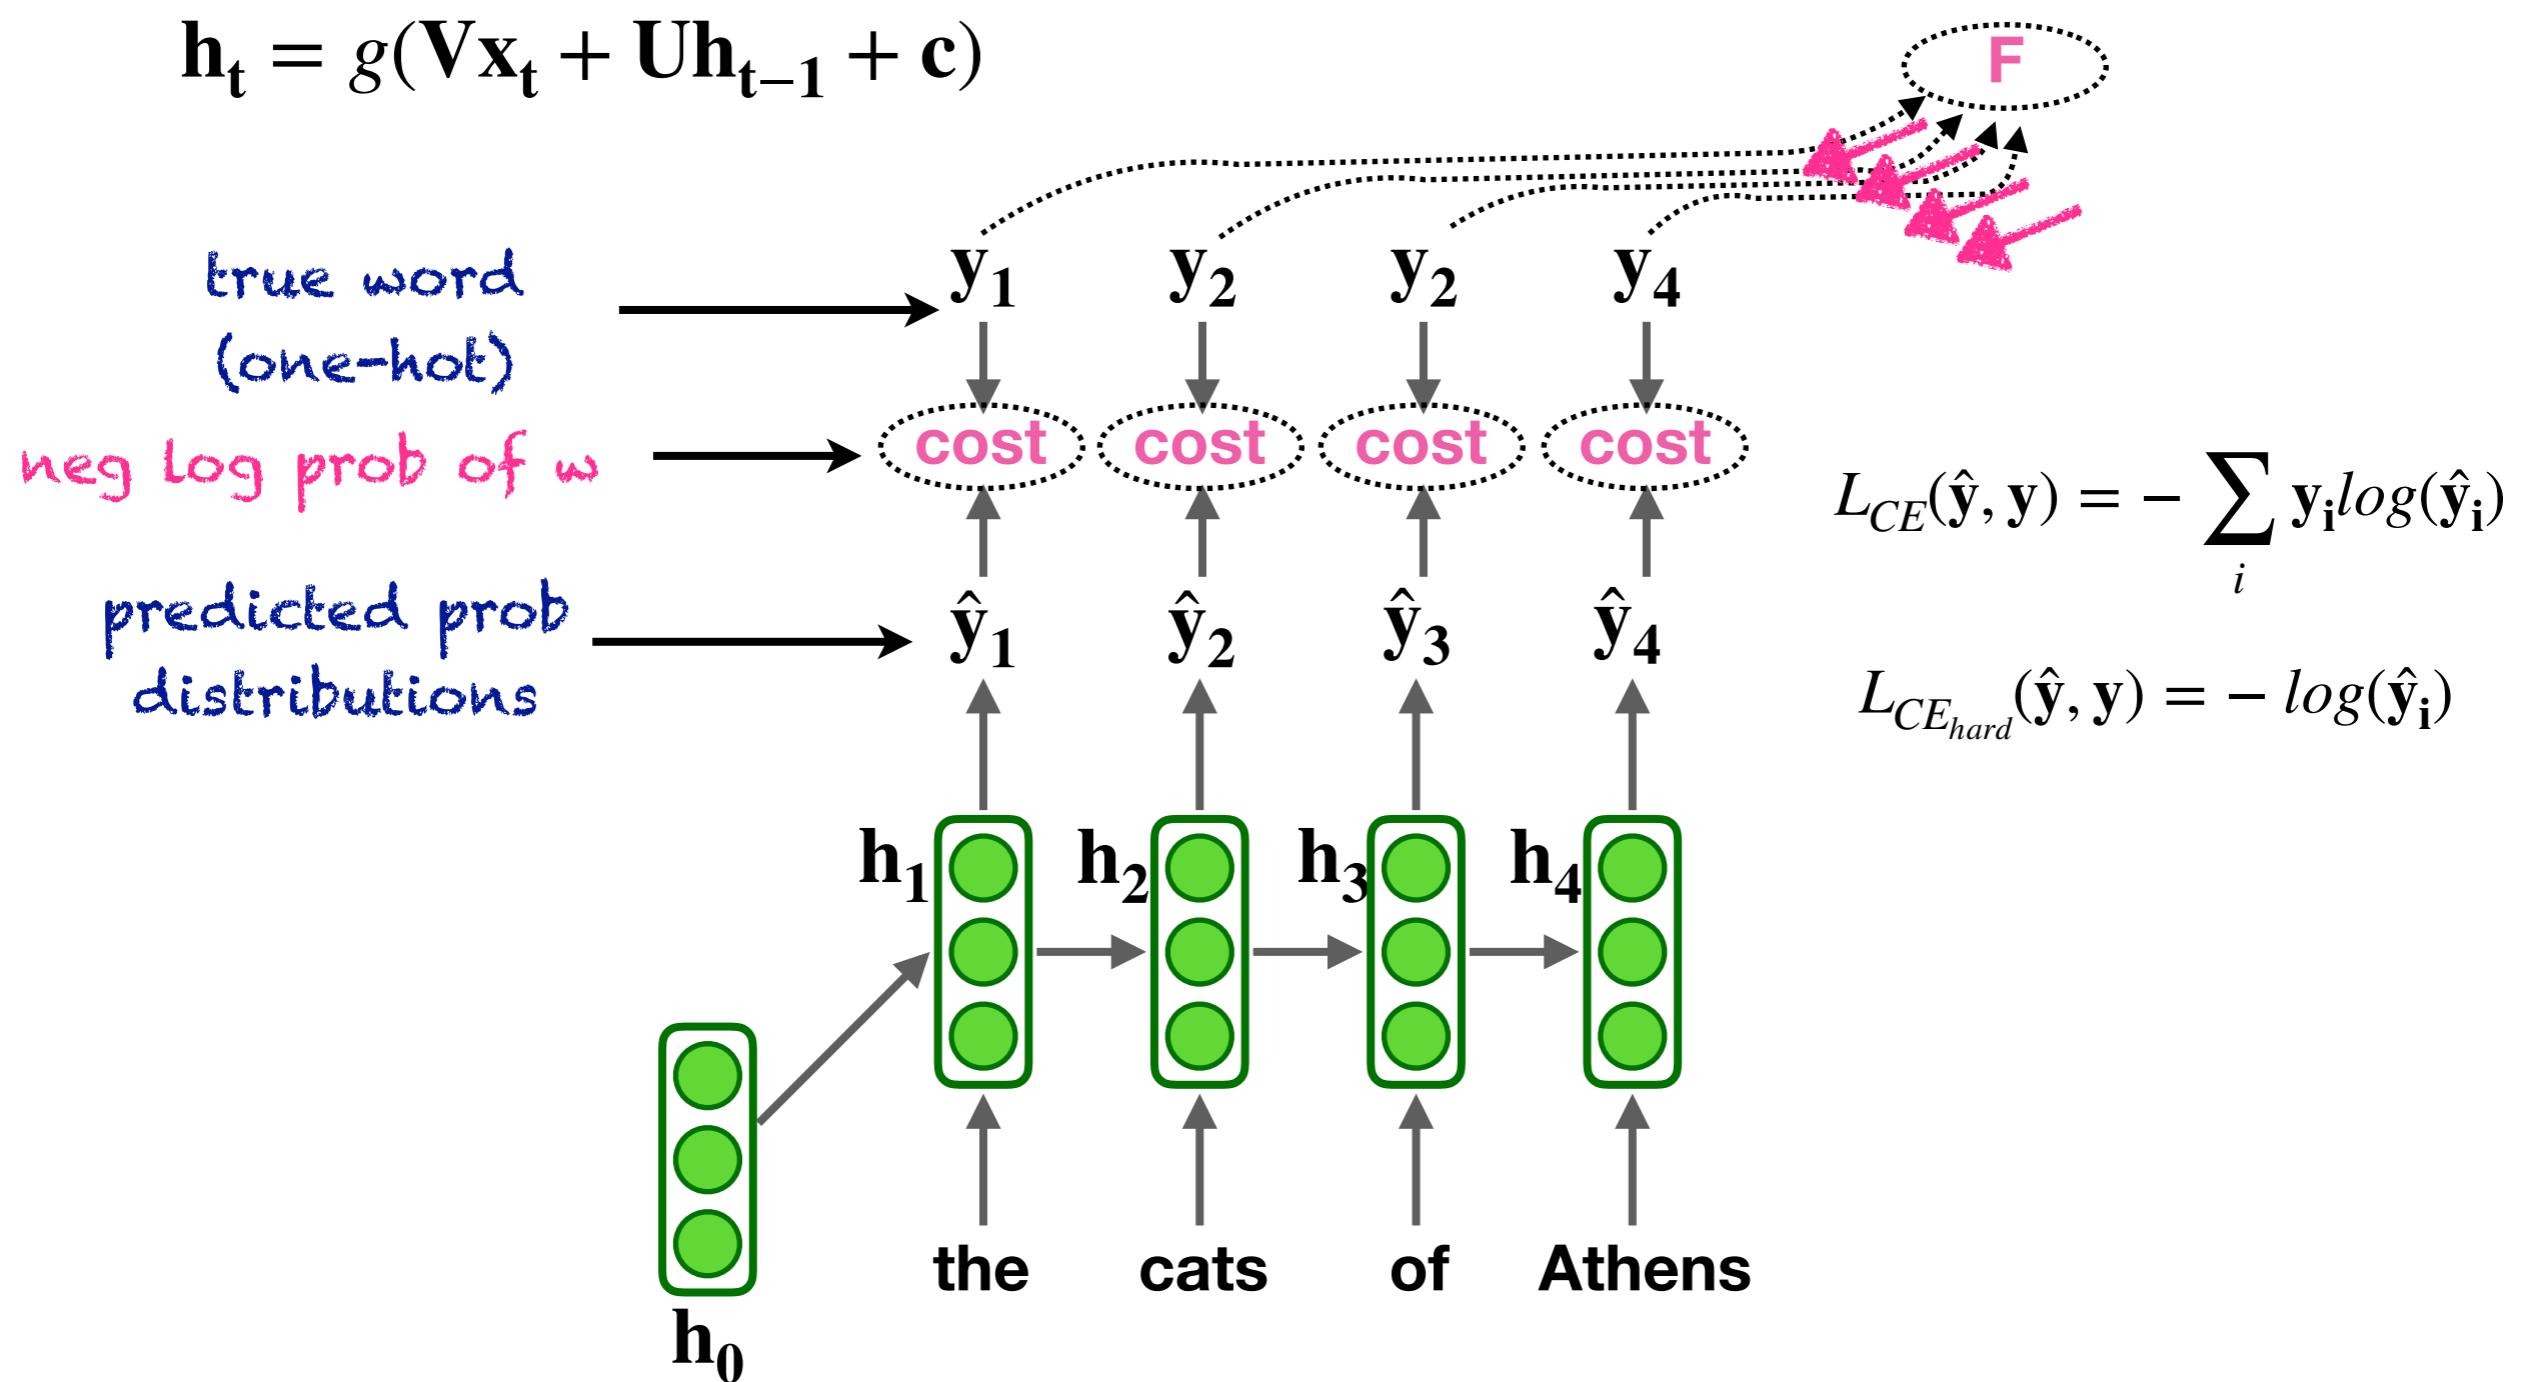


# RNN Language Model

# Training a RNN LM

$$\hat{y}_t = \text{softmax}(\mathbf{W}\mathbf{h}_t + \mathbf{b}) \in \mathbb{R}^{|V|}$$

$$\mathbf{h}_t = g(\mathbf{V}\mathbf{x}_t + \mathbf{U}\mathbf{h}_{t-1} + \mathbf{c})$$



# What about these issues?

- ▶ Can it handle similar words?
  - ▶ she *bought* a bicycle
  - ▶ she *purchased* a bicycle
- ▶ Long-distance dependencies?
  - ▶ for *programming* she yesterday purchased her own brand new *laptop*
  - ▶ for *running* she yesterday purchased her brand new *sportswatch*



However, in practice the vanilla RNN  
has some trouble.. more soon



# Generate with a RNN LM - some fun!

## Generating Baby Names

### character-level RNN-LM

Lets try one more for fun. Lets feed the RNN a large text file that contains 8000 baby names listed out, one per line (names obtained from [here](#)). We can feed this to the RNN and then generate new names! Here are some example names, only showing the ones that do not occur in the training data (90% don't):

*Rudi Levette Berice Lussa Hany Mareanne Chrestina Carissy Marylen Hammie Janye Marlise Jacacie  
Hendred Romand Charienna Nenotto Ette Dorane Wallen Marly Darine Salina Elvyn Ersia Maralena Minoria Ellia  
Charmin Antley Nerille Chelon Walmor Evena Jeryly Stachon Charisa Allisa Anatha Cathanie Geetra Alexie Jerin  
Cassen Herbett Cossie Velen Daurenge Robester Shermond Terisa Licia Roselen Ferine Jayn Lusine Charyanne  
Sales Sanny Resa Wallon Martine Merus Jelen Candica Wallin Tel Rachene Tarine Ozila Ketia Shanne Arnande  
Karella Roselina Alessia Chasty Deland Berther Geamar Jackein Mellisand Sagdy Nenc Lessie Rasemy Guen  
Gavi Milea Anneda Margoris Janin Rodelin Zeanna Elyne Janah Ferzina Susta Pey Castina*

You can see many more [here](#). Some of my favorites include "Baby" (haha), "Killie", "Char", "R", "More", "Mars", "Hi", "Saddie", "With" and "Ahbort". Well that was fun. Of course, you can imagine this being quite useful inspiration when writing a novel, or naming a new startup :)

# Generate with a RNN LM - some fun!

## Algebraic Geometry (Latex)

The results above suggest that the model is actually quite good at learning complex syntactic structures.

Impressed by these results, my labmate ([Justin Johnson](#)) and I decided to push even further into structured territories and got a hold of [this book](#) on algebraic stacks/geometry. We downloaded the raw Latex source file (a 16MB file) and trained a multilayer LSTM. Amazingly, the resulting sampled Latex *almost* compiles. We had to step in and fix a few issues manually but then you get plausible looking math, it's quite astonishing:

For  $\bigoplus_{n=1,\dots,n} \text{where } \mathcal{L}_{n,\bullet} = 0$ , hence we can find a closed subset  $\mathcal{H}$  in  $\mathcal{H}$  and any sets  $\mathcal{F}$  on  $X$ ,  $U$  is a closed immersion of  $S$ , then  $U \rightarrow T$  is a separated algebraic space.

*Proof.* Proof of (1). It also start we get

$$S = \text{Spec}(R) = U \times_X U \times_X U$$

and the comparicoly in the fibre product covering we have to prove the lemma generated by  $\coprod Z \times_U U \rightarrow V$ . Consider the maps  $M$  along the set of points  $\text{Sch}_{fppf}$  and  $U \rightarrow U$  is the fibre category of  $S$  in  $U$  in Section ?? and the fact that any  $U$  affine, see Morphisms, Lemma ???. Hence we obtain a scheme  $S$  and any open subset  $W \subset U$  in  $\text{Sh}(G)$  such that  $\text{Spec}(R') \rightarrow S$  is smooth or an

$$U = \bigsqcup U_i \times_{S_i} U_i$$

which has a nonzero morphism we may assume that  $f_i$  is of finite presentation over  $S$ . We claim that  $\mathcal{O}_{X,x}$  is a scheme where  $x, x' \in S'$  such that  $\mathcal{O}_{X,x'} \rightarrow \mathcal{O}_{X',x'}$  is separated. By Algebra, Lemma ?? we can define a map of complexes  $\text{GL}_S(x'/S')$  and we win.  $\square$

To prove study we see that  $\mathcal{F}|_U$  is a covering of  $X'$ , and  $\mathcal{T}_i$  is an object of  $\mathcal{F}_{X/S}$  for  $i > 0$  and  $\mathcal{F}_p$  exists and let  $\mathcal{F}_i$  be a presheaf of  $\mathcal{O}_X$ -modules on  $\mathcal{C}$  as a  $\mathcal{F}$ -module. In particular  $\mathcal{F} = U/\mathcal{F}$  we have to show that

$$\tilde{\mathcal{M}}^* = \mathcal{I}^* \otimes_{\text{Spec}(k)} \mathcal{O}_{S,s} - i_X^{-1} \mathcal{F}$$

is a unique morphism of algebraic stacks. Note that

$$\text{ARROWS} = (\text{Sch}/S)^{opp}_{fppf}, (\text{Sch}/S)_{fppf}$$

and

$$V = \Gamma(S, \mathcal{C}) \longrightarrow (U, \text{Spec}(A))$$

is an open subset of  $X$ . Thus  $U$  is affine. This is a continuous map of  $X$  is the inverse, the groupoid scheme  $S$ .

*Proof.* See discussion of sheaves of sets.  $\square$

The result for prove any open covering follows from the less of Example ???. It may replace  $S$  by  $X_{\text{spaces},\text{etale}}$  which gives an open subspace of  $X$  and  $T$  equal to  $S_{\text{Zar}}$ , see Descent, Lemma ???. Namely, by Lemma ?? we see that  $R$  is geometrically regular over  $S$ .

**Lemma 0.1.** Assume (3) and (3) by the construction in the description.

Suppose  $X = \lim |X|$  (by the formal open covering  $X$  and a single map  $\text{Proj}_X(\mathcal{A}) = \text{Spec}(B)$  over  $U$  compatible with the complex

$$\text{Set}(\mathcal{A}) = \Gamma(X, \mathcal{O}_X \mathcal{O}_X).$$

When in this case of to show that  $Q \rightarrow \mathcal{C}_{Z/X}$  is stable under the following result in the second conditions of (1), and (3). This finishes the proof. By Definition ?? (without element is when the closed subschemes are catenary. If  $T$  is surjective we may assume that  $T$  is connected with residue fields of  $S$ . Moreover there exists a closed subspace  $Z \subset X$  of  $X$  where  $U$  in  $X'$  is proper (some defining as a closed subset of the uniqueness it suffices to check the fact that the following theorem

(1)  $f$  is locally of finite type. Since  $S = \text{Spec}(R)$  and  $Y = \text{Spec}(R)$ .

*Proof.* This is form all sheaves of sheaves on  $X$ . But given a scheme  $U$  and a surjective étale morphism  $U \rightarrow X$ . Let  $U \cap U = \coprod_{i=1,\dots,n} U_i$  be the scheme  $X$  over  $S$  at the schemes  $X_i \rightarrow X$  and  $U = \lim_i X_i$ .  $\square$

The following lemma surjective restrecomposes of this implies that  $\mathcal{F}_{x_0} = \mathcal{F}_{x_0} = \mathcal{J}_{X,\dots,g}$ .

**Lemma 0.2.** Let  $X$  be a locally Noetherian scheme over  $S$ ,  $E = \mathcal{F}_{X/S}$ . Set  $\mathcal{I} = \mathcal{J}_1 \subset \mathcal{I}'_n$ . Since  $\mathcal{I}^n \subset \mathcal{I}^n$  are nonzero over  $i_0 \leq p \leq n$  is a subset of  $\mathcal{J}_{n,p} \circ \tilde{A}_2$  works.

**Lemma 0.3.** In Situation ???. Hence we may assume  $q' = 0$ .

*Proof.* We will use the property we see that  $p$  is the most functor (??). On the other hand, by Lemma ?? we see that

$$D(\mathcal{O}_{X'}) = \mathcal{O}_X(D)$$

where  $K$  is an  $F$ -algebra where  $\delta_{n+1}$  is a scheme over  $S$ .  $\square$

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

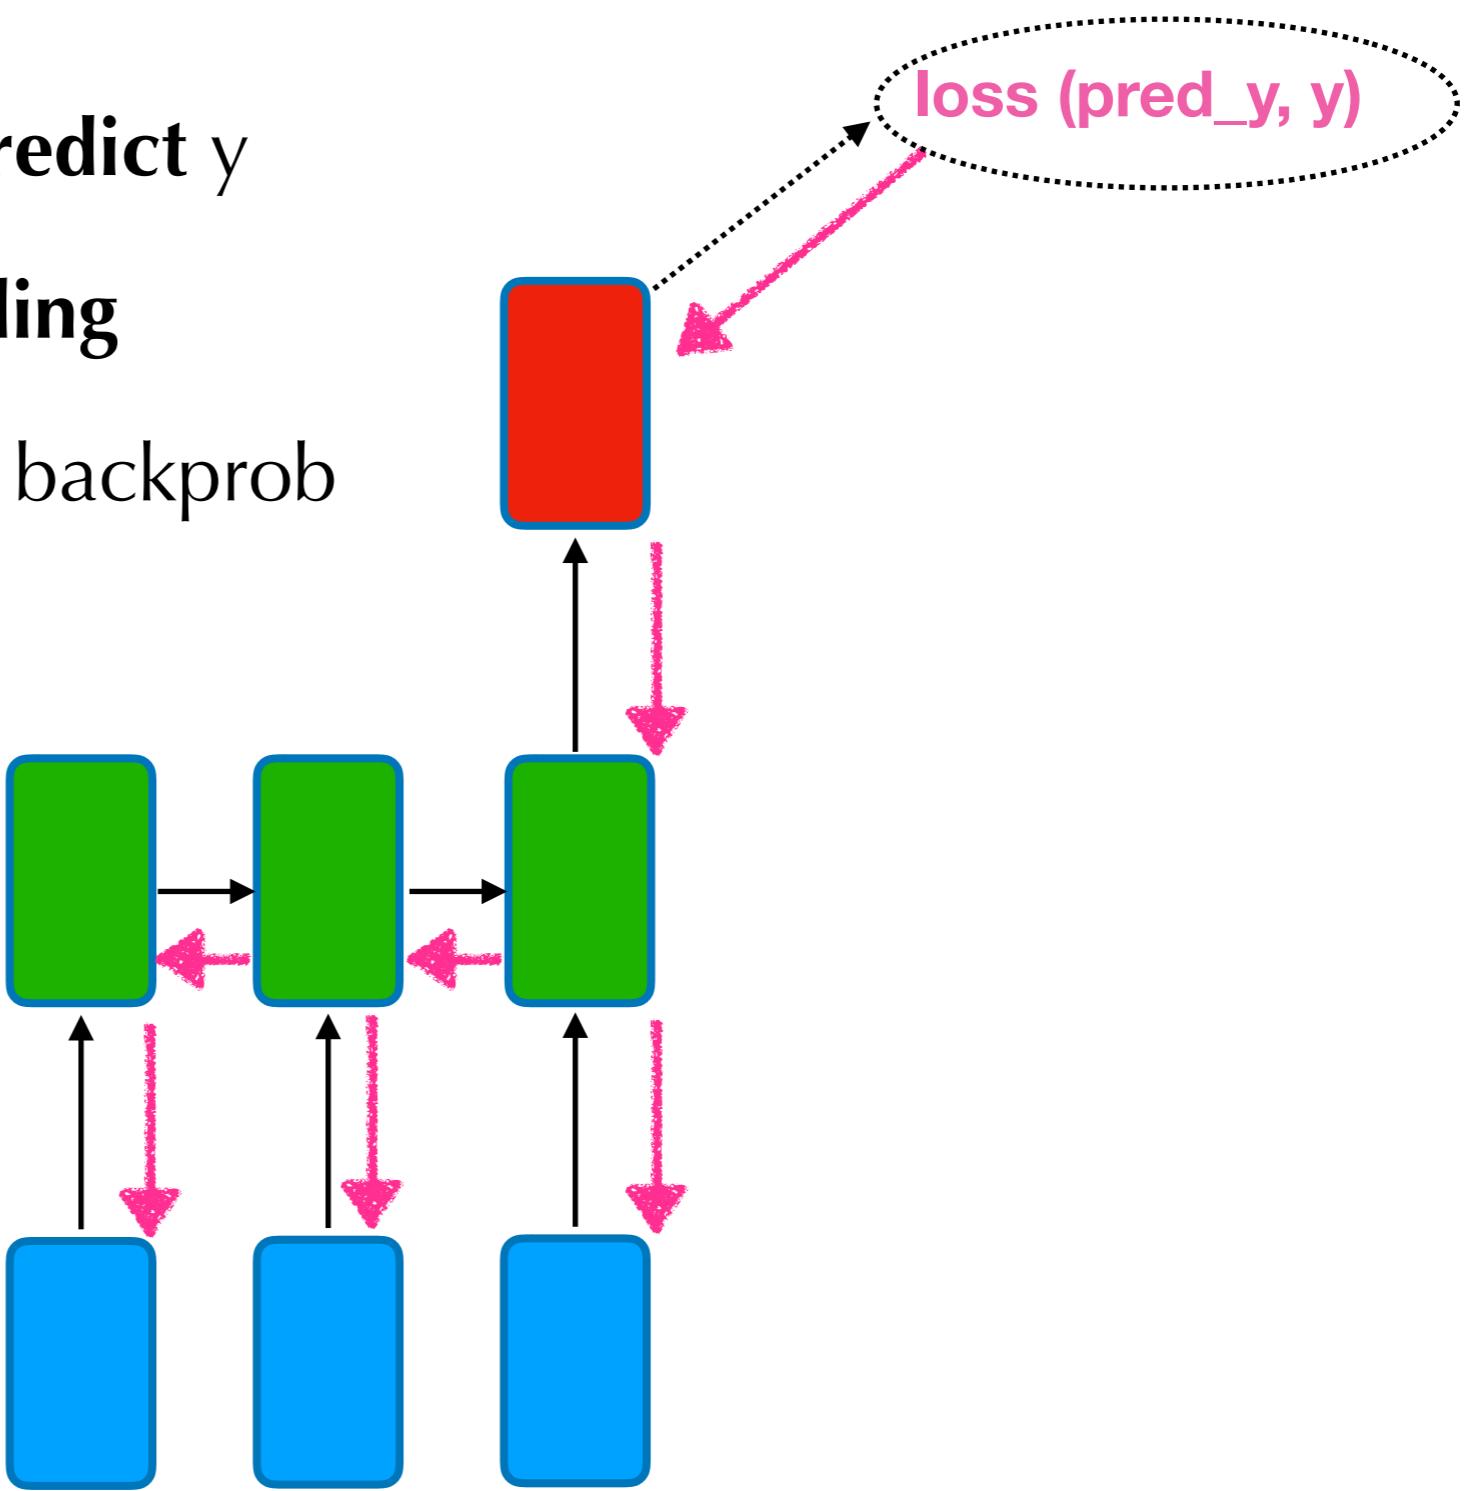
# RNNs - Interim summary

- ▶ **LM:** a model that predicts the next word
- ▶ **RNN:** a family of neural networks
  - ▶ to model sequential input of any length
  - ▶ apply the same parameters on each time step
  - ▶ can optionally produce output at each time step
- ▶ RNN's are great as LMs. But they can be used for much more!

# **Common Usage Patterns**

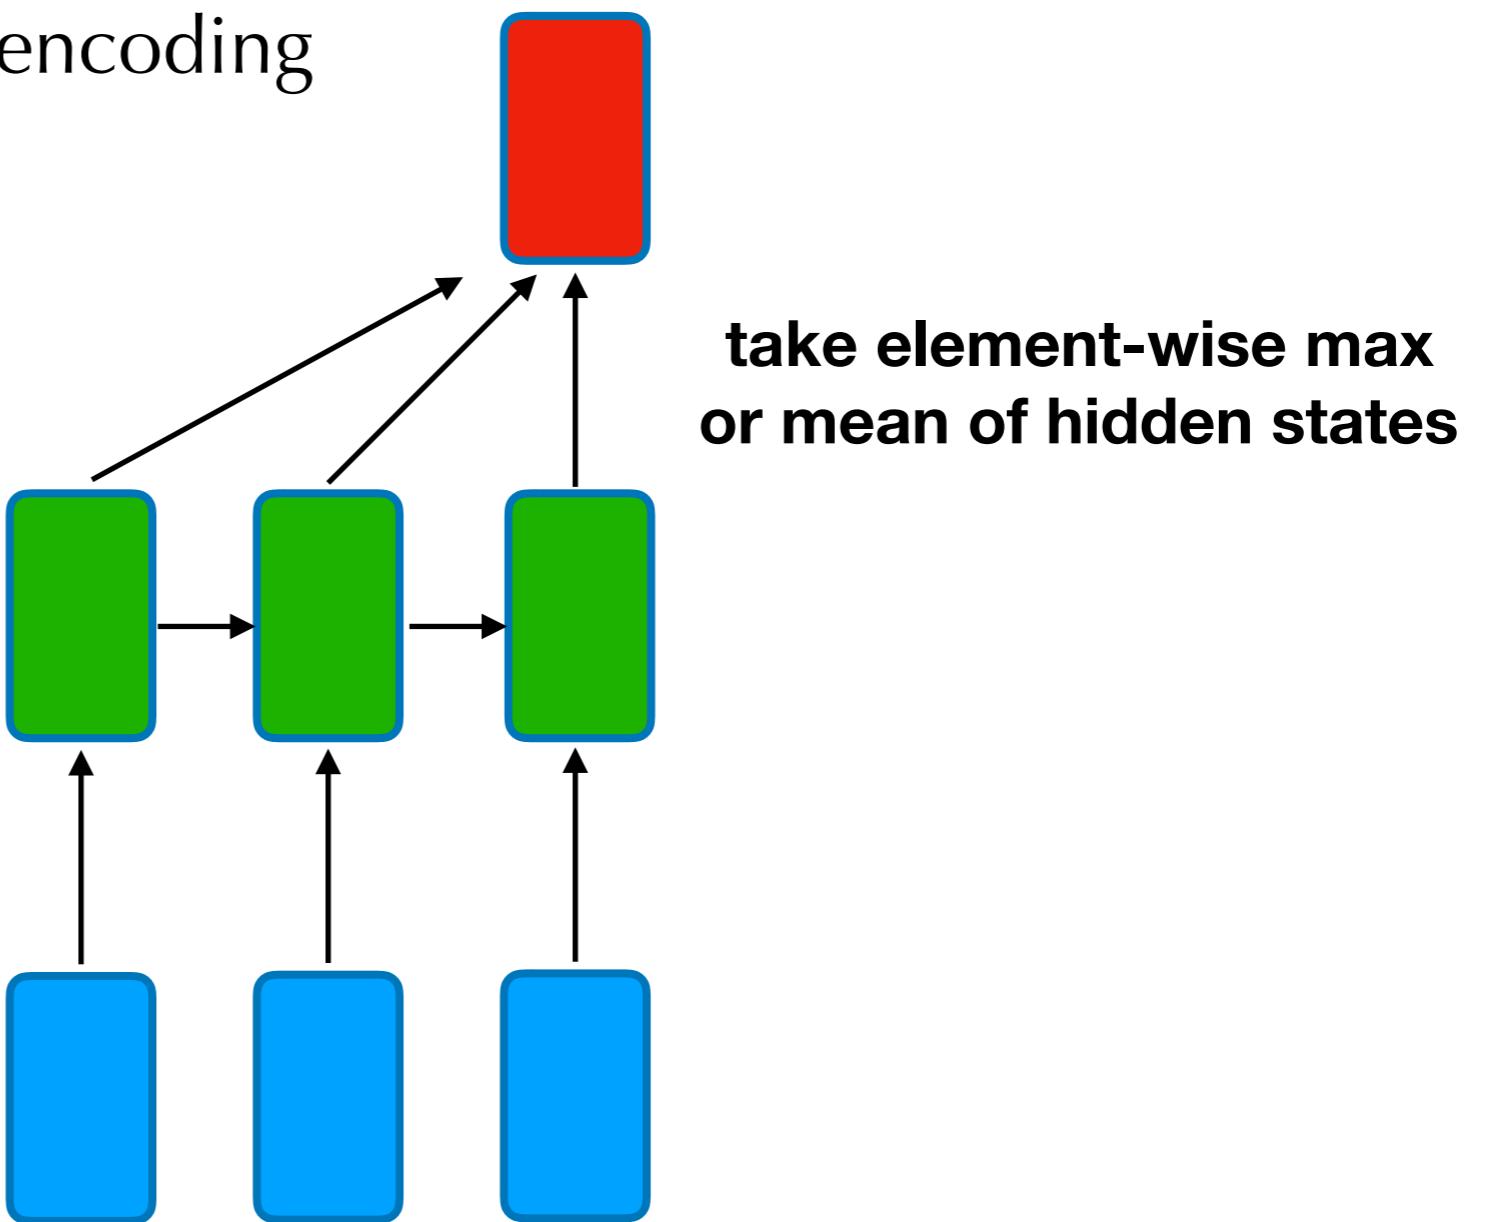
# Example: An RNN as acceptor

- ▶ Use **last state** to **predict y**
  - ▶ **sentence encoding**
- ▶ Calculate loss and backprob



# Example: An RNN as acceptor

- ▶ Use **average of states** to predict  $y$ 
  - ▶ other sentence encoding

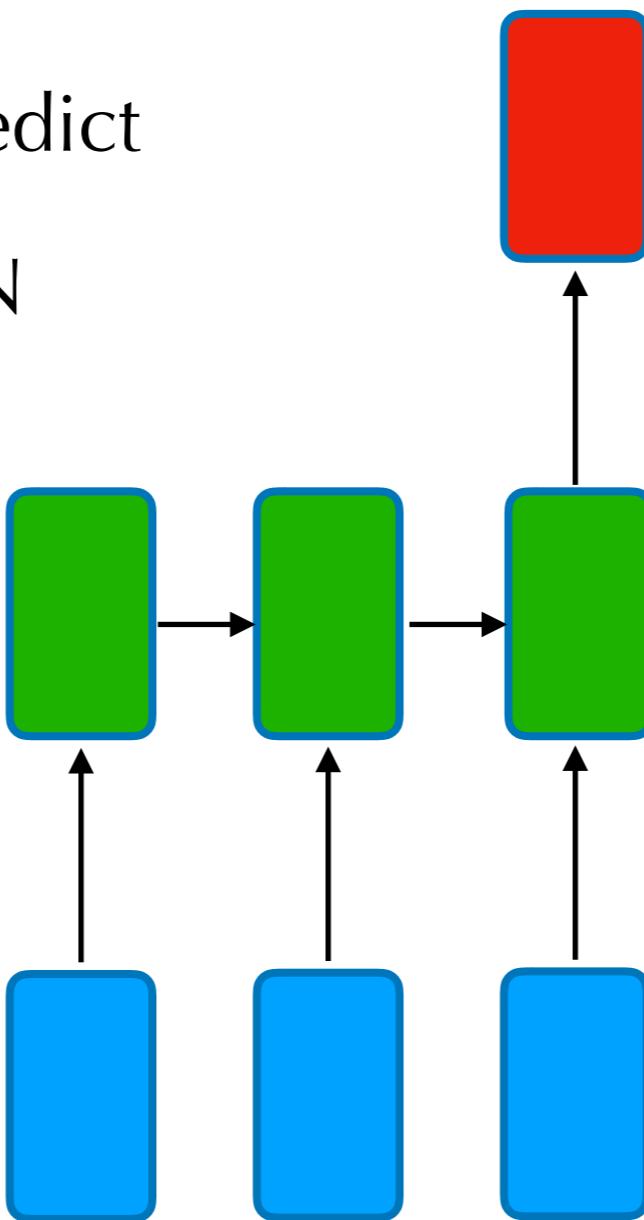


# Example: An RNN as encoder

- ▶ Use **last state** as encoding of the information in the sequence; use as “feature” in other NN

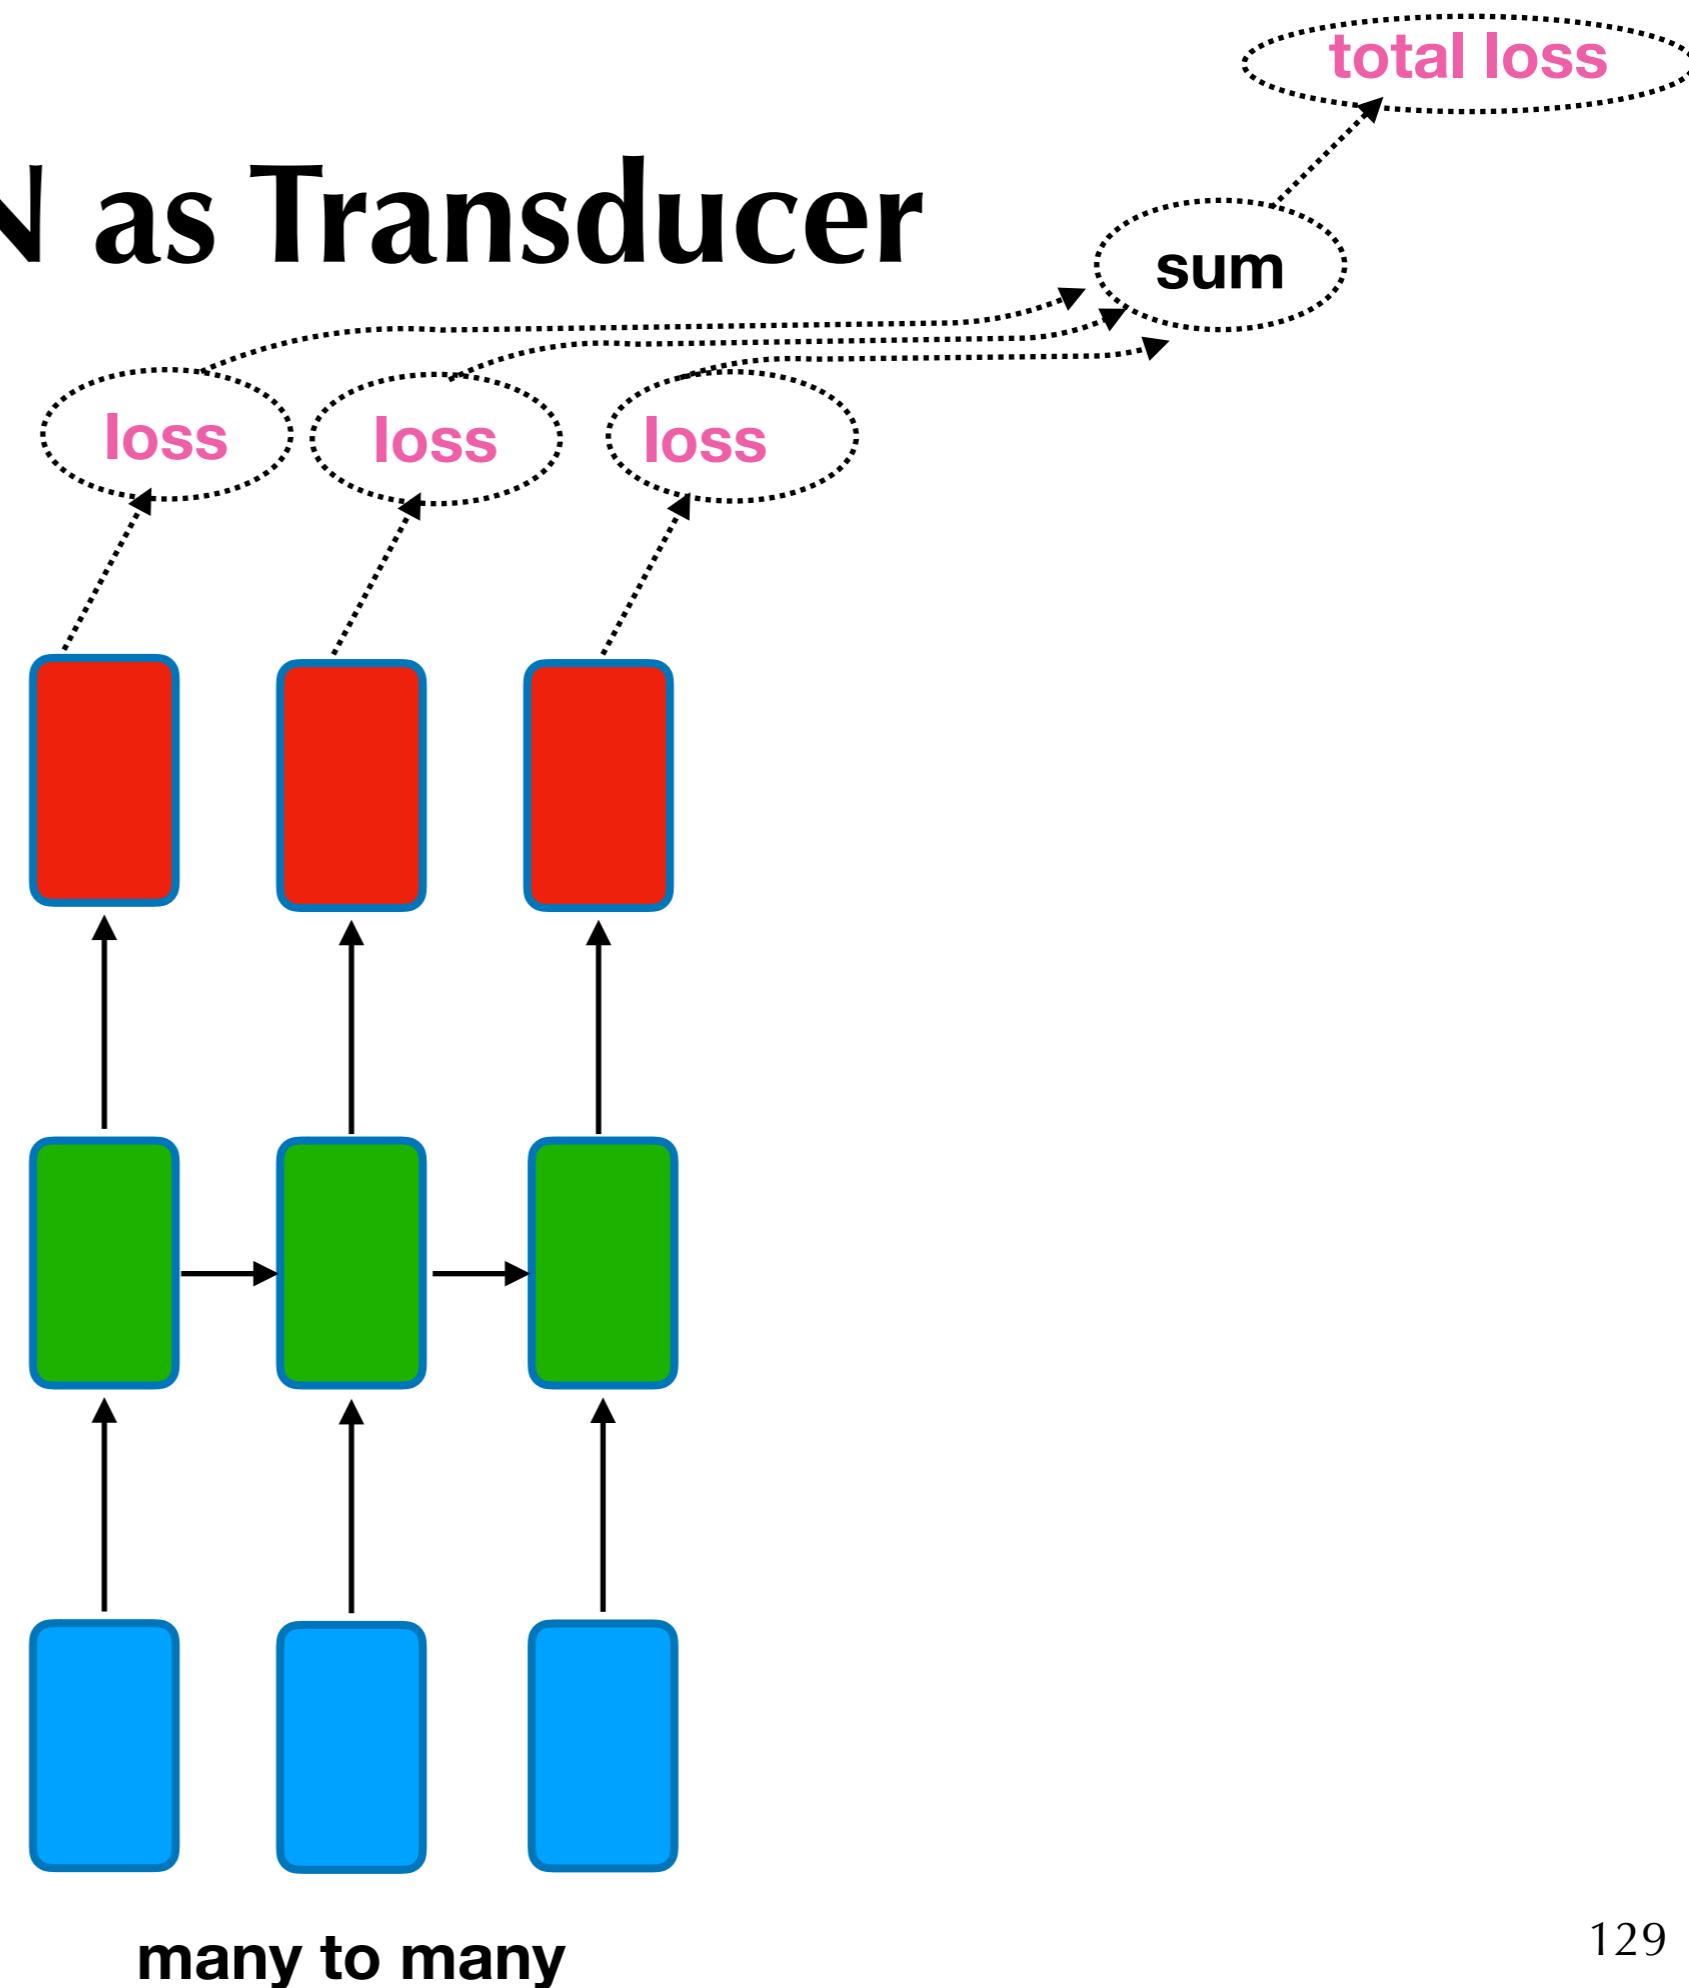
- ▶ encode, not predict

- ▶ E.g. character RNN



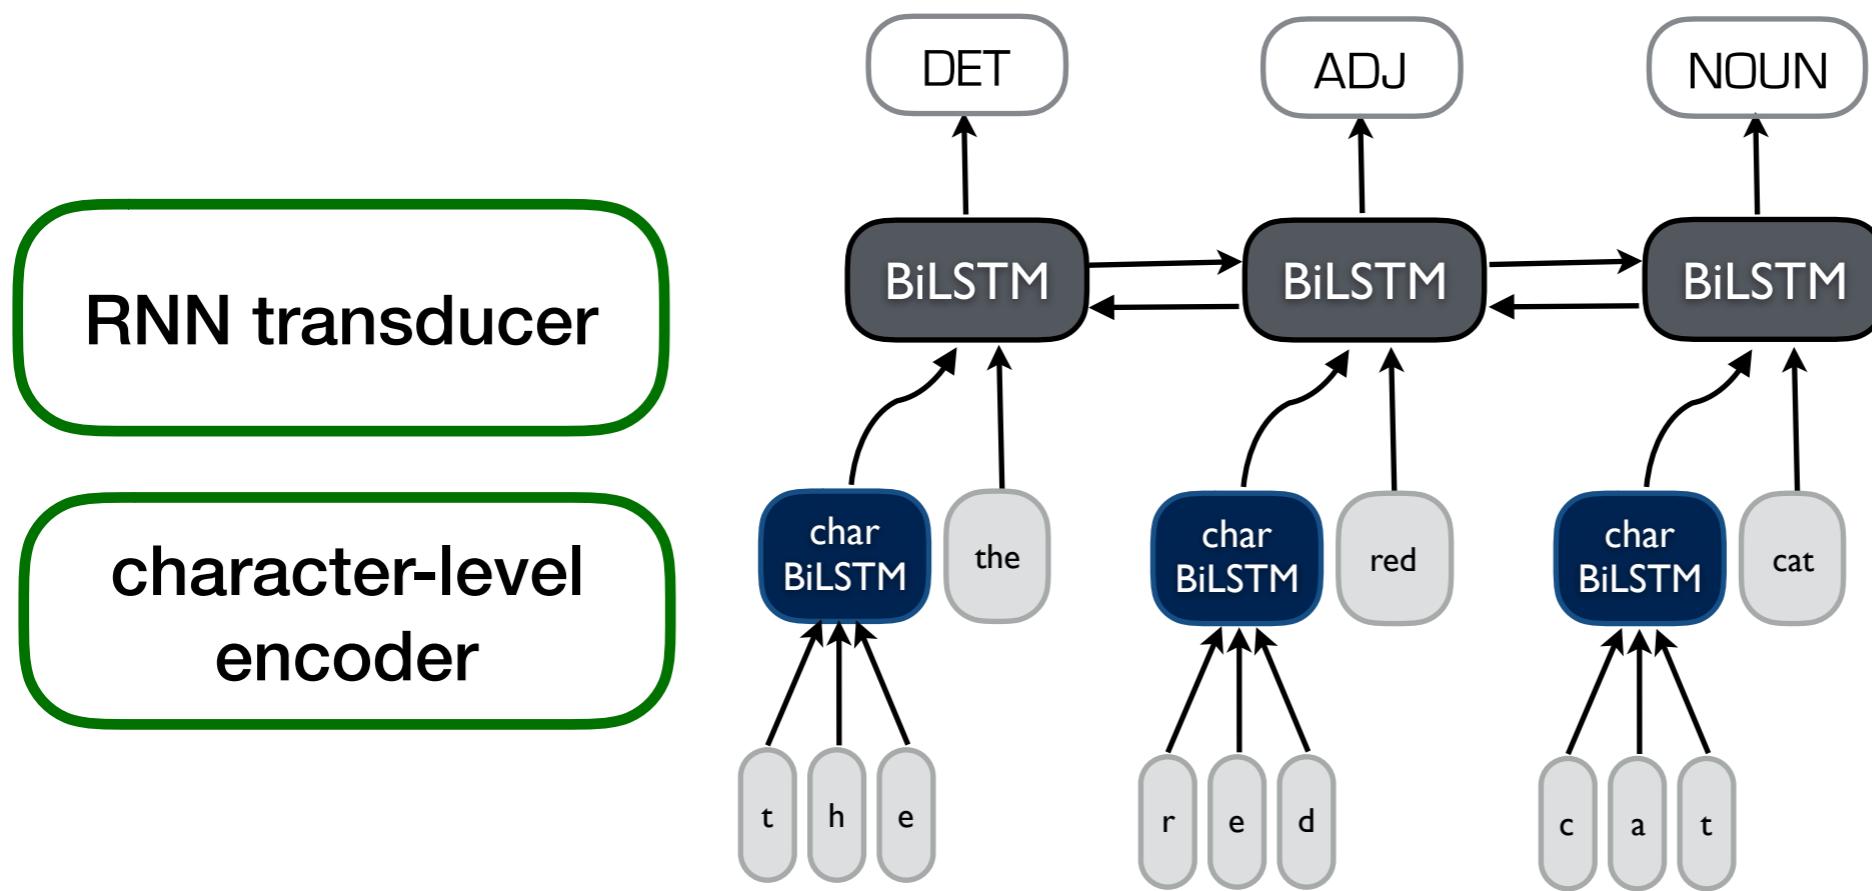
# RNN as Transducer

- ▶ predict an output for each time step  $t$
- ▶ E.g. Tagging (POS, NER)



# Combining them - a hierarchical RNN: Example for POS

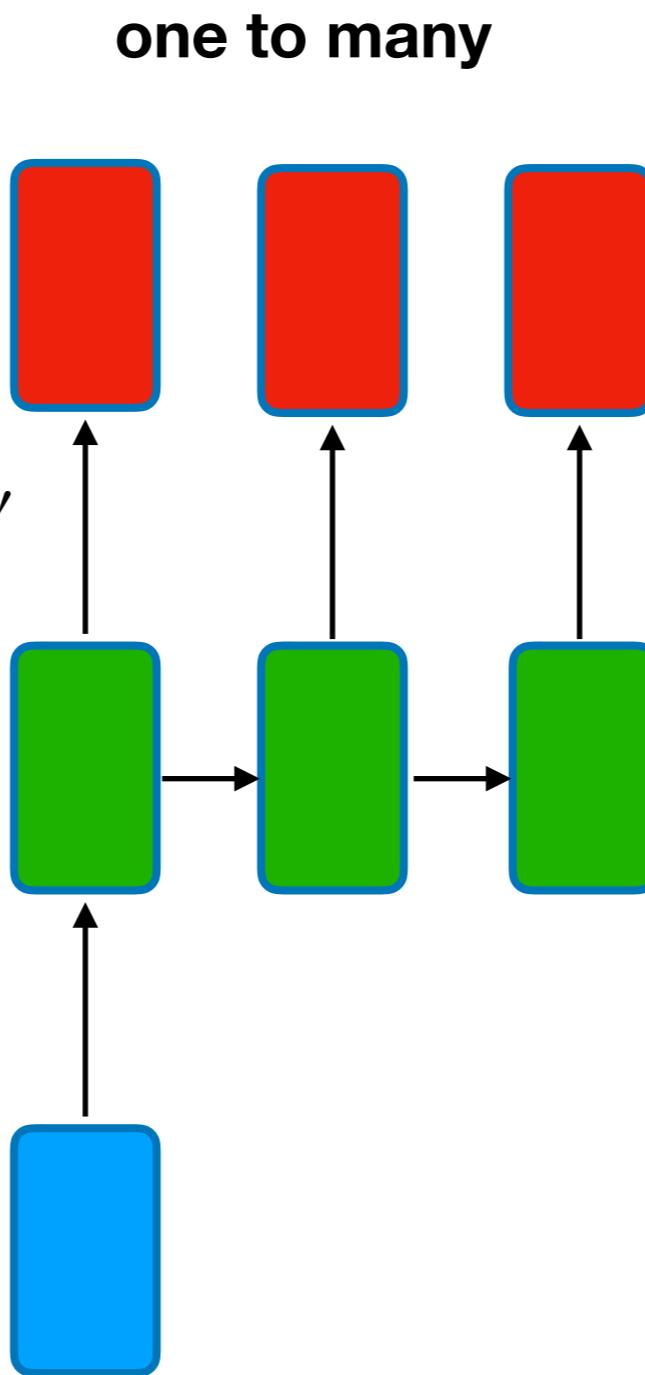
- ▶ Use **RNN transducer** and lower-level **RNN encoder** for **characters** (more in a second)



(Plank et al., 2016)

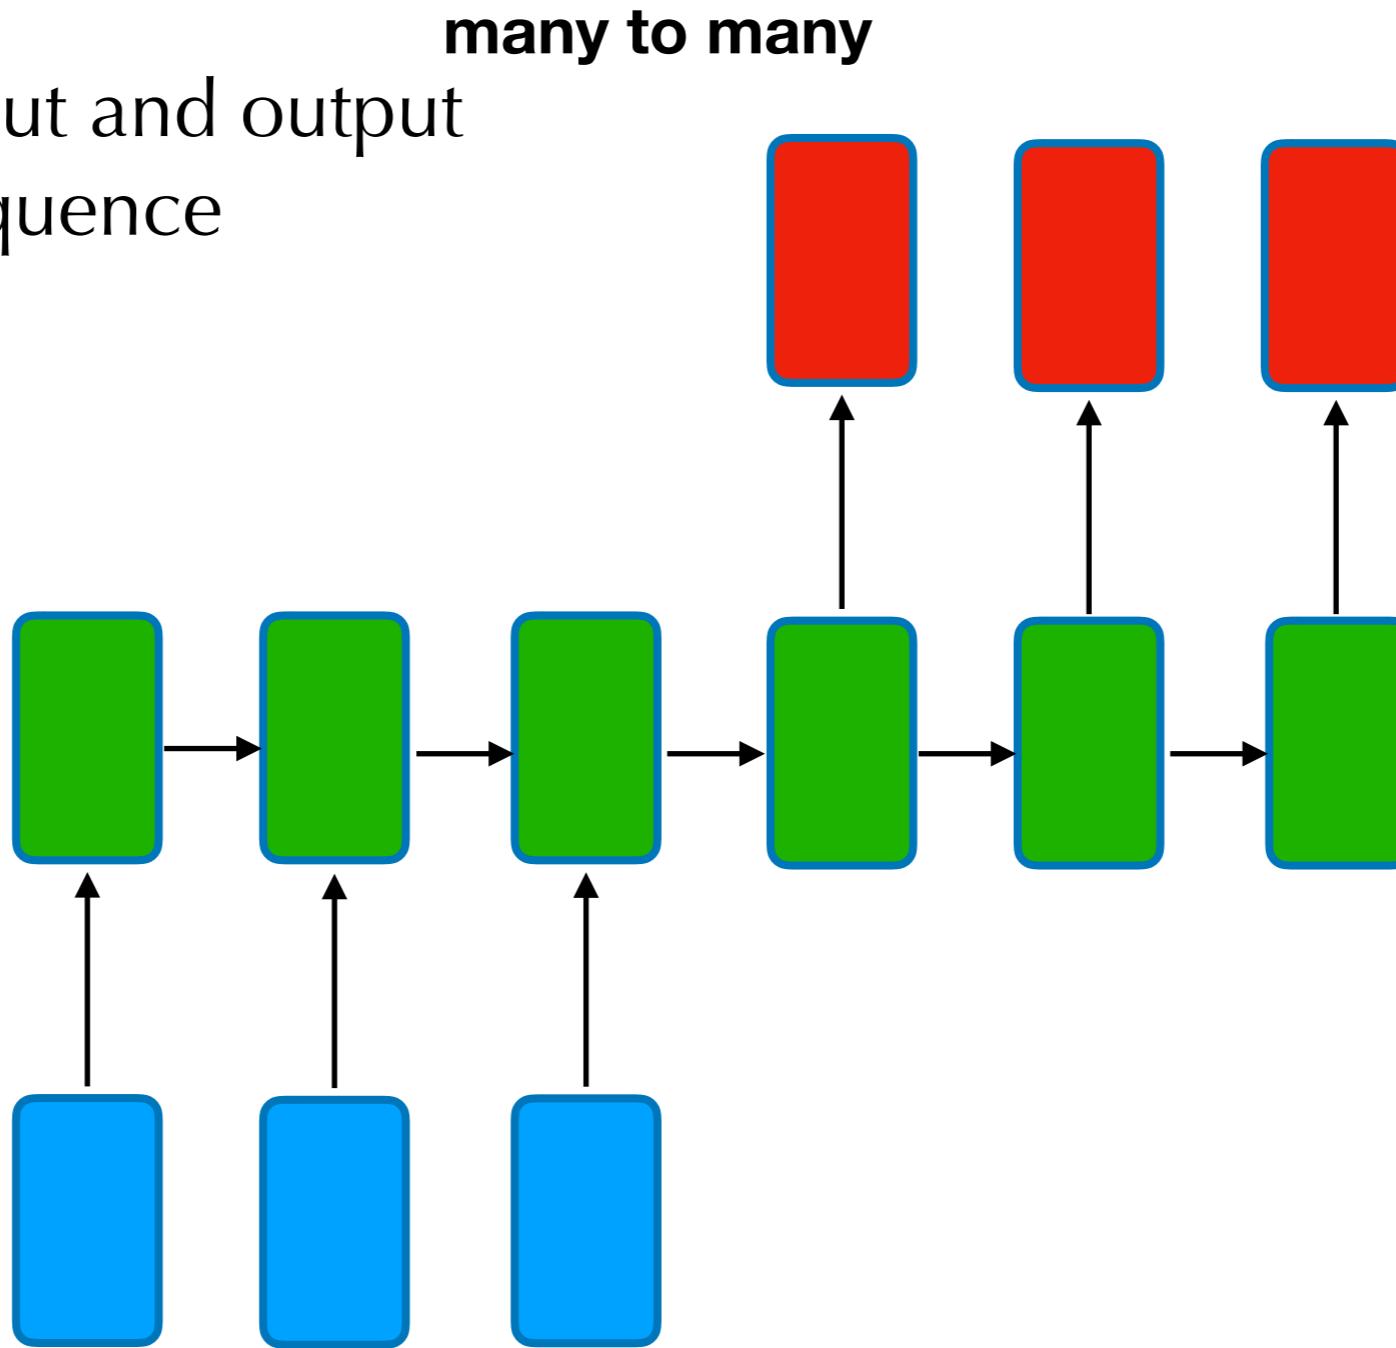
# RNN as generator

- ▶ Conditional generation
- ▶ E.g. image caption generation, speech synthesis



# RNN encoder-decoder (seq2seq)

- ▶ Both input and output are a sequence

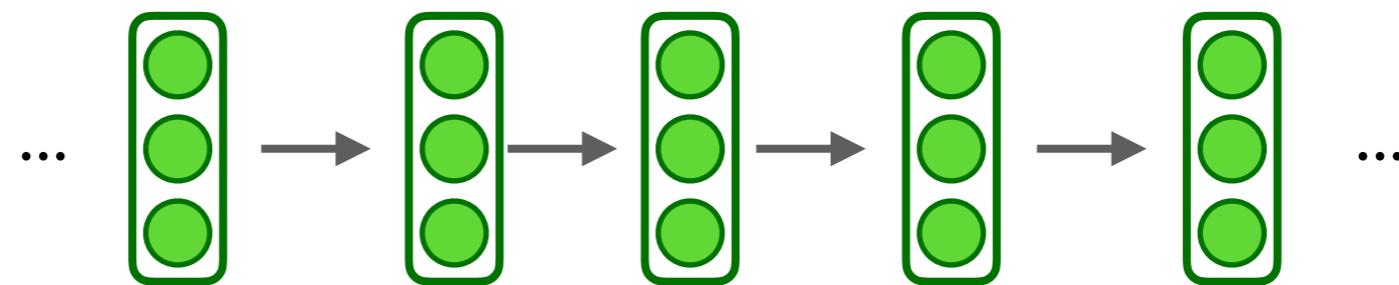


**Deeper, better  
models?**

# Only left to right?

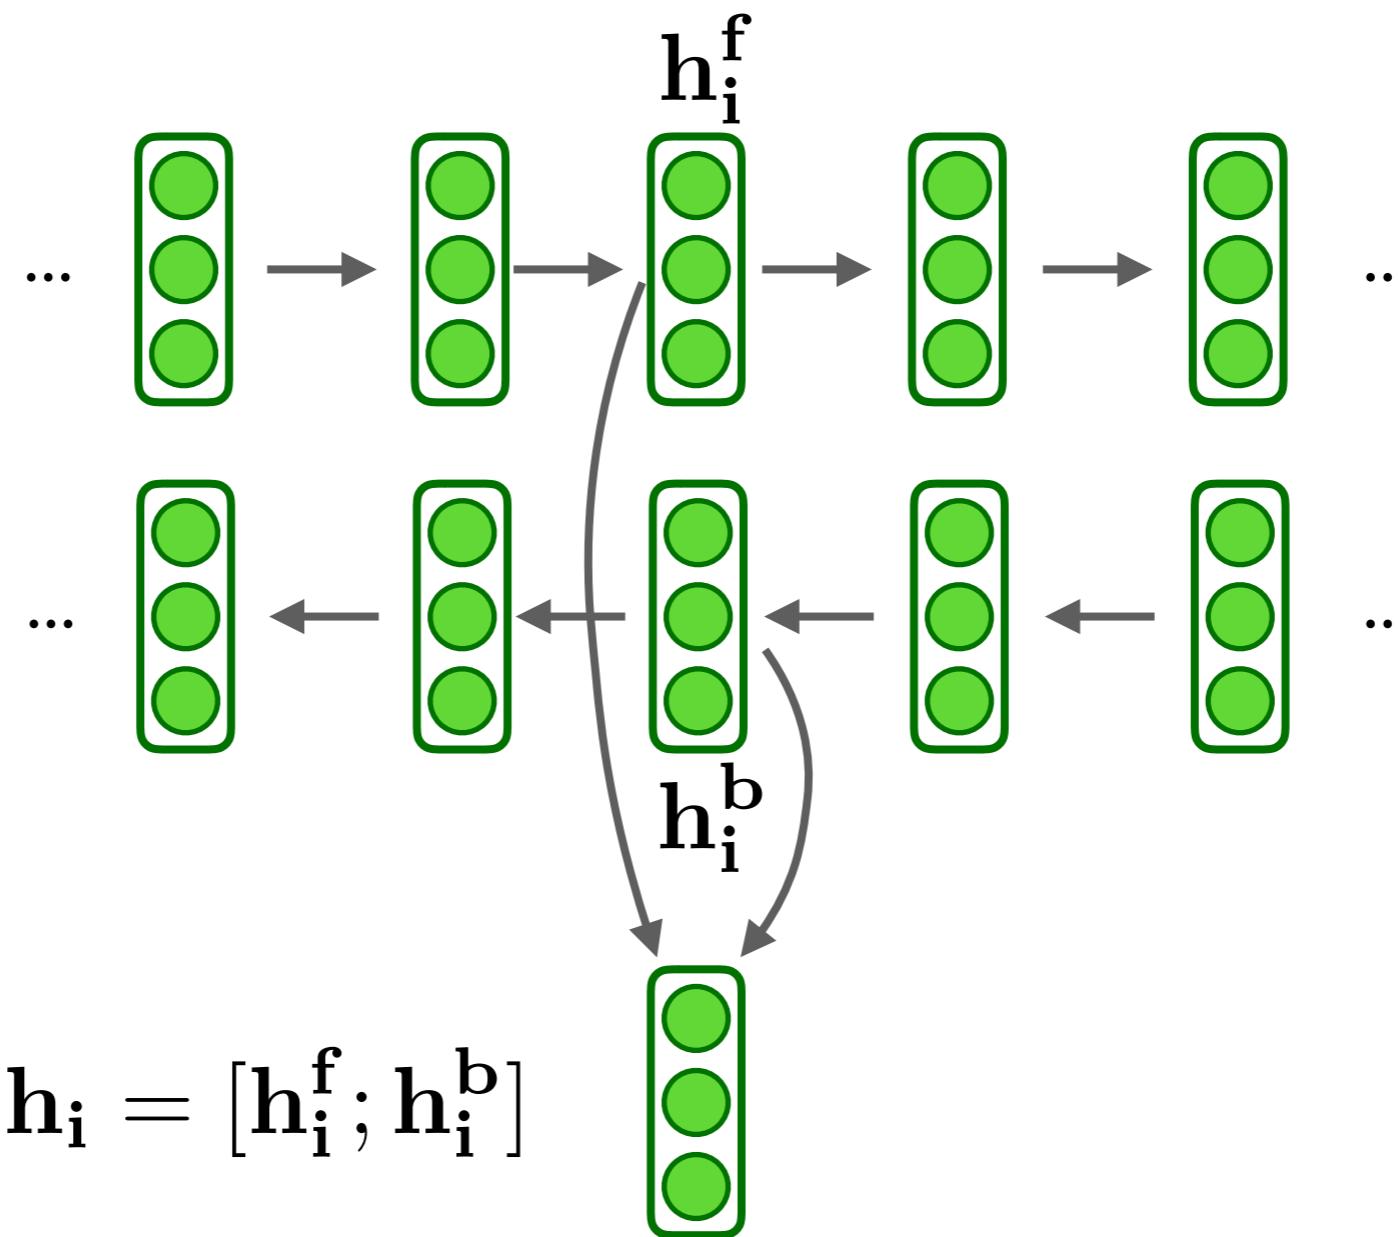
The person who hunts ducks out on the weekends

... person who hunts ducks out ...



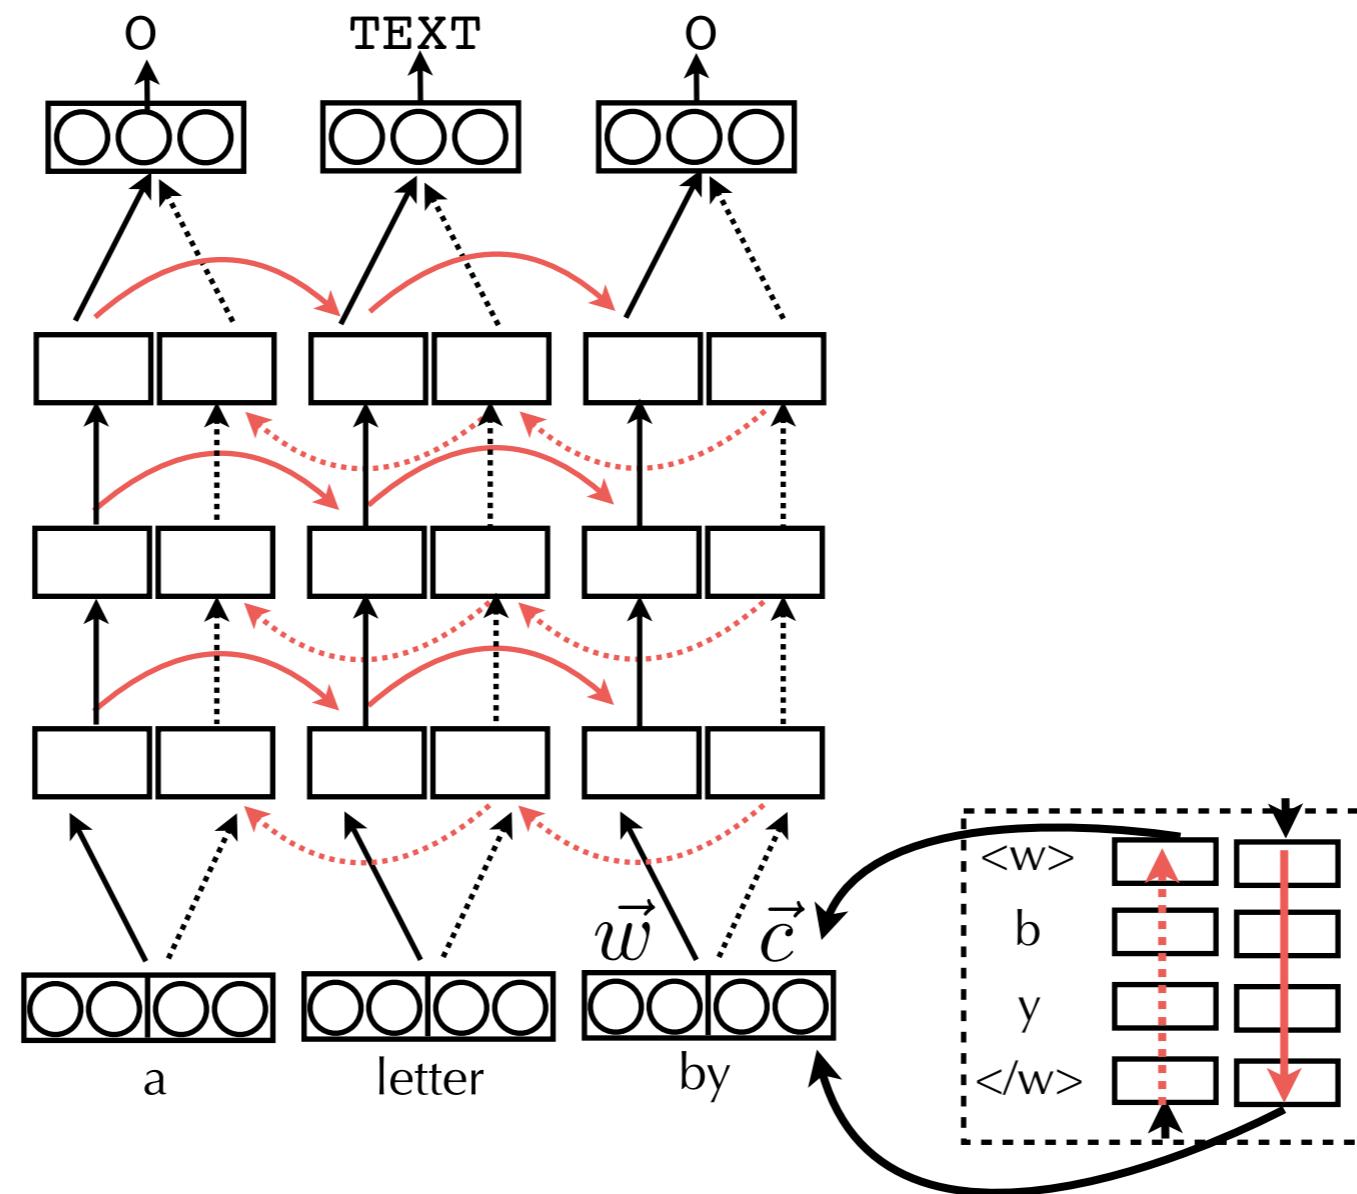
# Bidirectional RNNs

... person who hunts ducks out ...



# Stacked RNNs

- ▶ Multiple layers of RNNs, e.g., bi-RNNs



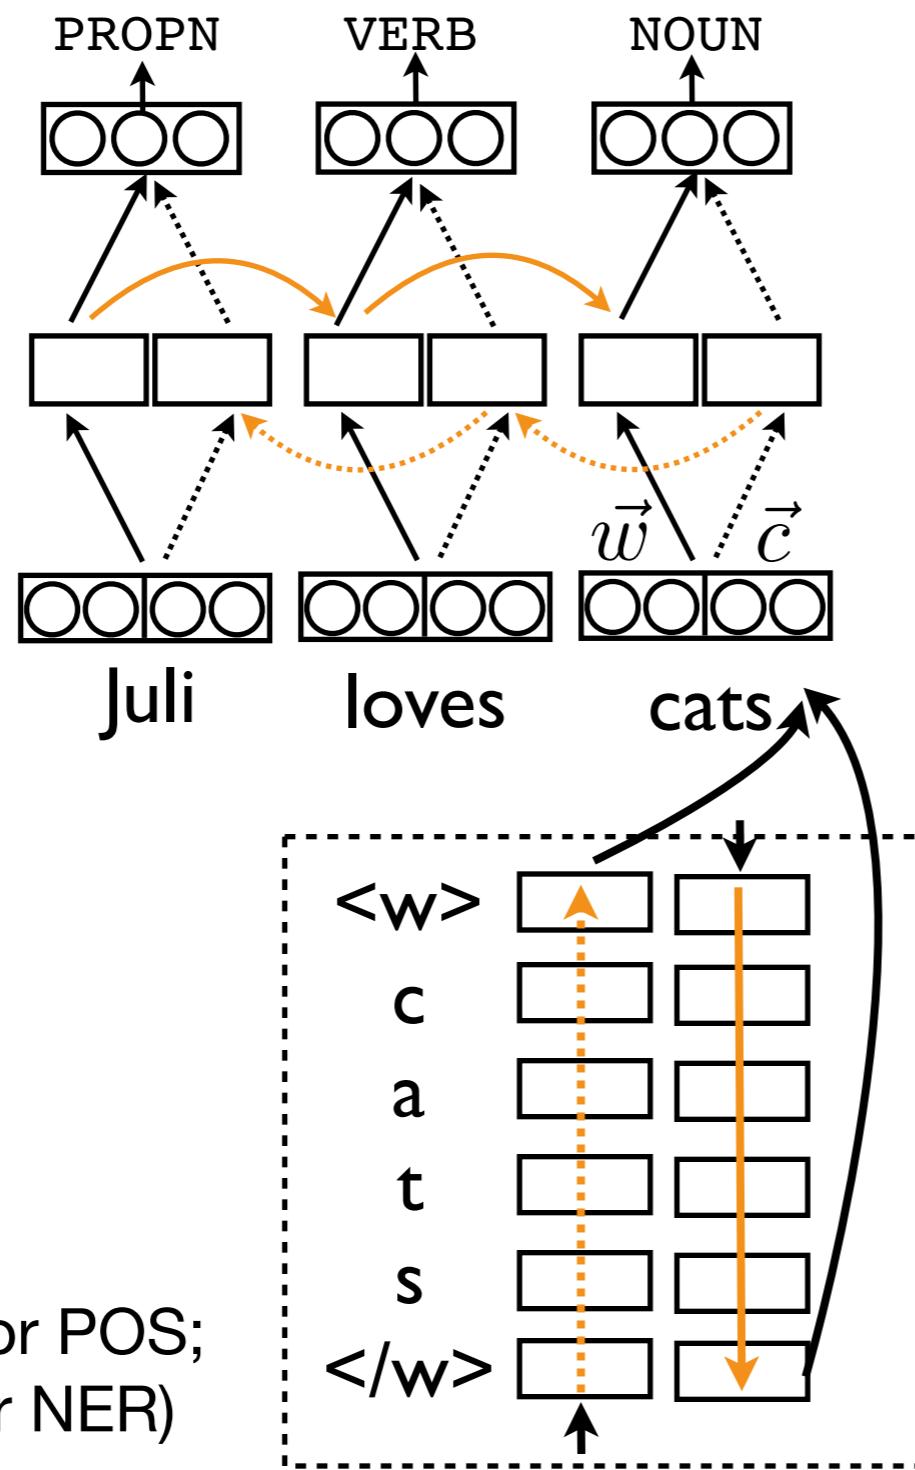
Guards against the long  
tail?

Subword representations

# OOVs (out-of-vocabulary) words

- ▶ So far we saw <UNK>
  - ▶ But that conflates a lot of information into a single <UNK> representations
- ▶ Are we better off modeling at the subword level?

# Subword representations: Characters



\*able (98% adj  
in WSJ)

bi\* (85% noun  
in Danish)

(Plank et al., 2016 for POS;  
Ling et al., 2015 for NER)

# How to model subwords?

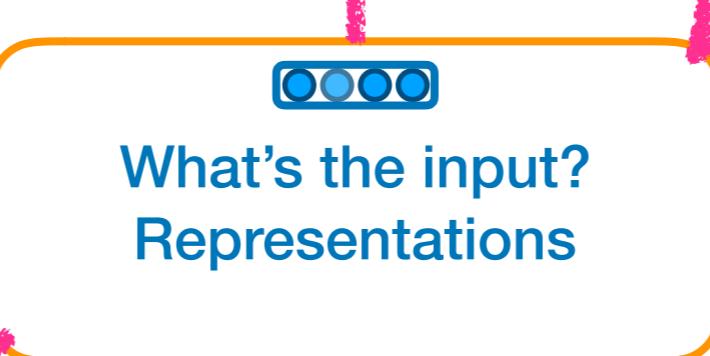
- ▶ Representations:
  - ▶ *Characters*
  - ▶ *Bytes* (e.g., Gillick et al., 2015; Plank et al., 2016)
  - ▶ *Byte-Pair Encoding (BPE)* (Sennrich et al., 2016)
- ▶ Modeling choices:
  - ▶ RNN-variants, CNNs, ...
  - ▶ How to leverage the representations (only char level, combine, ...)

# Overview

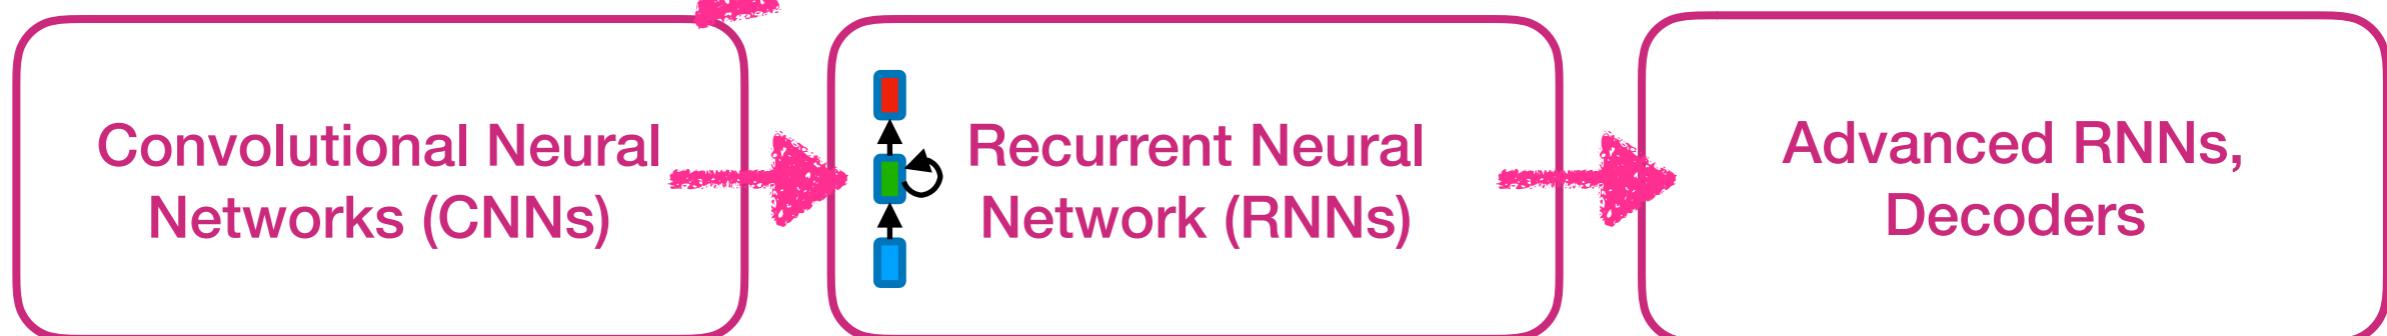
*foundations*



*representations*



*beyond FFNNs*



# A note on terminology

- ▶ RNN = “vanilla” RNN



- ▶ RNN flavors (=gated RNNs):

- ▶ GRU



and LSTMs



- ▶ Why? Problem of RNNs: Vanishing gradients!

# Gated RNN architectures

# Vanishing Gradient

- ▶ Example:
  - ▶ The cat, which ate a ...., was full
  - ▶ The cats, which ... , were full
- ▶ Backprop can have difficulties with long sequences:  
vanishing gradient problem
  - ▶ if the gradient becomes very close to zero:
    - ▶ is it because there is no dependency in the data?
    - ▶ or because of a wrong configuration of the parameters—the vanishing gradient condition?

# Exploding Gradients

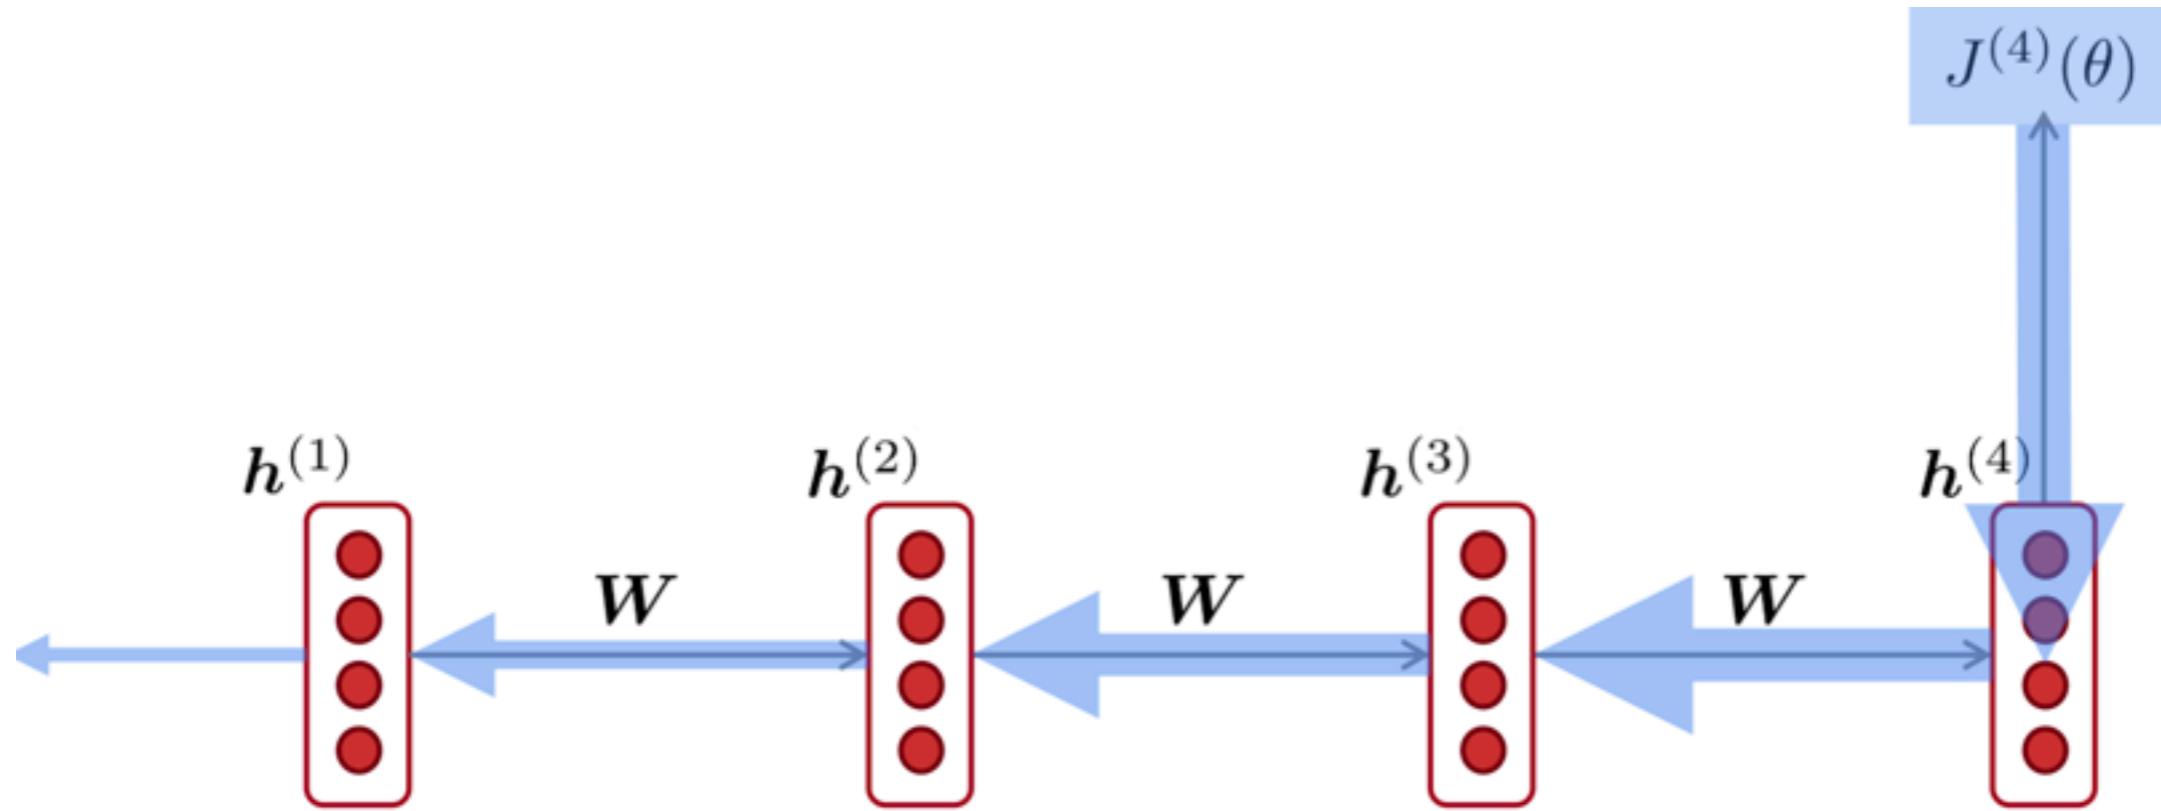
- ▶ Easier to catch. If the gradient becomes too big, then the SGD update becomes very large:

$$\theta^{new} = \theta^{old} - \overbrace{\alpha \nabla_{\theta} J(\theta)}^{\text{gradient}}$$

learning rate

- ▶ This might cause bad updates: too large updates, large loss
- ▶ In the worst case, you might get NaNs or Infs
- ▶ Solution: **gradient clipping** (scale down before update)

# Vanishing Gradient

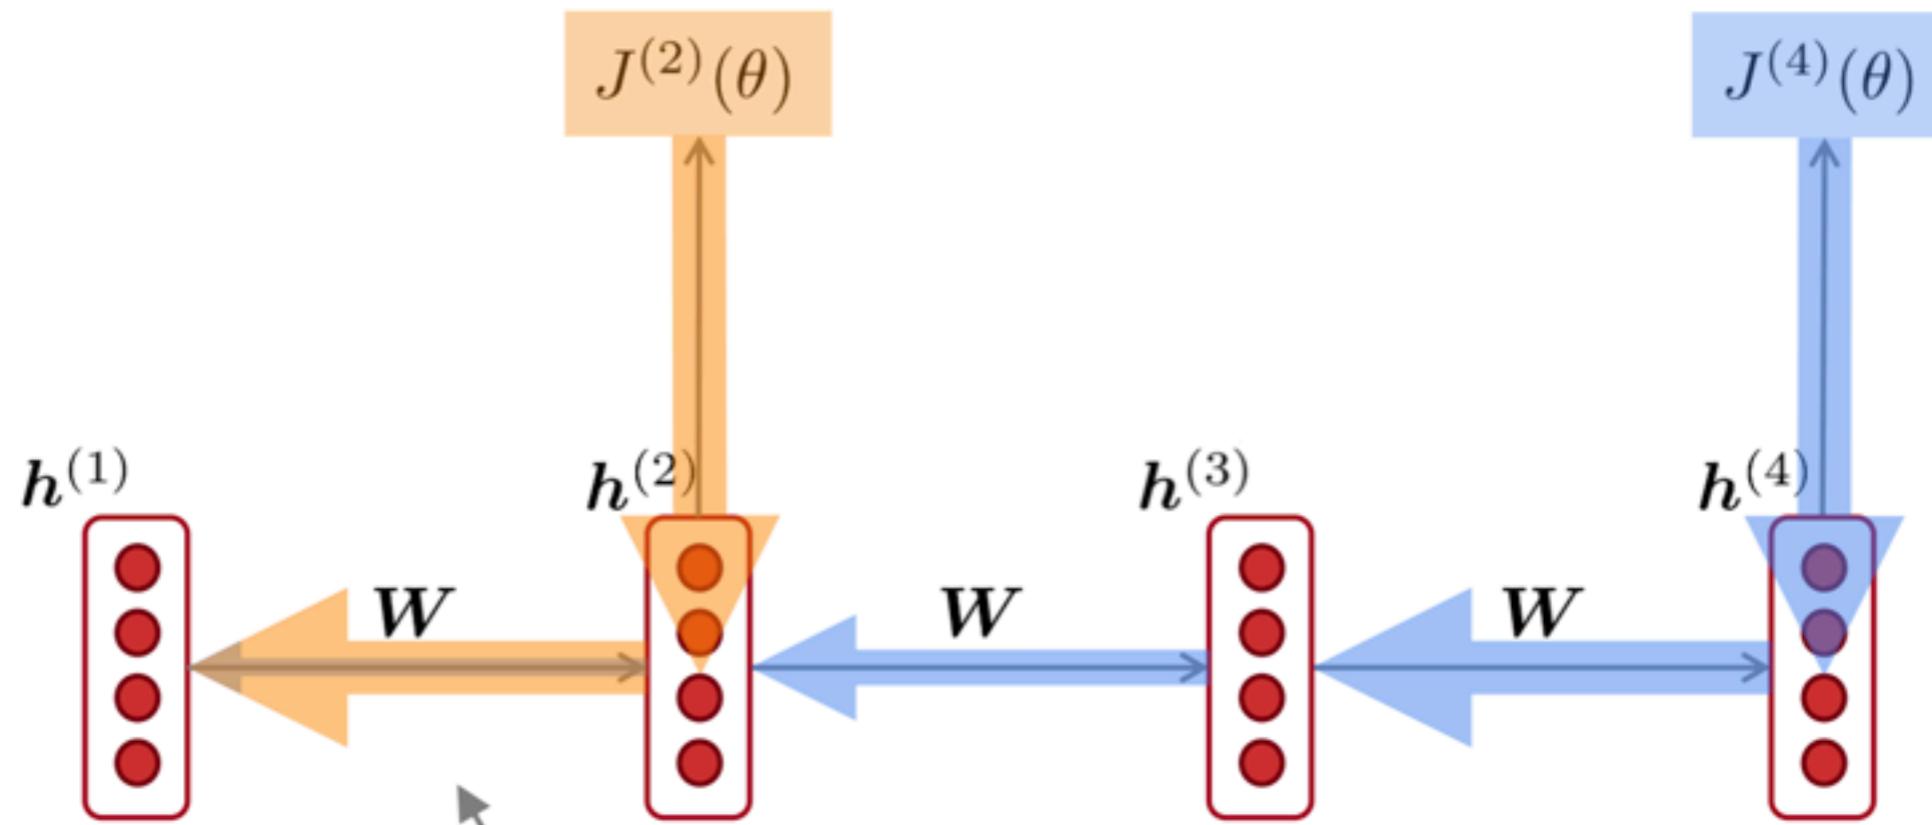


$$\frac{\partial J^{(4)}}{\partial h^{(1)}} = \boxed{\frac{\partial h^{(2)}}{\partial h^{(1)}}} \times \boxed{\frac{\partial h^{(3)}}{\partial h^{(2)}}} \times \boxed{\frac{\partial h^{(4)}}{\partial h^{(3)}}} \times \frac{\partial J^{(4)}}{\partial h^{(4)}}$$

What happens if these are small?

Vanishing gradient problem:  
When these are small, the gradient signal gets smaller and smaller as it backpropagates further

# Why is Vanishing Gradient a problem?



Gradient signal from faraway is lost because it's much smaller than gradient signal from close-by.

So model weights are only updated only with respect to near effects, not long-term effects.

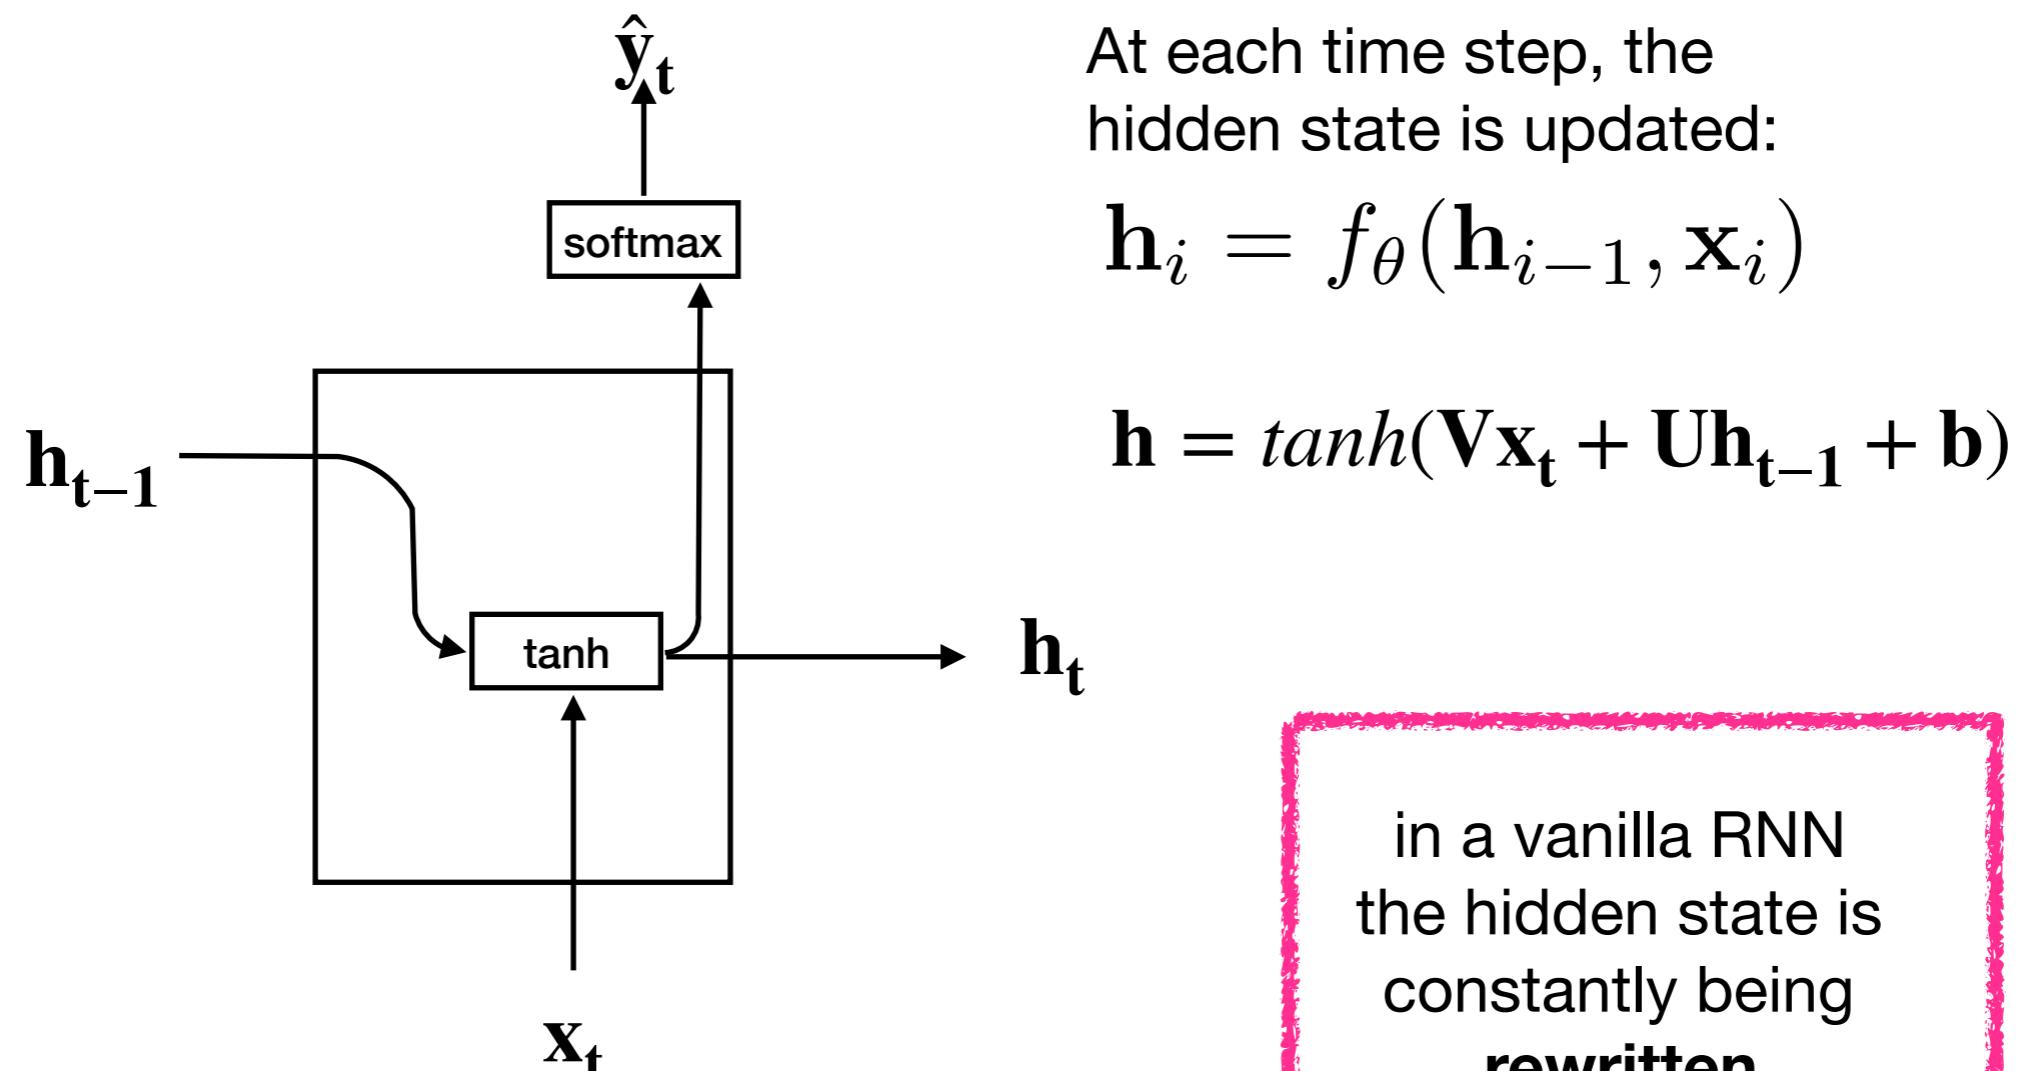
# Effect of vanishing gradient

- ▶ LM task:

“When she bought her laptop, she found that the keyboard layout was Danish. She went back to the shop to ask if the owner of the shop had another keyboard layout. Unfortunately this was not the case, so she kept the \_\_\_\_\_”

- ▶ Model needs to learn a dependency to “laptop”
- ▶ But if the gradients are small, the model won’t learn this
- ▶ RNNs are better at syntactic **recency** [Linzen et al., 2016]

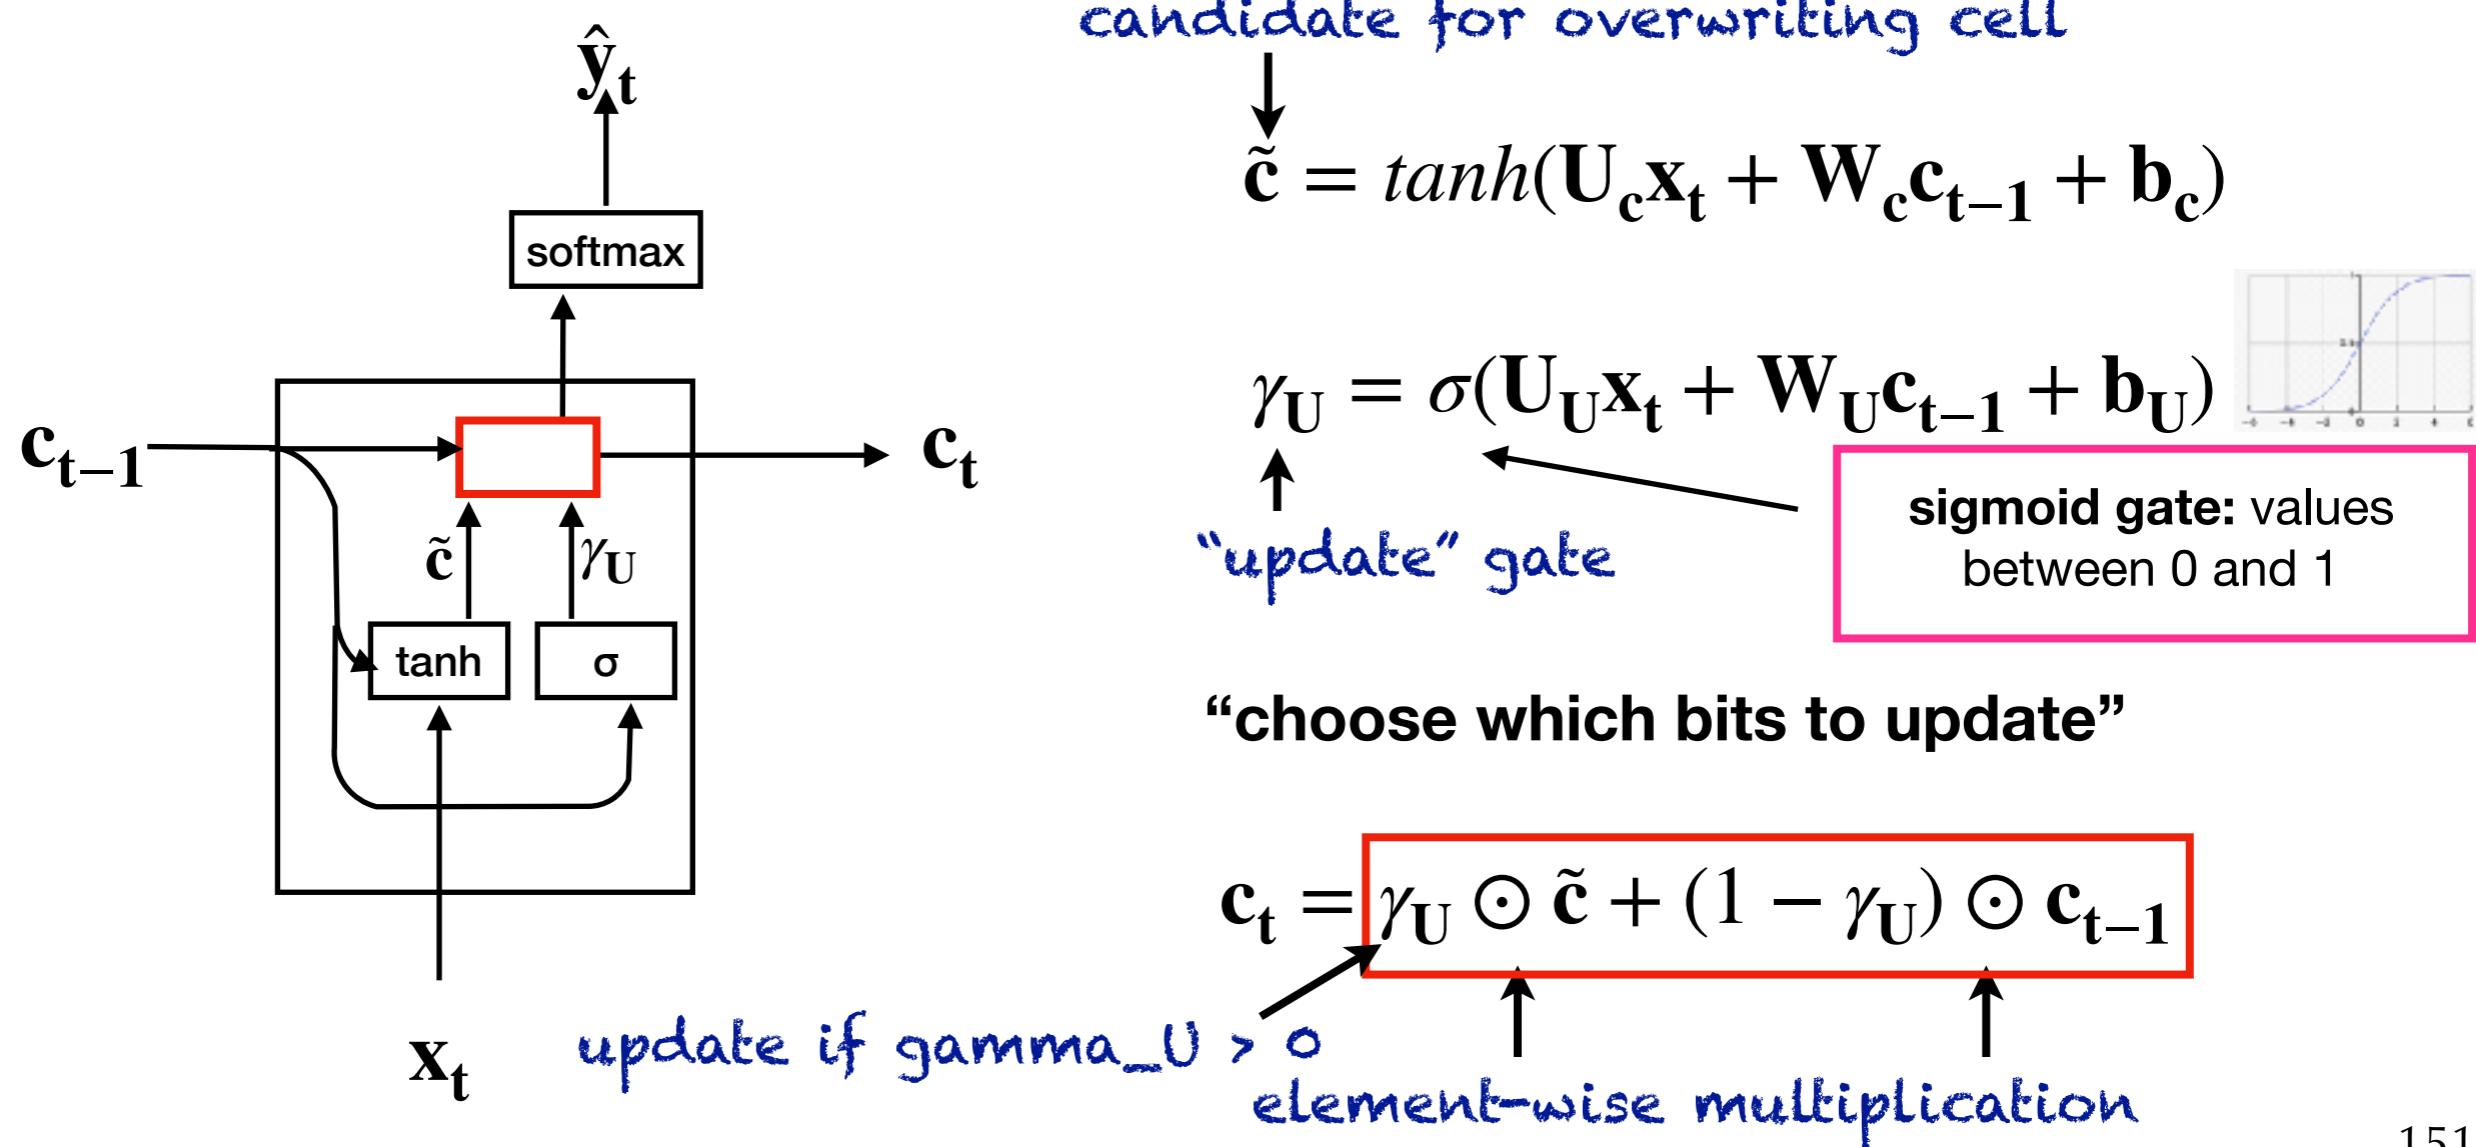
# RNN unit



**Gated RNN architectures:  
RNN flavors with a  
separate memory**

# GRU (Gated recurrent Unit) - simplified

- Cho et al. (2014) - key idea: dynamic memory update  $c$  ( $h=c$ )
- at every step  $t$ , consider overwriting candidate memory  $c$



# GRU (Gated recurrent Unit) - full

- ▶ GRU: creates “adaptive” connections
- ▶ perhaps prune some unnecessary connections adaptively

**Update gate:** controls what parts of the hidden state are updated vs preserved

**Reset gate:** controls what parts of the previous hidden state are used to compute new content

$$\begin{aligned}\tilde{\mathbf{c}} &= \tanh(\mathbf{U}_c \mathbf{x}_t + \mathbf{W}_c (\gamma_R \odot \mathbf{c}_{t-1}) + \mathbf{b}_c) \\ \gamma_U &= \sigma(\mathbf{U}_U \mathbf{x}_t + \mathbf{W}_U \mathbf{c}_{t-1} + \mathbf{b}_U) \\ \gamma_R &= \sigma(\mathbf{U}_R \mathbf{x}_t + \mathbf{W}_R \mathbf{c}_{t-1} + \mathbf{b}_R) \\ \mathbf{c}_t &= \gamma_U \odot \tilde{\mathbf{c}} + (1 - \gamma_U) \odot \mathbf{c}_{t-1}\end{aligned}$$

all vectors of same size

**How does this help the vanishing gradient problem?**  
GRUs make it easier to retain info long-term (e.g. by not updating bits)

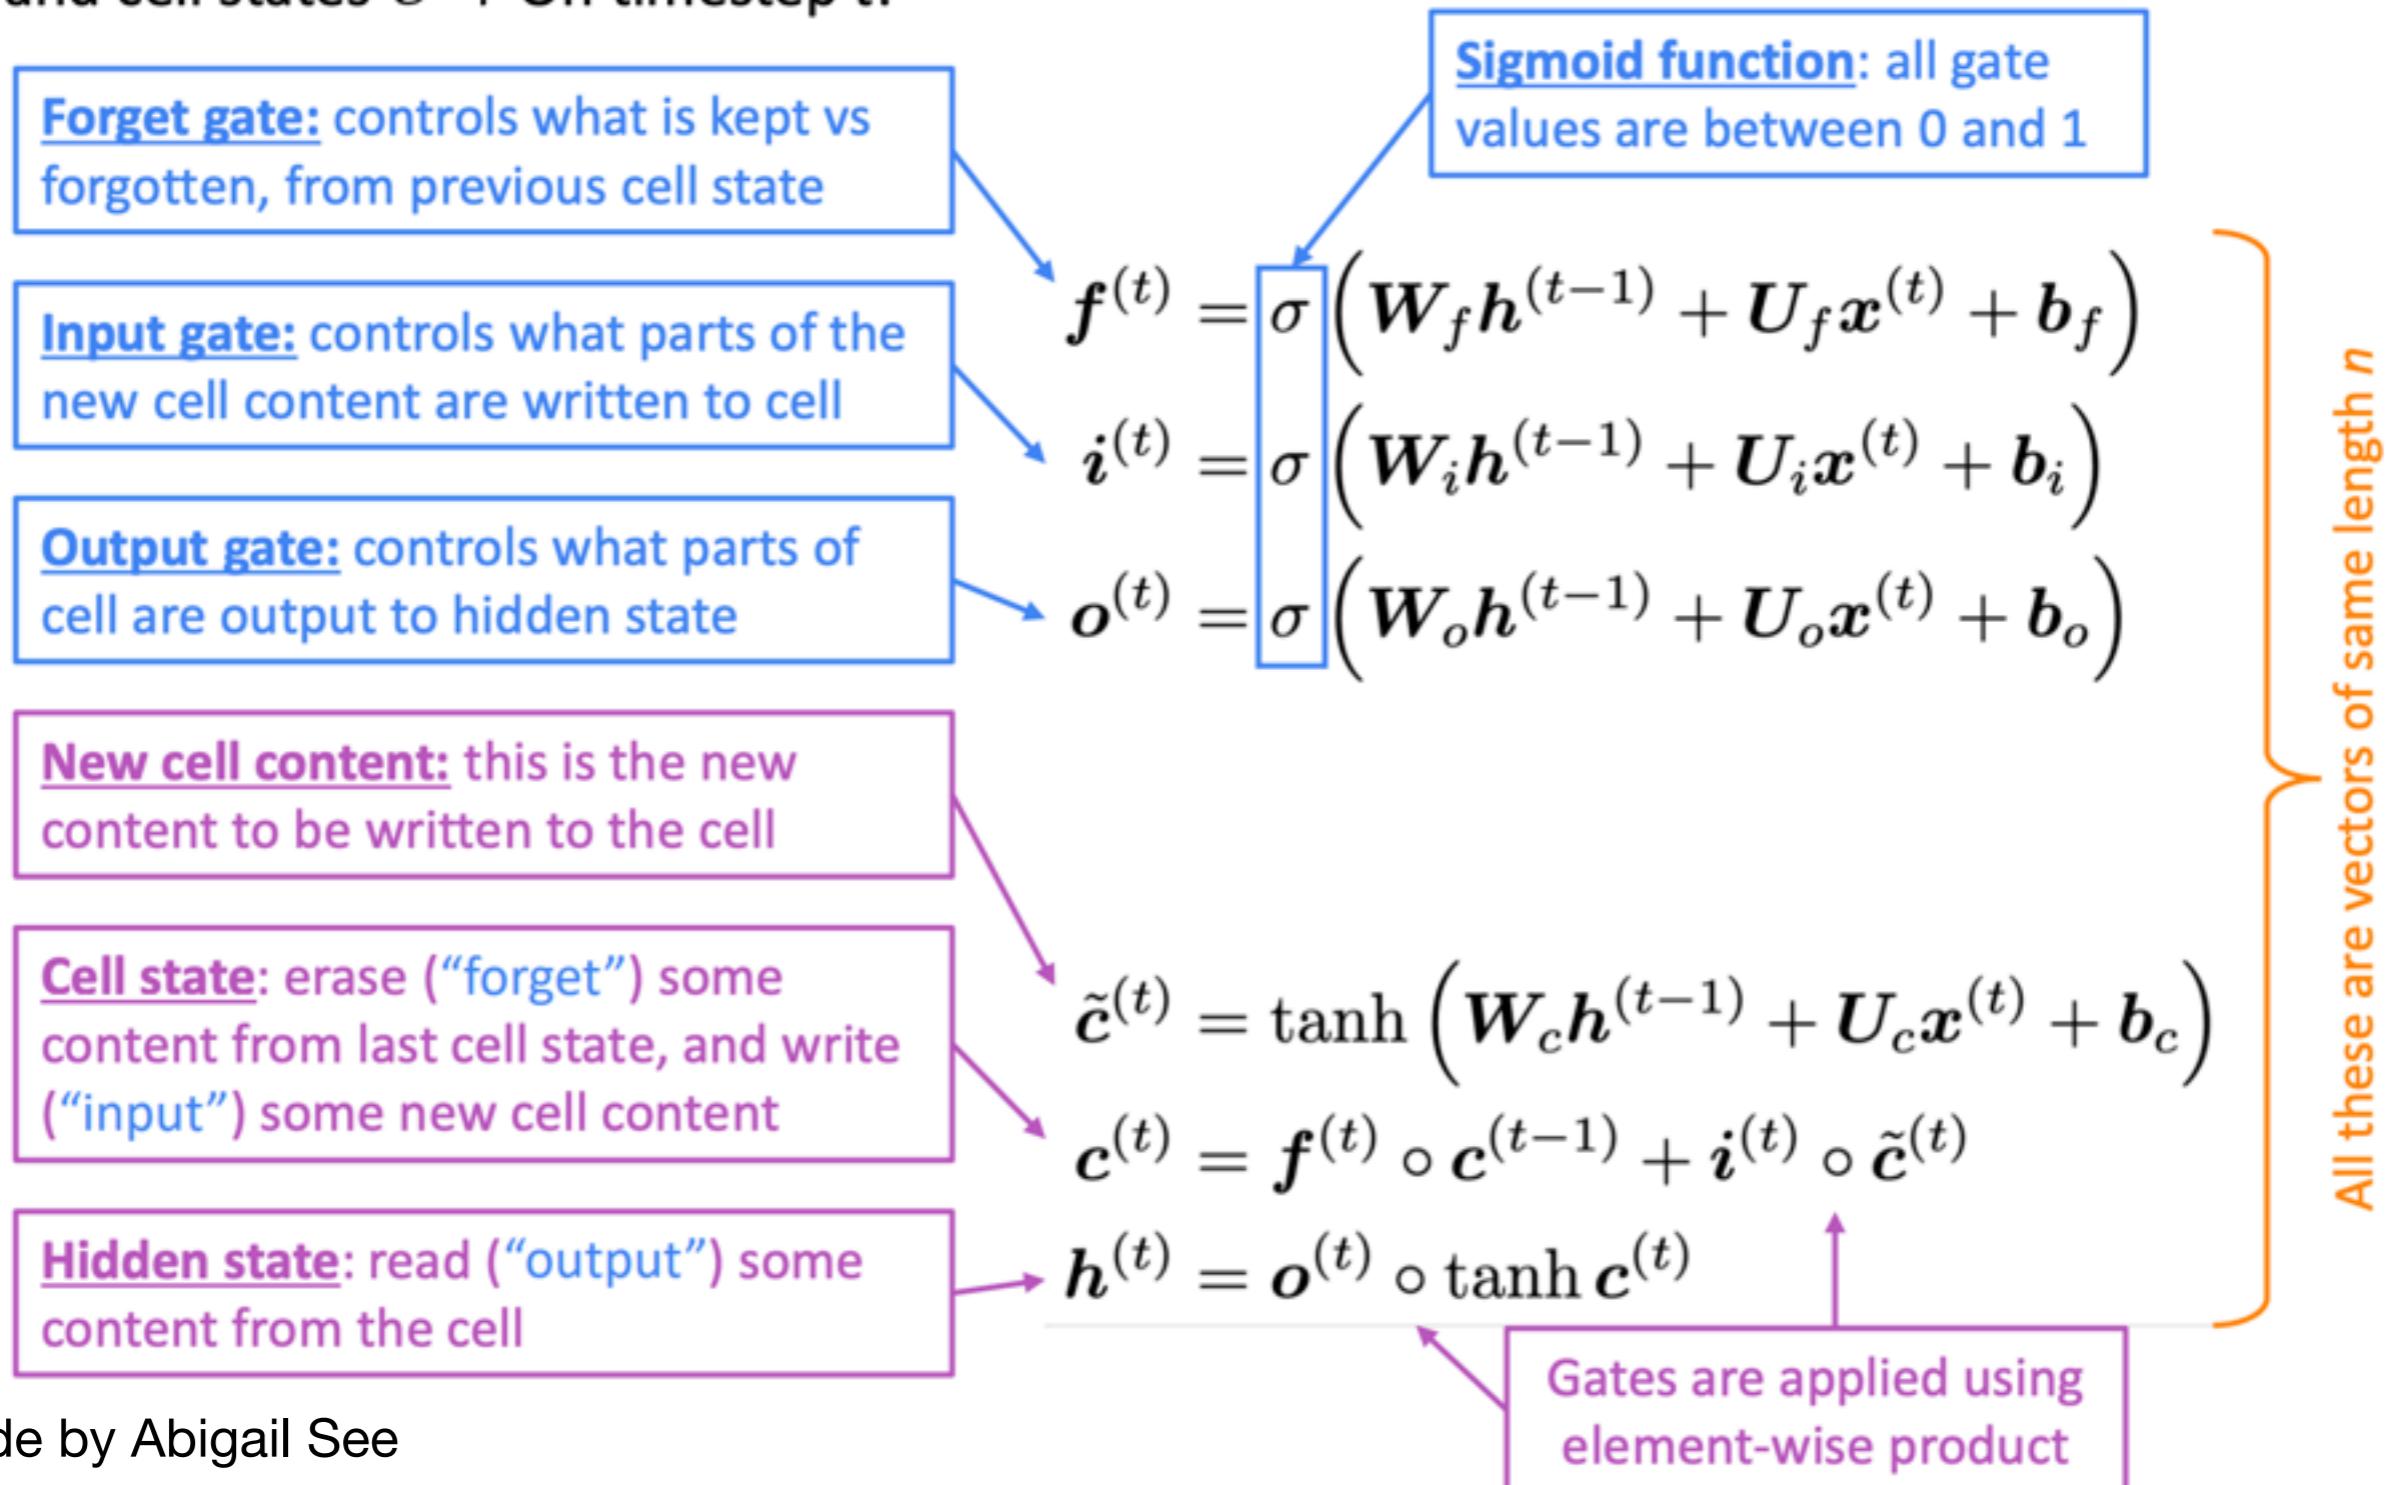
Slide inspired by Abigail See

# LSTM (Long-Short Term Memory)

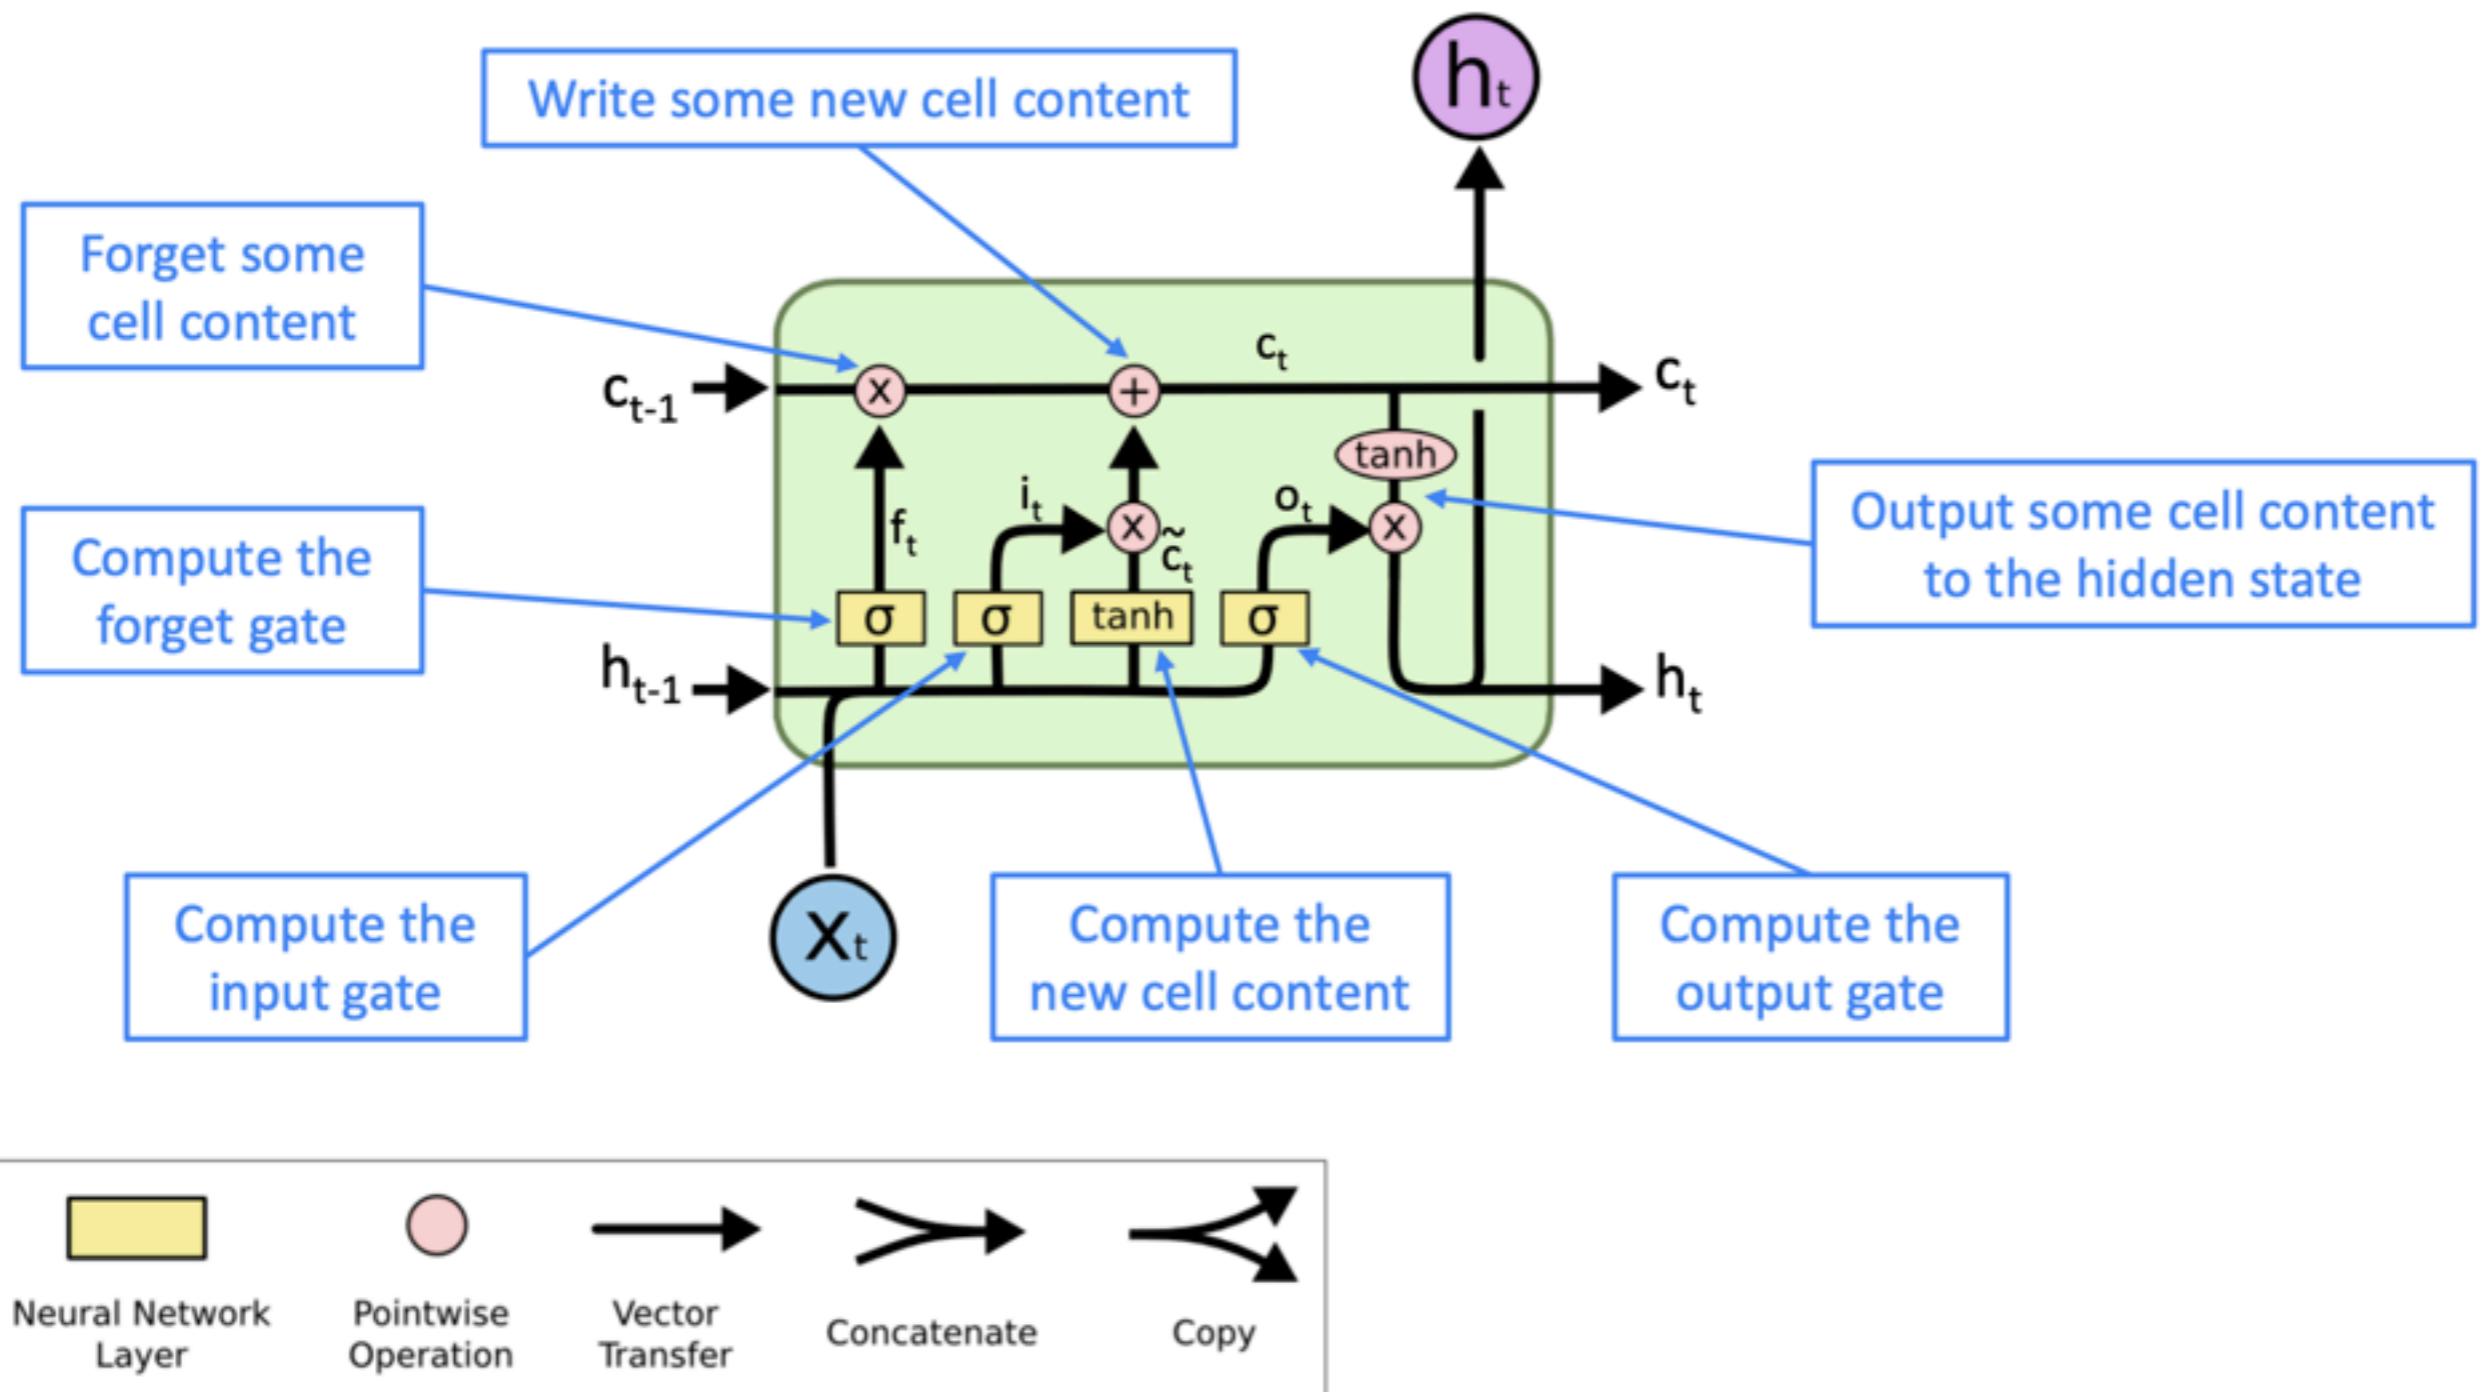
- ▶ Introduced by Hochreiter & Schmidhuber 1997
- ▶ Separate memory cell  $c$  and hidden state  $h$
- ▶ Three gates:
  - ▶ **forget gate:** controls what is kept and forgotten from previous cell state
  - ▶ **input gate:** controls what part of the new cell content are written to the cell
  - ▶ **output gate:** controls what part of the new cell content are written to the hidden state

# LSTM (Long-Short Term Memory)

We have a sequence of inputs  $x^{(t)}$ , and we will compute a sequence of hidden states  $h^{(t)}$  and cell states  $c^{(t)}$ . On timestep  $t$ :



# LSTM (Long-Short Term Memory)

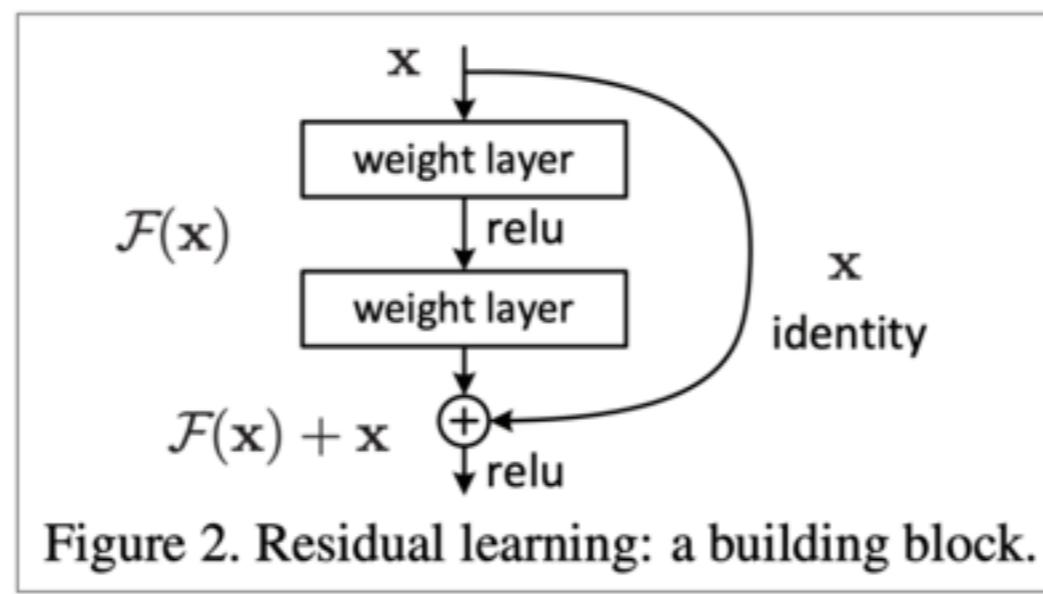


# GRU vs LSTM

- ▶ GRU is more efficient to learn (fewer parameters)
- ▶ Which is better?
  - ▶ No conclusive evidence that one is always superior to the other
- ▶ LSTM is typically a good starting choice
- ▶ Suggestion: switch to GRU if you want a more efficient model

# Residual connections

- ▶ Is the vanishing gradient problem specific to RNNs?
  - ▶ No! Also for deep FFNN and ConvNets
- ▶ **Solution:** add direct “skip” connections (ResNet, residual connections) - proposed by He et al., (2015)
  - ▶ i.e. add  $F(x) + x$ , instead of  $F(x)$
  - ▶ allows for training deeper models

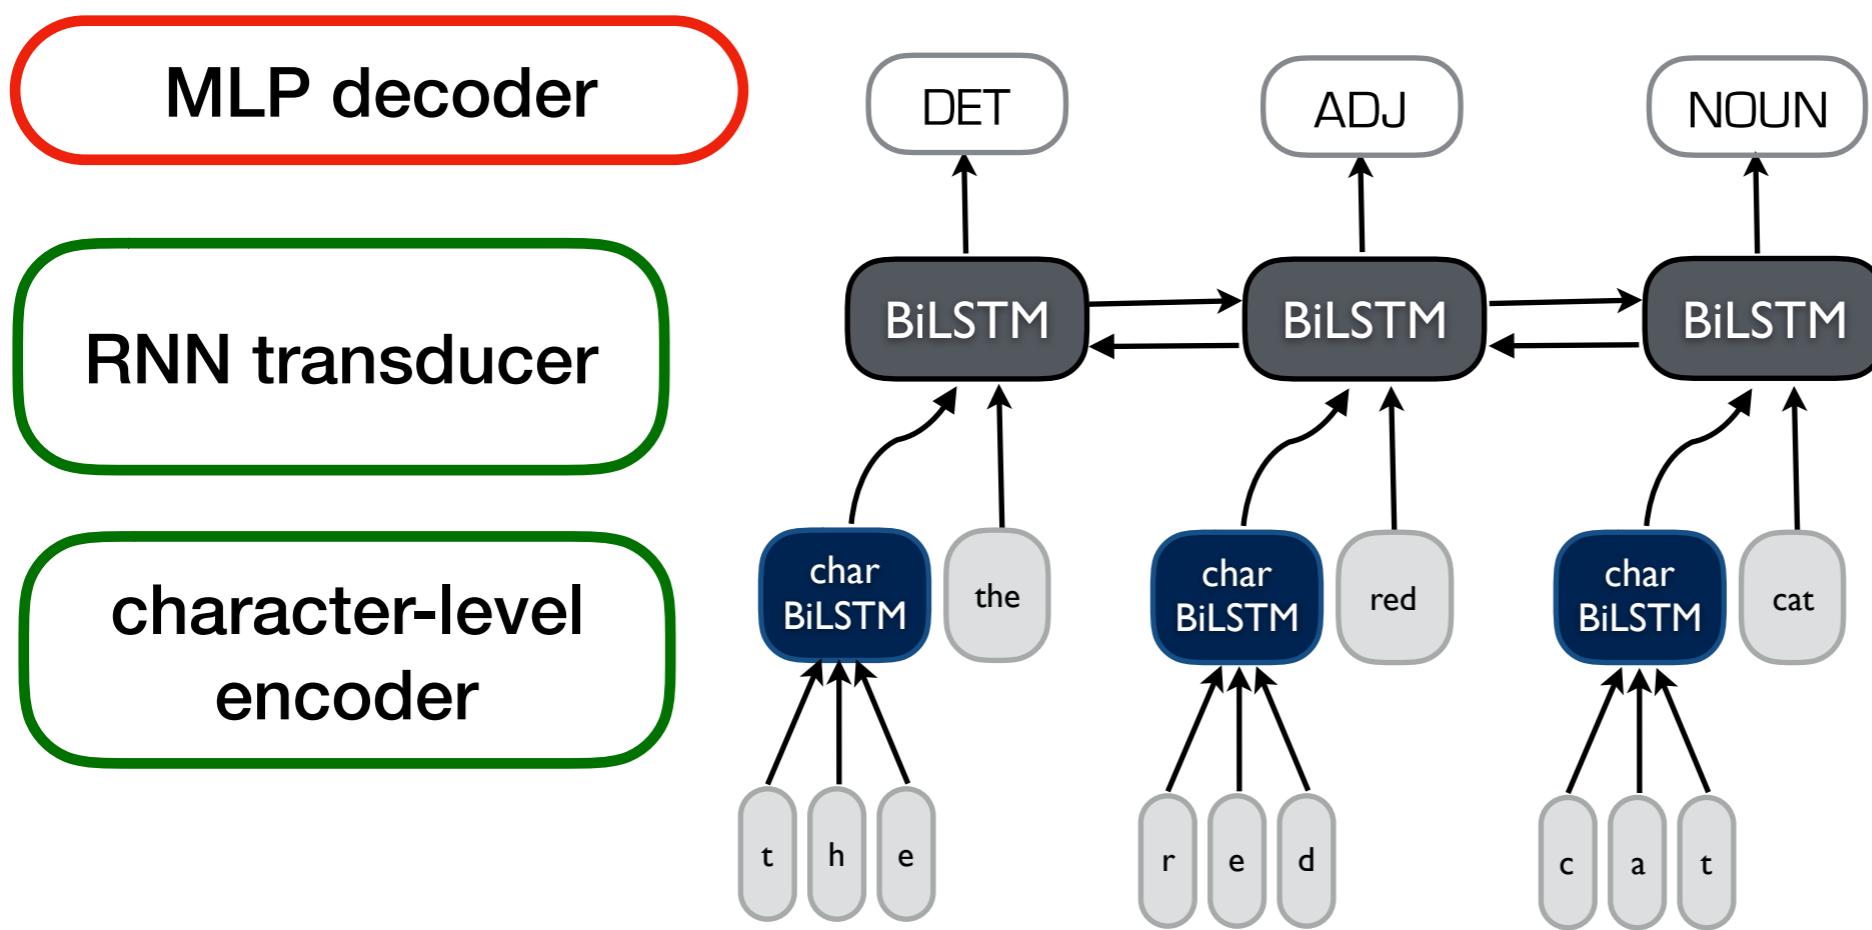


**LSTMs are everywhere...**

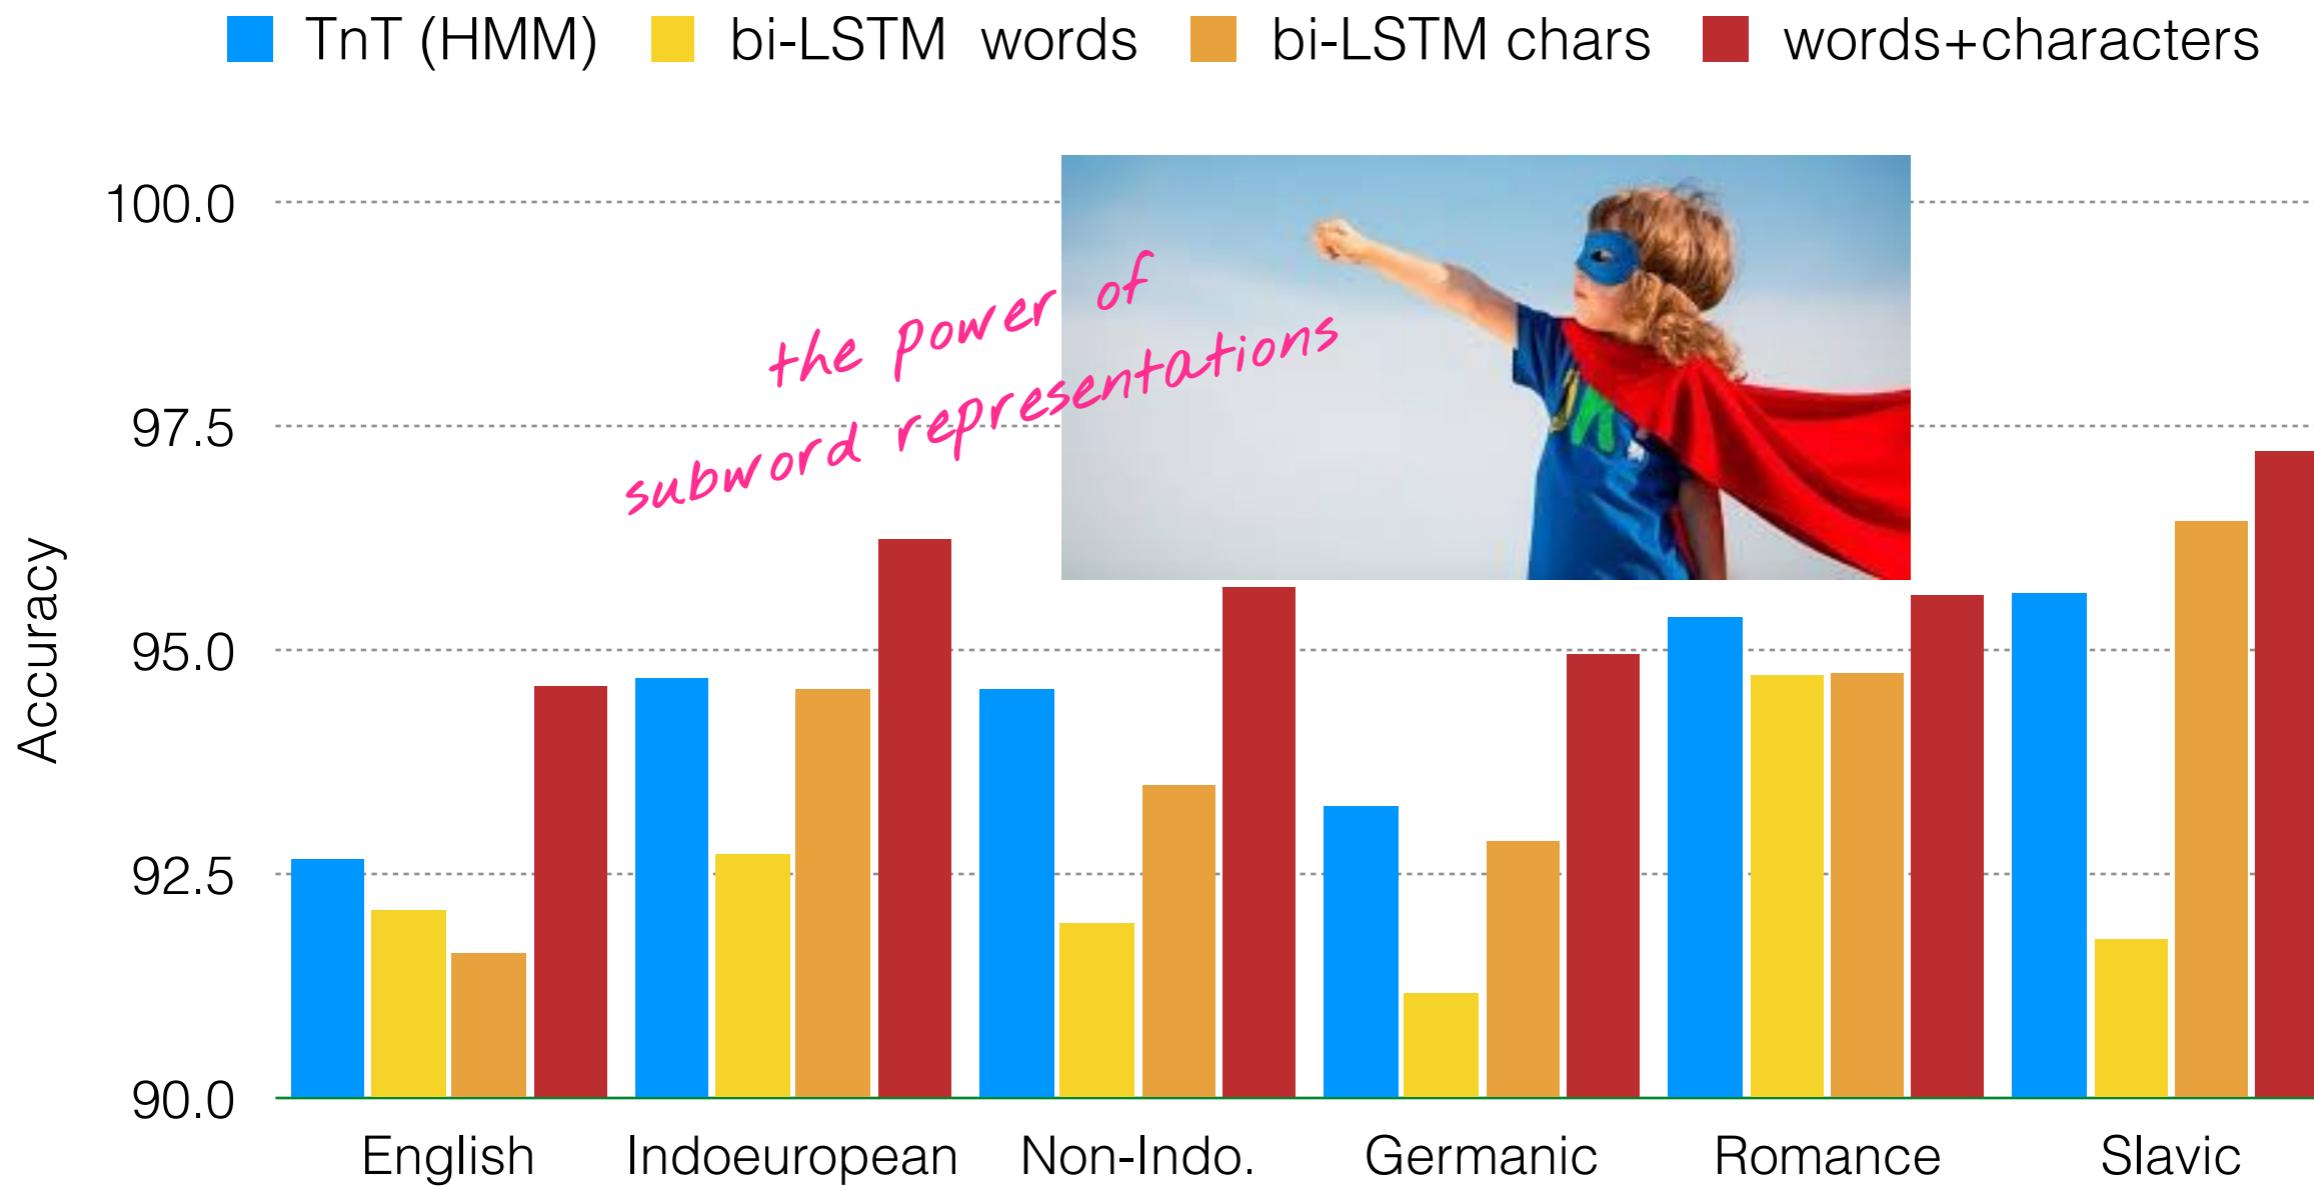
**Let's look briefly at different  
decoders via examples**

# A very common POS tagger

- ▶ Use **bi-LSTM transducer** with a lower-level **bi-LSTM encoder for characters** and a softmax decoder

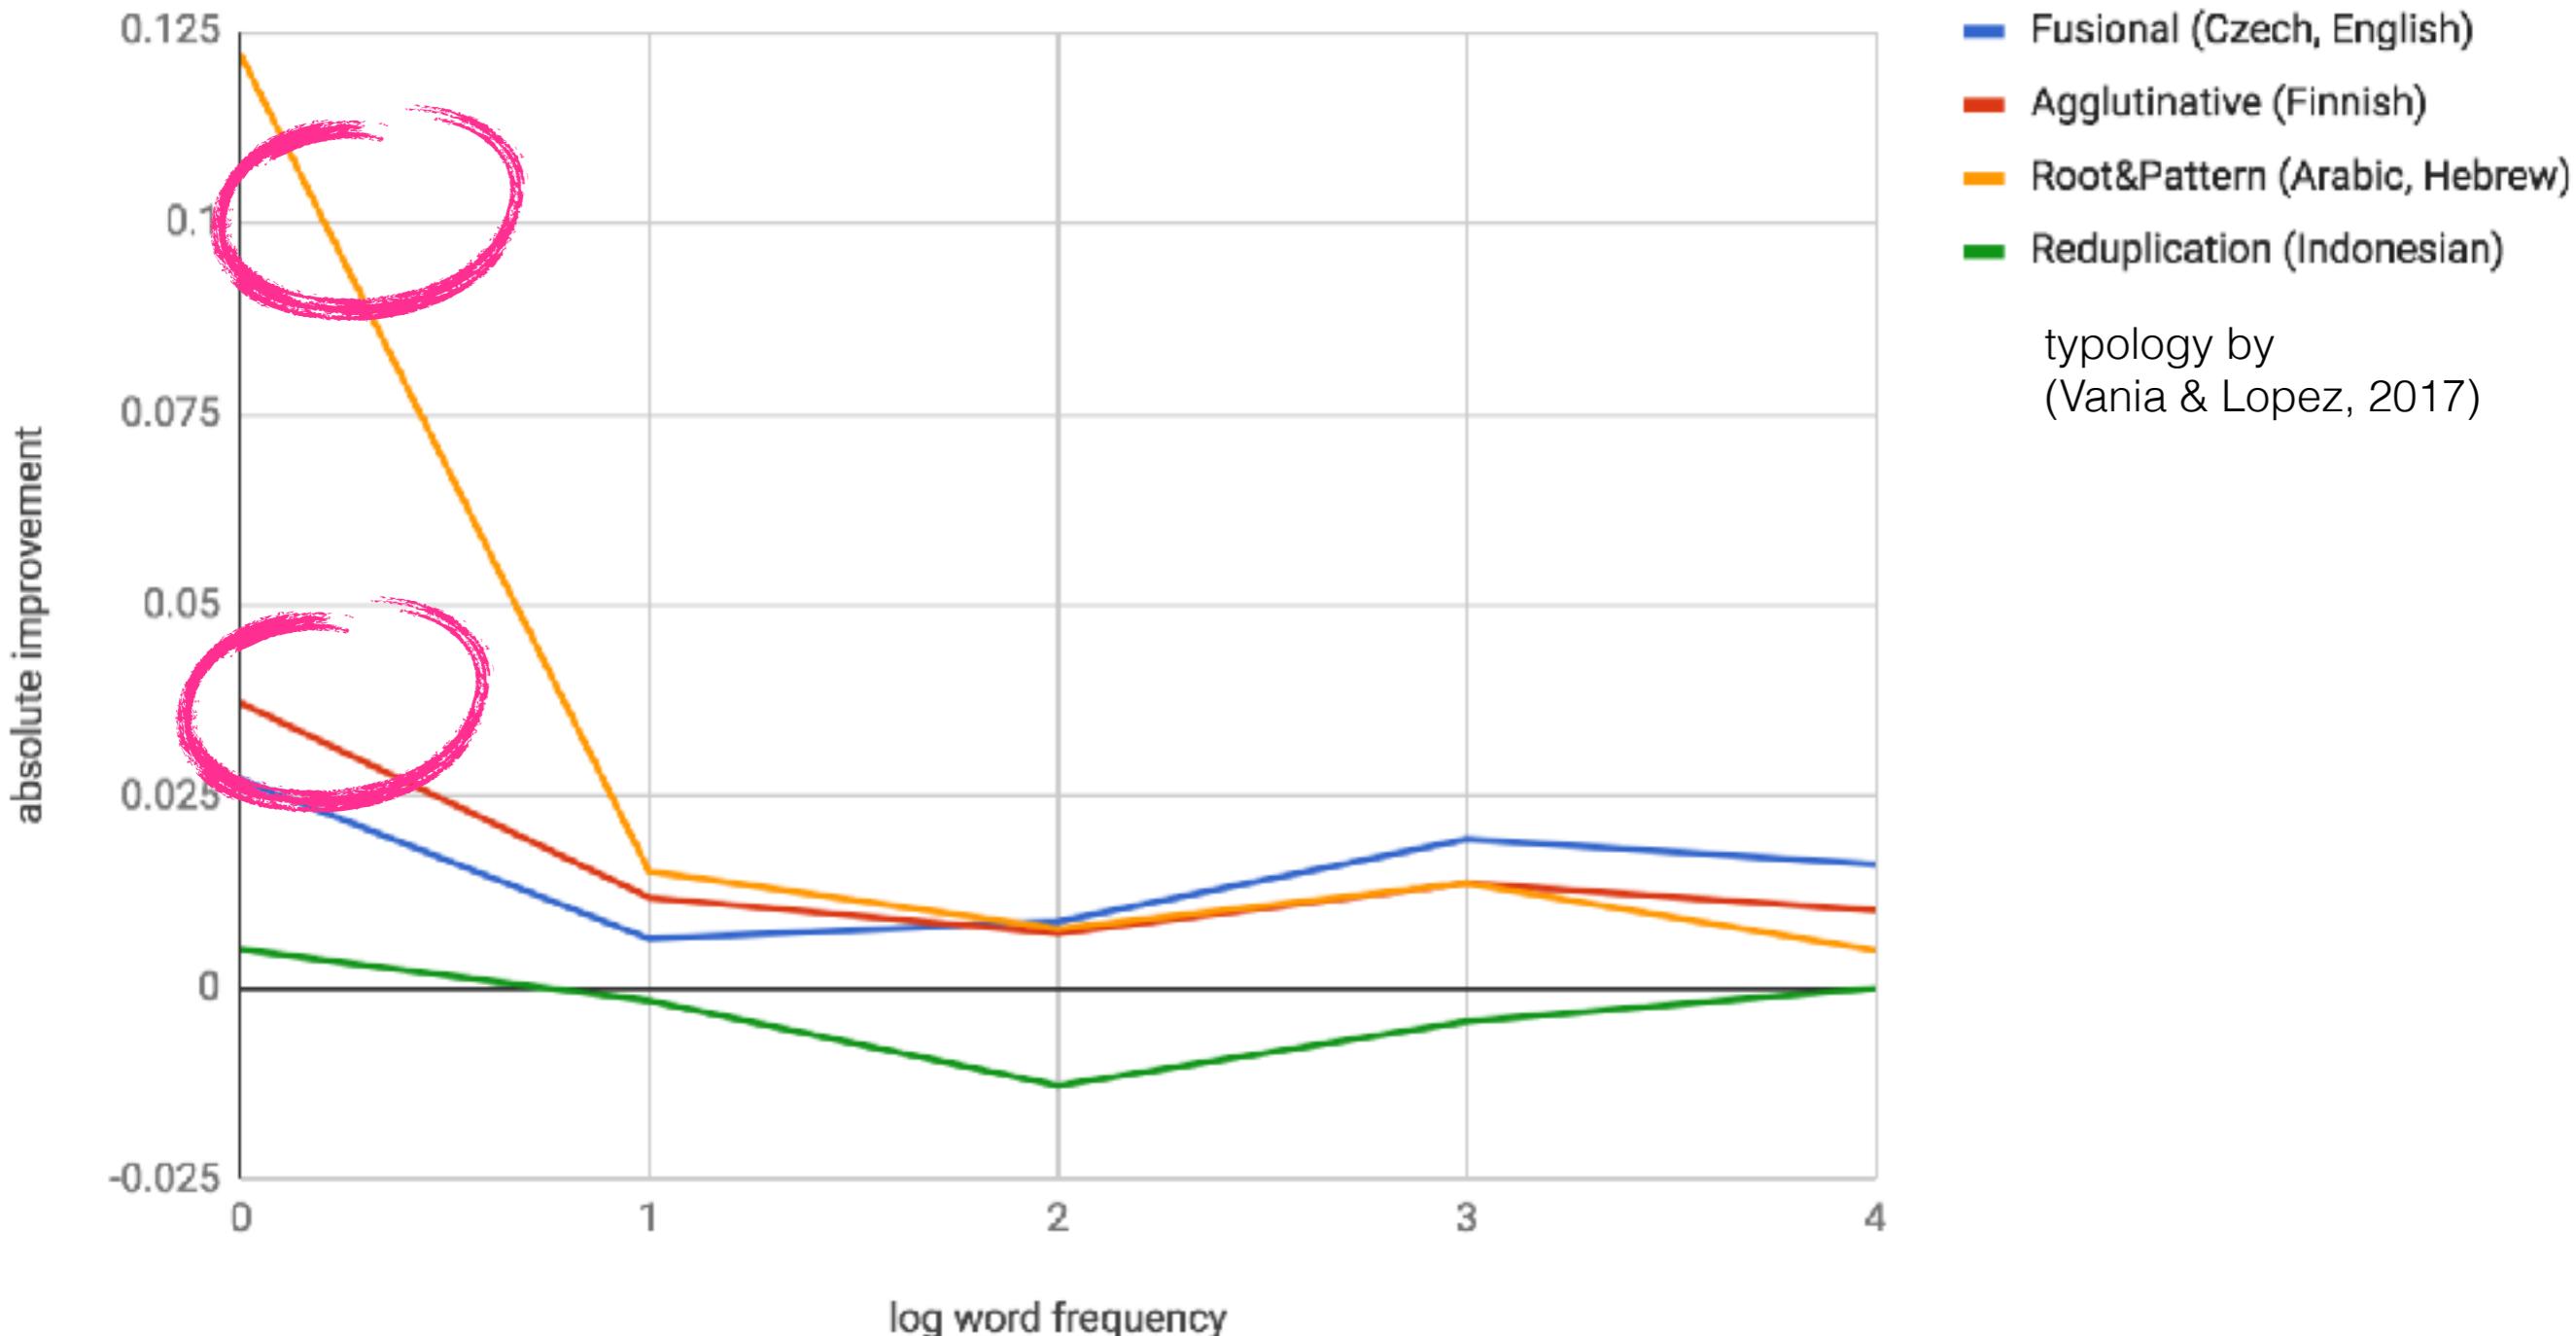


# POS tagging on many languages



17 coarse POS tags,  
experiments over 22 languages of UD 1.2, (Plank et al., 2016)

# A closer look at non-IE languages



— Fusional (Czech, English)  
— Agglutinative (Finnish)  
— Root&Pattern (Arabic, Hebrew)  
— Reduplication (Indonesian)

typology by  
(Vania & Lopez, 2017)

# Named Entity Recognition (NER)

- ▶ Example: Bill B-PER lives in Athens B-LOC
- ▶ (Huang et al., 2015): from RNN to bidirectional LSTM-CRF

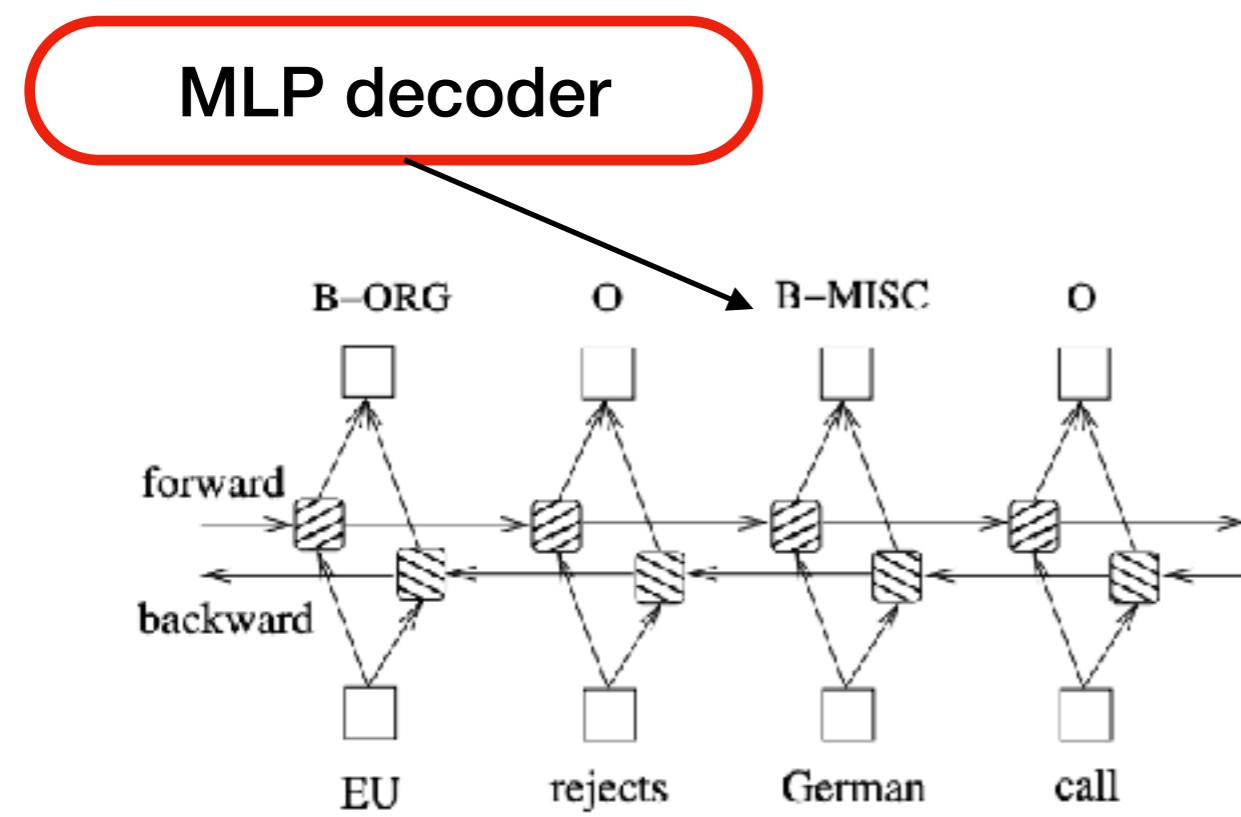


Figure 4: A bidirectional LSTM network.

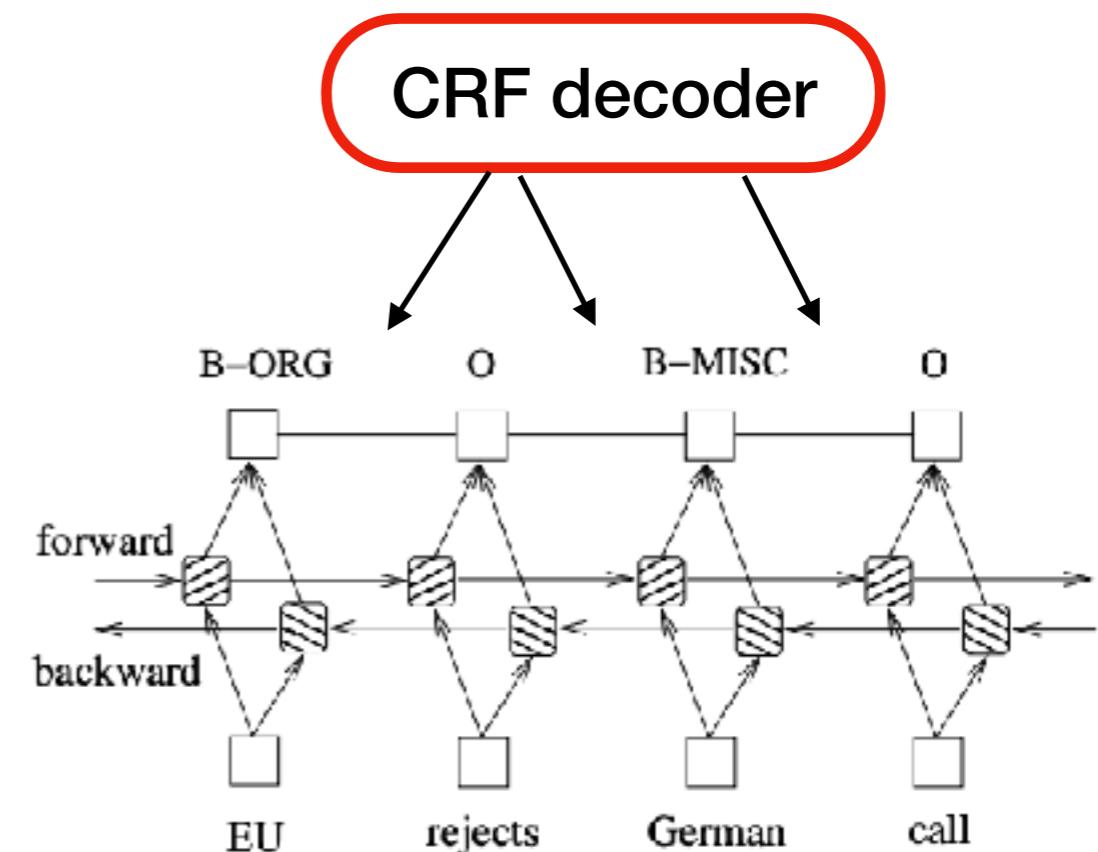
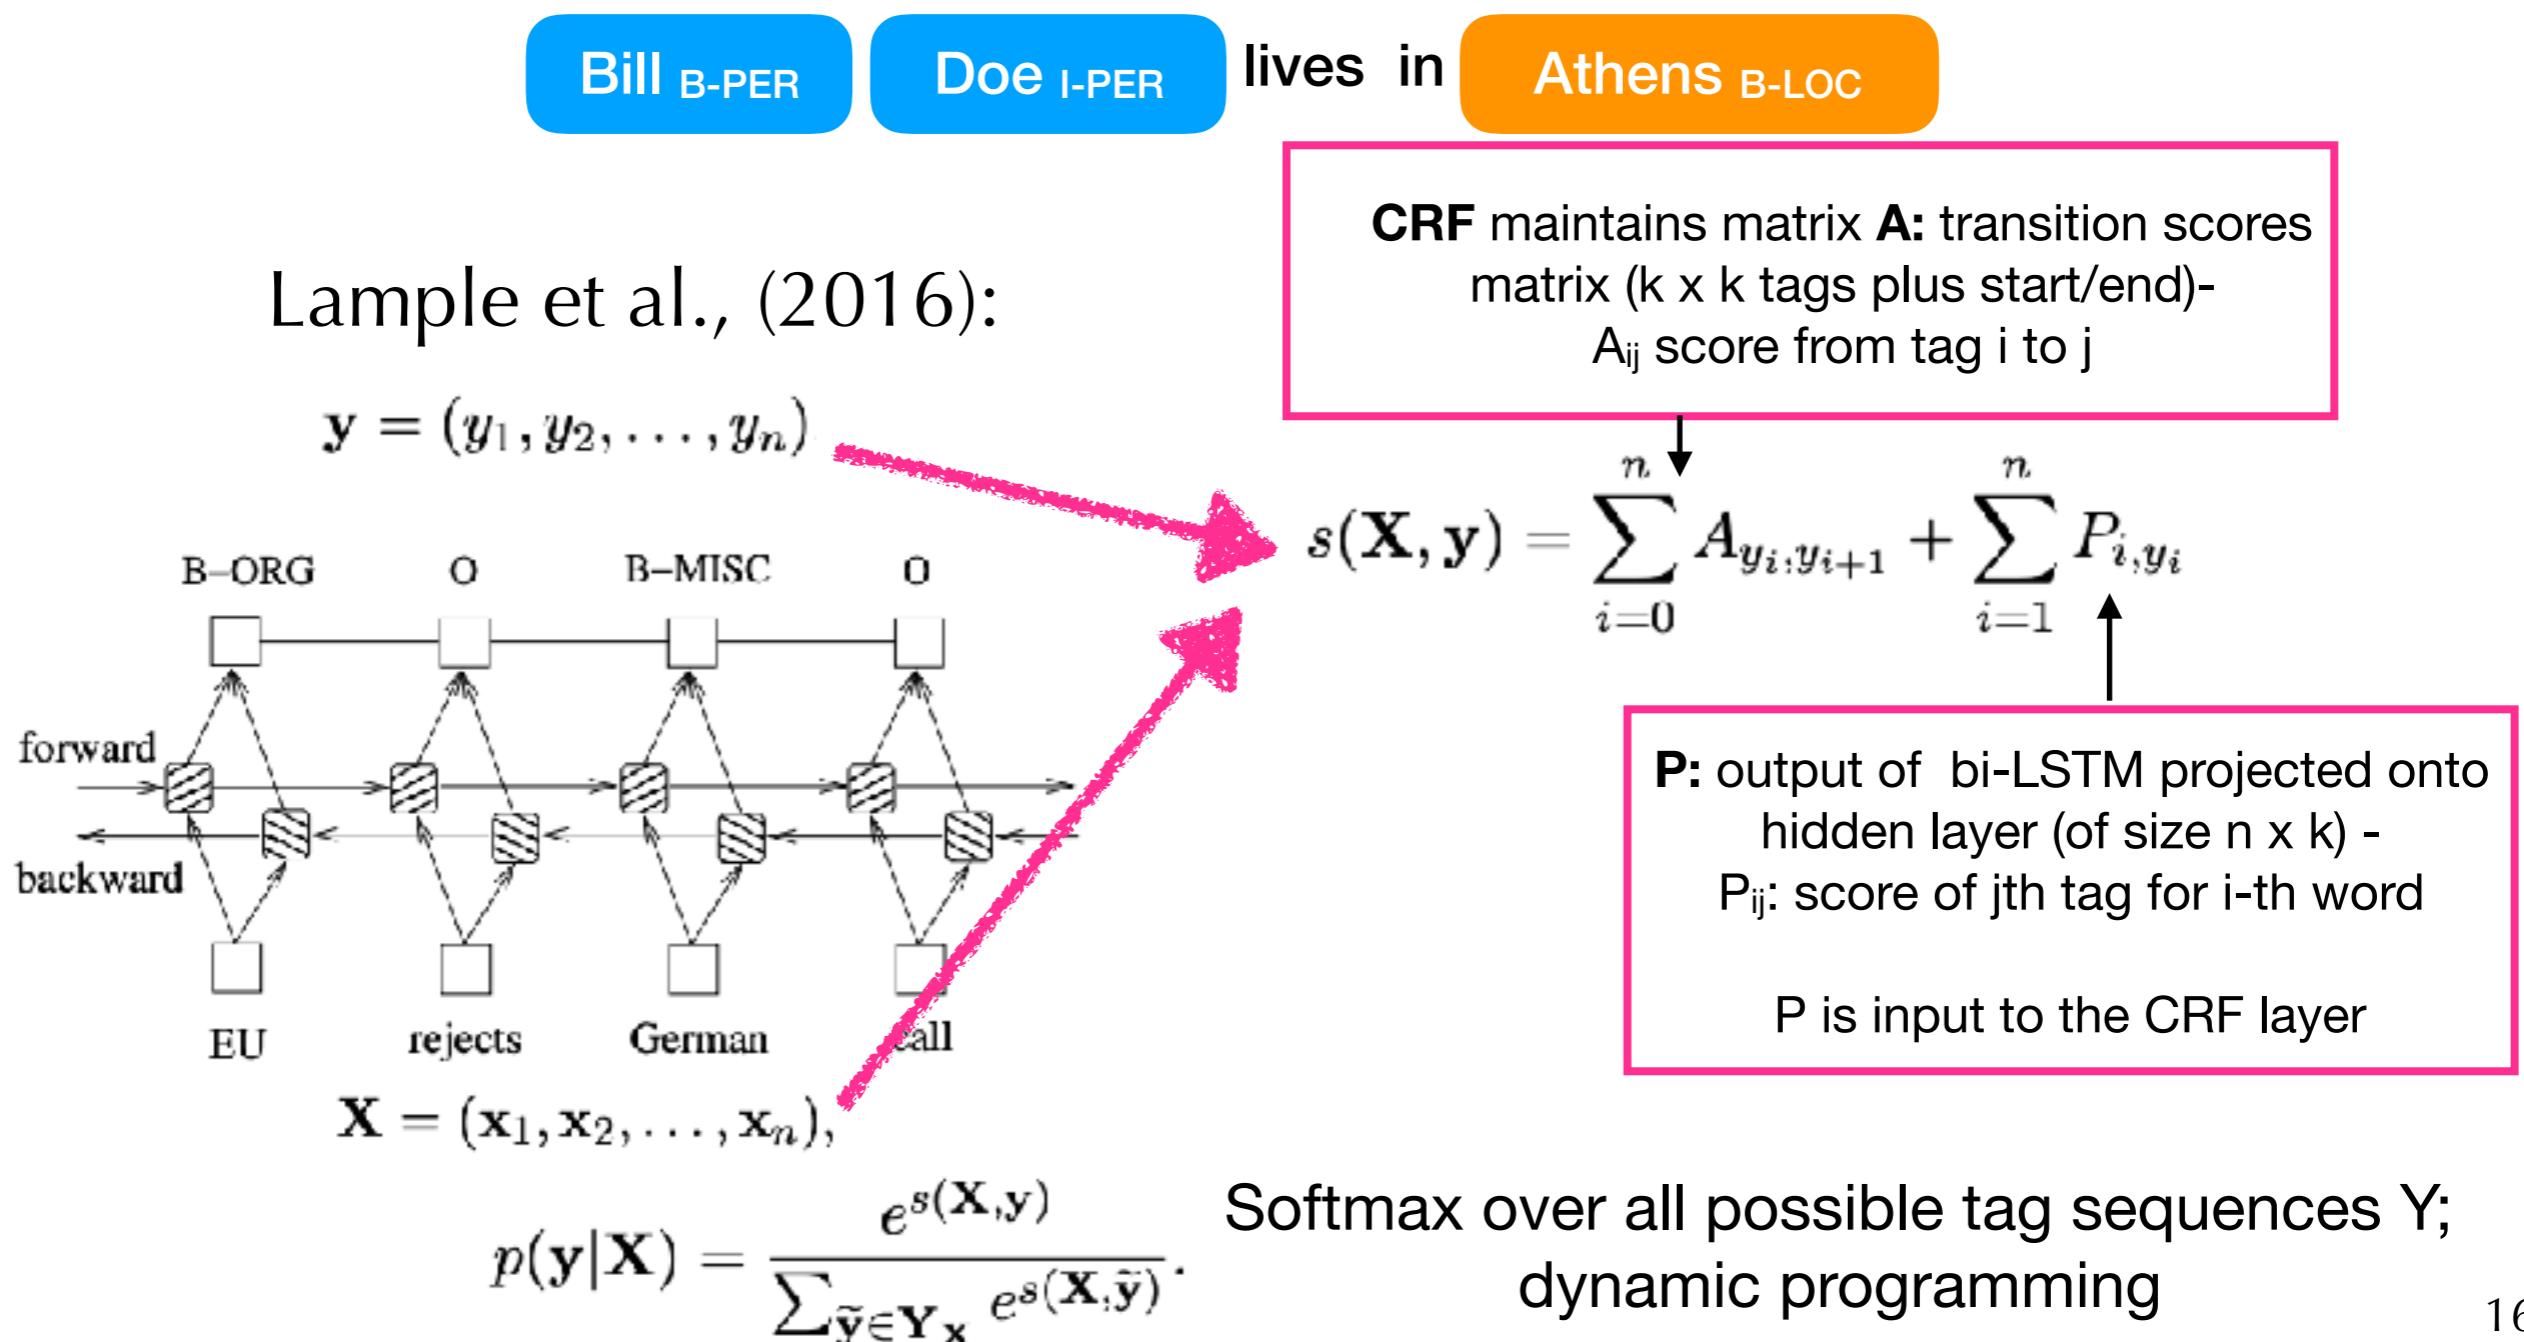


Figure 7: A BI-LSTM-CRF model.

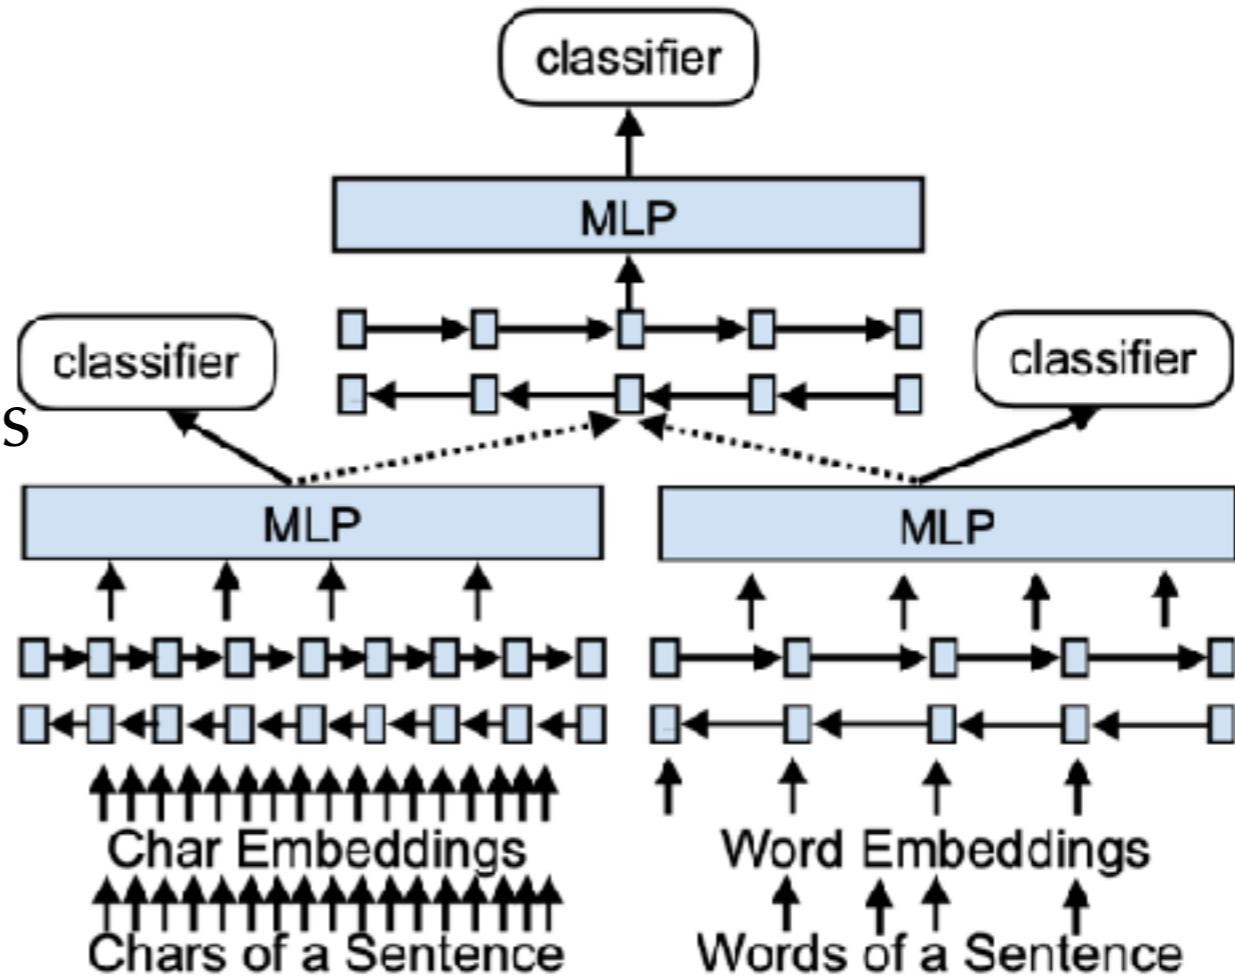
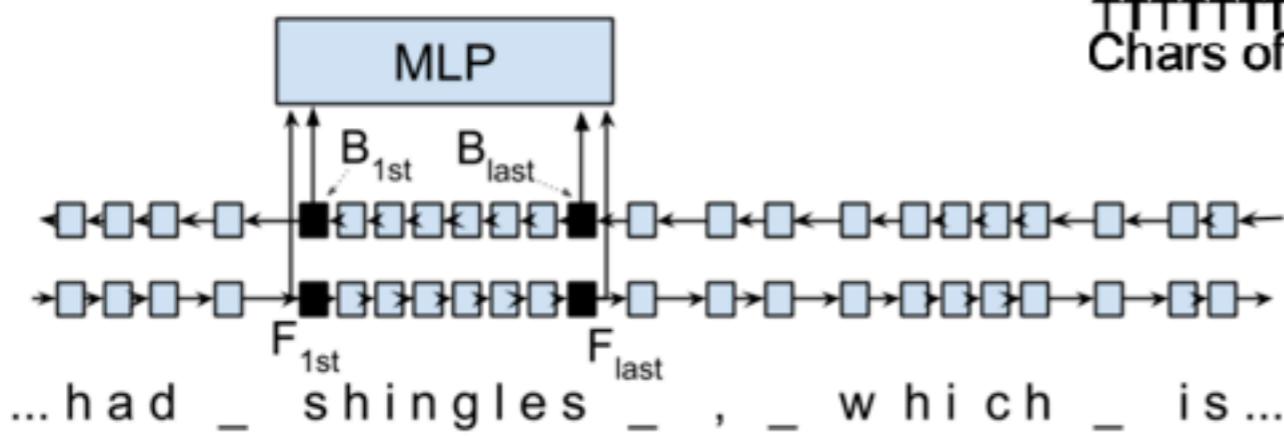
# CRF decoder

- Stronger sequential nature (e.g., I-PER after B-PER)



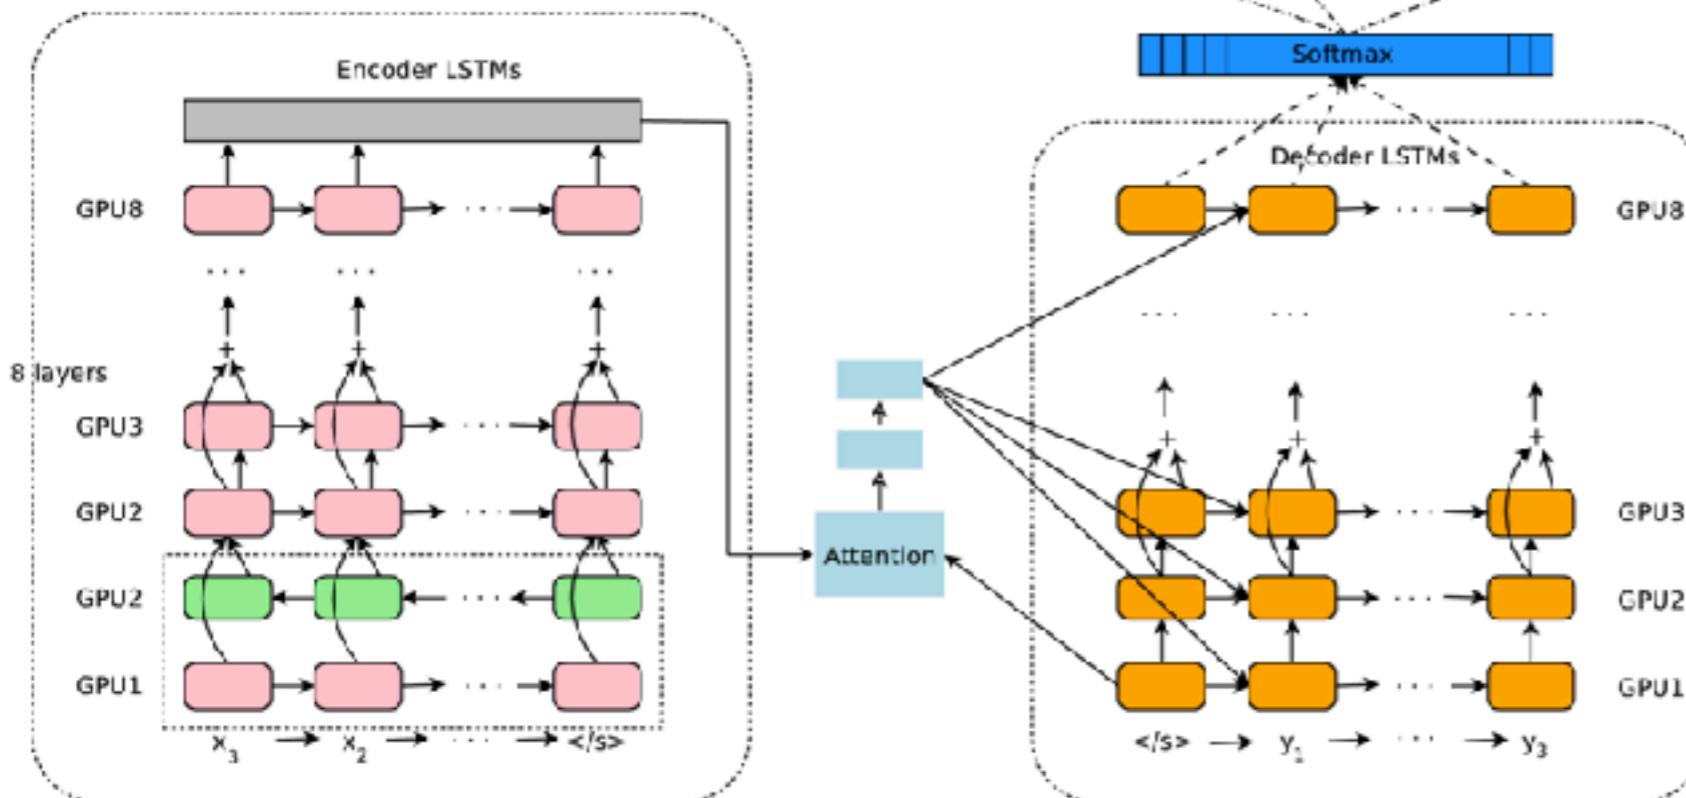
# Meta-BiLSTM

- ▶ The model so far is restricted to subwords to within words
- ▶ Recent SOTA model
  - ▶ context-sensitive character and word representations
  - ▶ Models trained synchronously & then combined



(Bohnet et al., 2018)

# Google's Neural MT System (Wu et al., 2016)



- ▶ **deep bidirectional LSTM** (stacked) with **residual connections** and **attention**
- ▶ huge improvements in MT quality

<https://arxiv.org/pdf/1609.08144.pdf>

- ▶ Now (**2019** onwards): other approaches have become dominant for certain NLP tasks (e.g. the **Transformer**) - see more on Monday (Arianna)

# Interim summary

- ▶ RNNs:

- ▶ Two fancy variants: **LSTM and GRU** to address the vanishing gradient problem
- ▶ Deep RNNs (stacking)
- ▶ Residual connections
- ▶ Two more concepts to cover:
  - ▶ beyond static word embeddings
  - ▶ gluing it all together: attention!

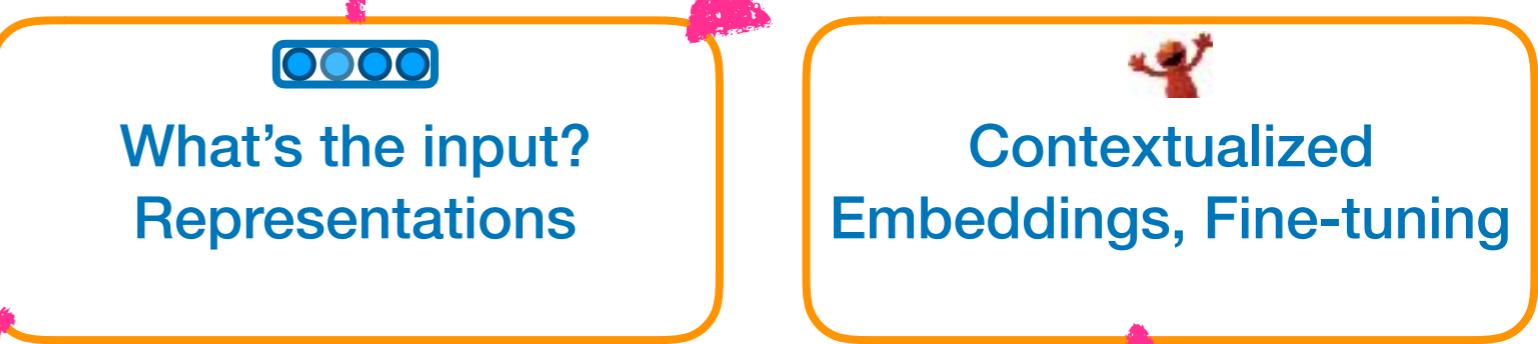


# Overview

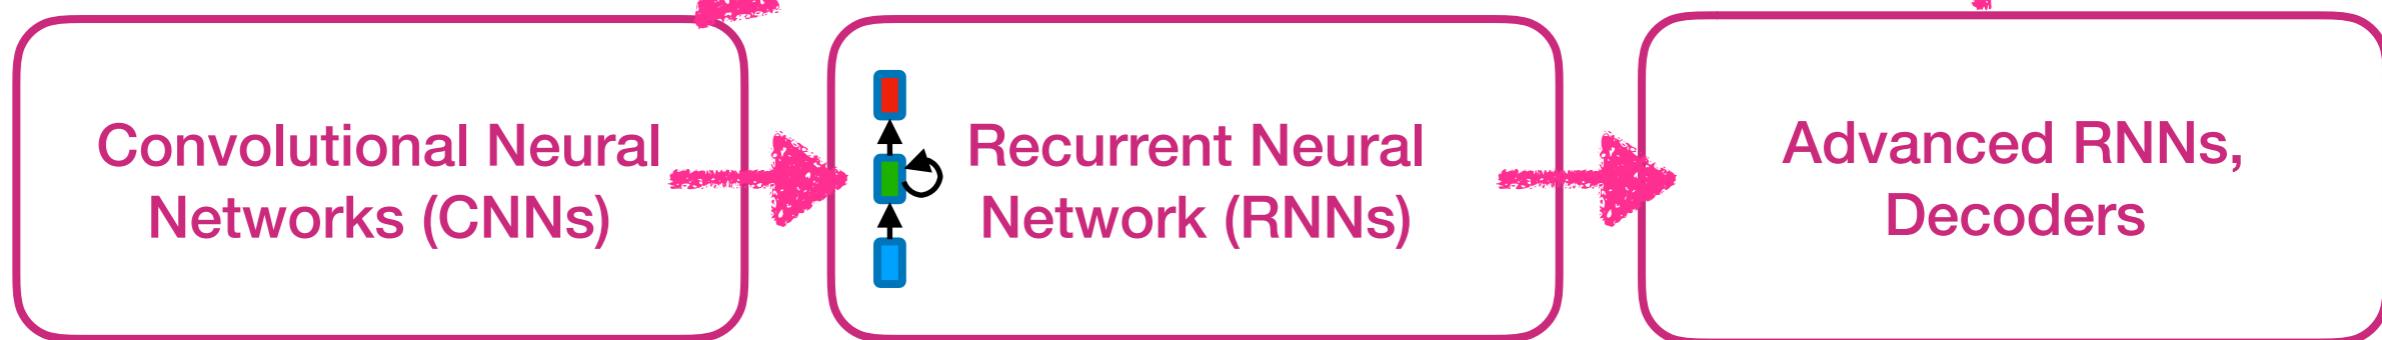
*foundations*



*representations*



*beyond FFNNs*



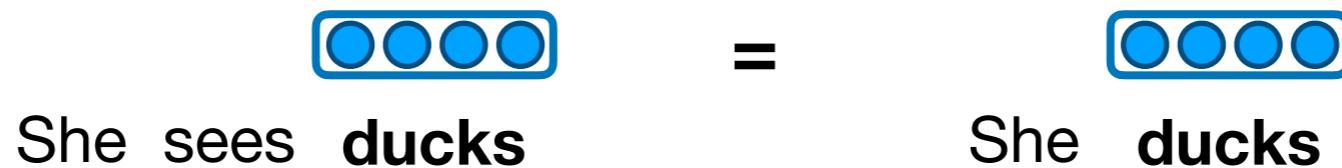
# Traditional (“static”) word embeddings

compress all contexts into a single vector

# Contextualized word embeddings

# Representing a word as vector so far

- ▶ Problem:
  - ▶ It is a **type-based** representation: always the same vector for a word regardless of its context (e.g. ‘ducks’)
  - ▶ Polysemy is not handled



- ▶ Solution: Contextualized embeddings
  - ▶ Learn a vector that depends on the context

# Language Models to the rescue!

Deep contextualized word representations

Matthew E. Peters<sup>†</sup>, Mark Neumann<sup>†</sup>, Mohit Iyyer<sup>†</sup>, Matt Gardner<sup>†</sup>,  
{matthewp, markn, mohiti, mattg}@allenai.org

Christopher Clark\*, Kenton Lee\*, Luke Zettlemoyer<sup>†\*</sup>  
{csquared, kentonl, lsz}@cs.washington.edu

\*Paul G. Allen School of Computer Science & Engineering, University of Washington  
<sup>†</sup>Allen Institute for Artificial Intelligence

Peters et al., NAACL 2018



Src: Wikipedia

# ELMo: Embeddings as Language Models

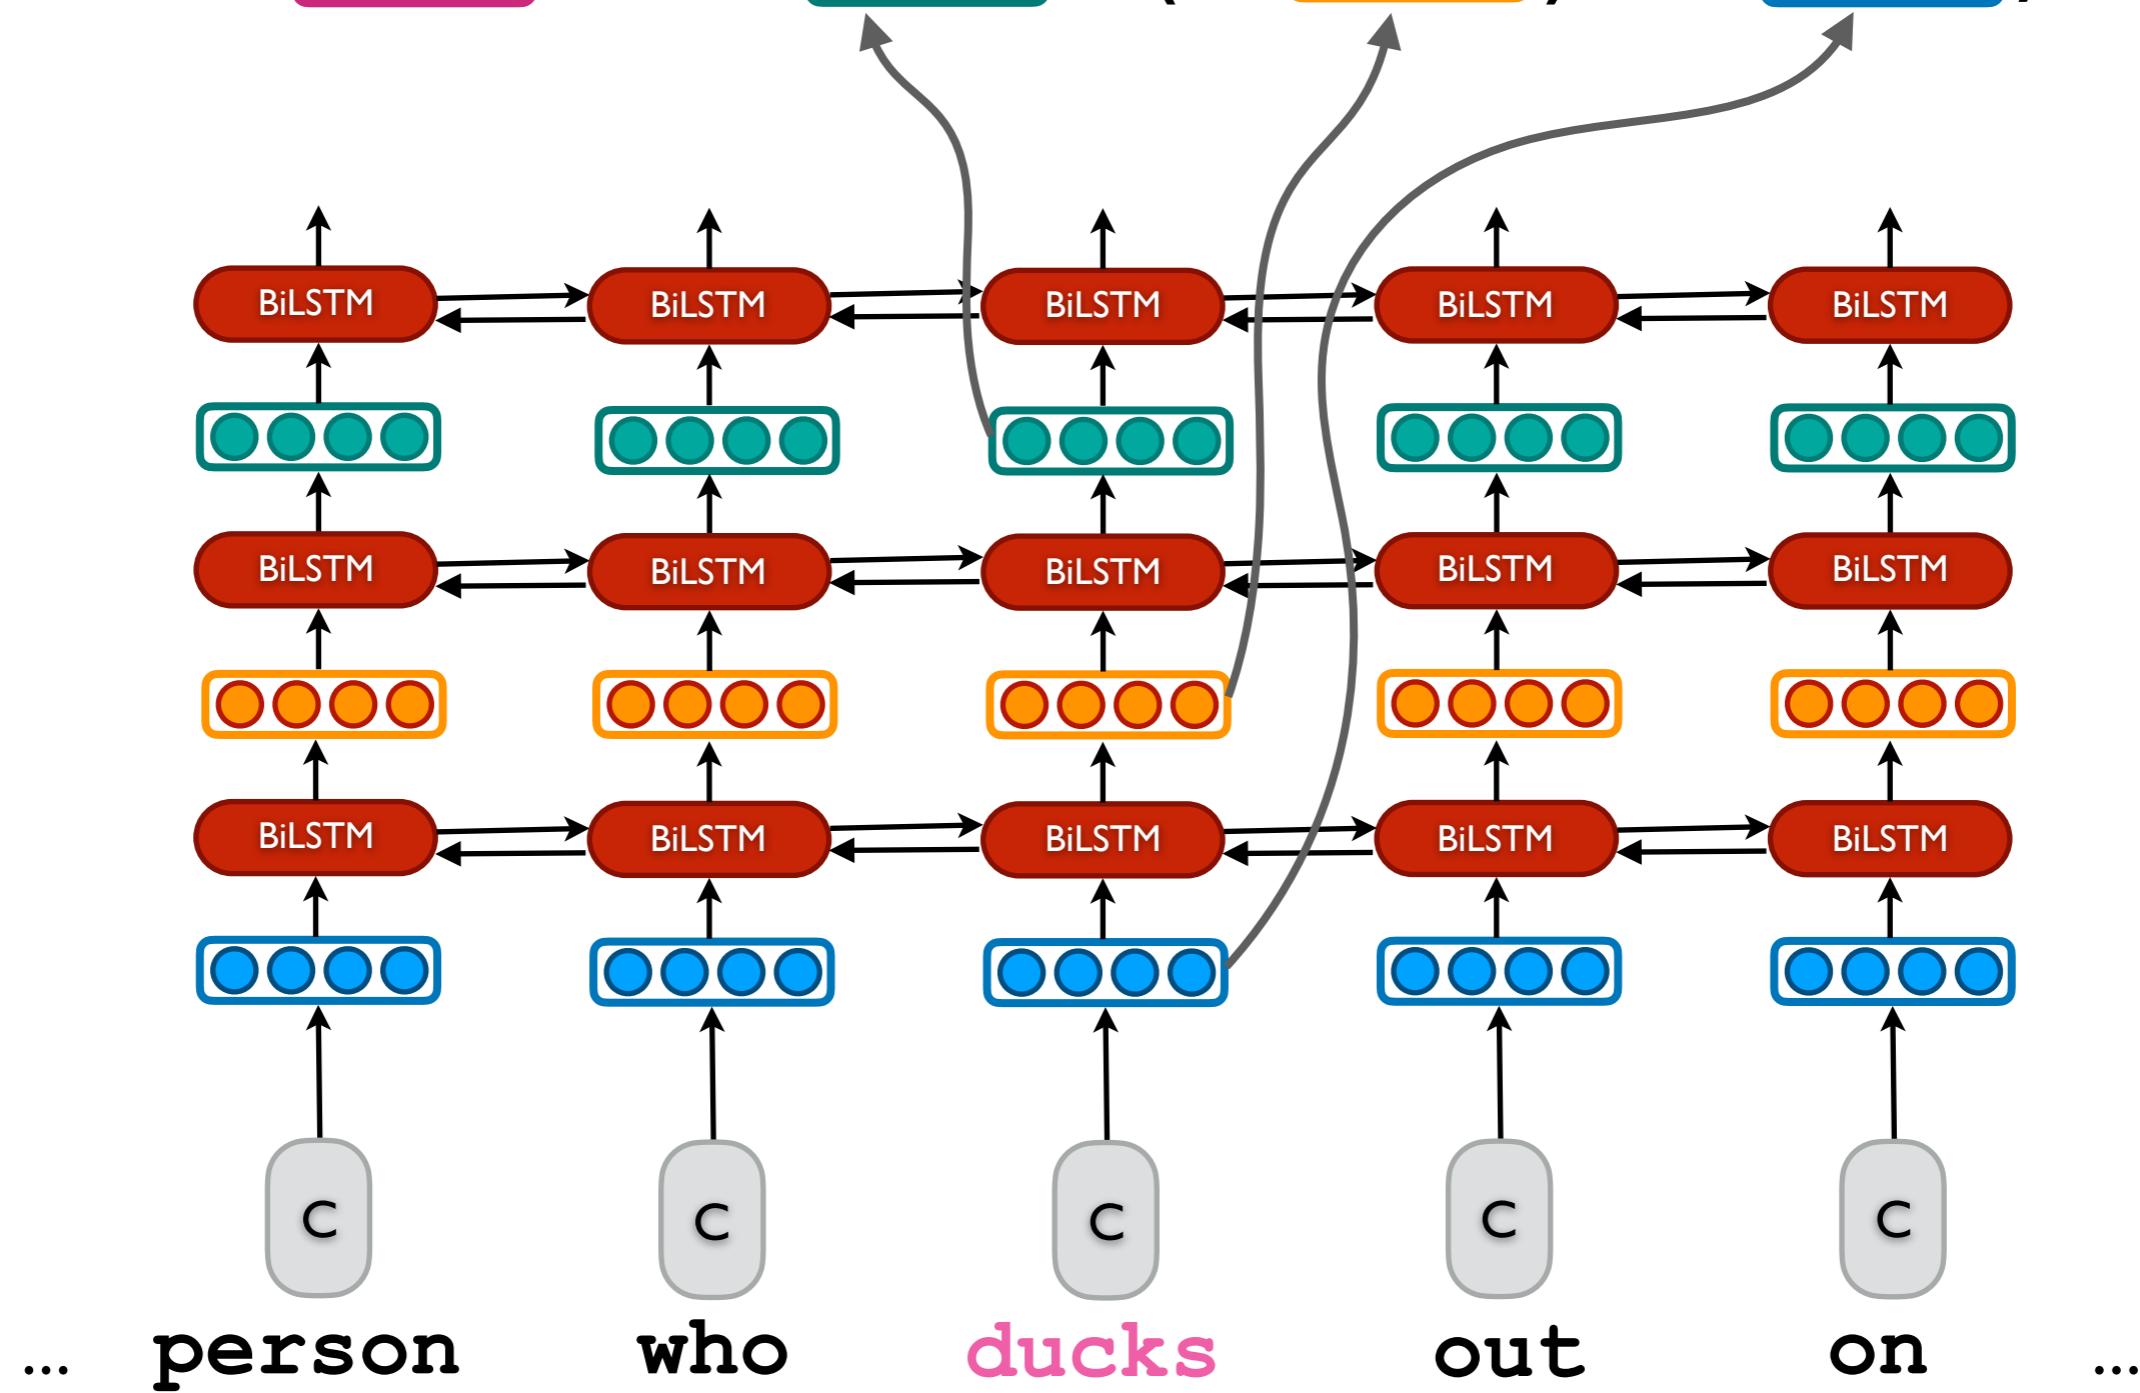
- ▶ Neural LMs embed the left and right context of a word
- ▶ We can use a bi-directional LM with the forward and the backward LSTM states

$$\sum_{k=1}^N \left( \log p(t_k | t_1, \dots, t_{k-1}; \Theta_x, \vec{\Theta}_{LSTM}, \Theta_s) + \log p(t_k | t_{k+1}, \dots, t_N; \Theta_x, \overleftarrow{\Theta}_{LSTM}, \Theta_s) \right).$$

- ▶ **Key Idea:** Learn word token vectors (not type!) using long contexts (not only context *windows*)
- ▶ ELMo uses a “deep” model to get different encodings (or “views”) from stacked RNNs

# Embeddings from Language Models

**ELMo**  = (  $\lambda_1^*$   ) + (  $\lambda_2^*$   ) + (  $\lambda_3^*$   )



# ELMo - Details

- ▶ ELMo: every token is assigned a representation that is a function of the entire input sentence ( $L = \#$  stacked layers)

$$\begin{aligned} R_k &= \{\mathbf{x}_k^{LM}, \overrightarrow{\mathbf{h}}_{k,j}^{LM}, \overleftarrow{\mathbf{h}}_{k,j}^{LM} \mid j = 1, \dots, L\} \\ &= \{\mathbf{h}_{k,j}^{LM} \mid j = 0, \dots, L\}, \end{aligned}$$

- ▶ This gives  $2L+1$  representations - Which to use?
  - ▶ Just the top layer (similar to TagLM; Peters et al., 2017)
  - ▶ Include all  $L+1$  layers, average
  - ▶ All layers, weighted average (best)

$$\mathbf{ELMo}_k^{task} = E(R_k; \Theta^{task}) = \gamma^{task} \sum_{j=0}^L s_j^{task} \mathbf{h}_{k,j}^{LM}$$

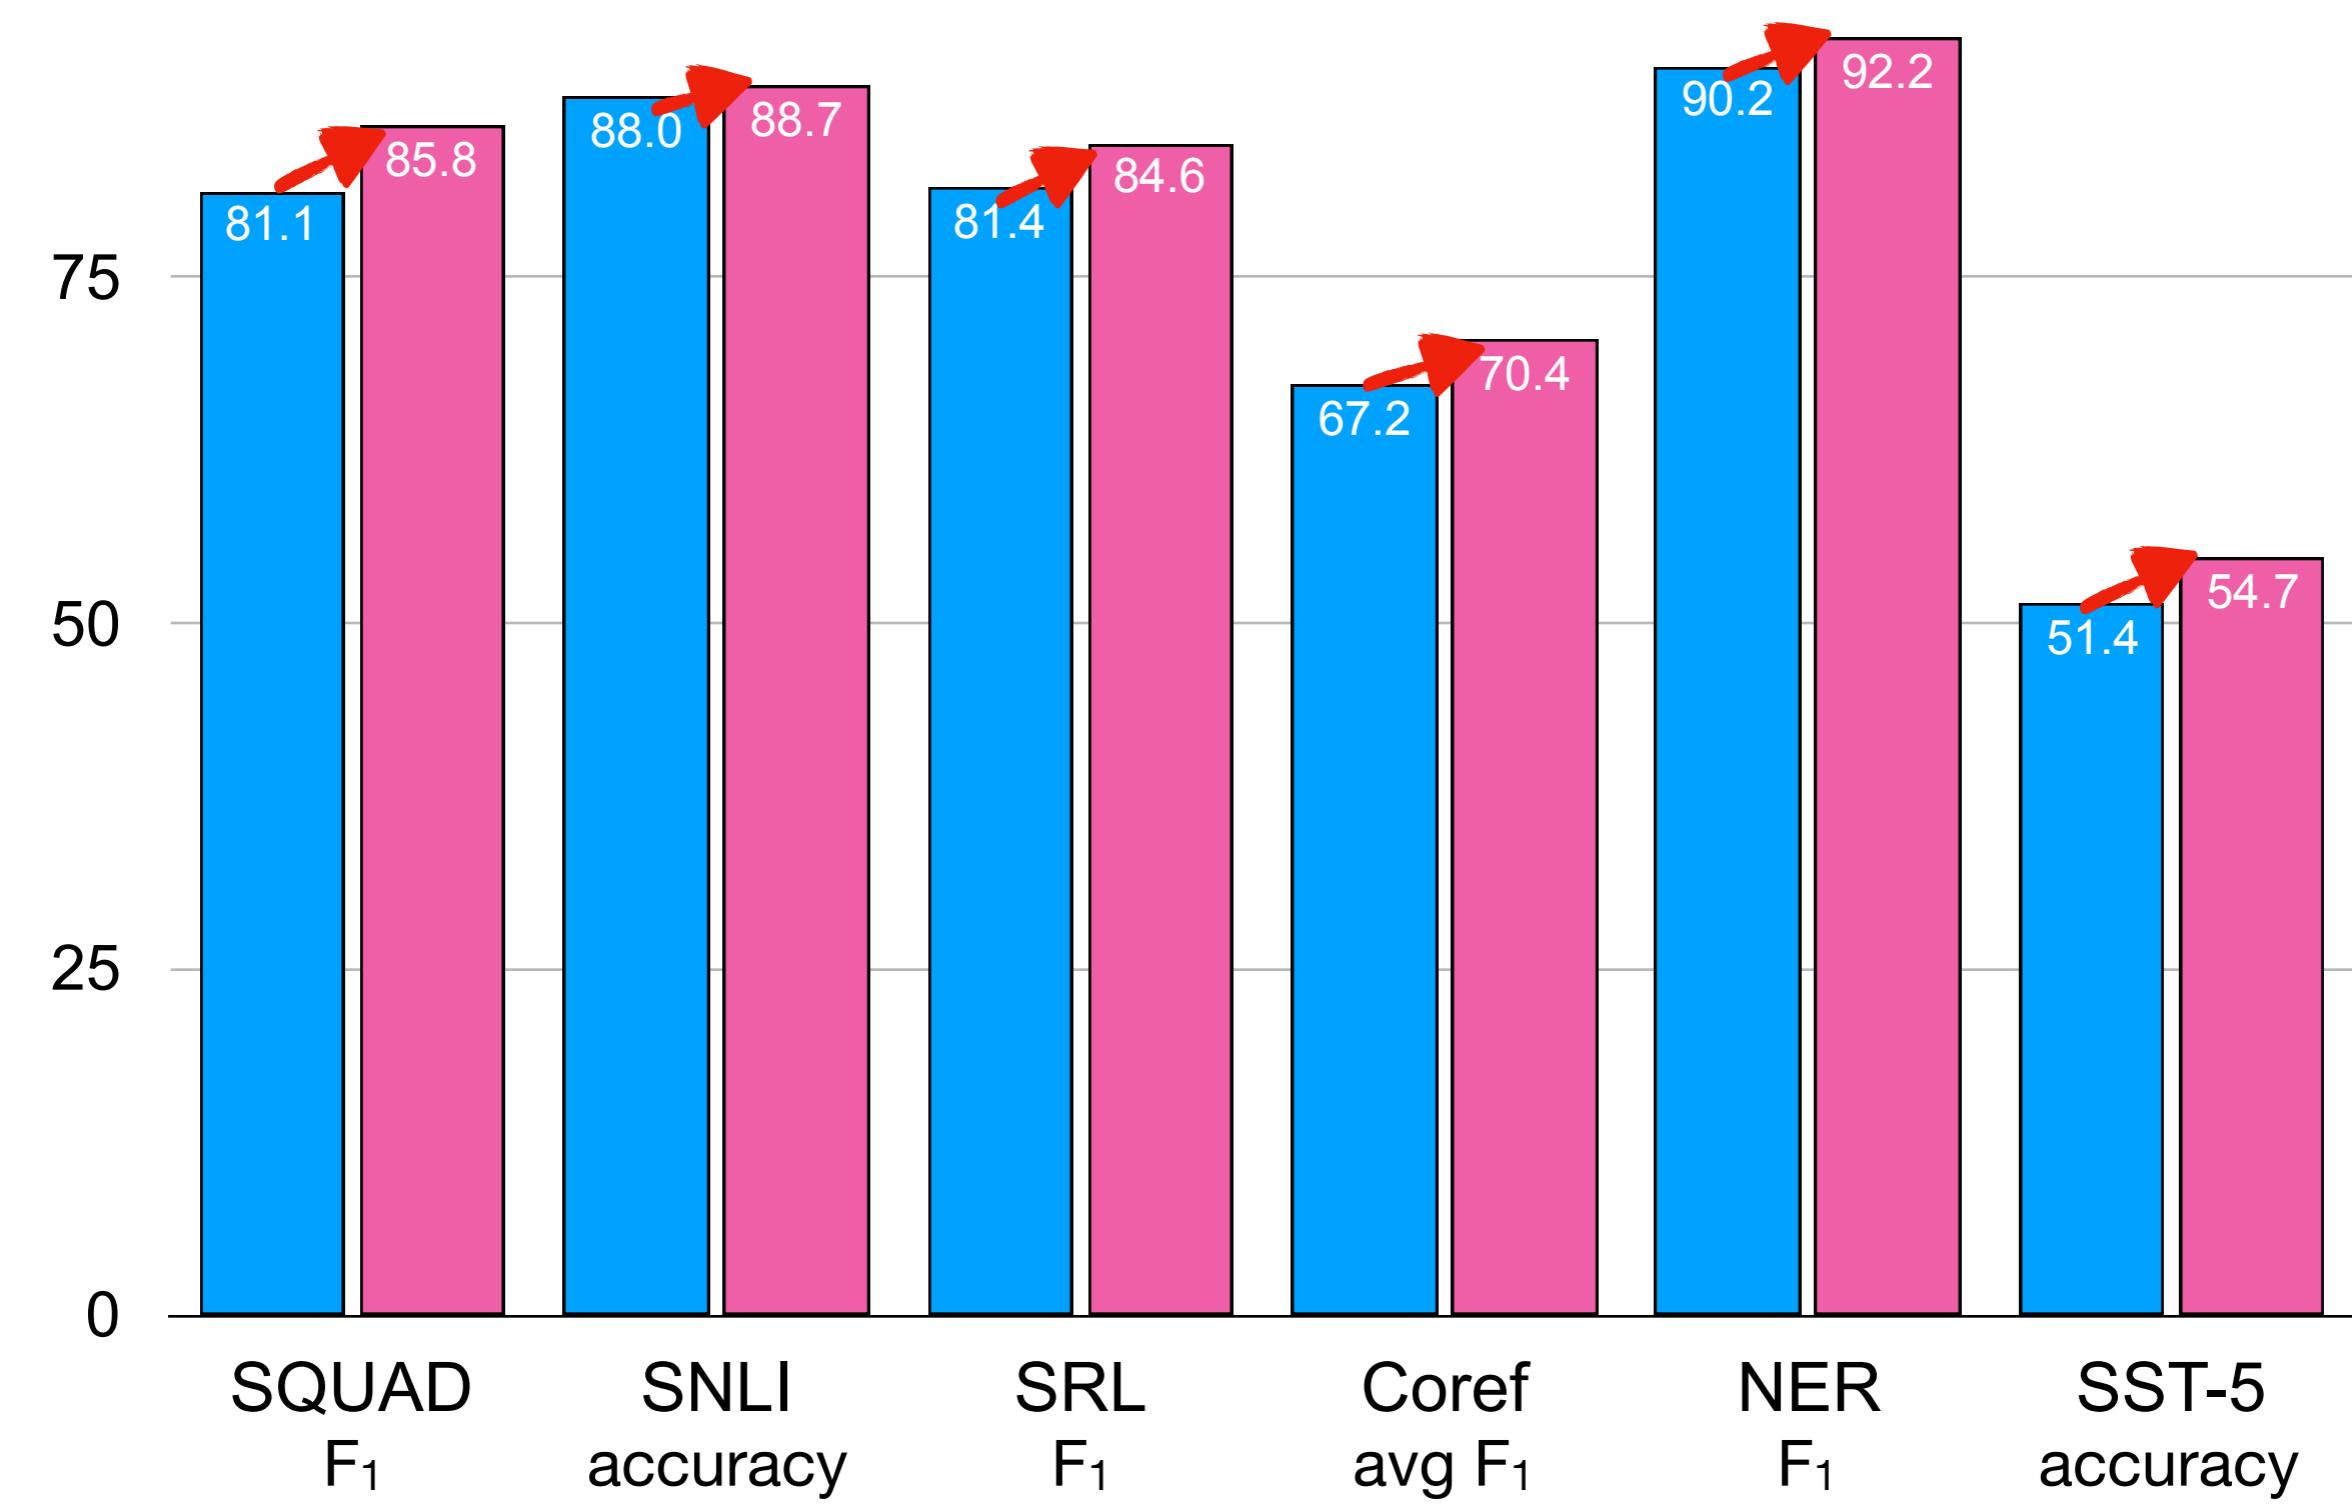
# How to use ELMo for your task?

- ▶ Recipe: For a given instance
  - ▶ Run biLM to get the representations for each word
  - ▶ Concatenate ELMo embeddings into task-specific model, e.g.,
    - ▶ as additional input to static word embeddings
    - ▶ as additional hidden representation
    - ▶ ... many choices, best might depend on end task



# Results over 6 NLP benchmarks

Peters et al., NAACL 2018



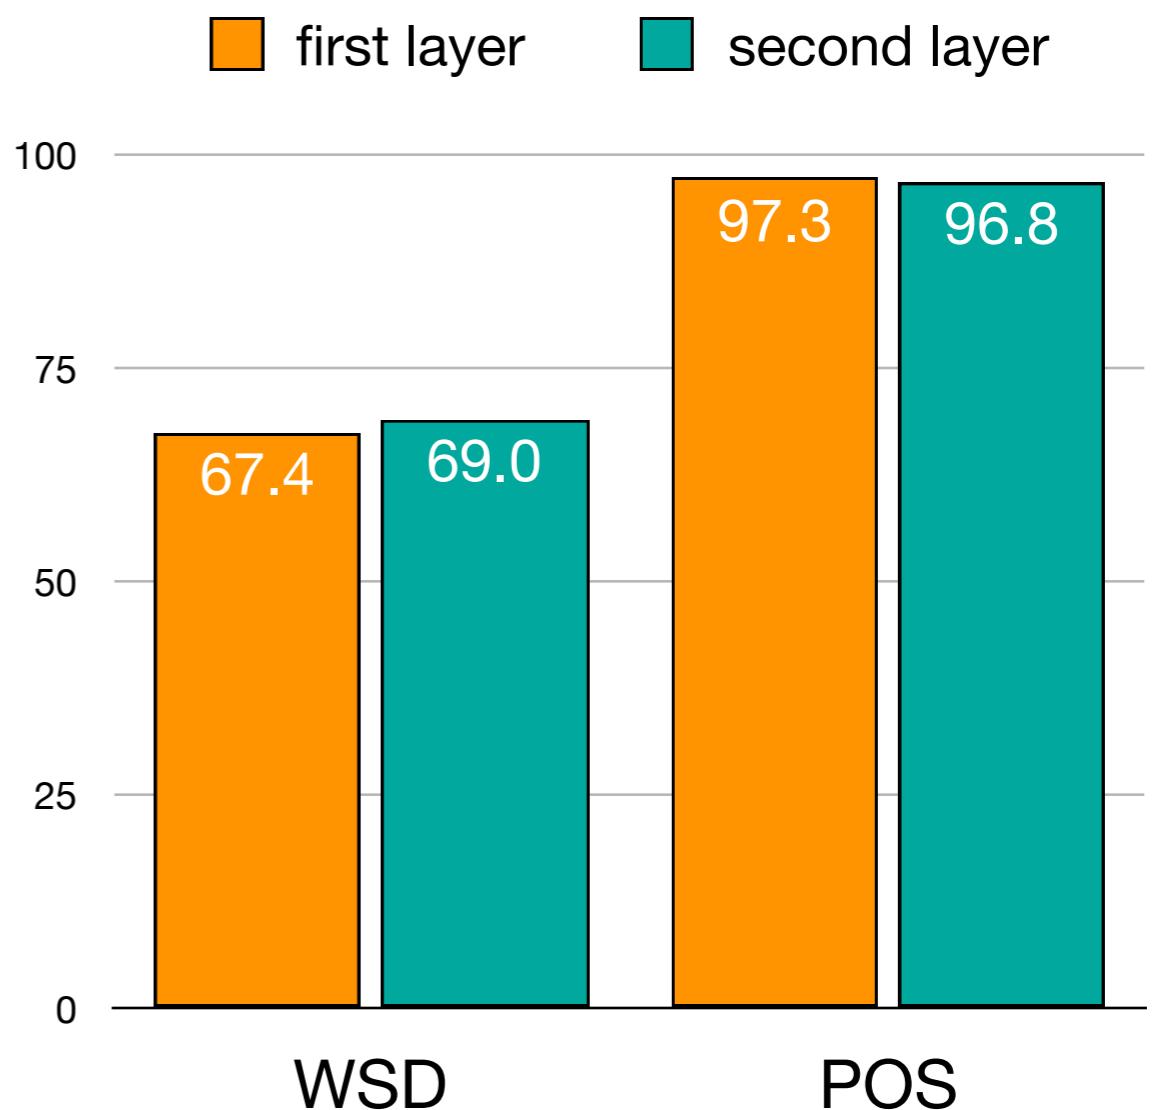
# Is ELMo the first such model? No!

- ▶ ELMo is *deeper* compared to an earlier model by Peters et al., 2017 ACL (**TagLM**)
- ▶ It doesn't require parallel data (as **CoVe** does) by McCann et al., 2017 NeurIPS
  - ▶ CoVe: use NMT as encoder (translation is meant to preserve meaning, so why not use it to provide context?)
- ▶ It obtained a new SOTA on 6 benchmarks

**What's in a  
representation?**

# Probing ELMo representations

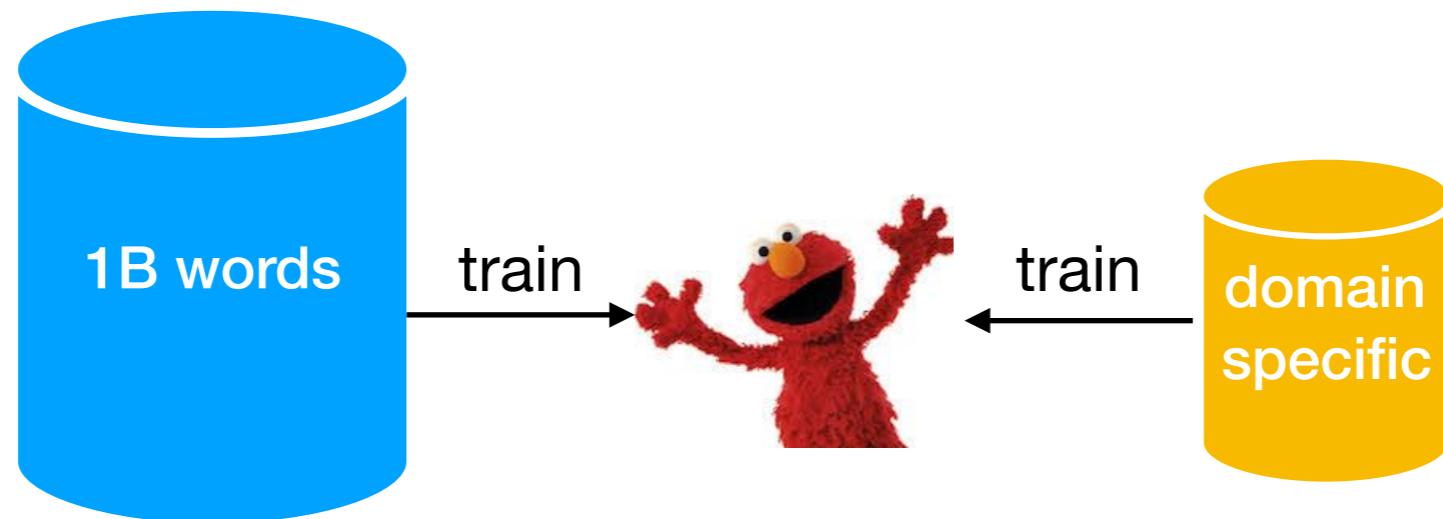
- ▶ What do ELMo representations capture?
  - ▶ Word Sense Disambiguation (WSD)
  - ▶ Part-of-Speech tagging (POS)
- ▶ Finding: Different layers encode different kinds of syntactic and semantic information



(Selected related work): [Tenney et al., 2019 ACL](#); [Liu et al., 2019 NAACL](#)  
[Belinkov & Glass, 2019 TAC](#)

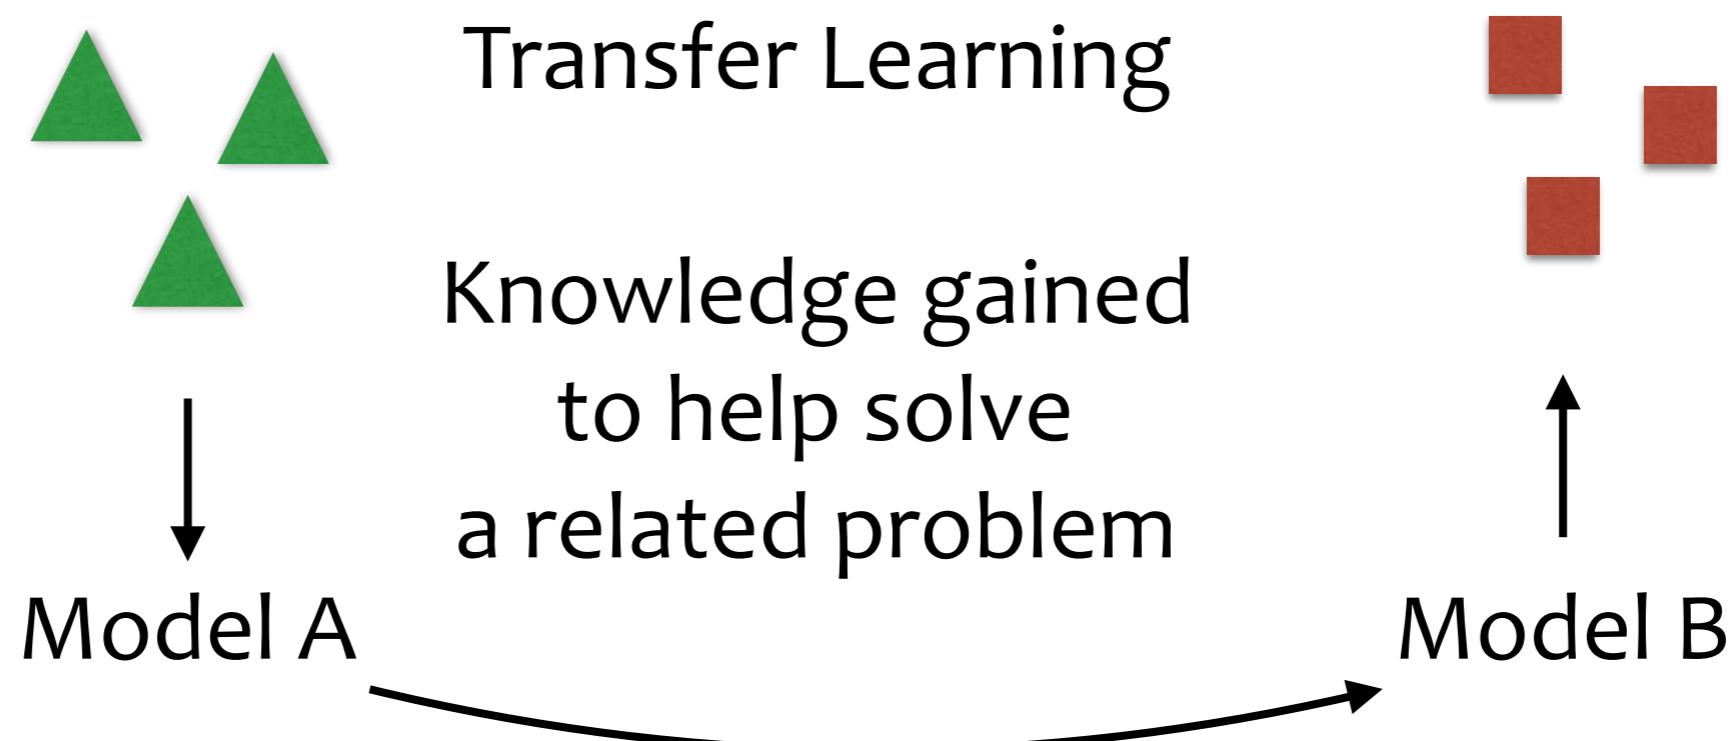
# On what was ELMo trained?

- ▶ A news corpus of 1B words: the 1-billion word language modeling benchmark (Chelba et al., 2014)
- ▶ ELMo can compute representations for any task
- ▶ In some cases, **fine-tuning** ELMo on domain-specific data leads to increased downstream performance



Fine-tuning: train on large data, continue training on small data (reuse weights)

# Fine-tuning: One way of Transfer Learning



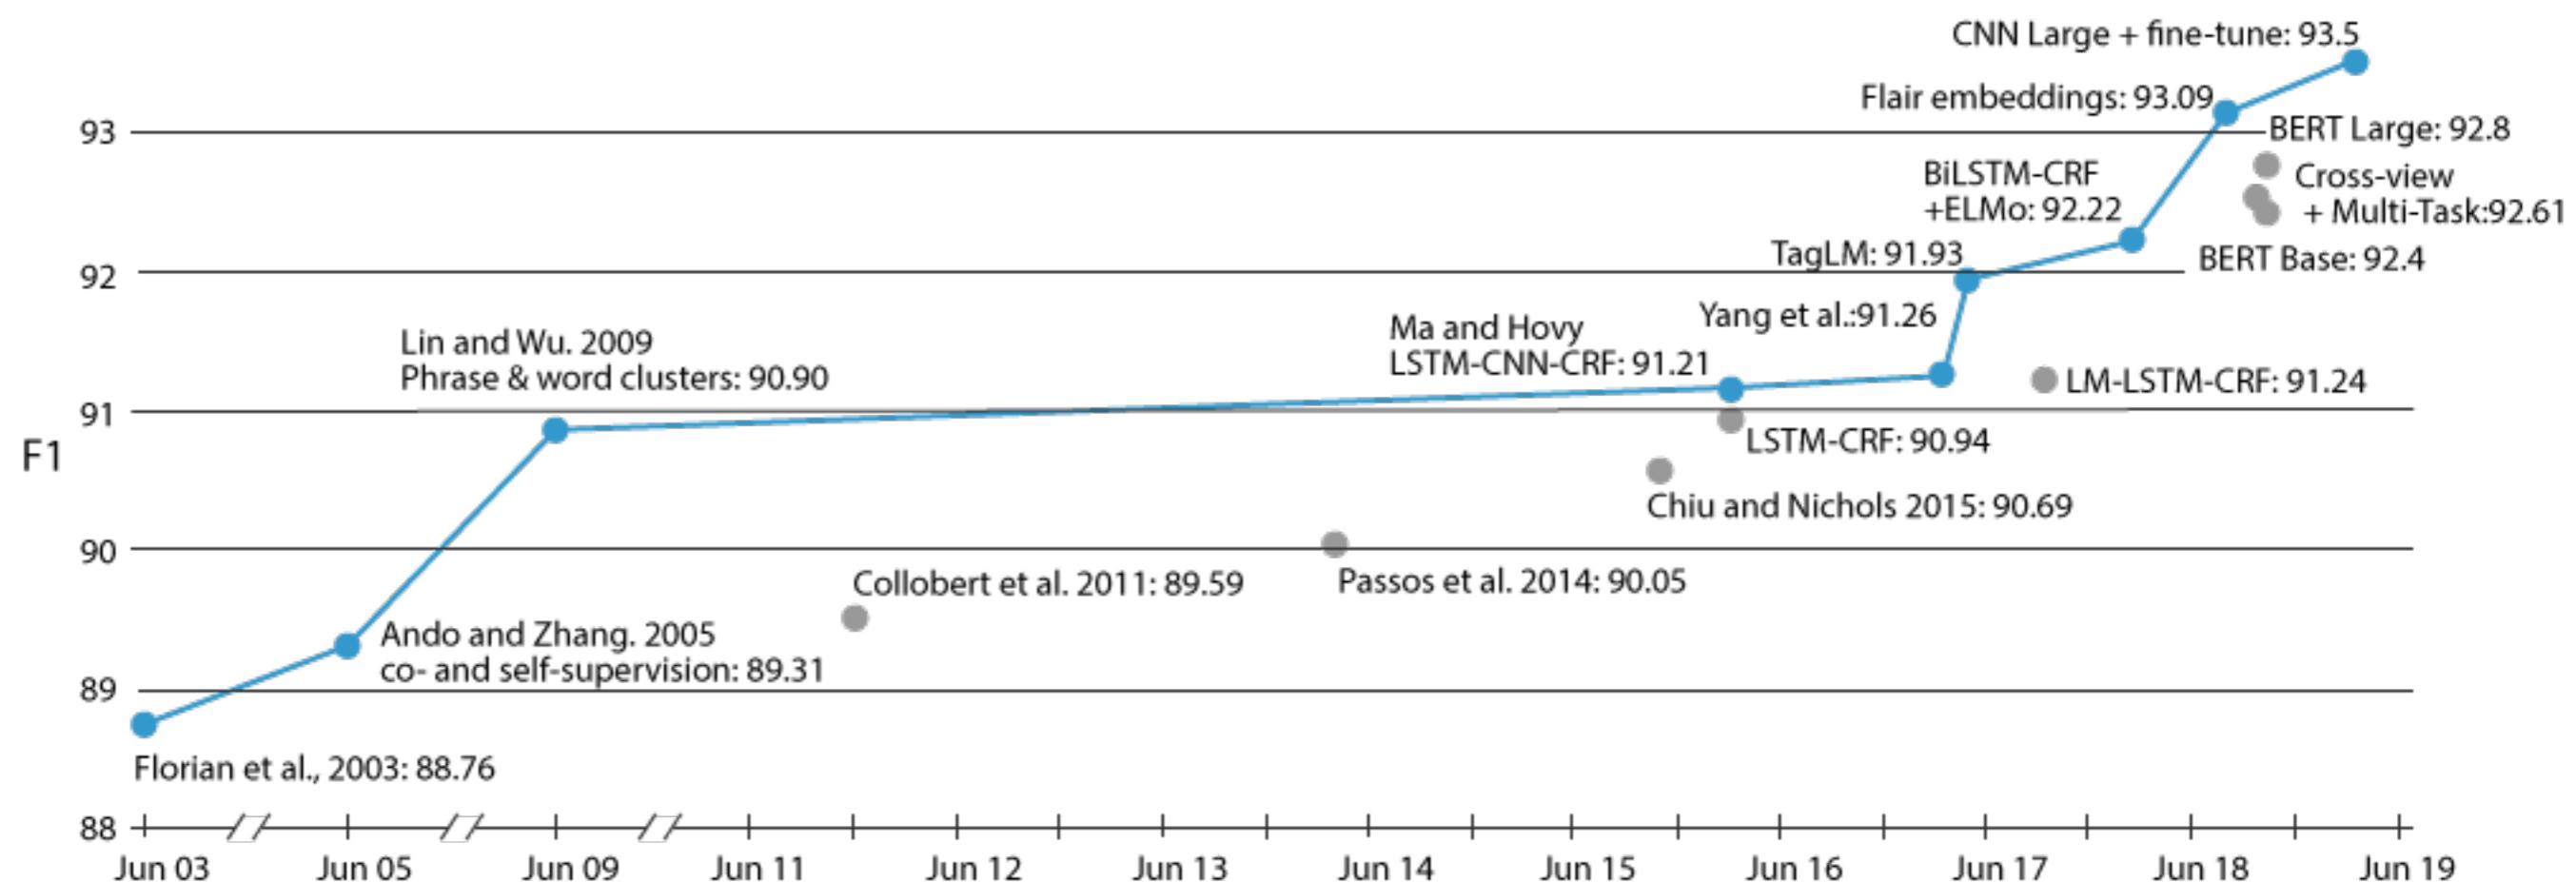
Language models learn  
transferable contextual  
representations

# To sum up: ELMo properties

- ▶ unsupervised
- ▶ contextual
- ▶ deep
- ▶ character-based
- ▶ extremely versatile (new type of word representation)
  
- ▶ Many follow-up words, most of which rely on the transformer model (Lecture 4), e.g., BERT

# NLP Progress on NER

- ▶ From Ruder et al.'s 2019 NAACL tutorial



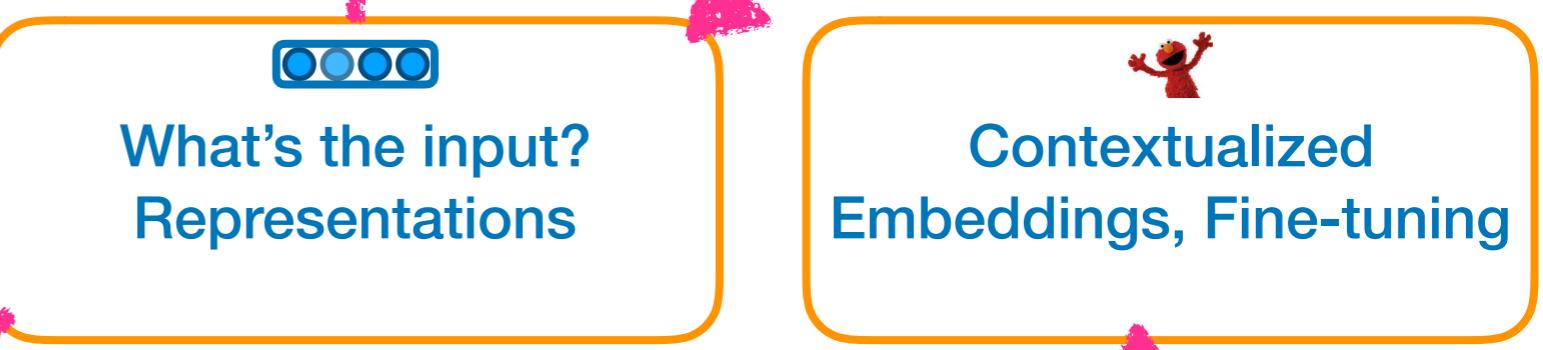
Performance on Named Entity Recognition (NER) on CoNLL-2003 (English) over time

# Overview

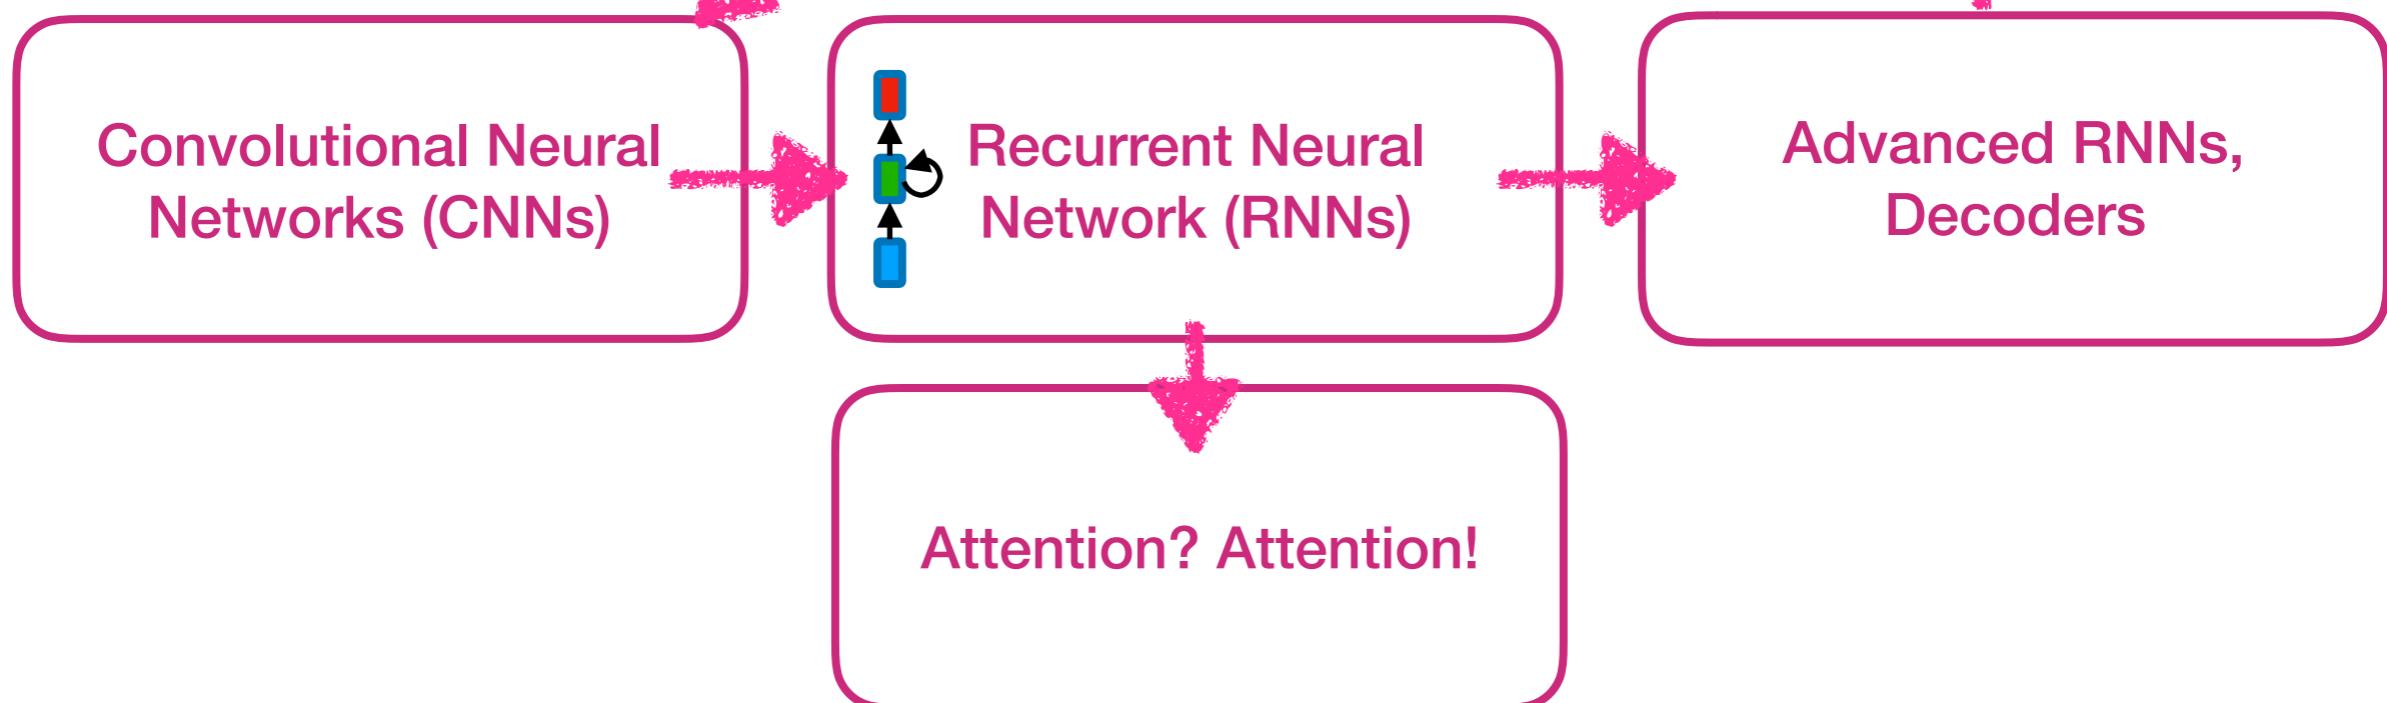
*foundations*



*representations*



*beyond FFNNs*

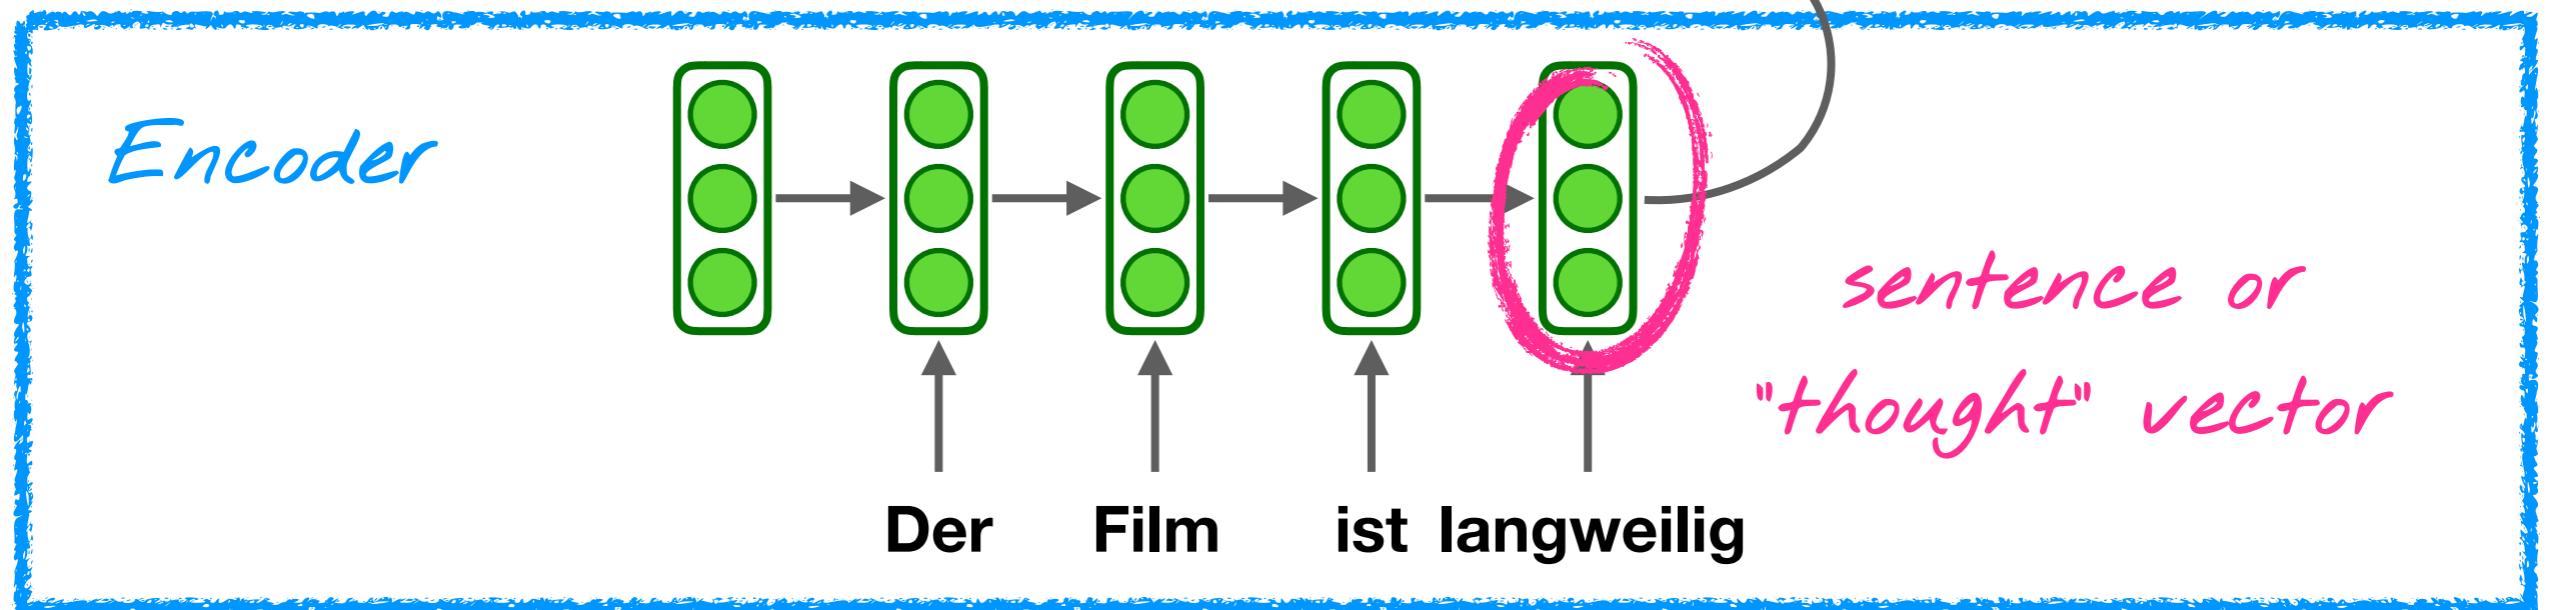
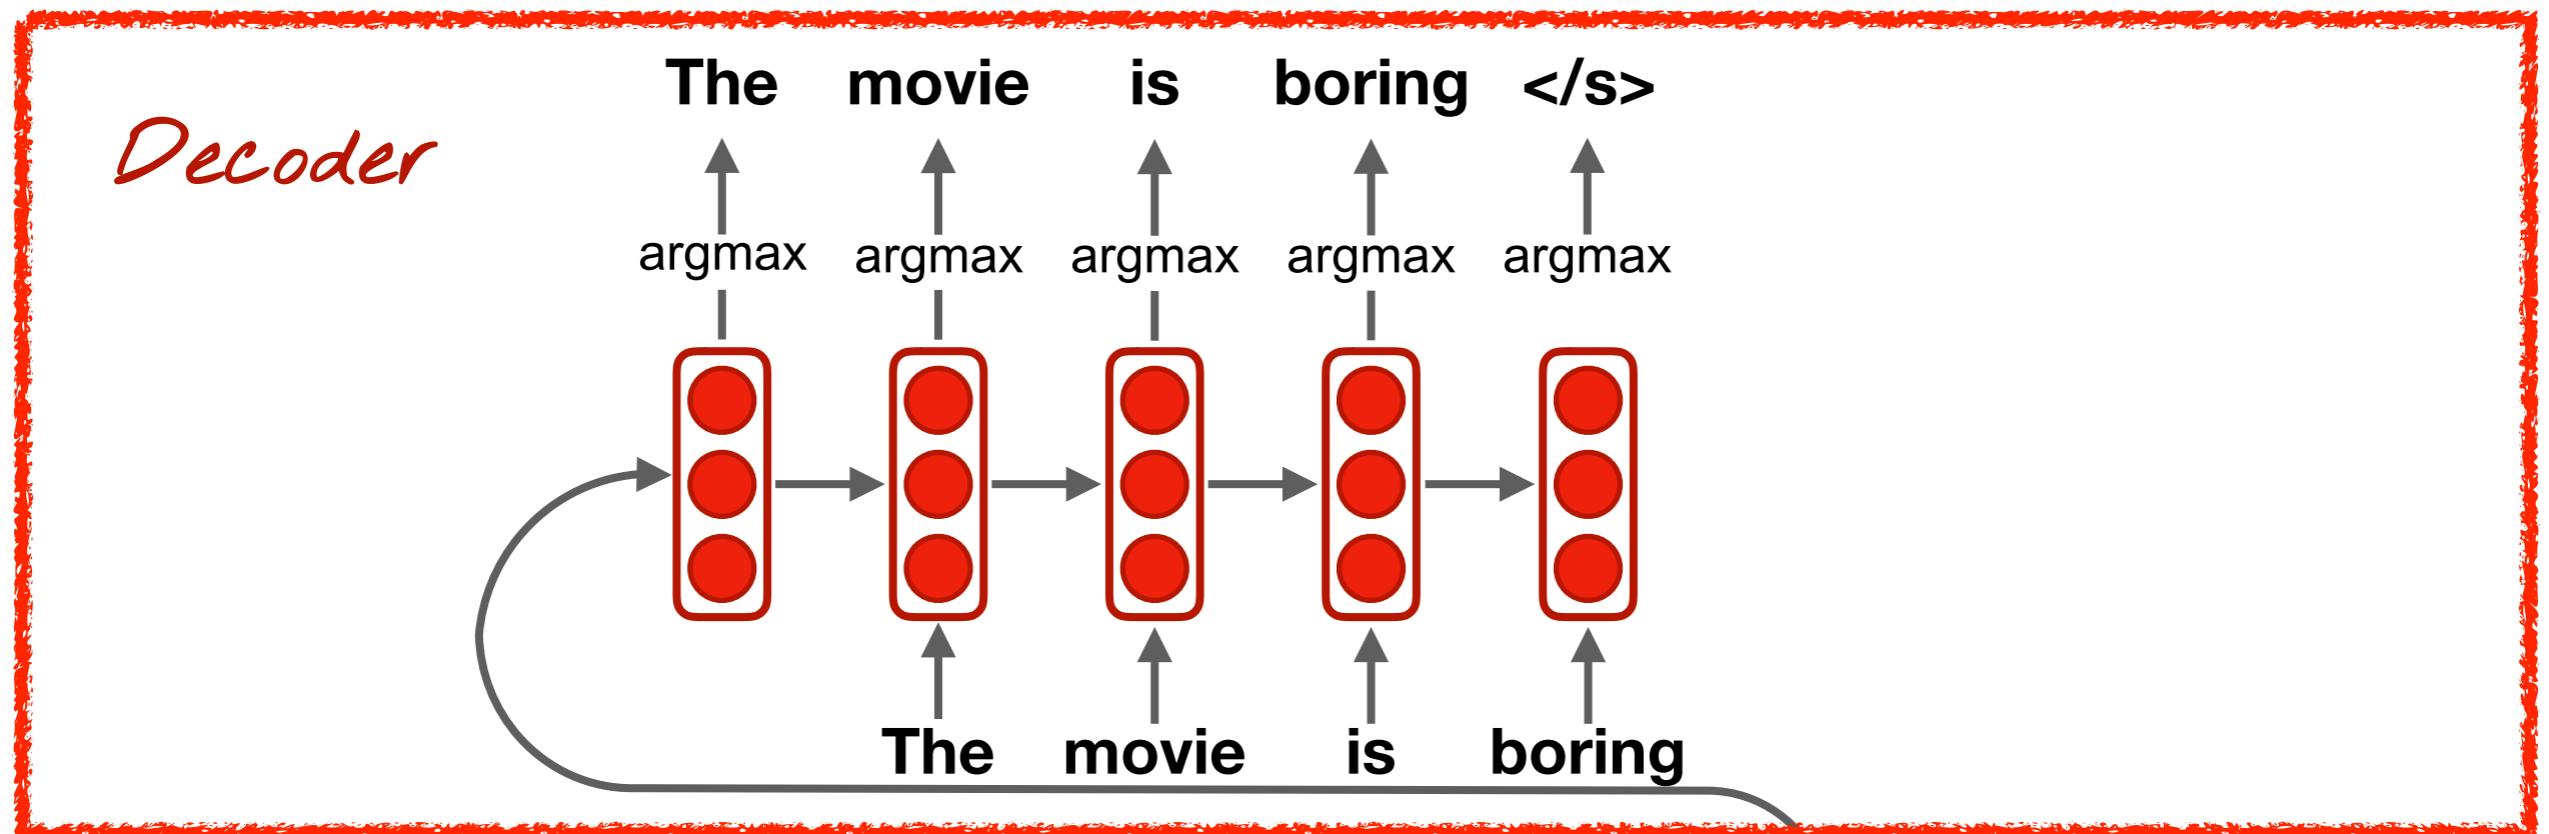


# Attention? Attention!

Many thanks to Lilian Weng for an awesome tutorial (<https://lilianweng.github.io/lil-log/2018/06/24/attention-attention.html>) and Graham Neubig's NN for NLP class (<http://www.phontron.com/class/nlp2019/>)

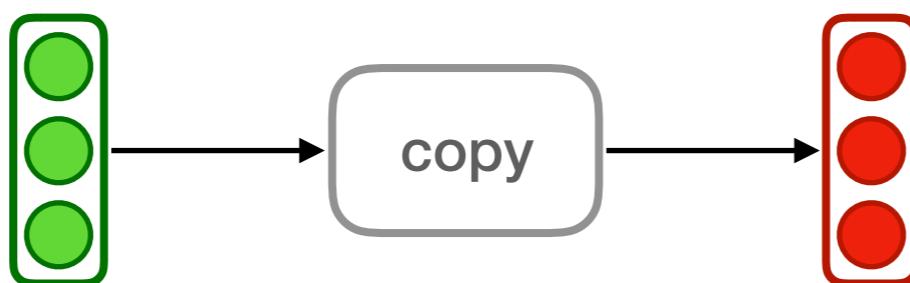
# Motivation: Encoder-decoder model

(Sutskever et al., 2014; Cho et al., 2014)

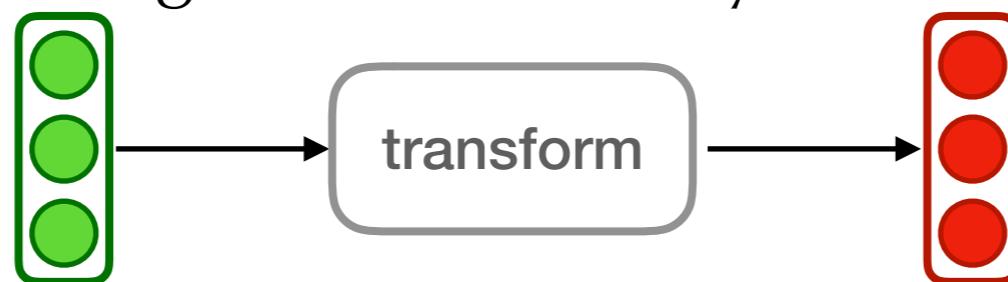


# How to pass the sentence vector?

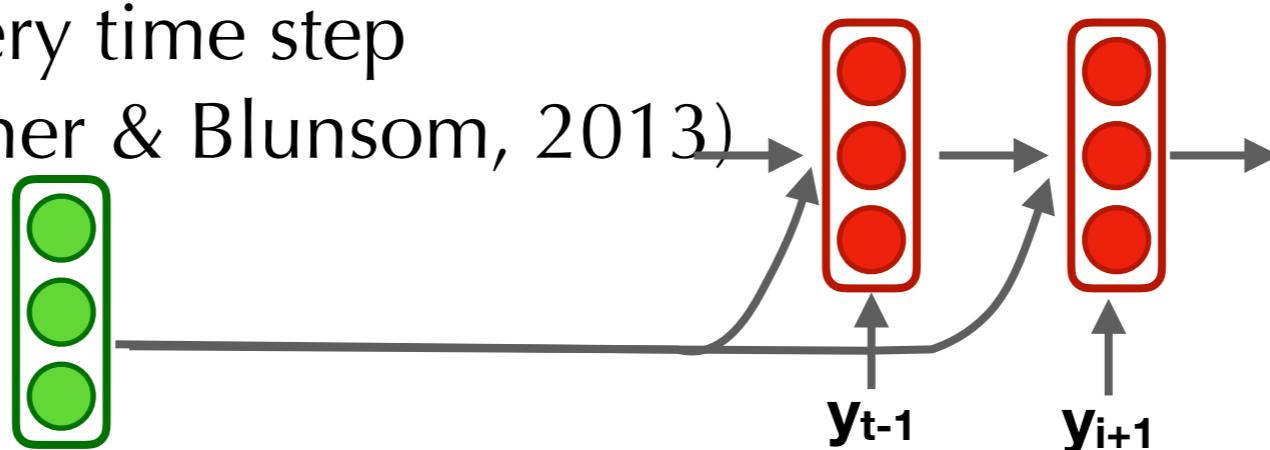
- ▶ Initialize decoder with encoder representation (Sutskever et al., 2014)



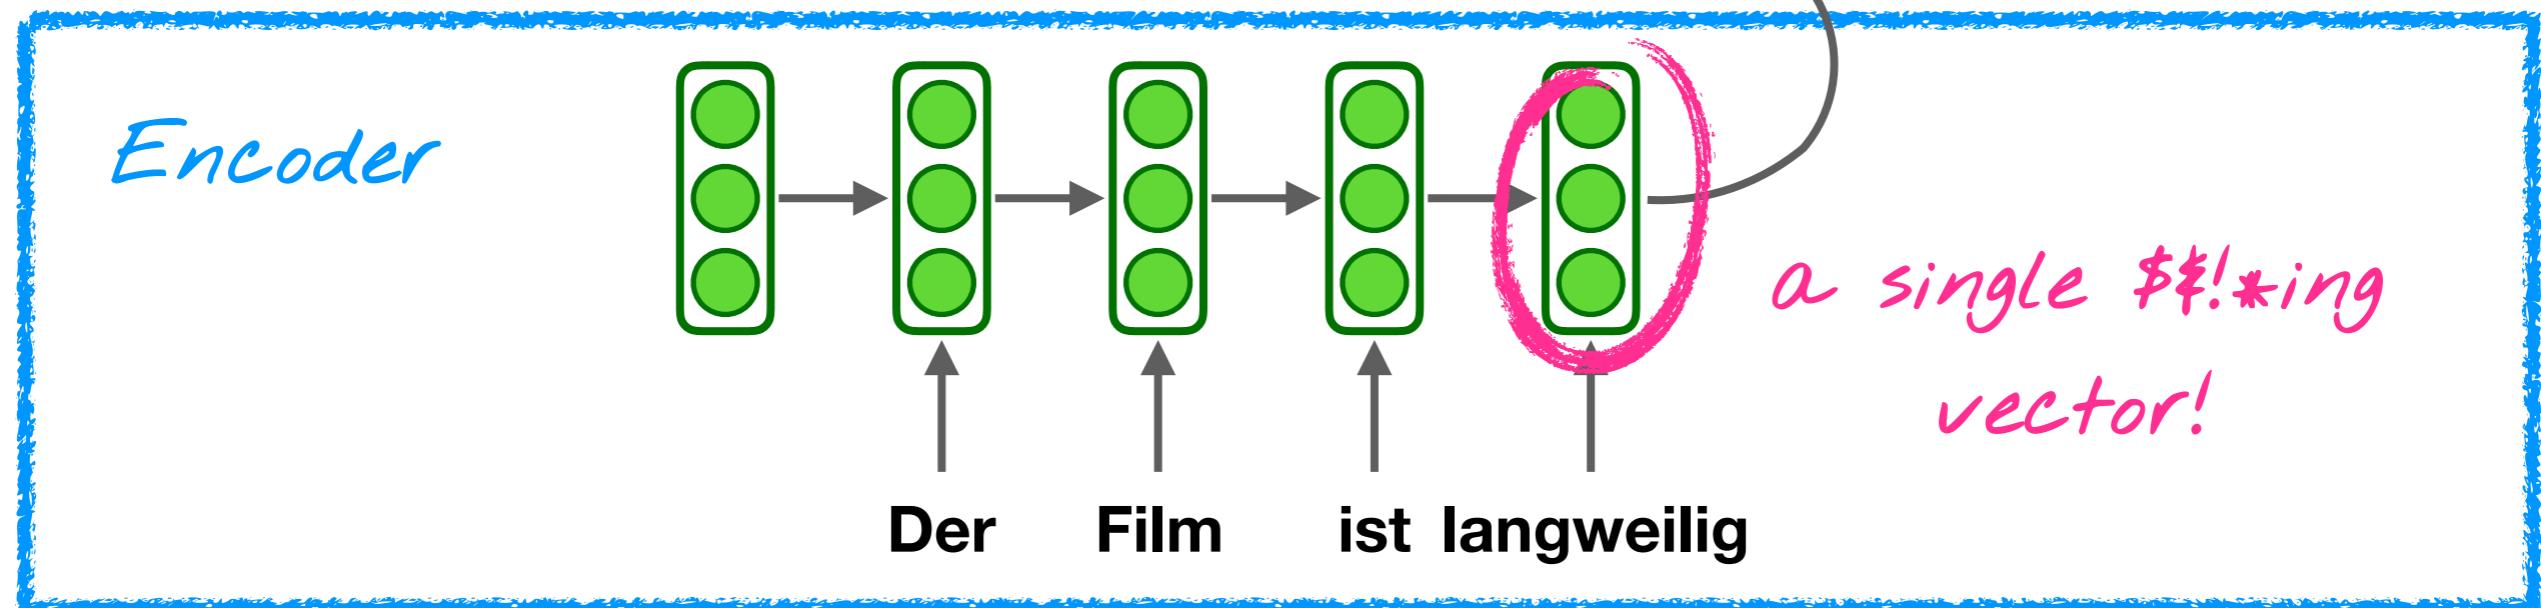
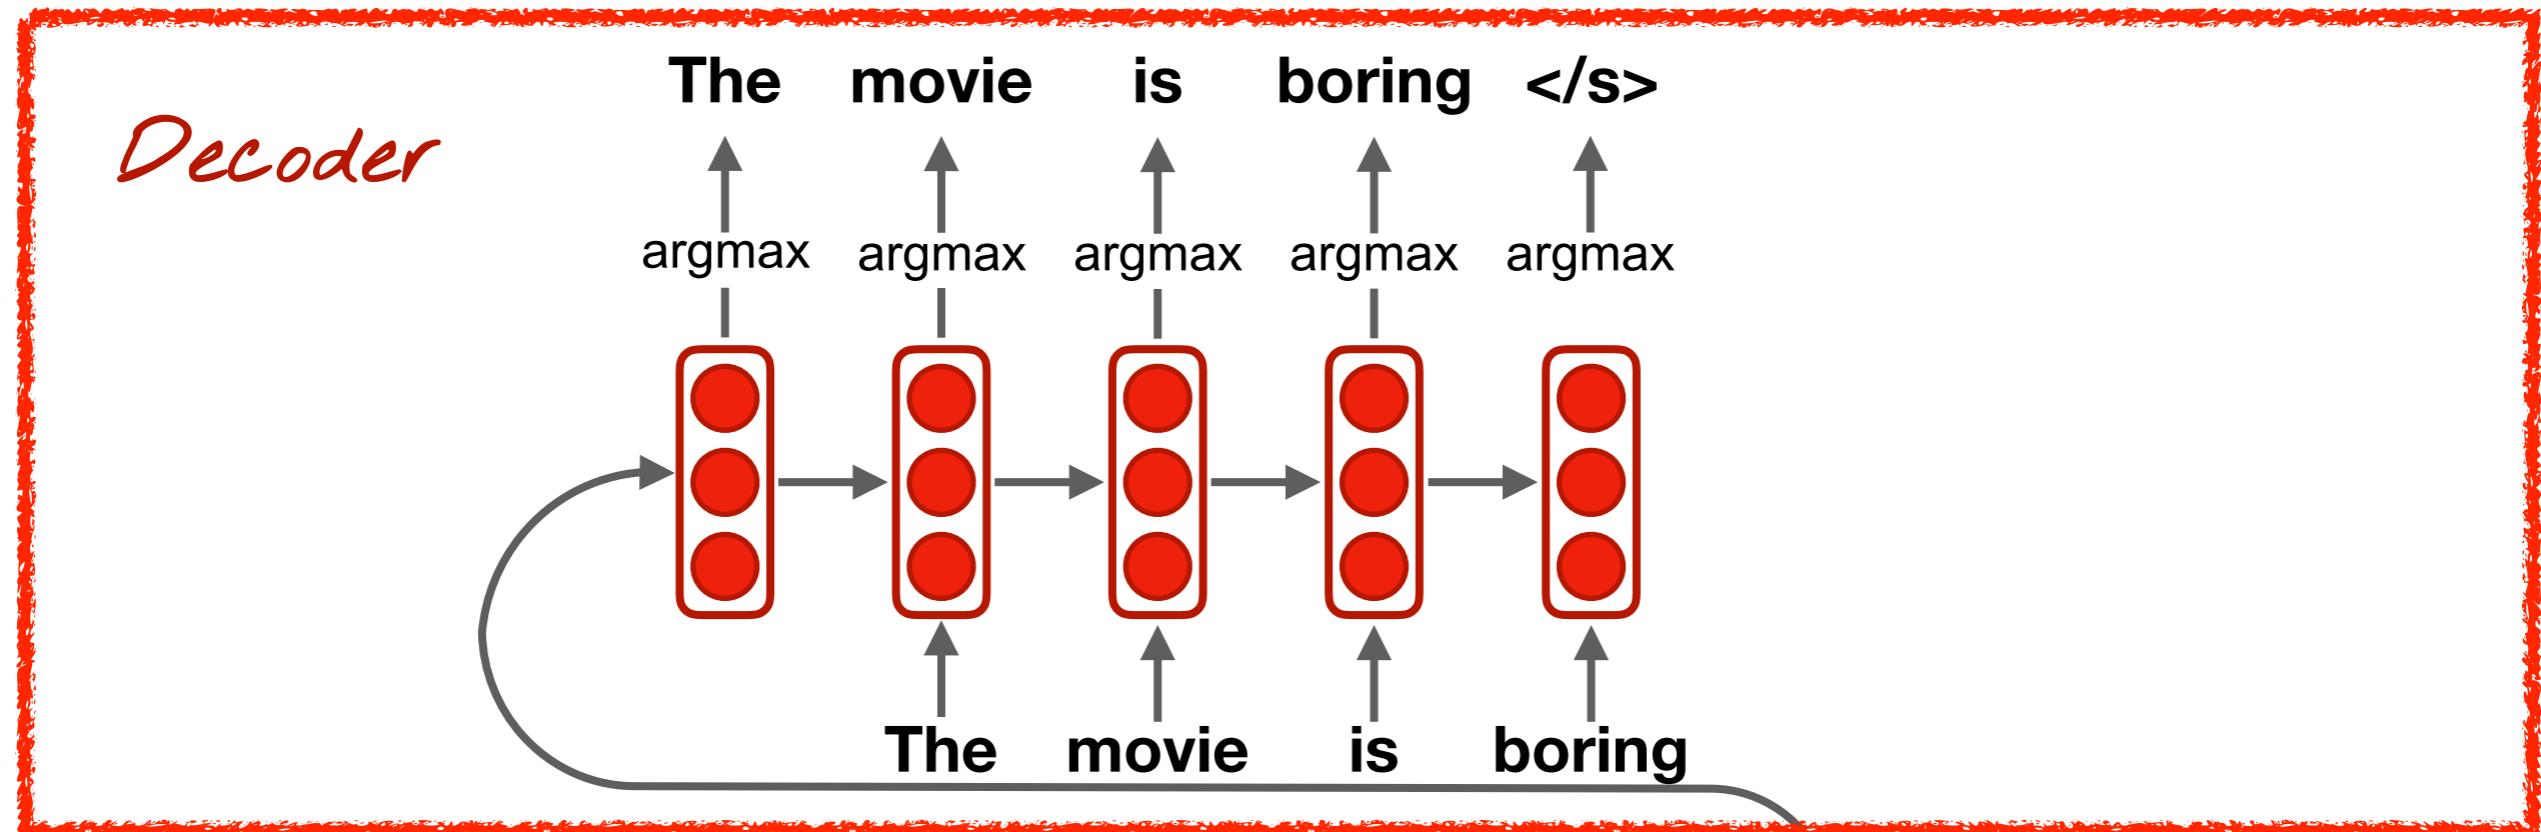
- ▶ Transform (change dimensionality)



- ▶ Input at every time step  
(Kalchbrenner & Blunsom, 2013)



# But: we're cramming it all into..

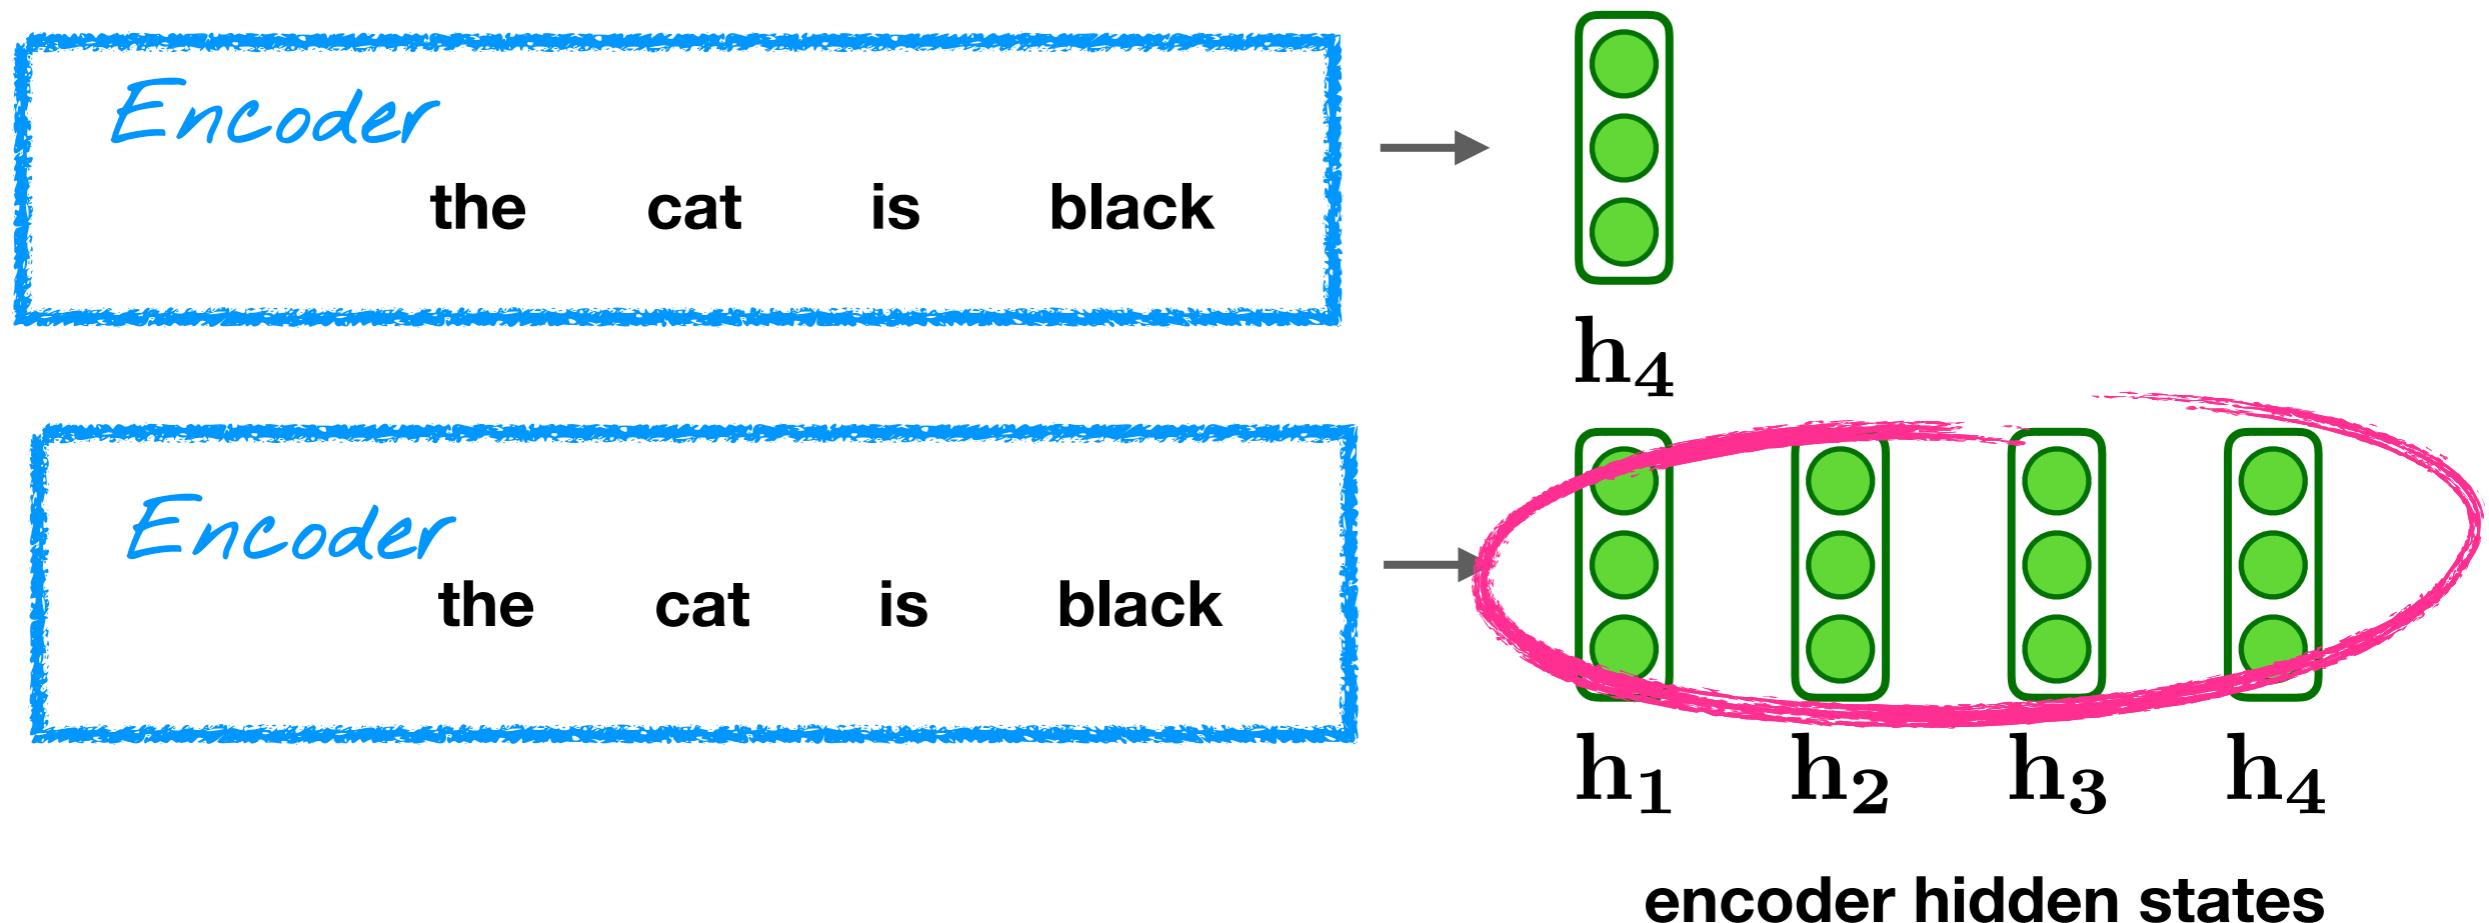


# Problem

- ▶ The encoder compresses the sentence into a single fixed-size vector. This representation is expected to be a good summary of the entire sentence.
- ▶ Disadvantage: incapability of remembering longer sequences.
- ▶ **“You can’t cram the meaning of a whole %&!\$ing sentence into a single \$&!\*\$ing vector!” — Ray Mooney**

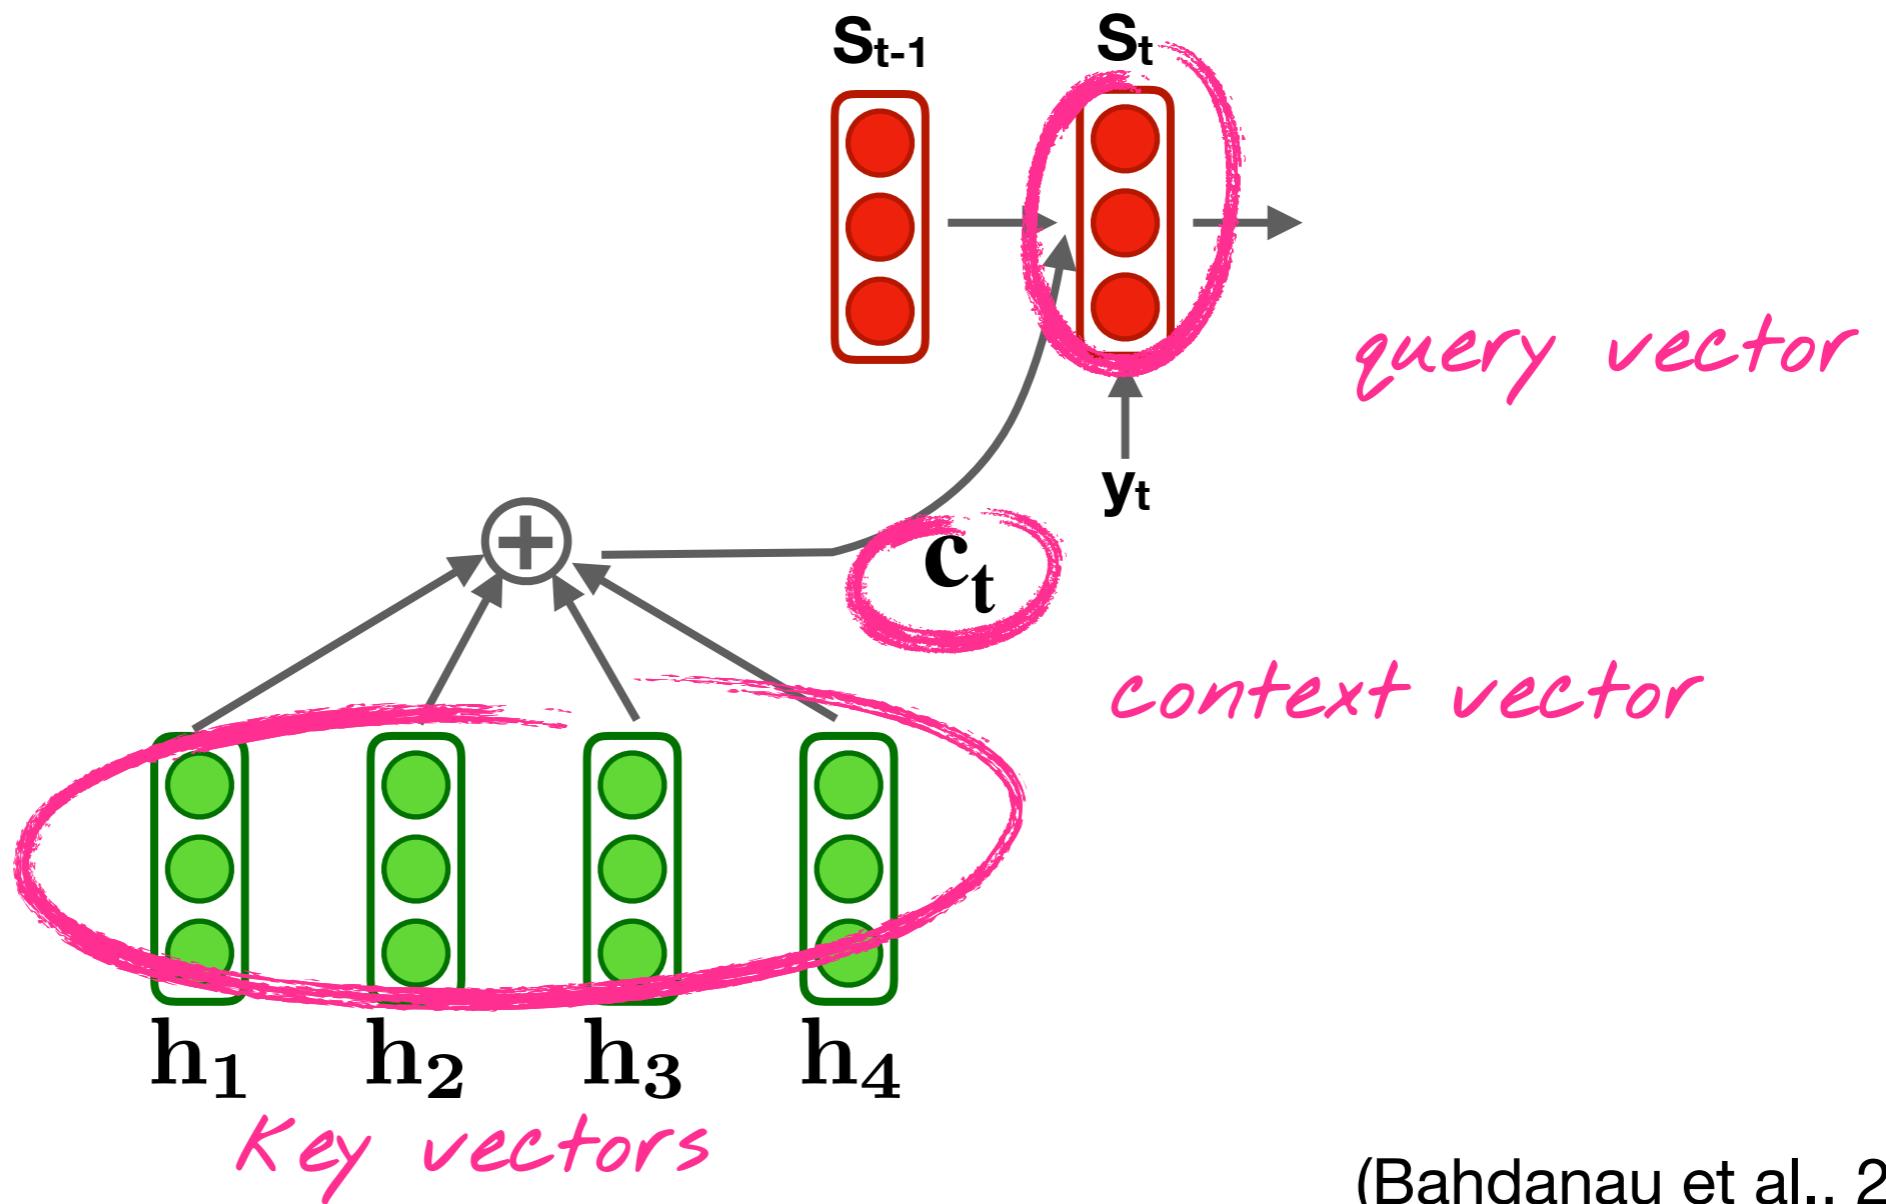
# Beyond a single static “crammed” vector

- ▶ What if we could use several vectors, based on the length of the input sequence?
- ▶ Idea: when we generate the next word in MT, perhaps we can learn to **attend** to the **relevant** source words



# Attention: Core Idea

- When decoding, perform a linear combination of the encoded input vectors, weighted by “attention weights”

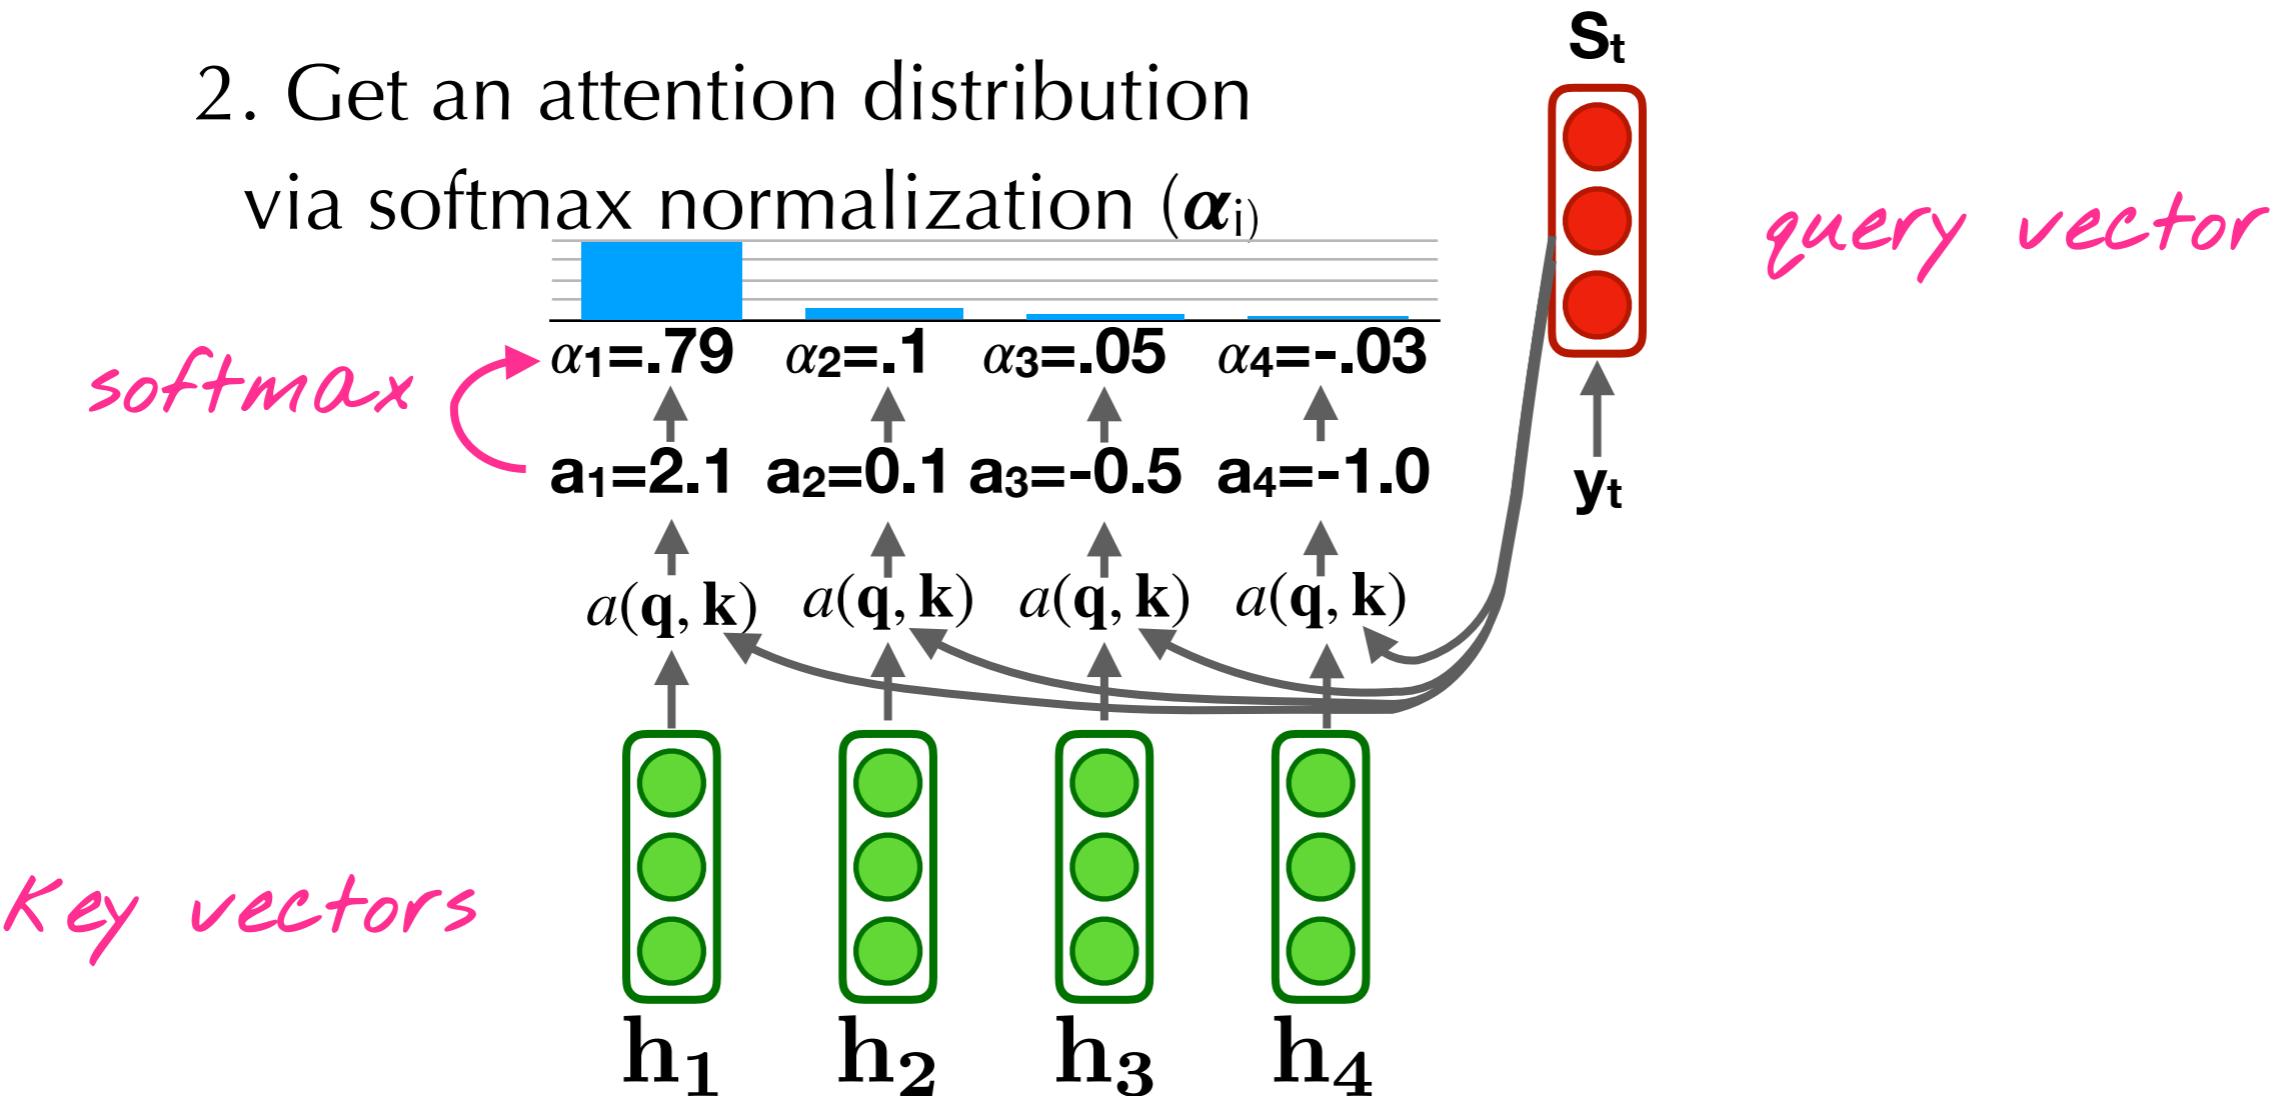


# Calculating attention (1/2): Attention weights $\alpha$

1. For each query-key pair, calculate an **attention score** ( $a_{ij}$ )

2. Get an attention distribution

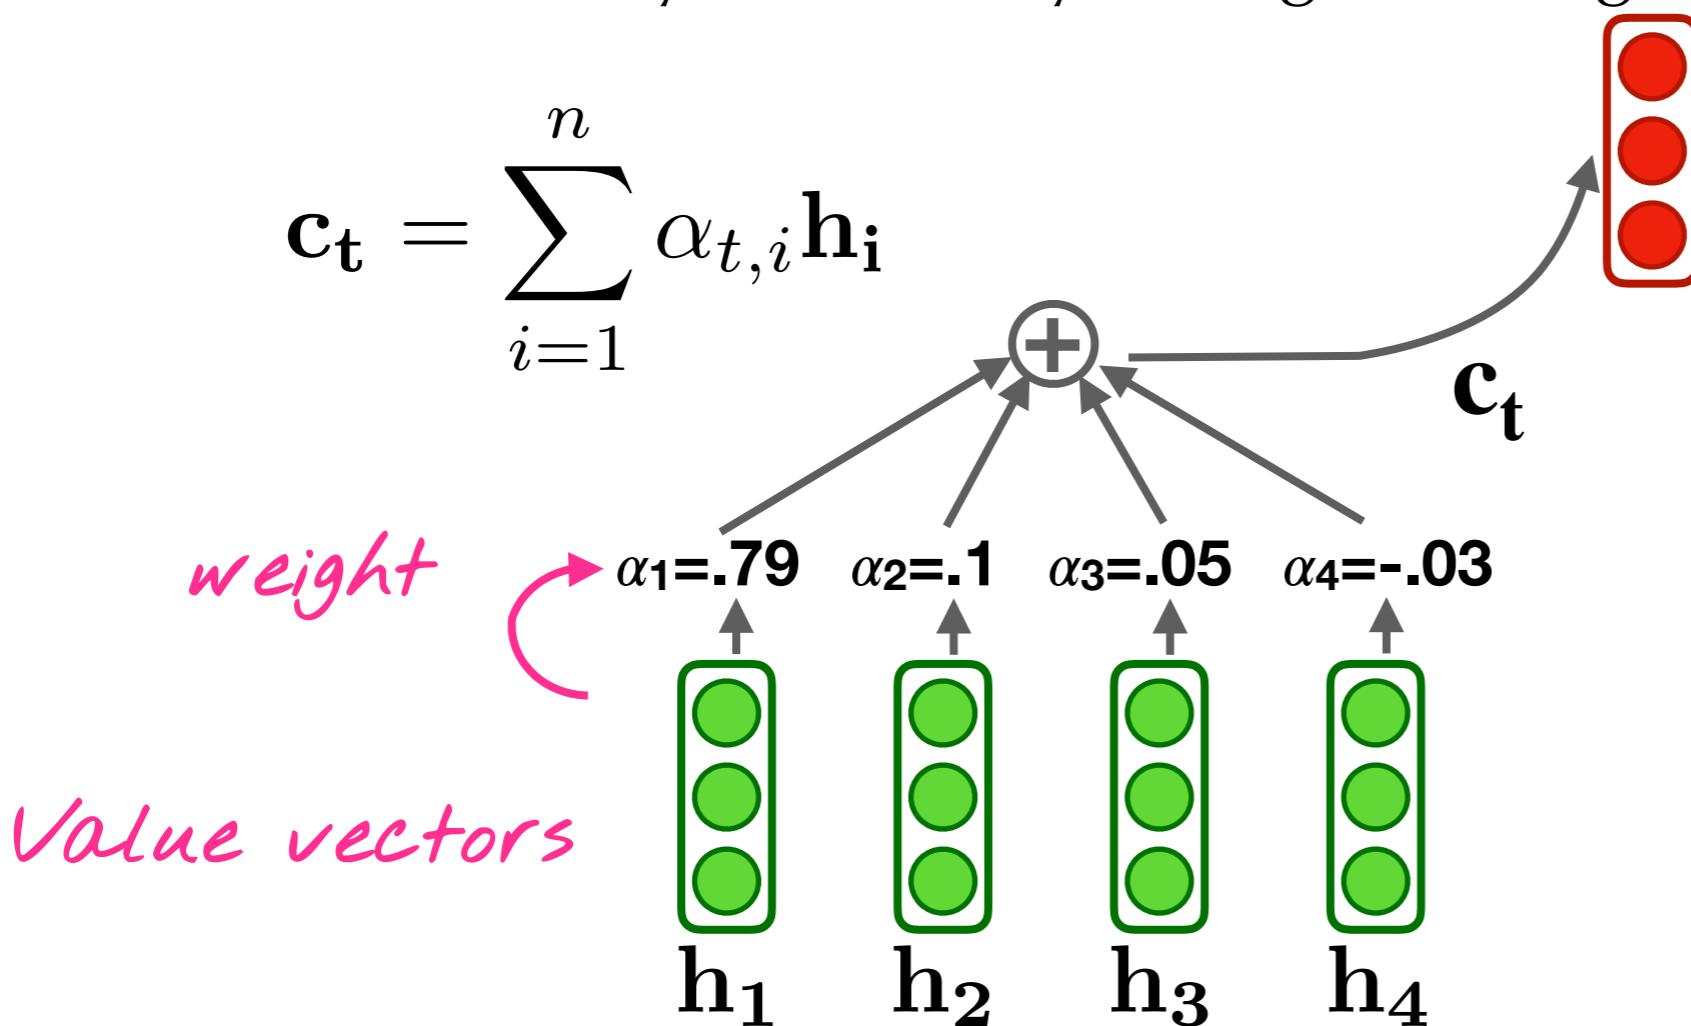
via softmax normalization ( $\alpha_i$ )



(Bahdanau et al., 2015)

# Calculating attention (1/2): Attention weights $\alpha$

3. Combine together **value vectors** (can be the encoder states, like the key vectors) by taking the weighted sum to get  $\mathbf{c}$



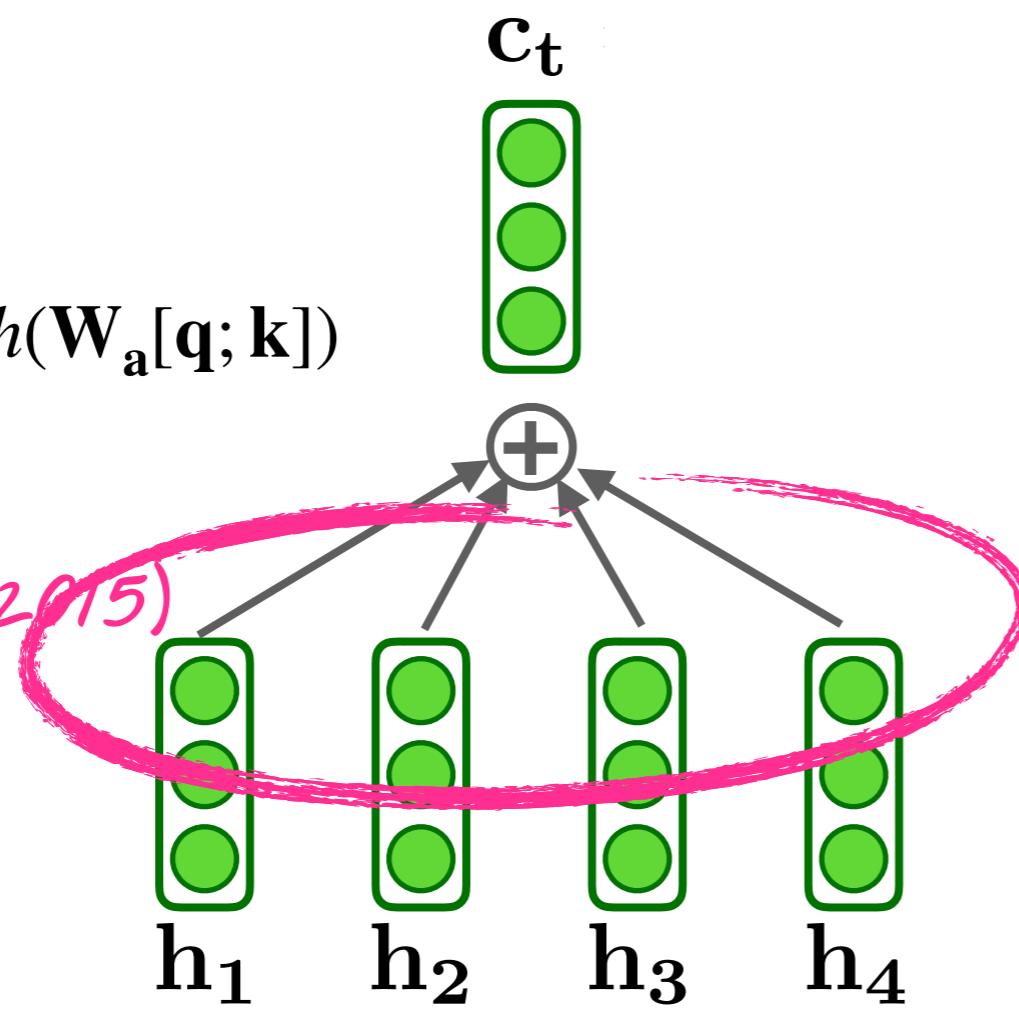
(Bahdanau et al., 2015)

# Summary: Additive attention (Bahdanau, 2015)

$$\mathbf{c}_t = \sum_{i=1}^n \alpha_{t,i} \mathbf{h}_i$$

$$a(\mathbf{q}, \mathbf{k}) = \mathbf{v}_a^T \tanh(\mathbf{W}_a[\mathbf{q}; \mathbf{k}])$$

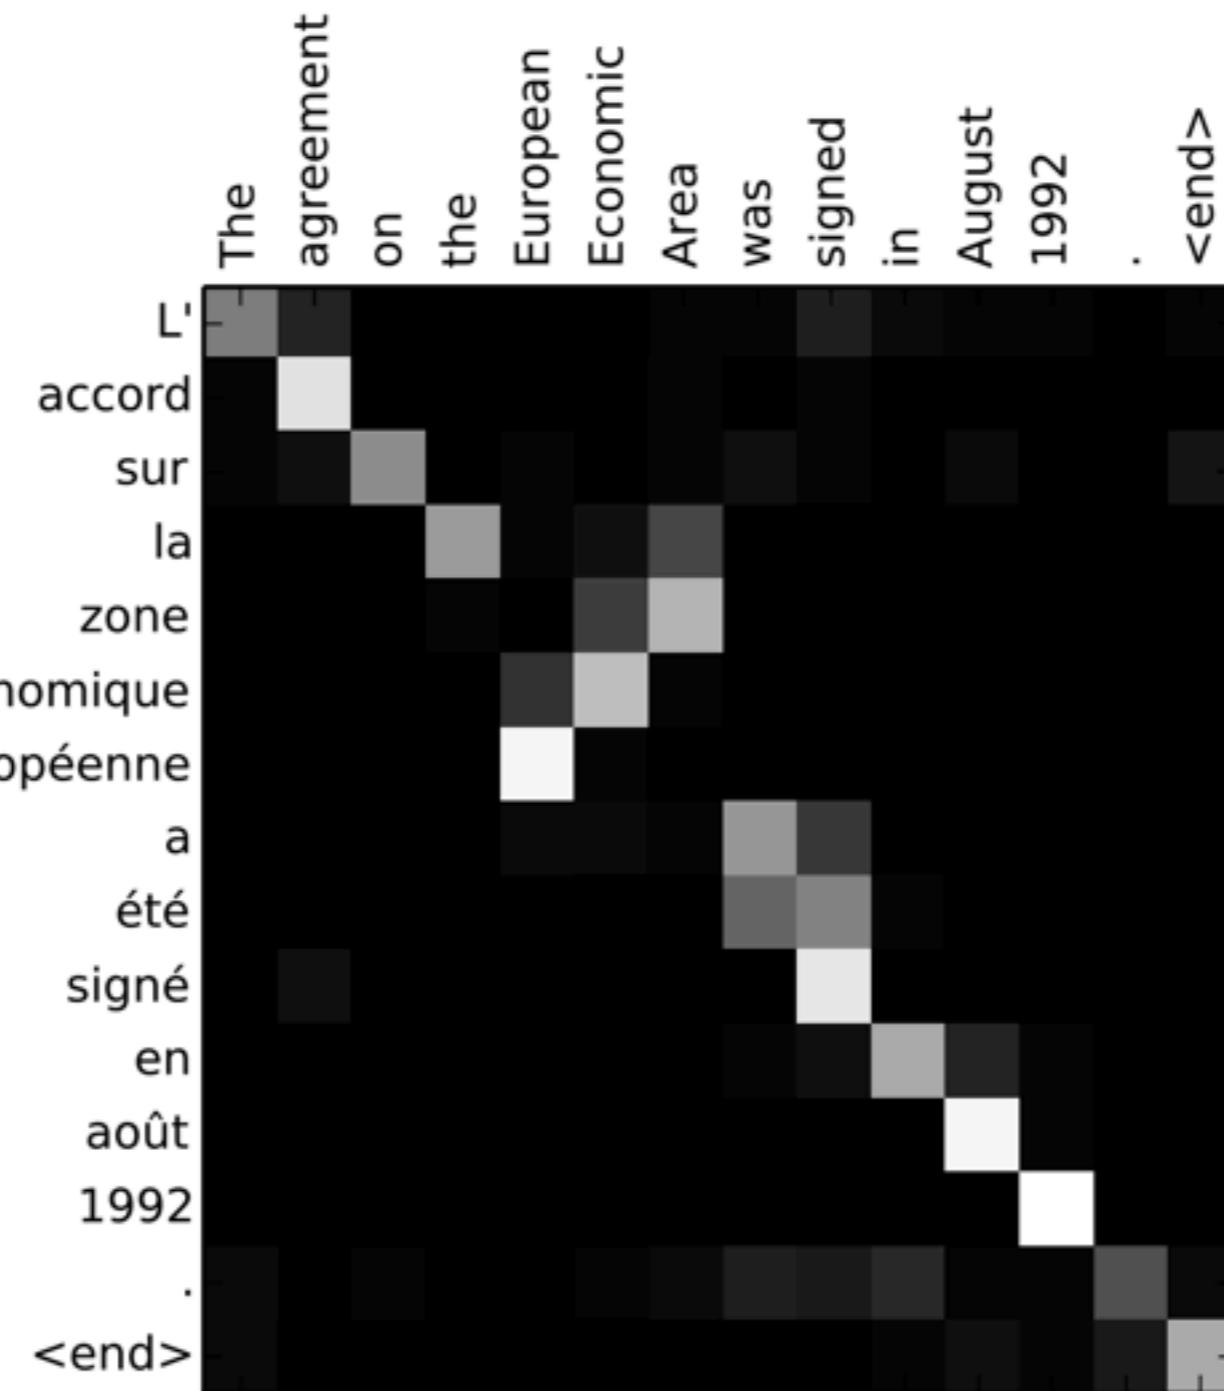
FFNN!  
(Bahdanau, 2015)



$$\alpha_1=.79 \quad \alpha_2=.1 \quad \alpha_3=.05 \quad \alpha_4=-.03$$

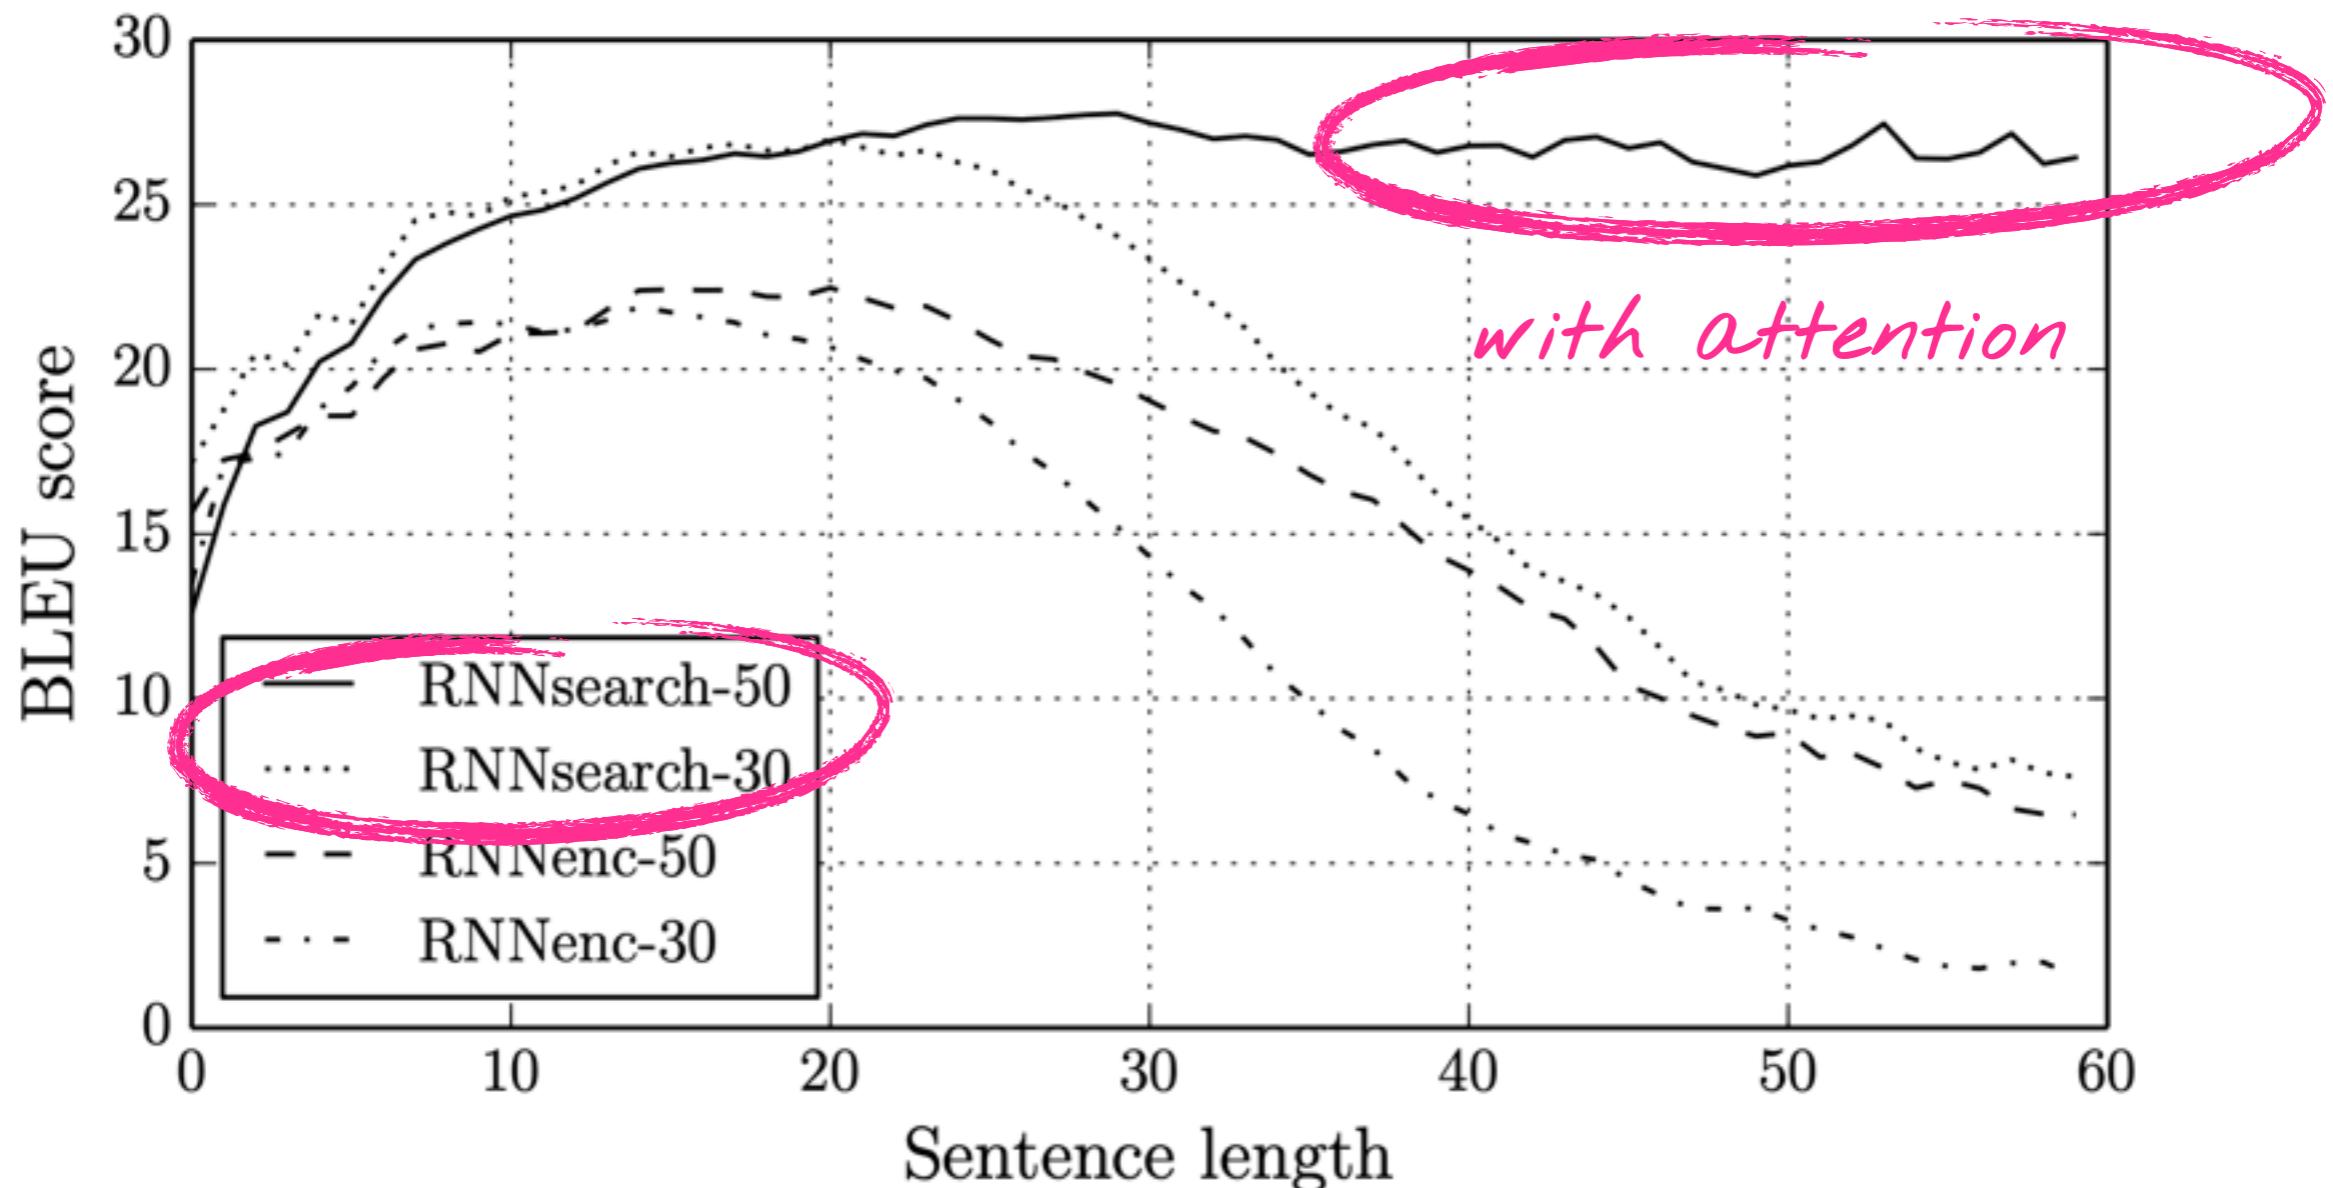
1. For each query-key pair, calculate weight  $\alpha_i$
2. Normalize via softmax
3. Combine together value vectors via weighted sum to get  $\mathbf{c}_t$
4. Use in your model in any part you like

# Alignment matrix (Bahdanau, 2015)



# Enc-dec performance deteriorates rapidly as input sentence length increases

Cho et al., (2014); Bahdanau et al. (2015)



**Different forms of  
attention are available  
(e.g., Luong et al., 2015)**

# Alignment Functions: What's $a(\mathbf{q}, \mathbf{k})$ ?

- ▶ In **Bahdanau et al.**, (2015): the alignment score function is a single **FFNN** (MLP) with a single hidden layer:
  - ▶ 
$$a(\mathbf{q}, \mathbf{k}) = \mathbf{v}_a^T \tanh(\mathbf{W}_a[\mathbf{q}; \mathbf{k}])$$
  - ▶ both  $\mathbf{v}_a$  and  $\mathbf{W}_a$  are trained with the network

# More alignment functions

- ▶ **Dot product** (Luong et al., 2015)
  - ▶  $a(\mathbf{q}, \mathbf{k}) = \mathbf{q}^T \mathbf{k}$
  - ▶ requires same size; but has no parameters!
- ▶ **Bilinear** (Luong et al., 2015)
  - ▶  $a(\mathbf{q}, \mathbf{k}) = \mathbf{q}^T W \mathbf{k}$
- ▶ **Scaled dot product** (Vaswani et al., 2017)
  - ▶ 
$$a(\mathbf{q}, \mathbf{k}) = \frac{\mathbf{q}^T \mathbf{k}}{\sqrt{|\mathbf{k}|}}$$
  - ▶ fixes problem of dot product that scale of dot product increases as dimensions get larger

A little more on  
attention

# Self-attention

- ▶ Attend to sentence itself (Cheng, Dong, Lapata, 2016)

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

**What to attend to?  
Some more examples**

# Image caption generation

- ▶ Salient parts of the image (e.g., Xu et al., 2015)



A woman is throwing a frisbee in a park.

# Character-level attention

- E.g. Rei et al., 2016

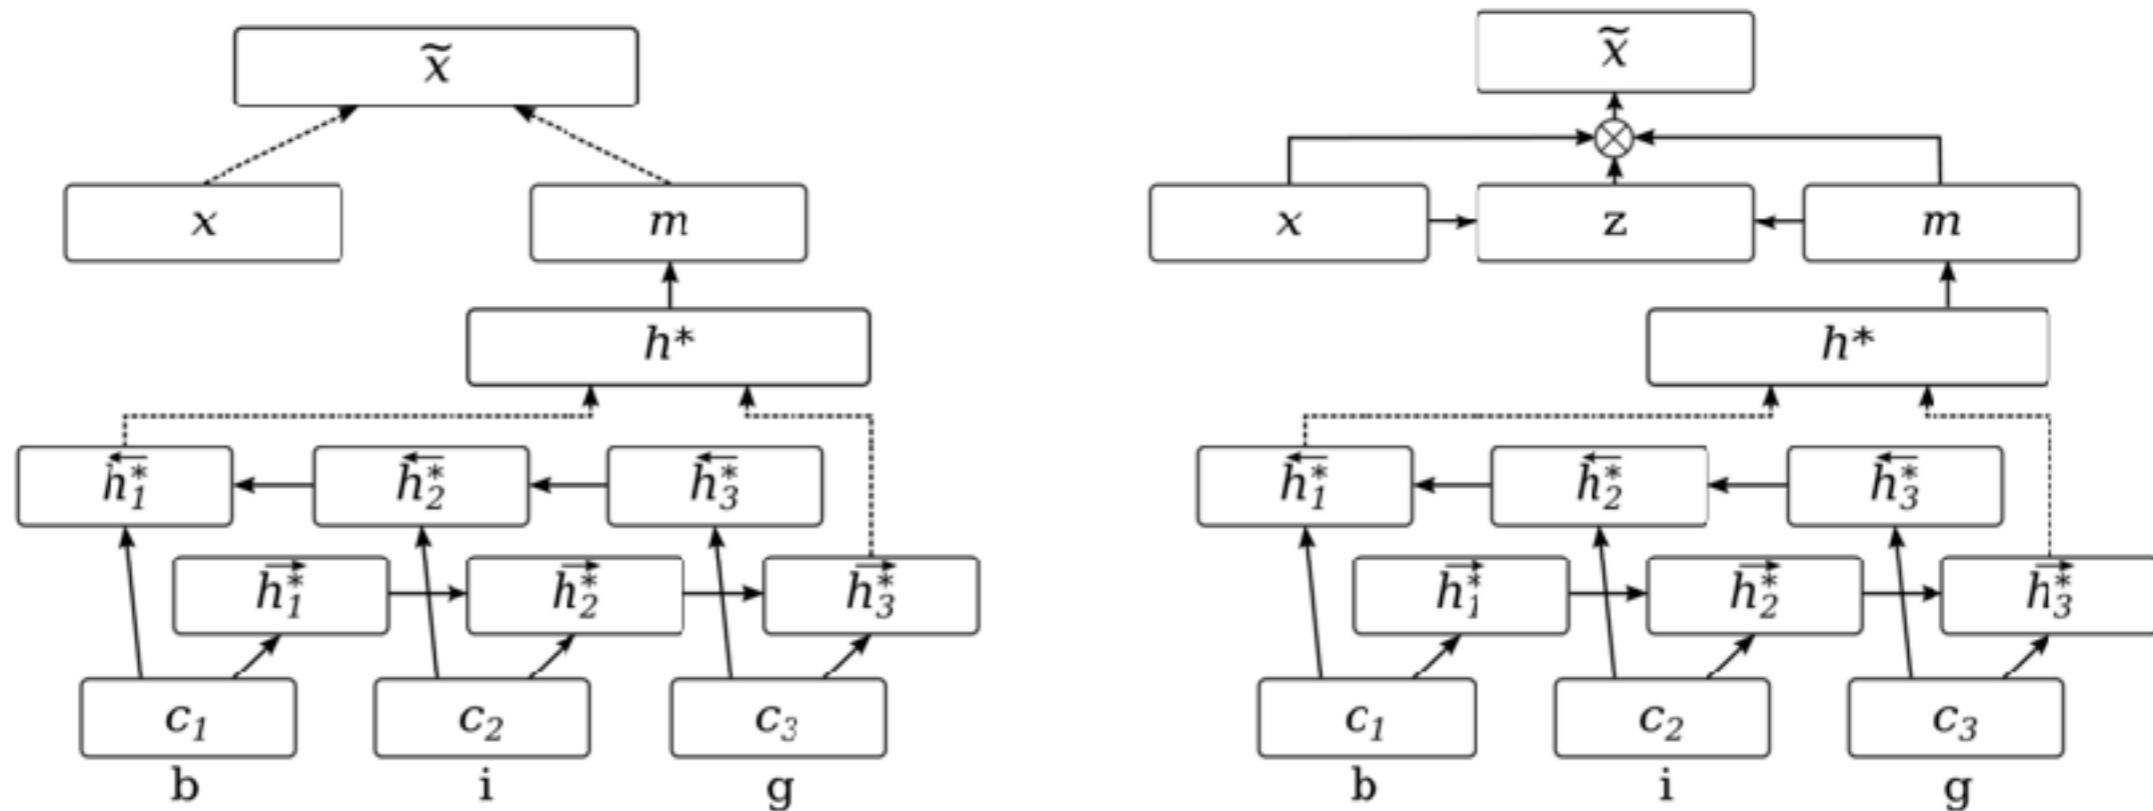


Figure 2: Left: concatenation-based character architecture. Right: attention-based character architecture. The dotted lines indicate vector concatenation.

**Attention is everywhere  
[and all you need?!] ->  
More on Monday :-)**



To summarize

# To sum up...

foundations

Motivation,  
Brief History, Overview

Back to the roots:  
Language Models

n-grams, Limitations

Feedforward NNs  
(FFNNs)

FFNN LM

representations

What's the input?  
Representations

n-hot & static  
word embeddings

ELMo

Contextualized  
Embeddings, Fine-tuning

beyond FFNNs

Convolutional Neural  
Networks (CNNs)

vanilla RNNs

Recurrent Neural  
Network (RNNs)

GRU, LSTM

Advanced RNNs,  
Decoders  
MLP, CRF

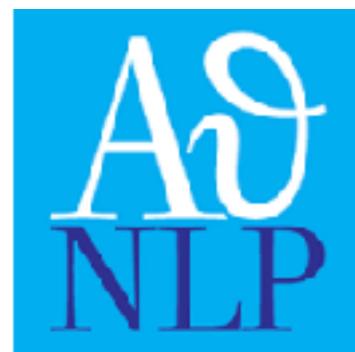
variable-size input  
& "soft" ngrams

Basis for Transformer

Attention? Attention!

# Questions? Thanks!

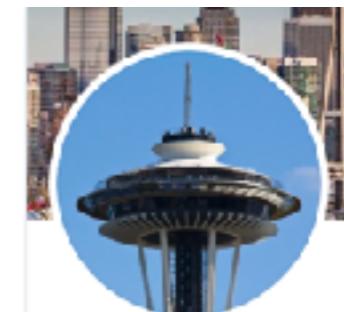
Thanks to all the  
organizers & sponsors of:



[bplank.github.io](https://bplank.github.io)  
<https://nlp.itu.dk/>

Barbara Plank  
@barbara\_plank  
ITU, Denmark

Follow us:



Research is supported by:



# References (incomplete)

- ▶ Jurasky & Martin textbook, chapter 3 (n-gram LMs), chapter 7 (neural LMs)
- ▶ Graham Neubig (2018): Language Models 4: Recurrent Neural Network Language Models
- ▶ Yoav Goldberg (2015): A Primer on Neural Network Models for Natural Language Processing
- ▶ Chris Manning & Abigail See (2018) Stanford class